

Kubernetes



Agenda

- Introduction
- Kubernetes Architecture
- Cluster
- Resources / Objects
- HPA
- Namespace
- Helm

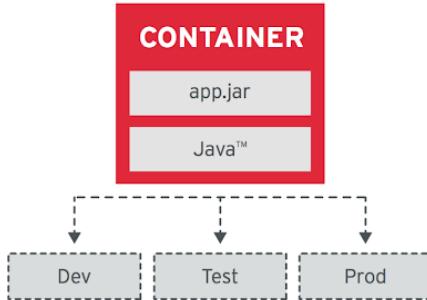


Introduction



Principles of Container-based

Image Immutability Principle



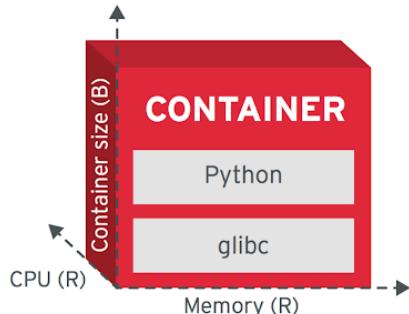
High Observability Principle



Process Disposability Principle



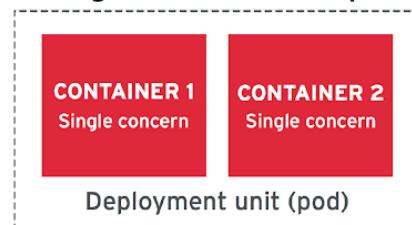
Runtime Confinement Principle



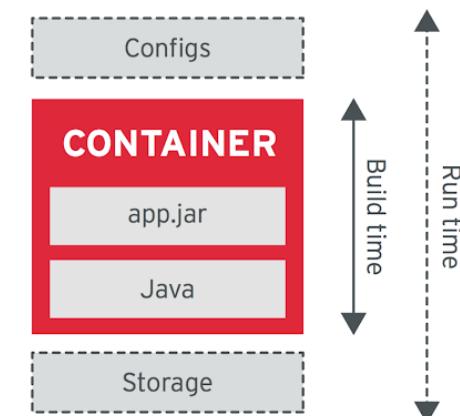
Lifecycle Conformance Principle



Single Concern Principle



Self-Containment Principle



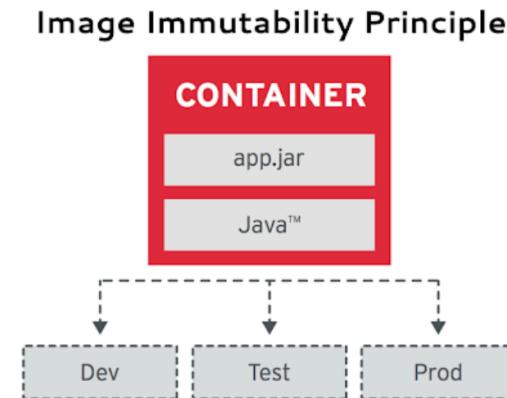
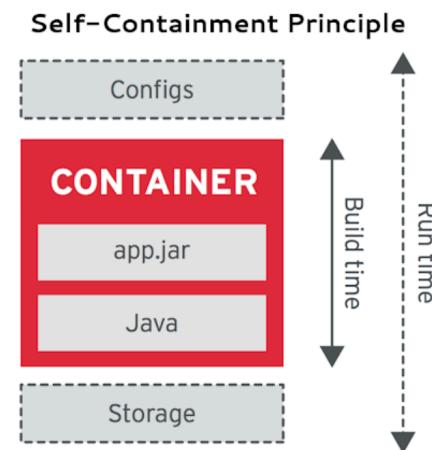
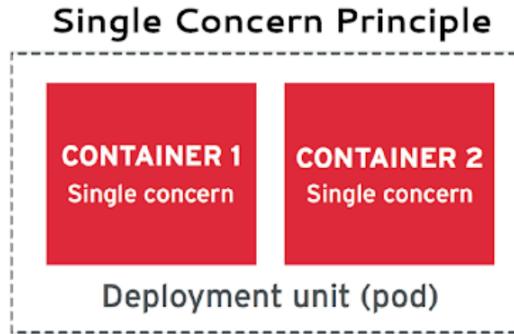
Principles of Container-based

Build Time

Single Concern: Each container addresses a single concern and does it well.

Self-Containment: A container relies only on the presence of the Linux kernel. Additional libraries are added when the container is built.

Image Immutability: Containerized applications are meant to be immutable, and once built are not expected to change between different environments.



Principles of Container-based

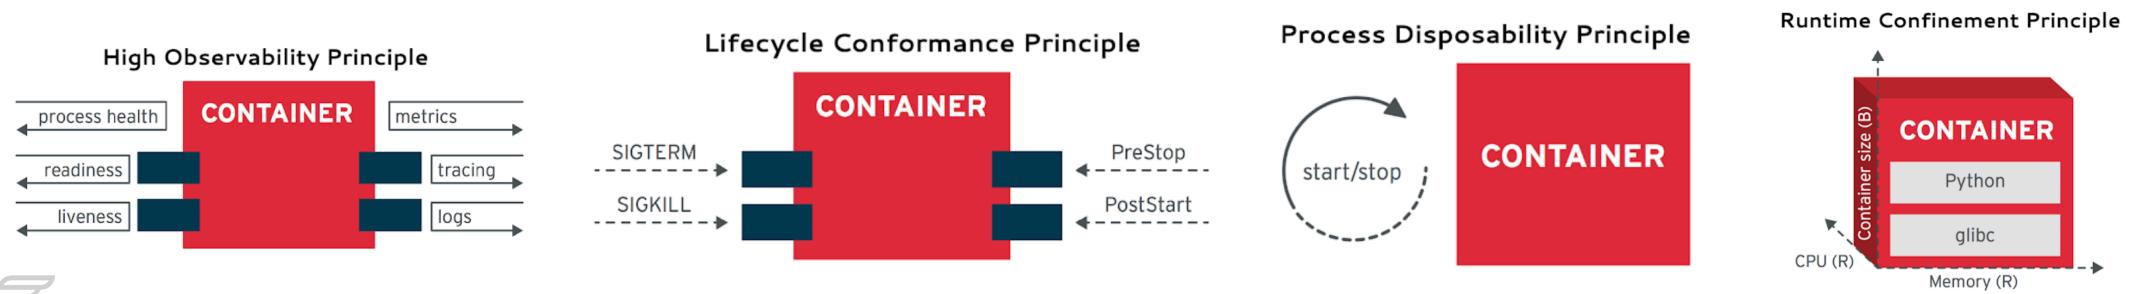
Run Time

High Observability: Every container must implement all necessary APIs to help the platform observe and manage the application in the best way possible.

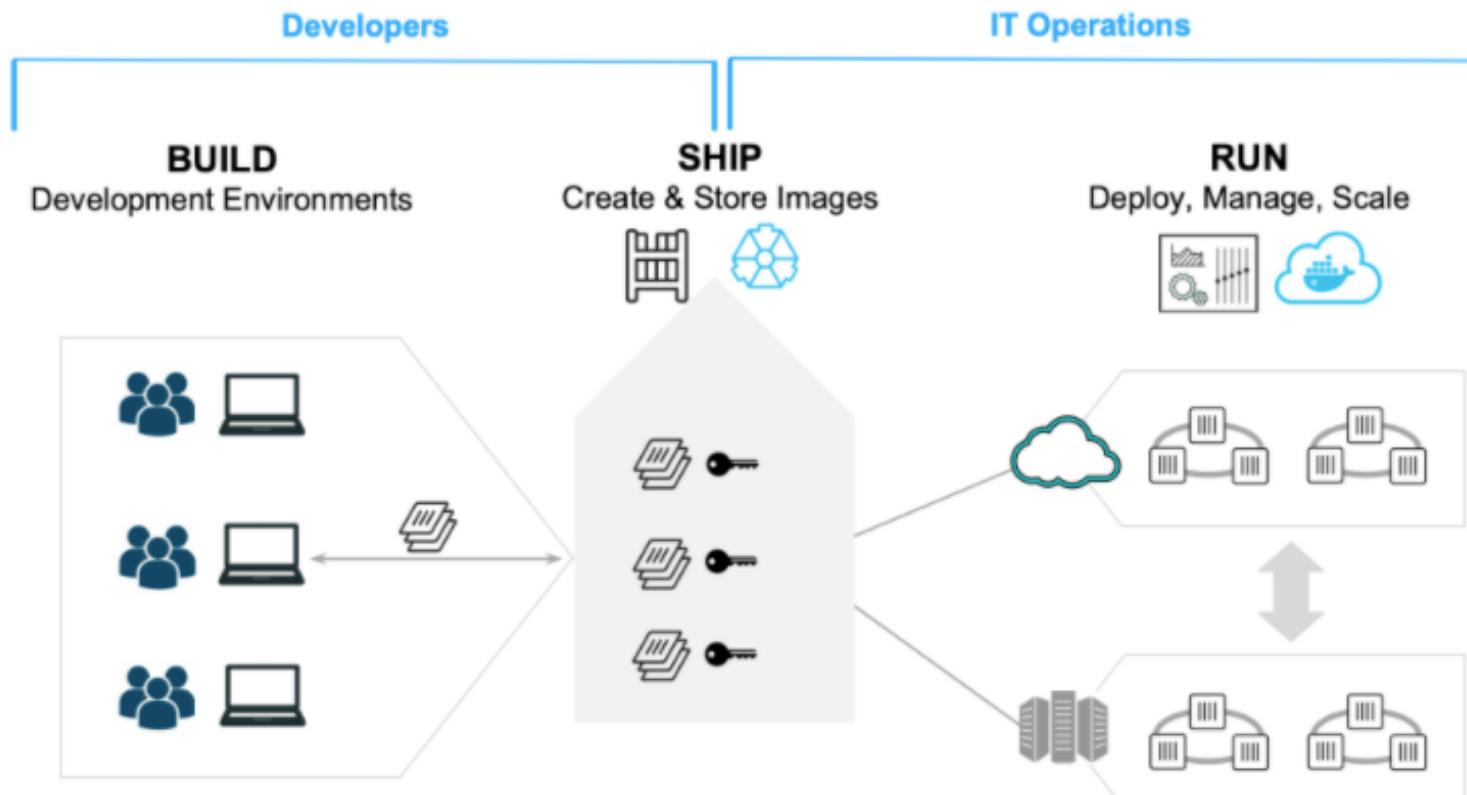
Lifecycle Conformance: A container must have a way to read events coming from the platform and conform by reacting to those events.

Process Disposability: Containerized applications must be as ephemeral as possible and ready to be replaced by another container instance at any point in time.

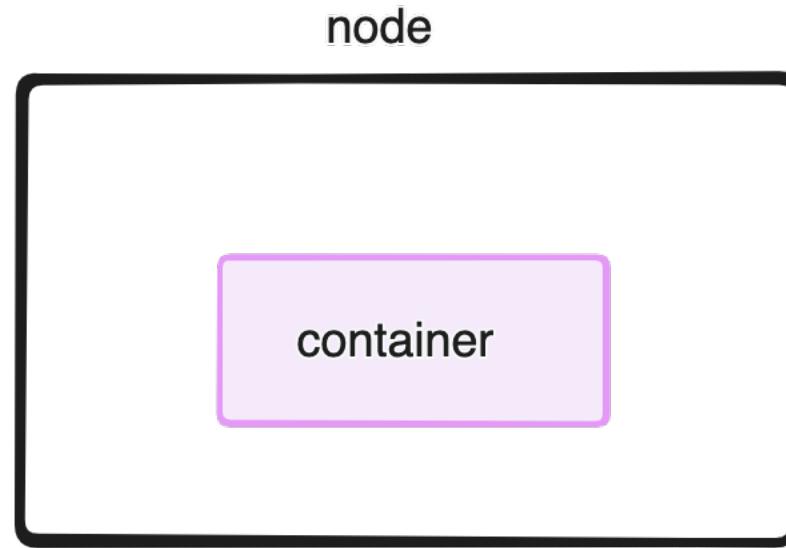
Runtime Confinement: Every container must declare its resource requirements and restrict resource use to the requirements indicated.



Using Docker: Build, Ship, Run



Docker standalone server

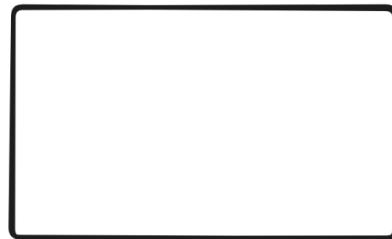


Docker multi server

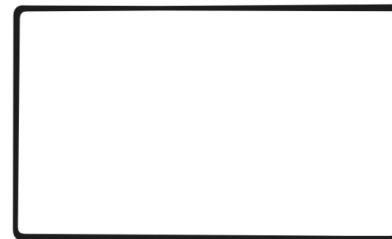
container

???

node



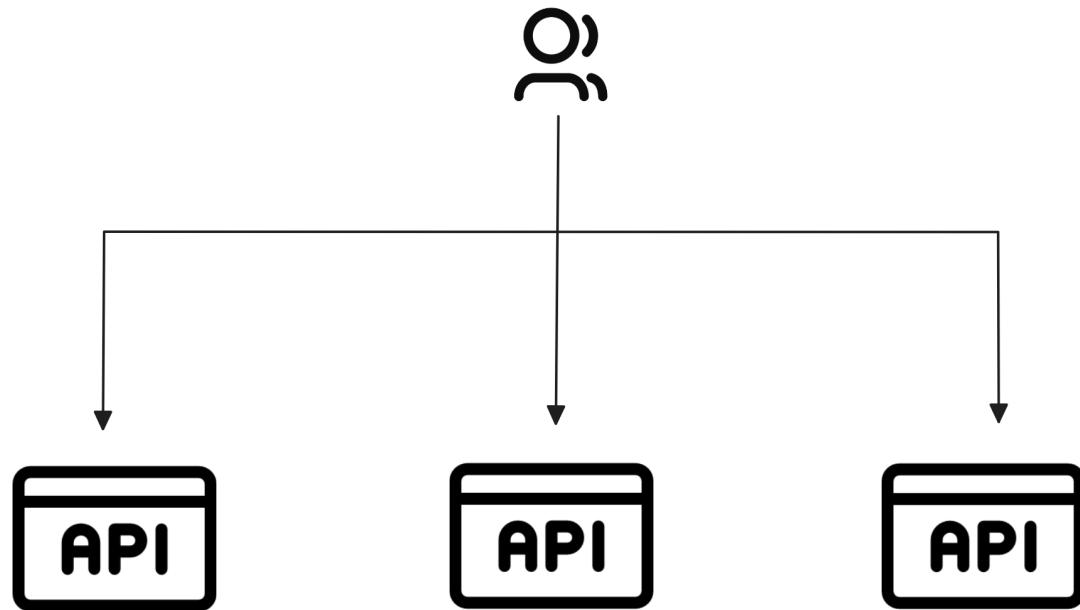
node



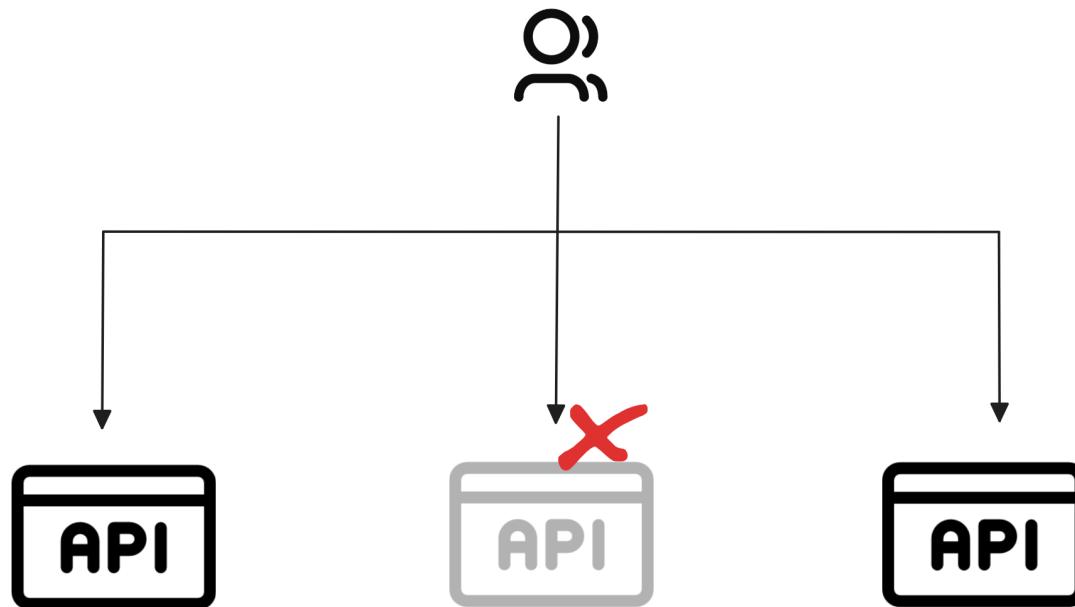
node



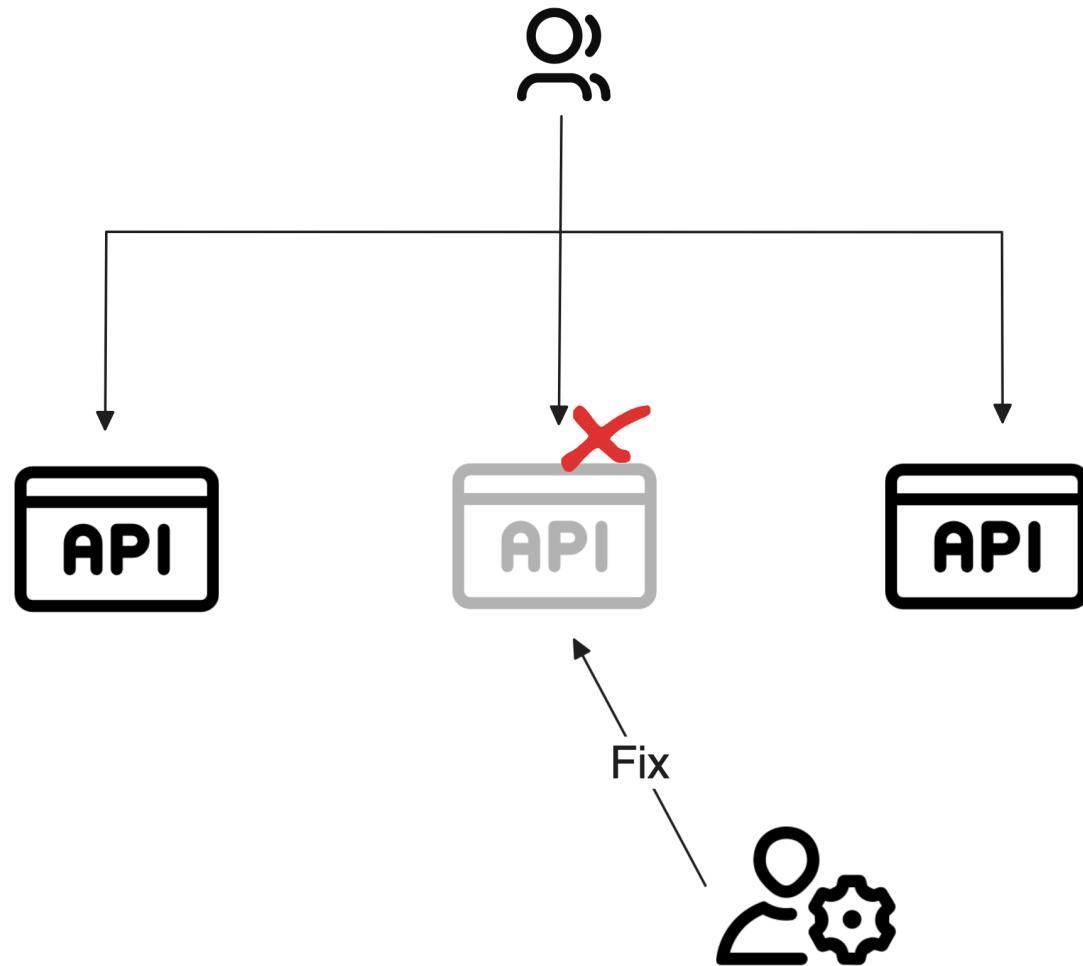
Applications problem



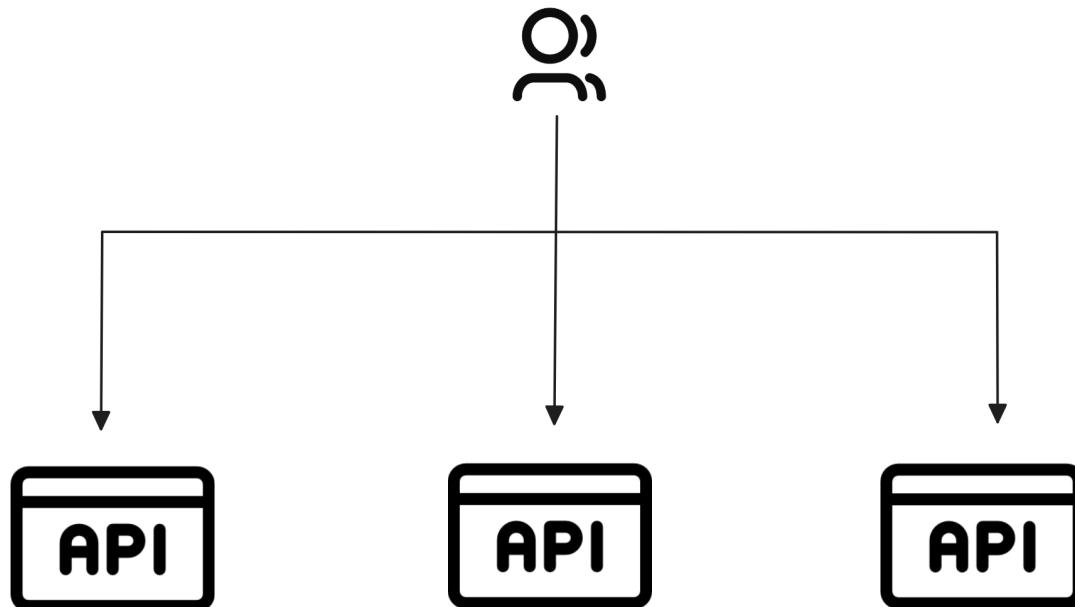
Applications problem



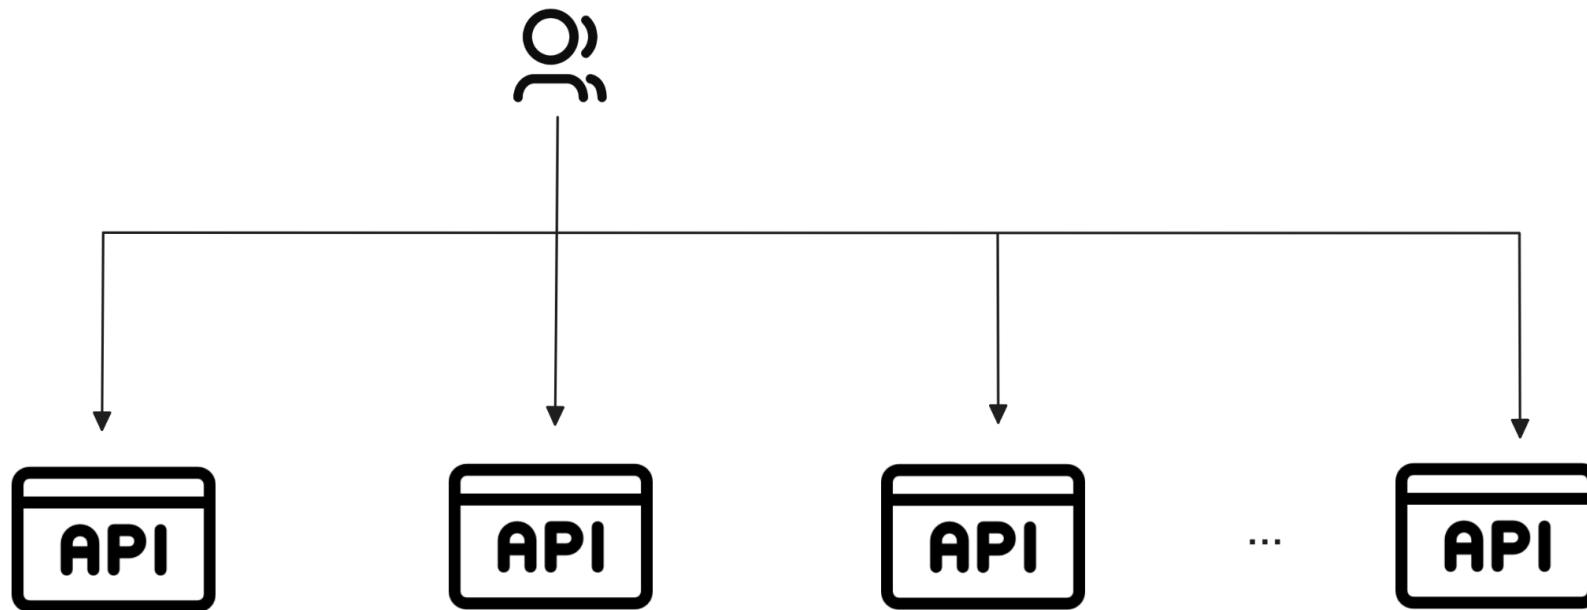
Applications problem



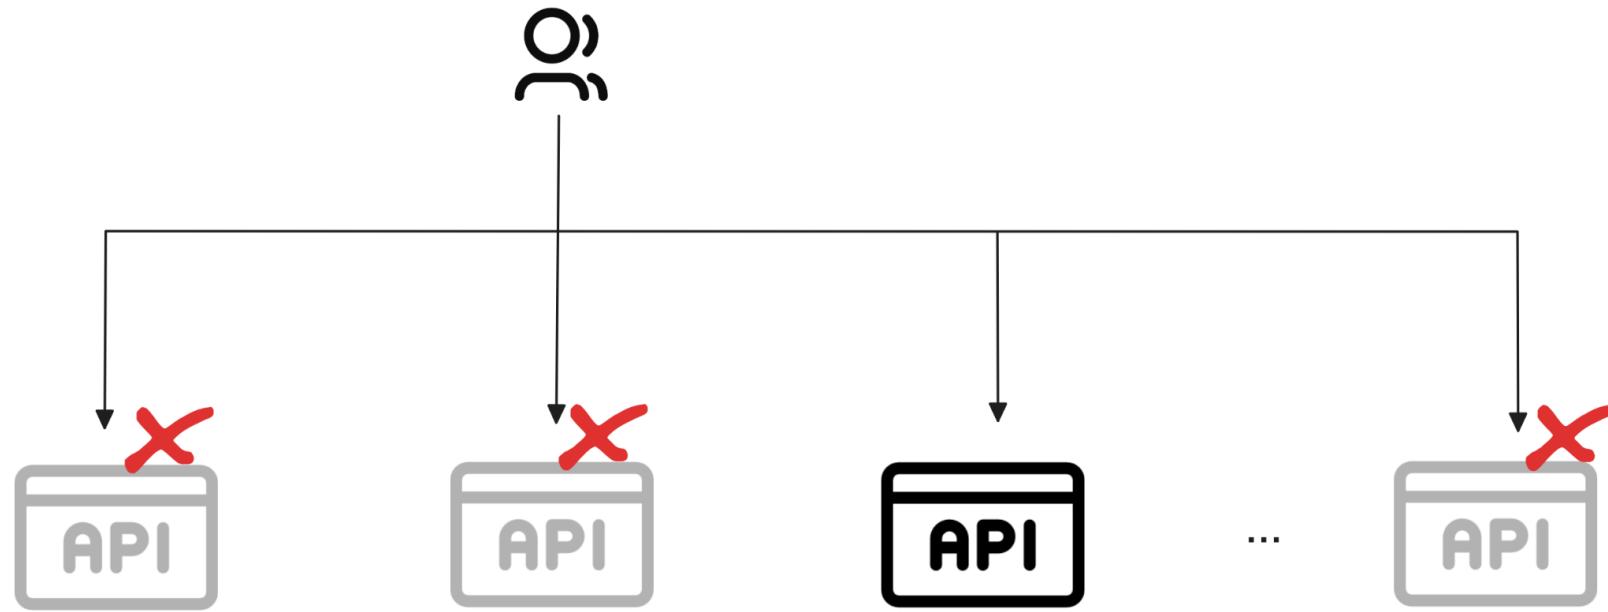
Applications problem



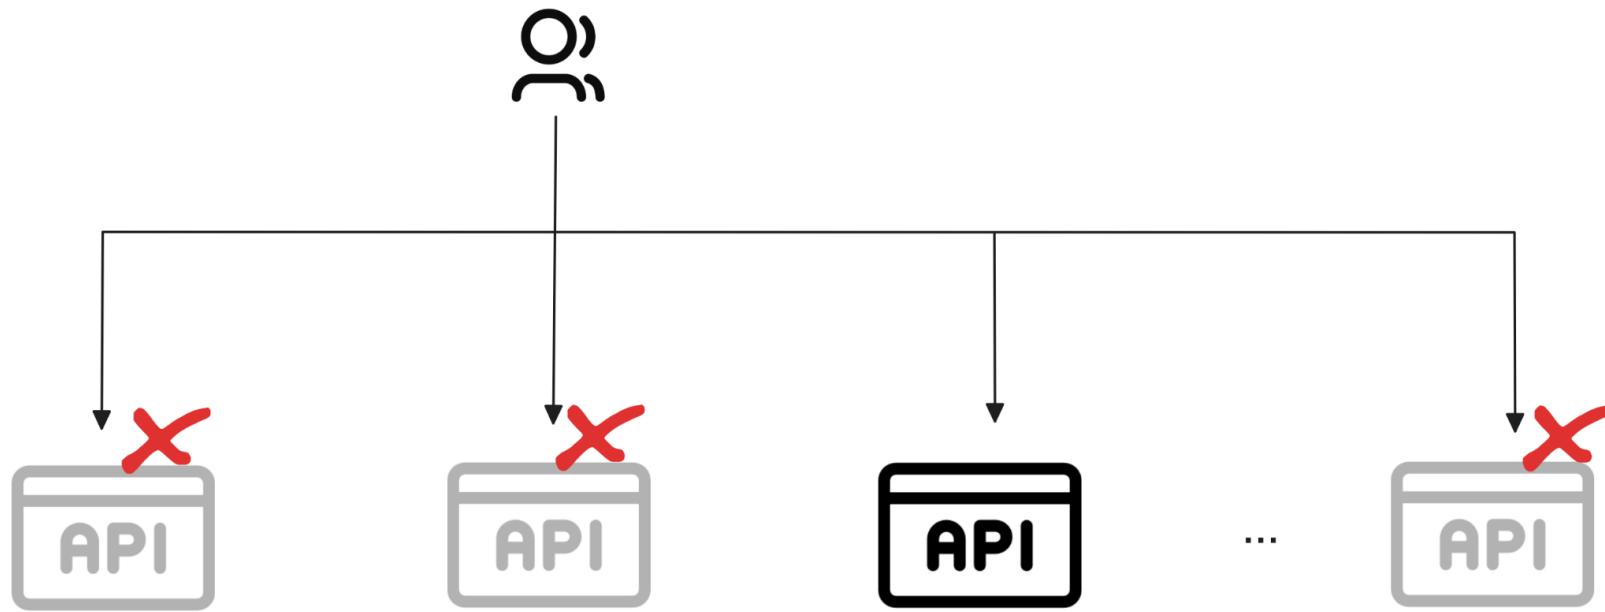
Applications problem



Applications problem



Applications problem



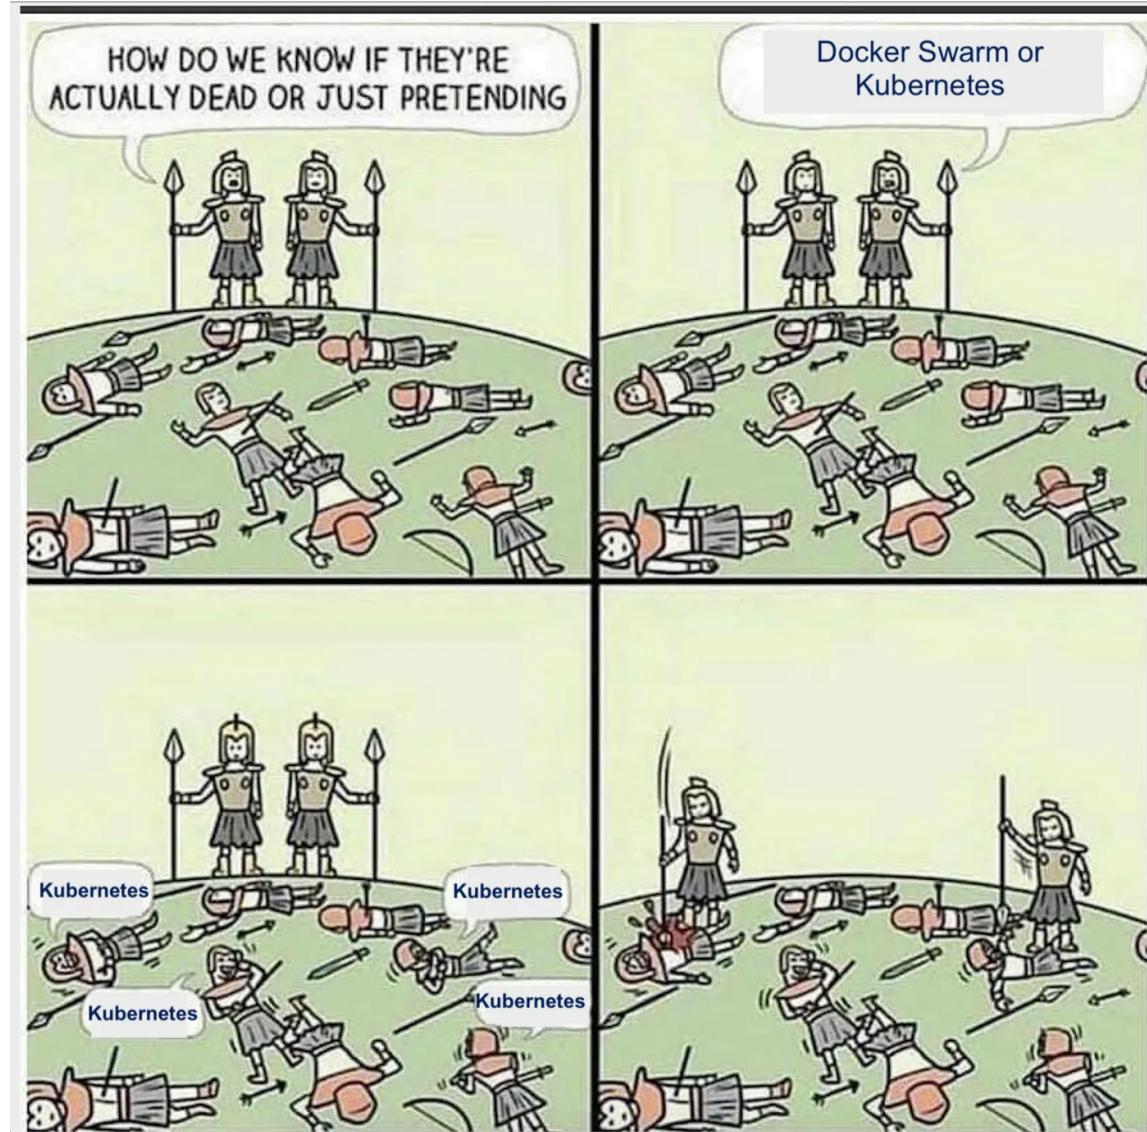
Fix !!!!



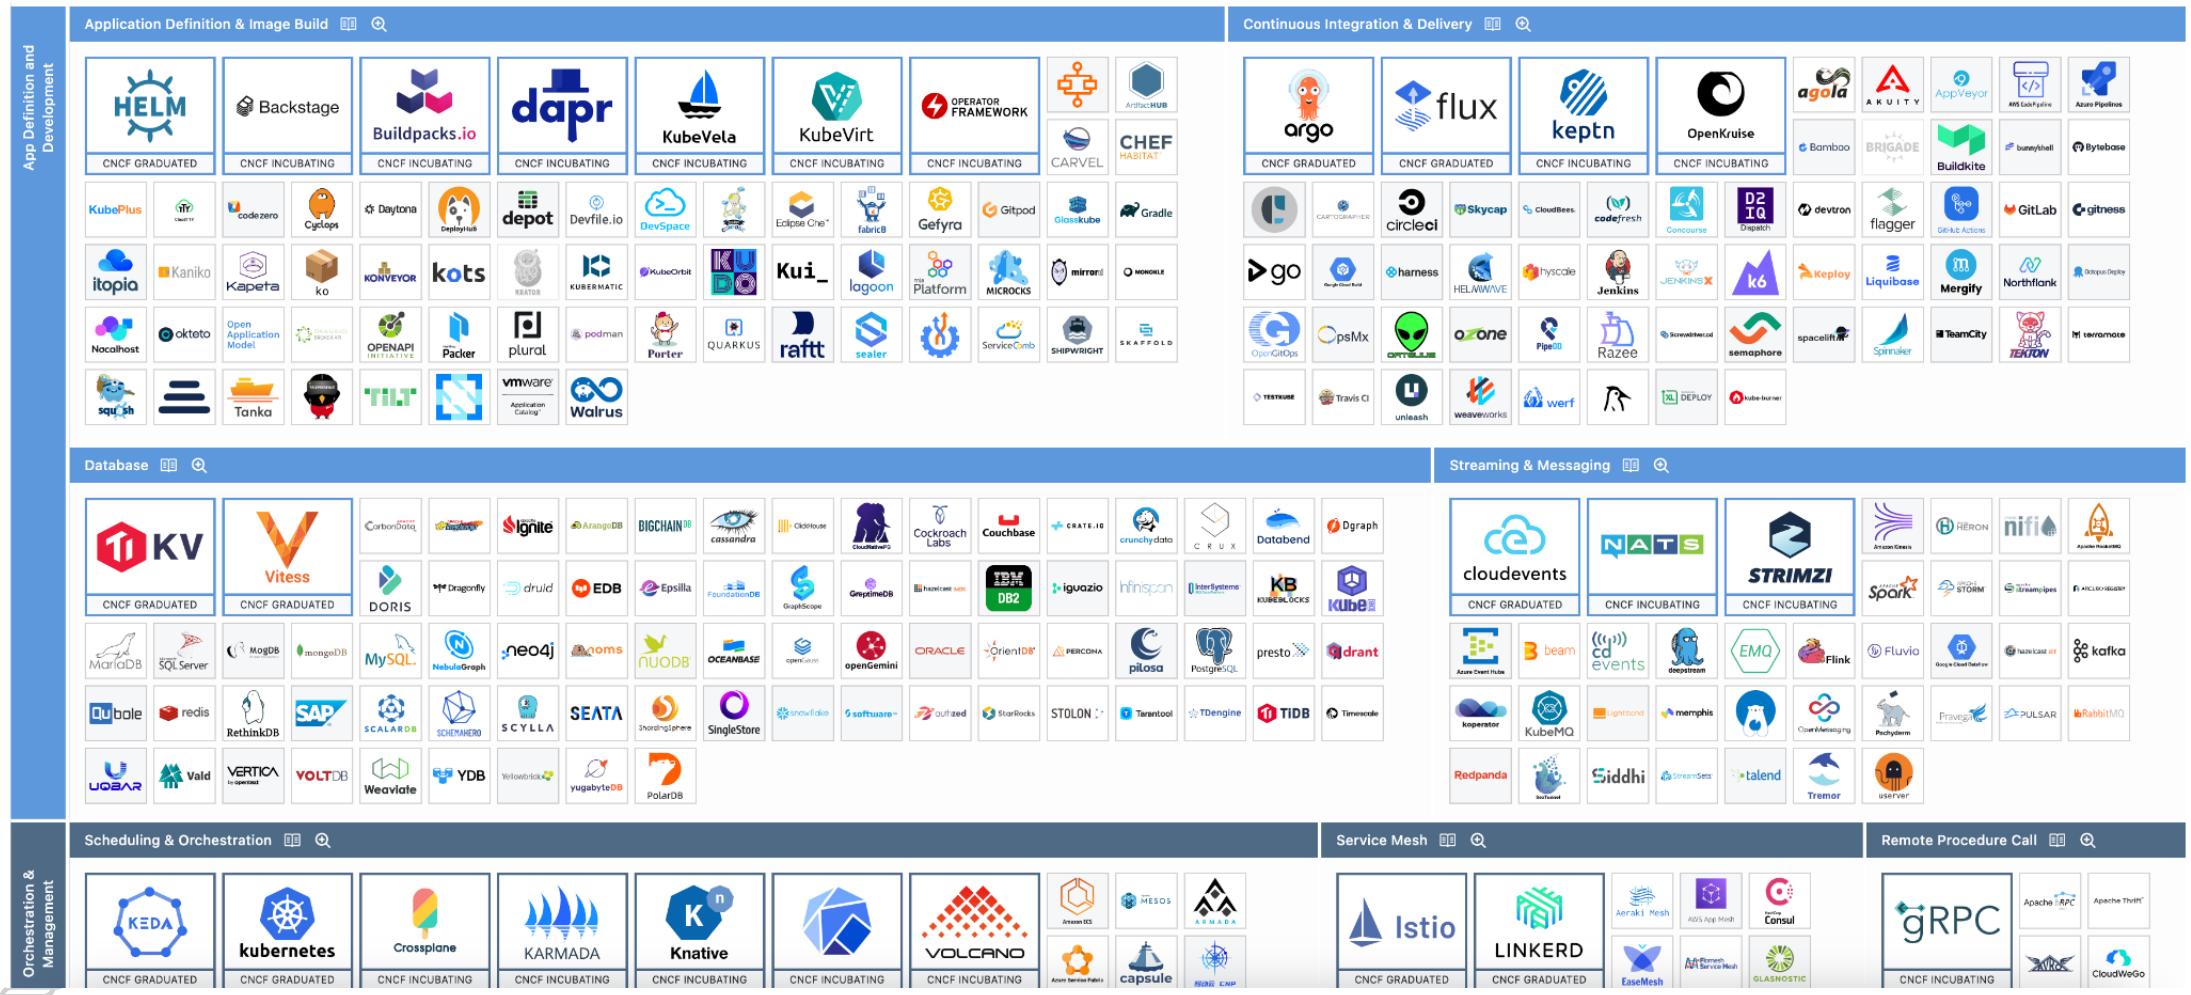
What is Kubernetes ?

- Kubernetes in Greek means the Helmsman
- Also known as k8s
- Inspired and informed by Google's experiences and internal systems
- 100% opensource, written in GO
- Container orchestrator, runs and manages containers
- Kubernetes v1.0 released on July 2015, join CNCF on March 2016
- Container choose Docker, Container Orchestrator choose Kubernetes





CNCF Landscape



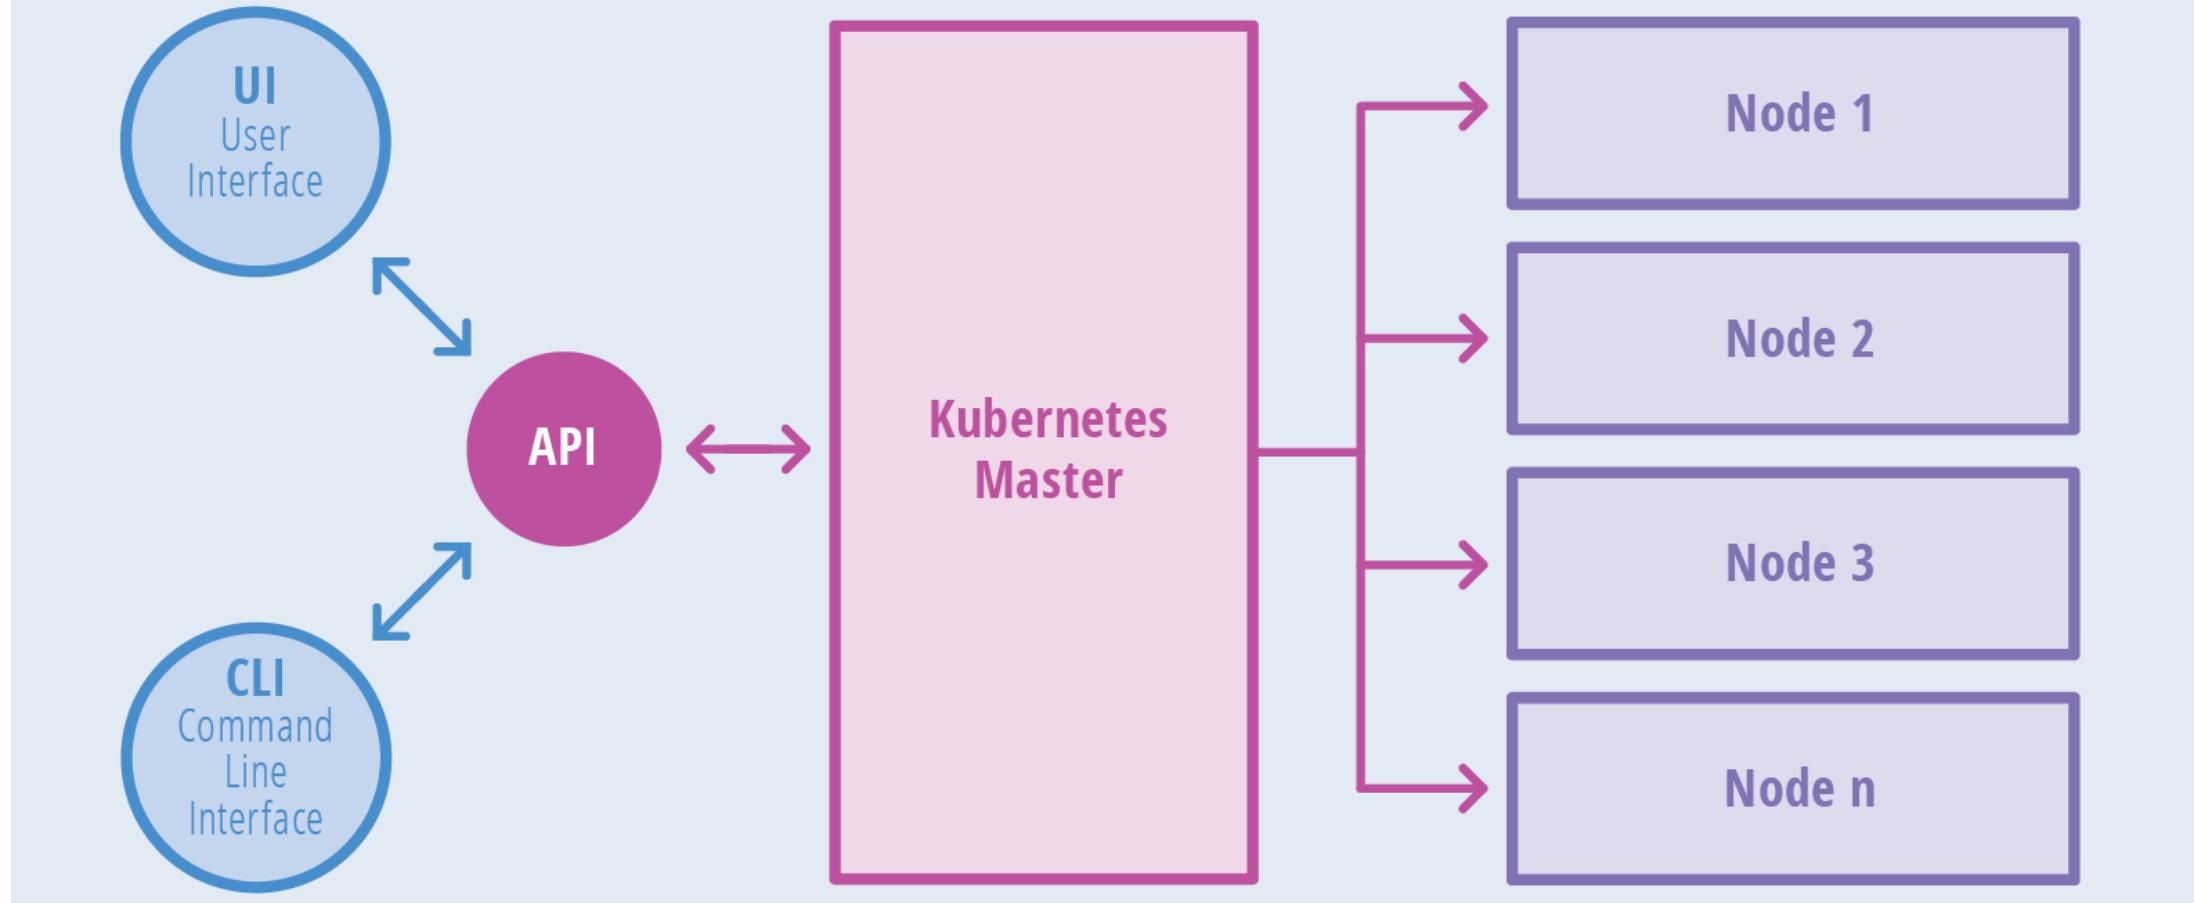
The features of Kubernetes

- Automatic bin packing
- Self healing
- Automated rollout & rollback
- Service Discovery & load Balancing
- Horizontal scaling
- Declarative Configuration



Kubernetes Architecture

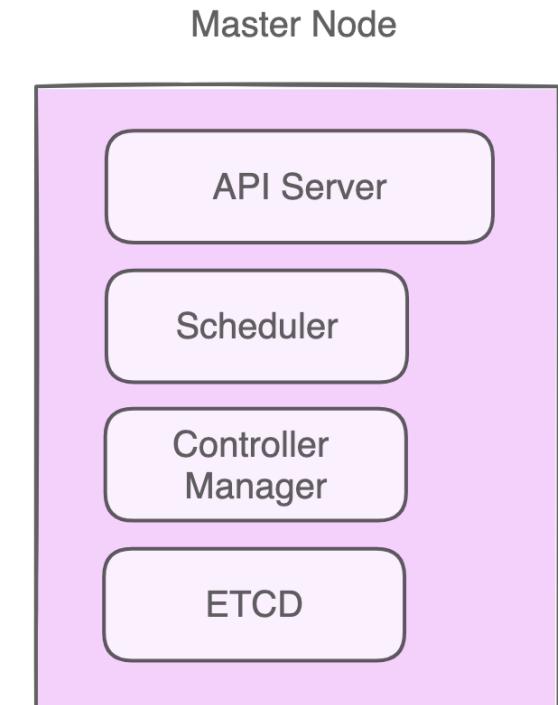
Kubernetes Architecture



Kubernetes Architecture

Master node

- Responsible for managing cluster
- Access master node via CLI, UI or API
- For fault tolerance, can be more than 1 master in cluster
- Include 4 components: API Server, Scheduler, Controller manager and ETCD

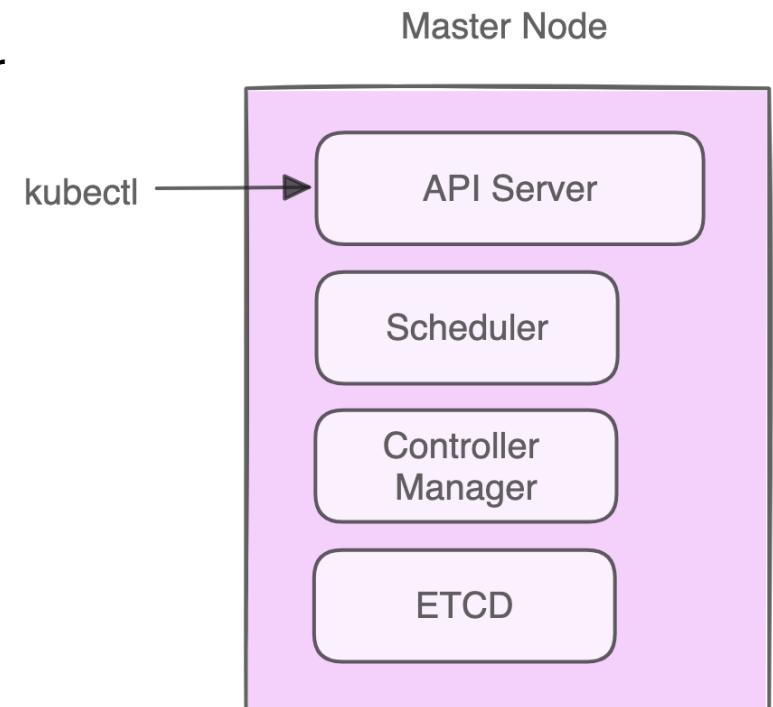


Kubernetes Architecture

Master node

API Server

- Centralized component where all the cluster components are communicated
- Execute and validate from user command

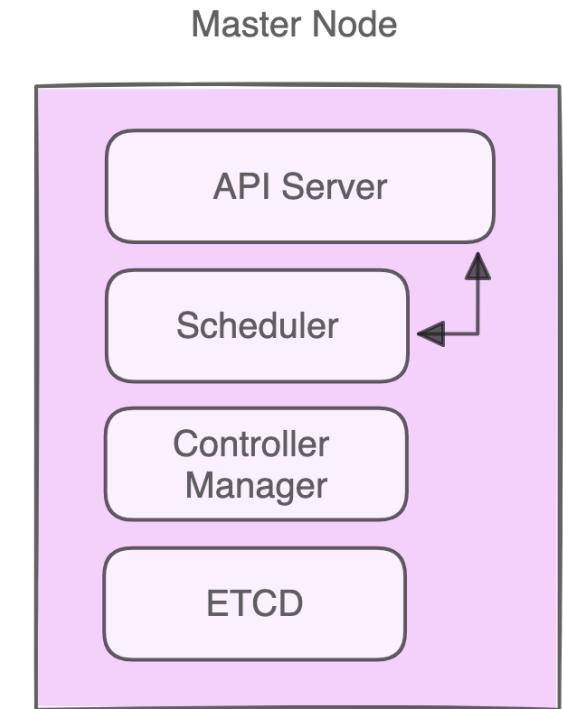


Kubernetes Architecture

Master node

Scheduler

- Assigning the application to worker node
- Automatically detect which pod to place on which node based on all factors

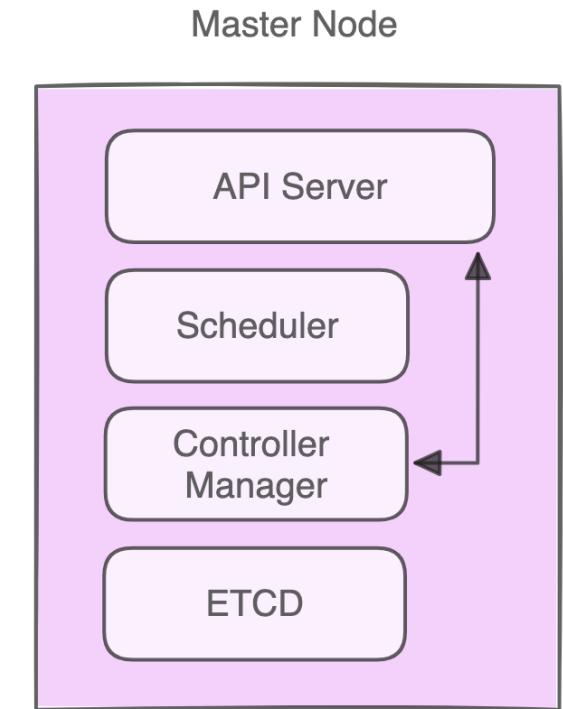


Kubernetes Architecture

Master node

Controller Manager

- The Controller Manager maintains the cluster
- Handles node failures, replicating components, maintaining the correct number of pods, etc.

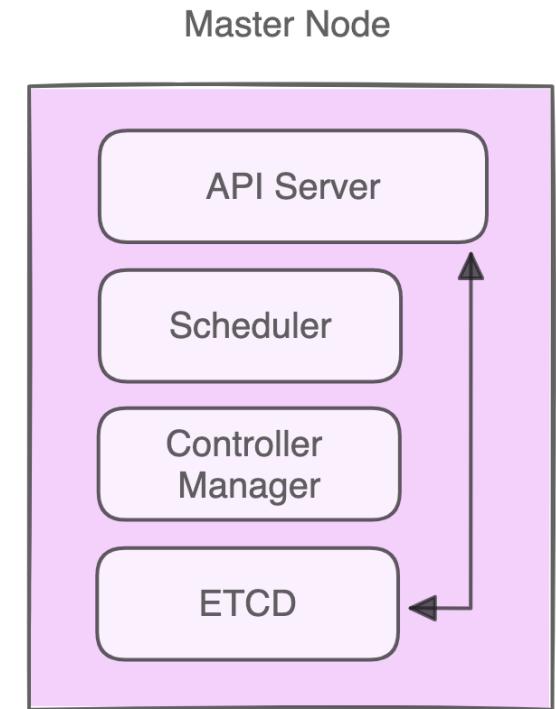


Kubernetes Architecture

Master node

ETCD

- ETCD is a distributed reliable key-value store used by Kubernetes to store all data used to manage the cluster
- When you have multi master node in cluster, ETCD stores all information in a distributed manner
- Application data doesn't store in ETCD

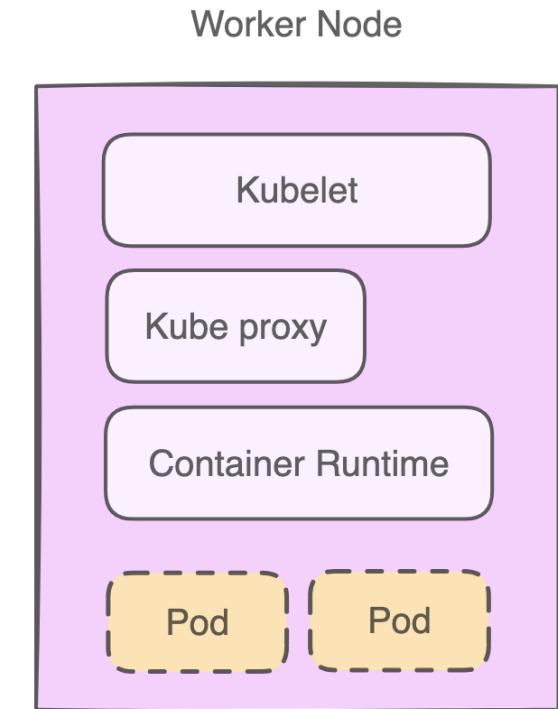


Kubernetes Architecture

Worker node

- Responsible about applications
- Include 3 components: **Container Runtime, Kubelet and Kube**

Proxy

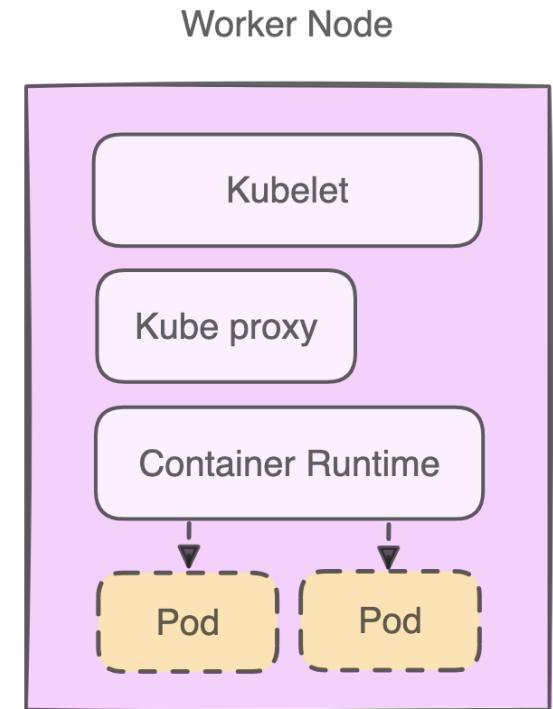


Kubernetes Architecture

Worker node

Container Runtime

- Container runtime which runs the containers like Docker, crio or containers
- Pulling the images and run the containers

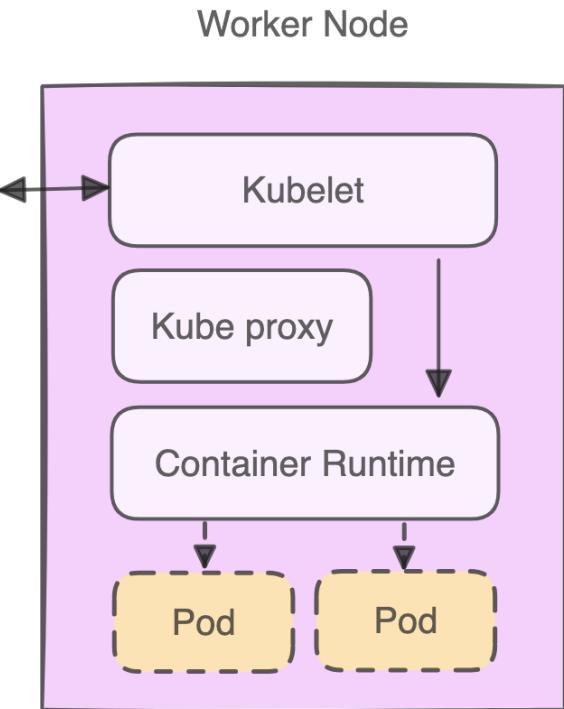


Kubernetes Architecture

Worker node

Kubelet

- Interacting with the master node
- Contact the container runtime for create pod depend on Po
- Provide health information of the worker node

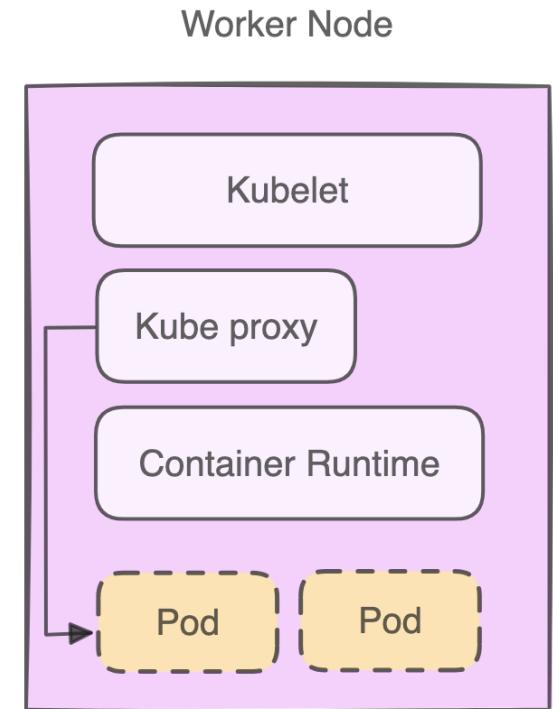


Kubernetes Architecture

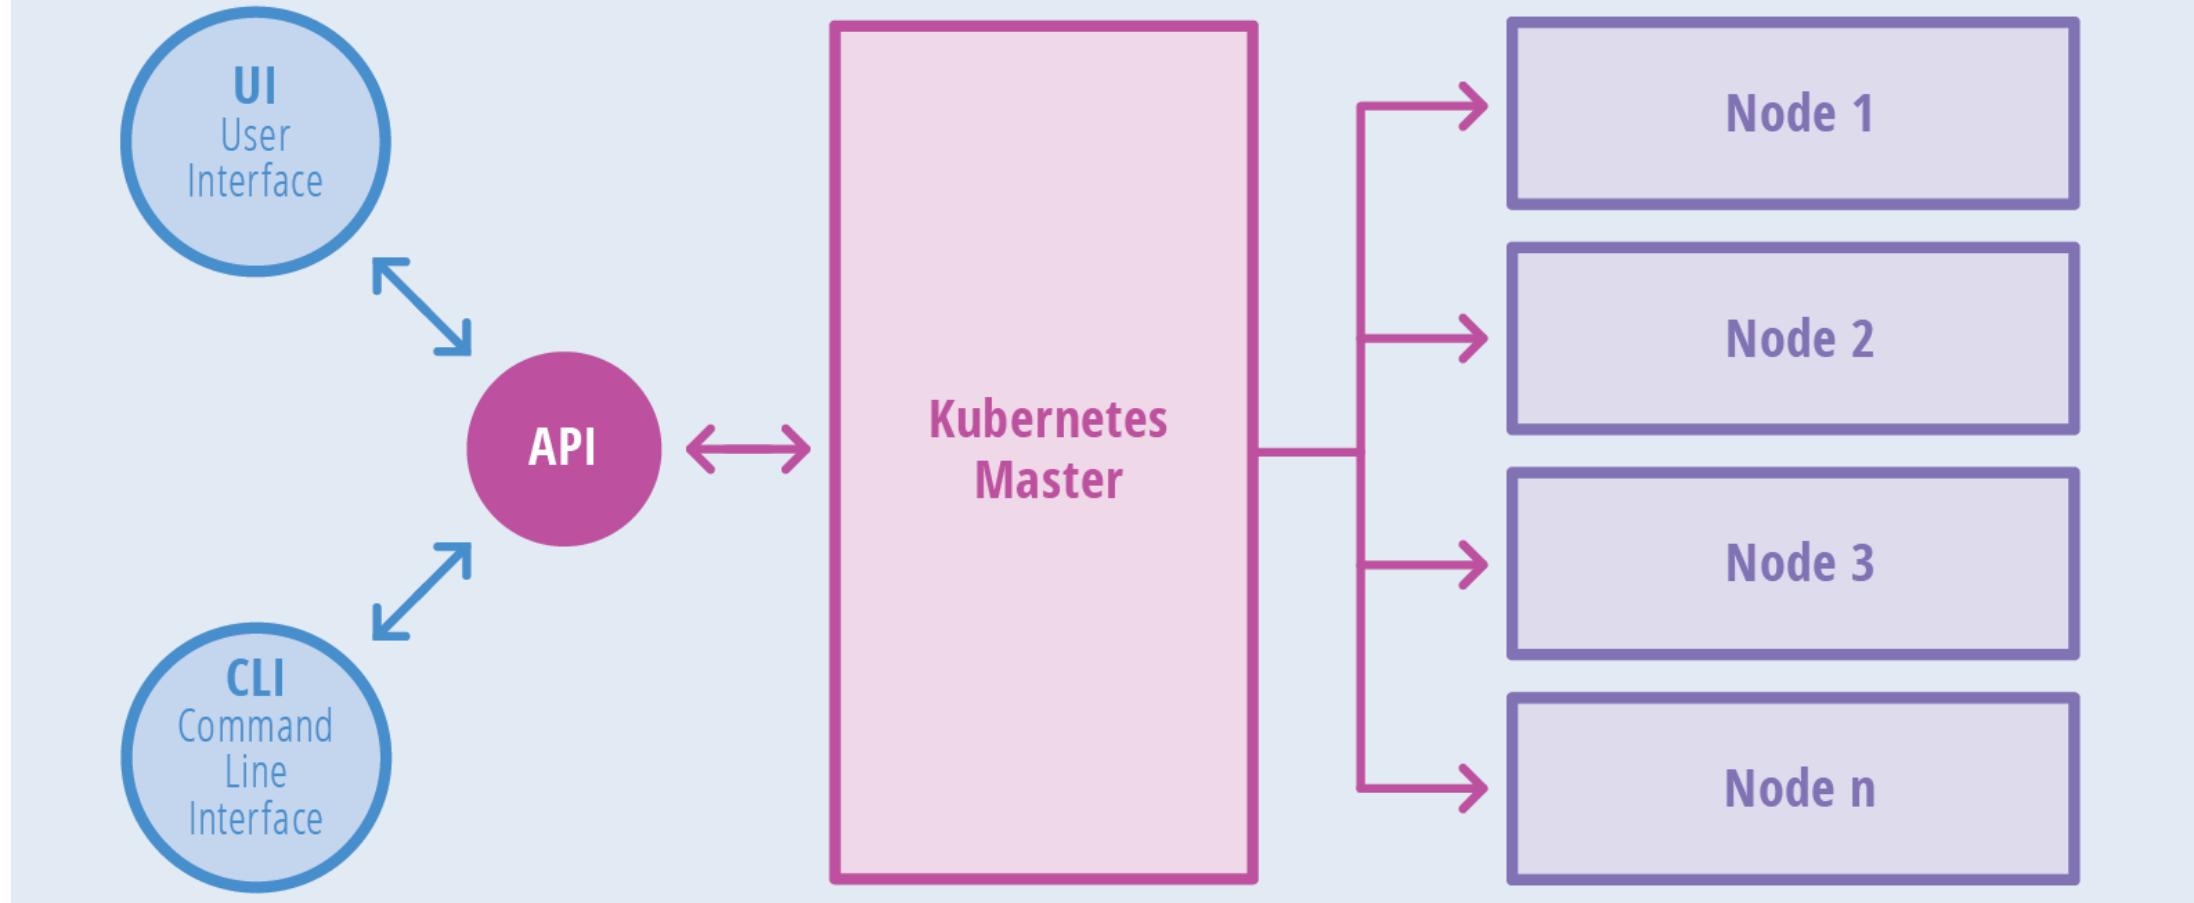
Worker node

Kube Proxy

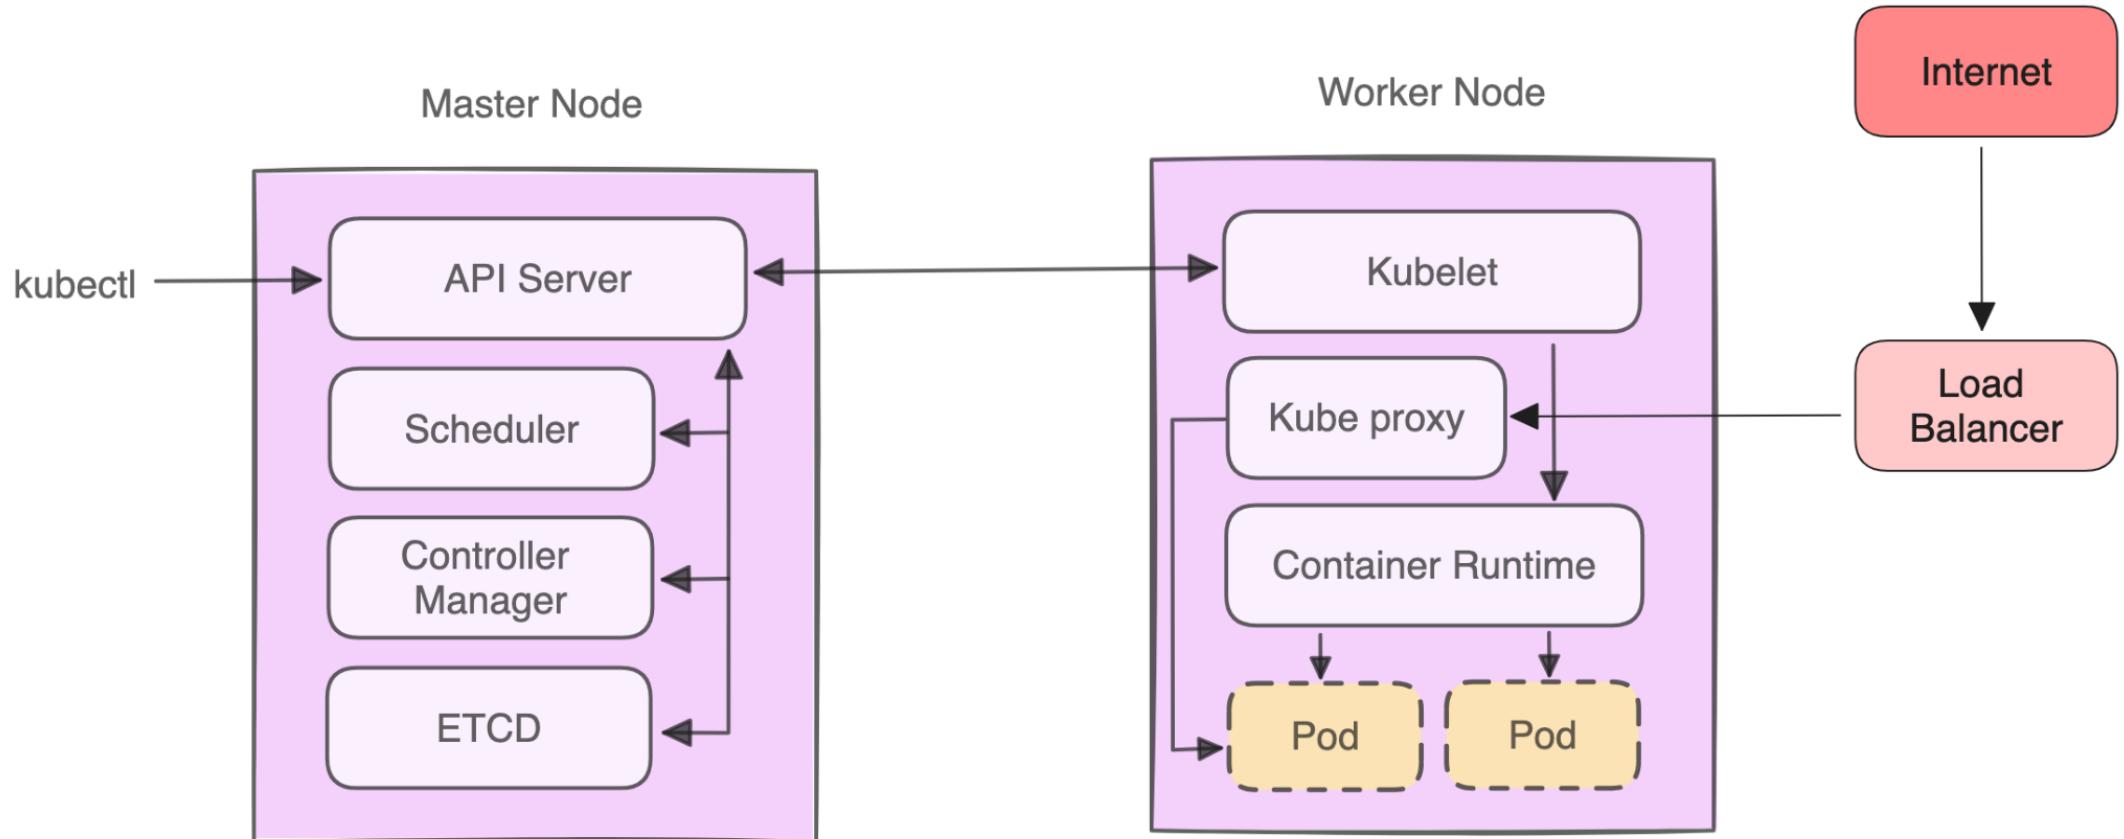
- Responsible for ensuring network traffic is routed properly to internal and external services
- Manage route for reduce latency time for example pod will communicate with another pod on same node first



Kubernetes Architecture



Kubernetes Architecture



Kubectl



Kubectl

- Command line
- Control the Kubernetes cluster manager
- CRUD Kubernetes resources



Kubectl example

- \$ kubectl apply
- \$ kubectl create
- \$ kubectl get
- \$ kubectl describe
- \$ kubectl delete



Cluster



Cluster

- On Cloud
- On Premise



Cluster on cloud



Amazon **EKS**



Google Kubernetes Engine



Azure Kubernetes Service (**AKS**)

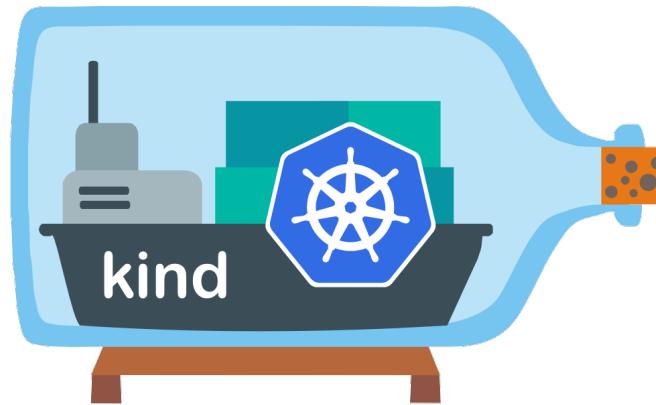
 Alibaba Cloud



Cluster on premise



Cluster on local development



Cluster comparing

Feature	Kubernetes on Cloud	Kubernetes On-Premise
Setup & Maintenance	<ul style="list-style-type: none">- Easy to set up with automated updates and patches- Minimal operational overhead	<ul style="list-style-type: none">- Requires manual setup- Ongoing maintenance and upgrades are the user's responsibility
Cost	<ul style="list-style-type: none">- High depend cloud provider- Costs tied to usage- No upfront infrastructure investment	<ul style="list-style-type: none">- Lower costs- High initial investment for hardware and infrastructure
Scalability	<ul style="list-style-type: none">- Highly scalable with near-infinite resources.- Automatic scaling options available	<ul style="list-style-type: none">- Limited by on-premise hardware resources.- Scaling may require additional hardware and setup.
Disaster Recovery	<ul style="list-style-type: none">- Built-in disaster recovery and backups.	<ul style="list-style-type: none">- Must design and implement your own
Customization	<ul style="list-style-type: none">- Limited by cloud provider's offerings and configurations.	<ul style="list-style-type: none">- Full control over configurations and infrastructure.



Cluster comparing

Feature	Kubernetes on Cloud	Kubernetes On-Premise
Security	- Security managed by cloud providers	- Full control over security configurations
Networking	- Advanced networking features available (e.g., VPCs, load balancers). - Simplified setup with integration to other cloud services	- Full control over network configurations. - Requires in-house expertise for complex networking setups
Upgrades & New Features	- Automatic access to the latest features and updates	- Full control over when and how to upgrade - Delayed to the latest features unless manually updated

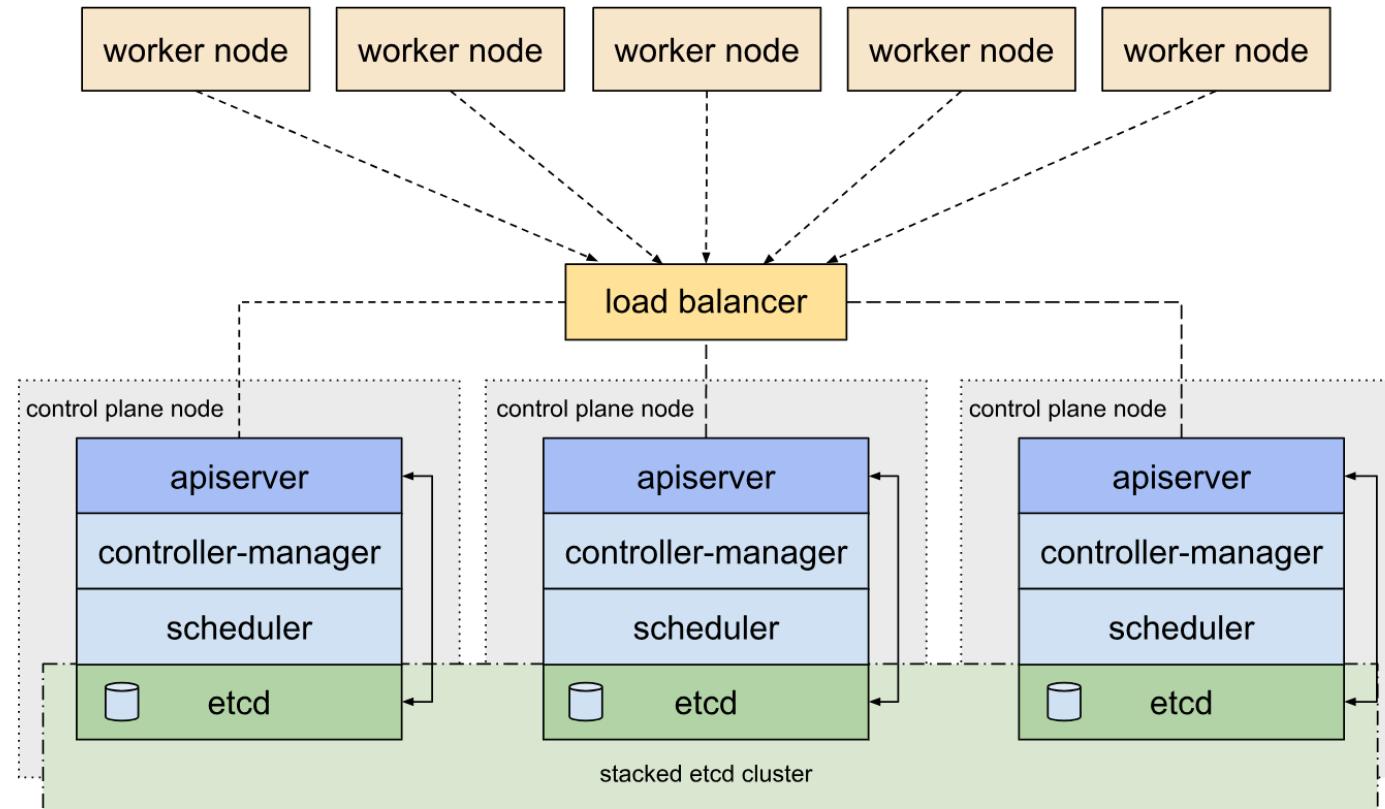


Options for Highly Available

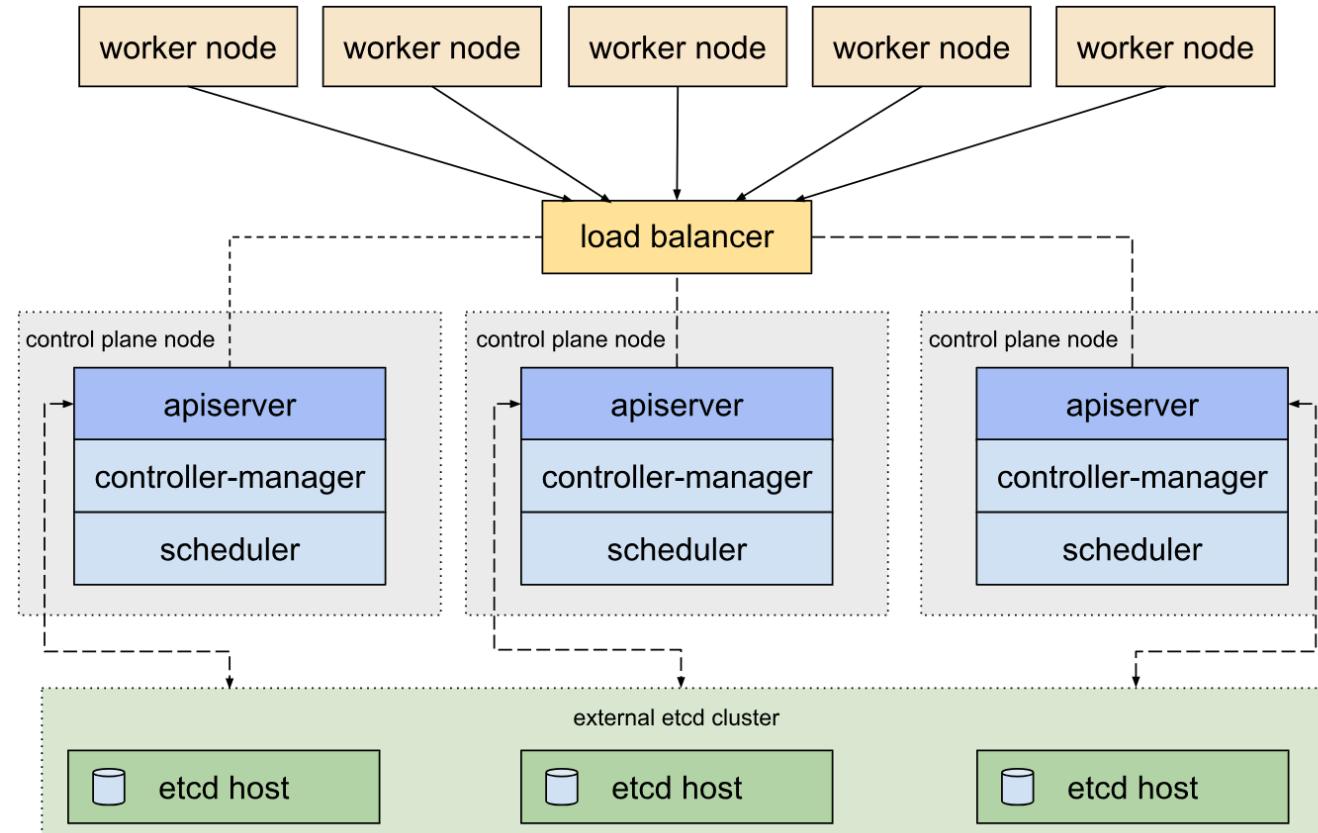
- Stacked etcd topology
- External etcd topology



Stacked etcd topology



External etcd topology



Fault Tolerance Table

Cluster Size	Majority	Failure Tolerance
1	1	0
2	2	0
3	2	1
4	3	1
5	3	2
6	4	2
7	4	3
8	5	3
9	5	4

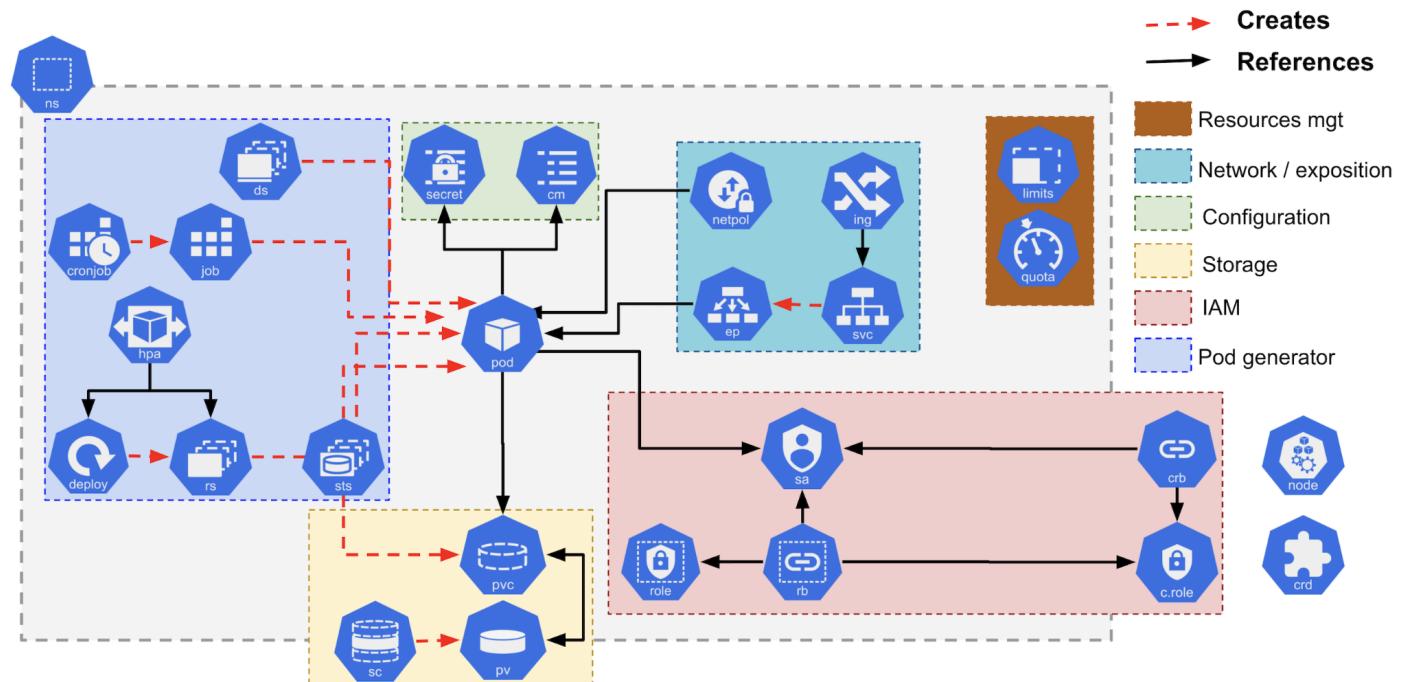


Resources / Objects

Resources / Objects

- Pod
- Deployment
- Service Ingress
- Config map
- Secret
- PV & PVC

Kubernetes Resources Map

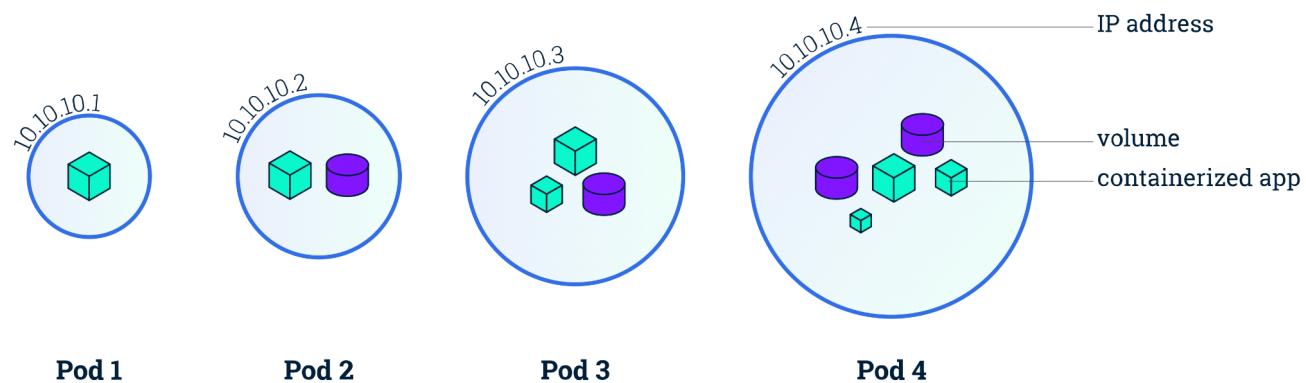


Pod

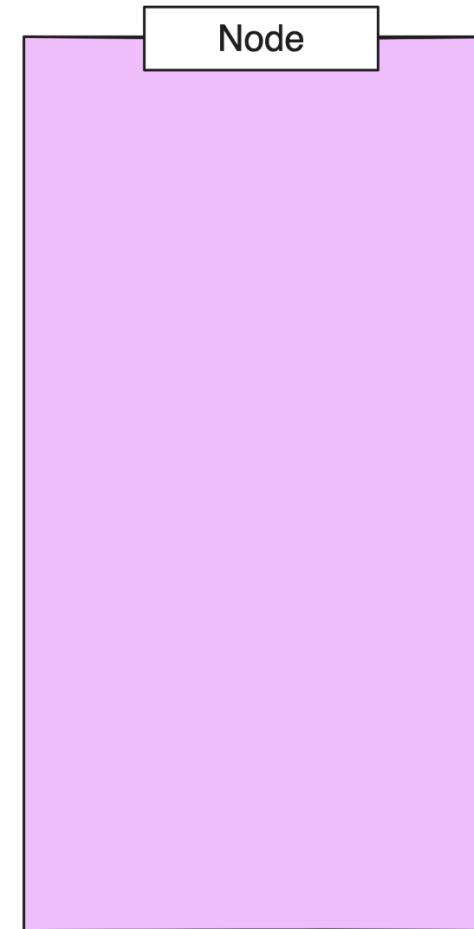
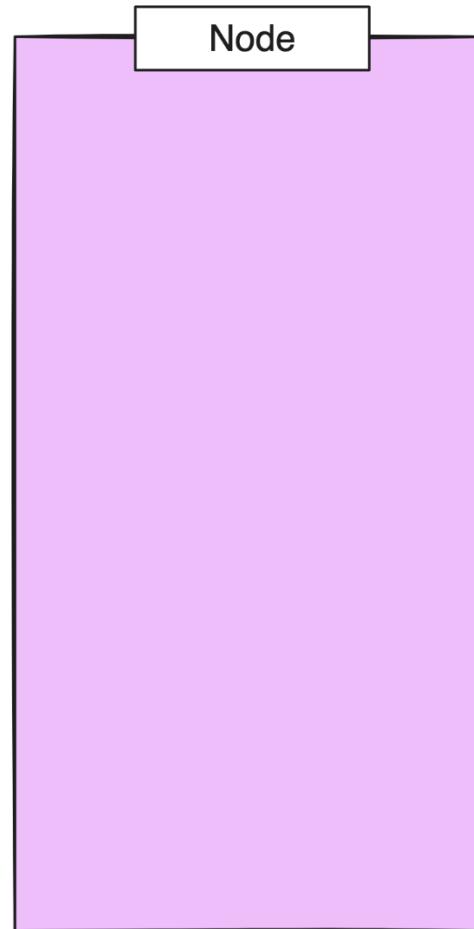
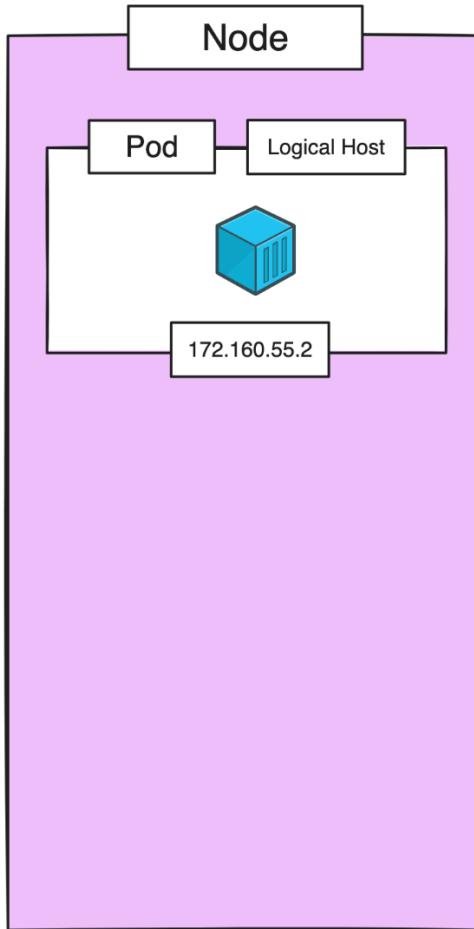


Pod

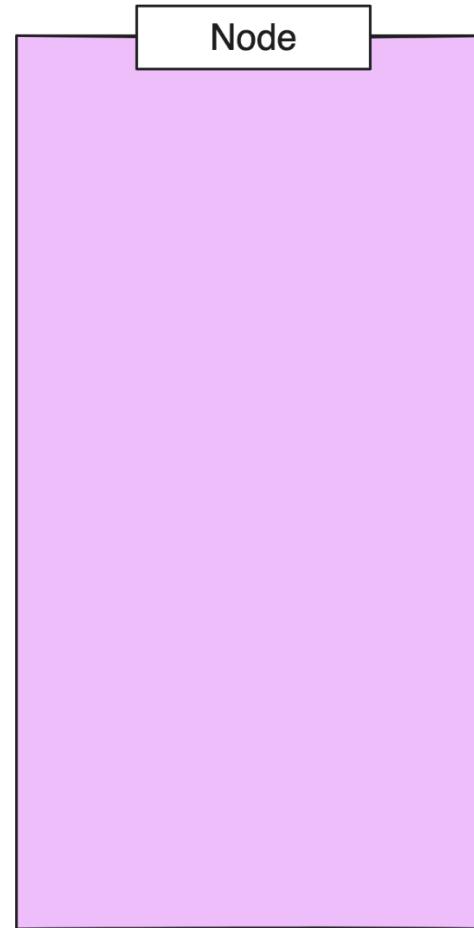
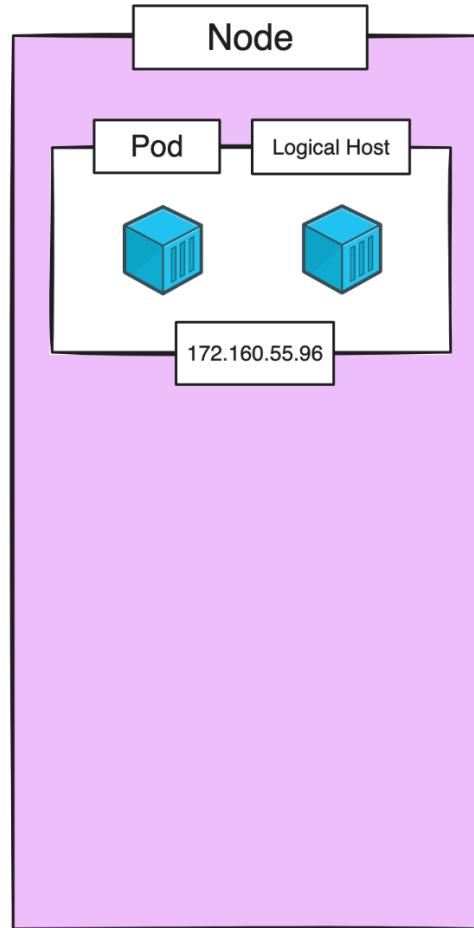
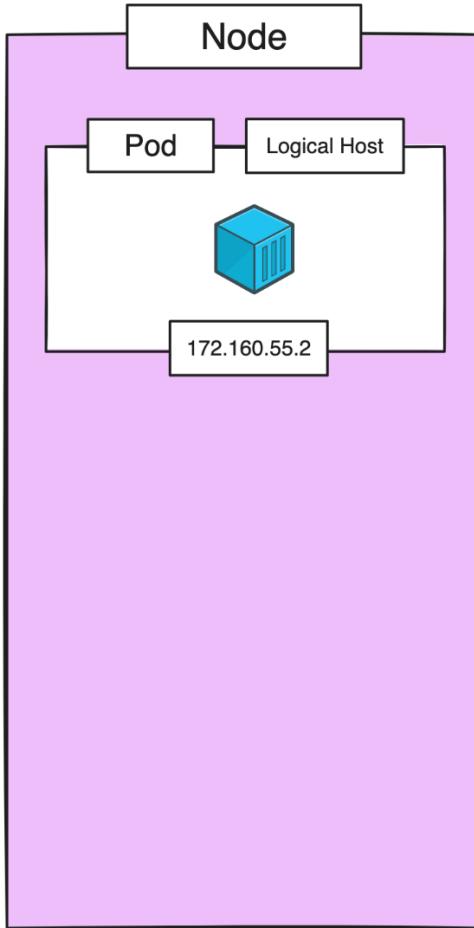
- Smallest deployable units of computing that you can create and manage in Kubernetes
- Logical Application
- One or more container in Pod (recommend 1 container per pod)
- Any containers in same pod will share storage volumes and network resources
- Each pod is assigned a unique ip



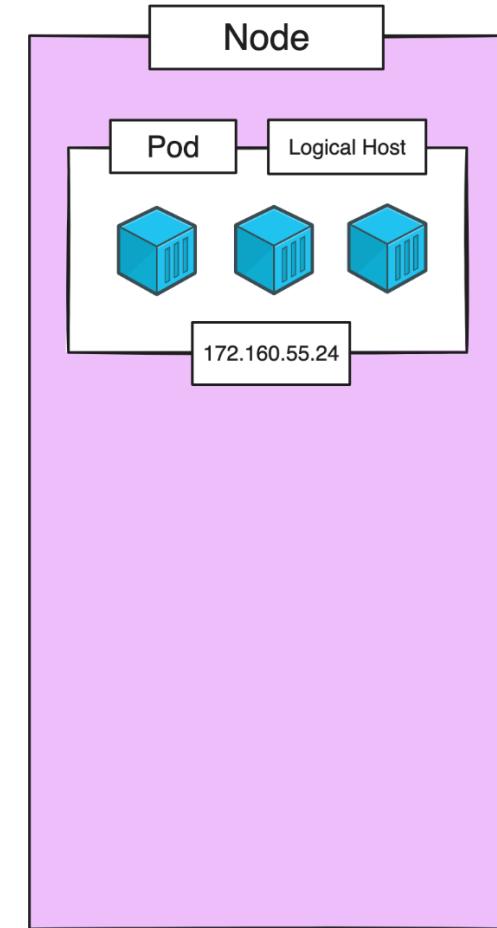
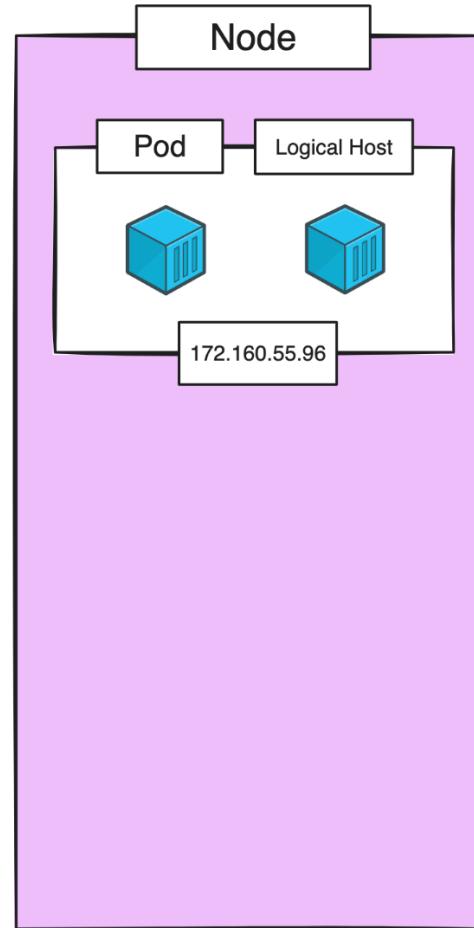
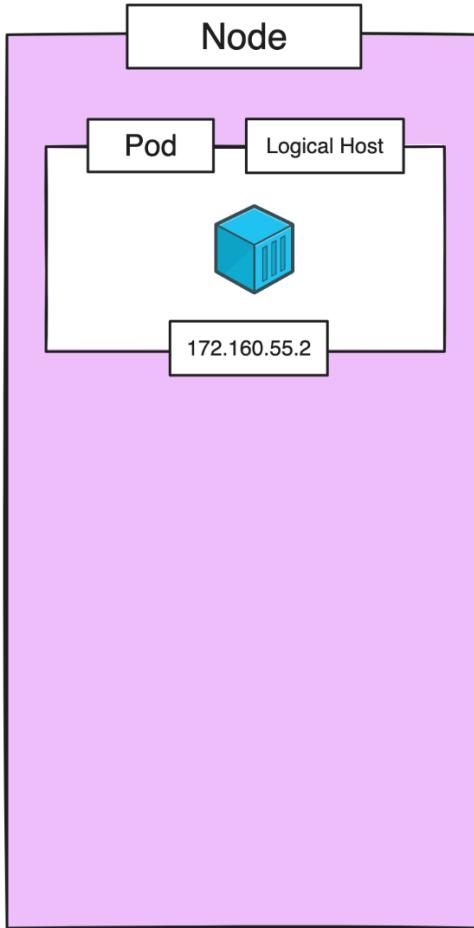
Node & Pod



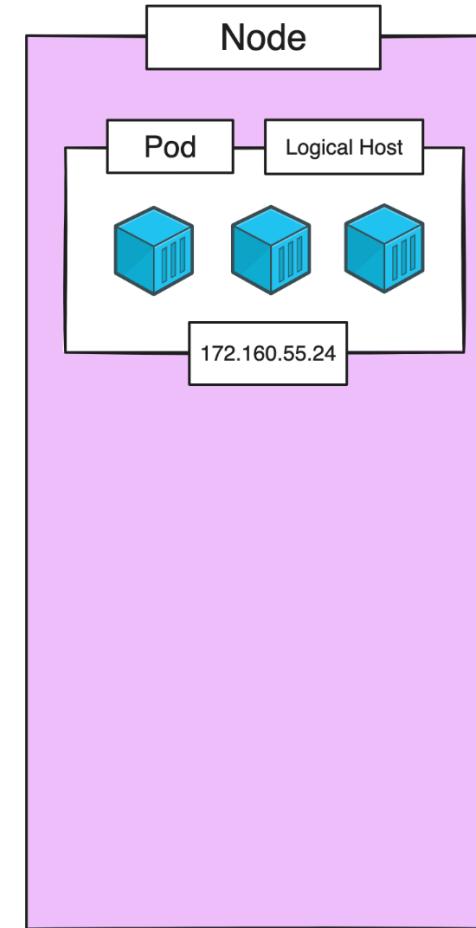
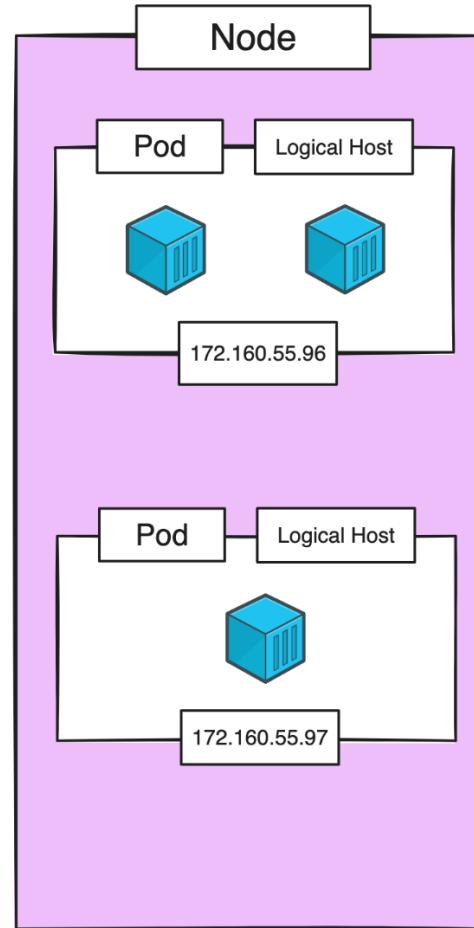
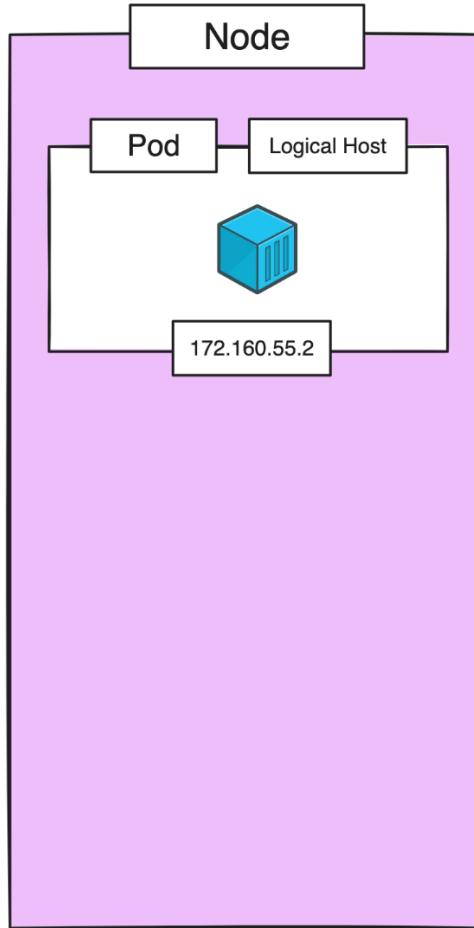
Node & Pod



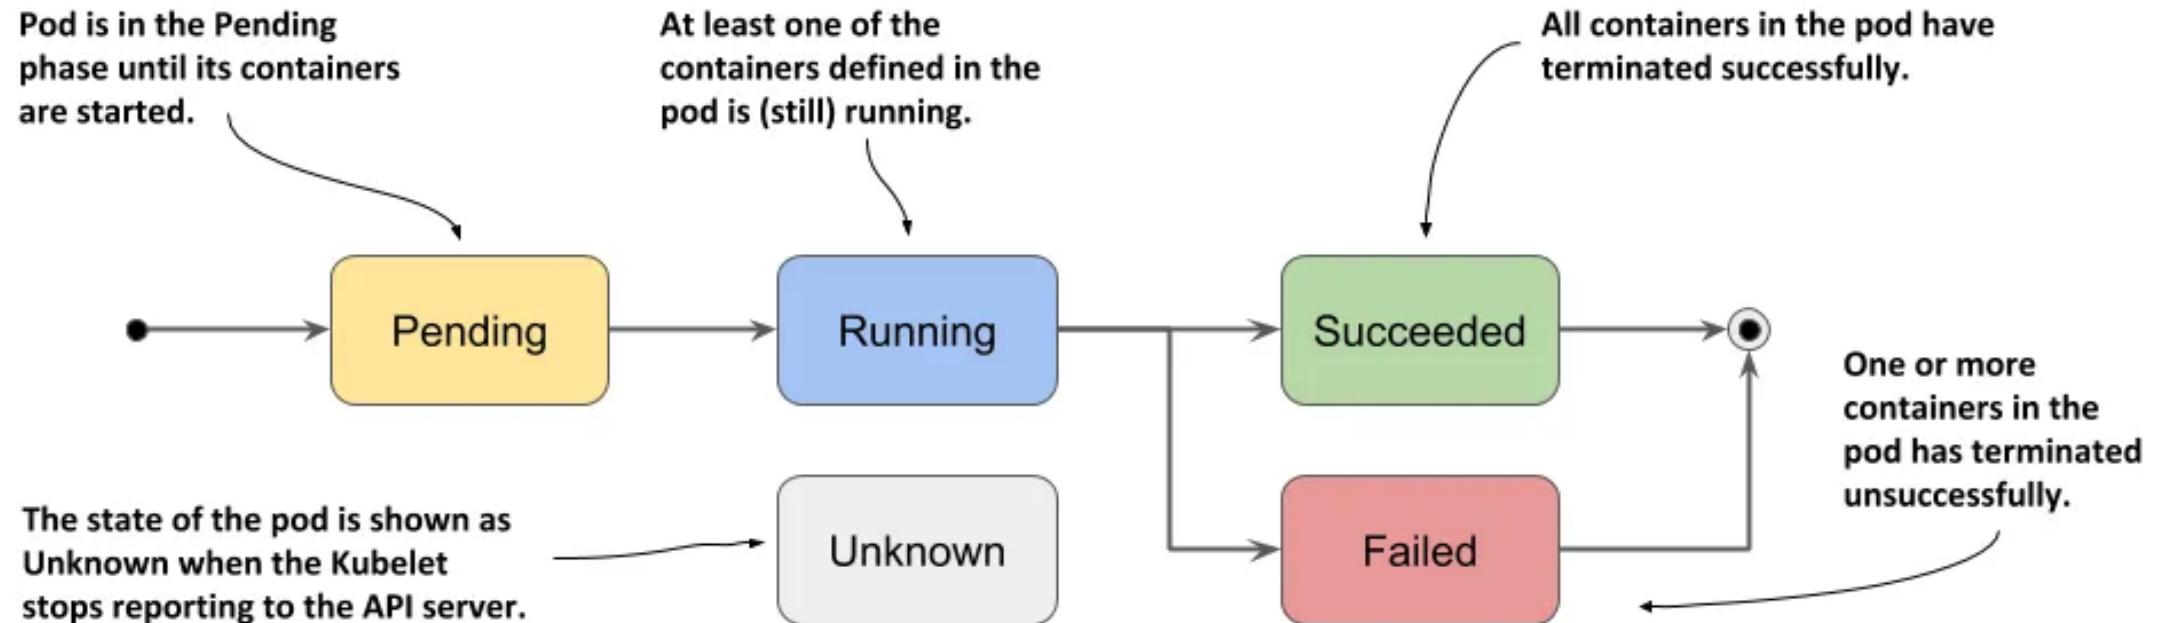
Node & Pod



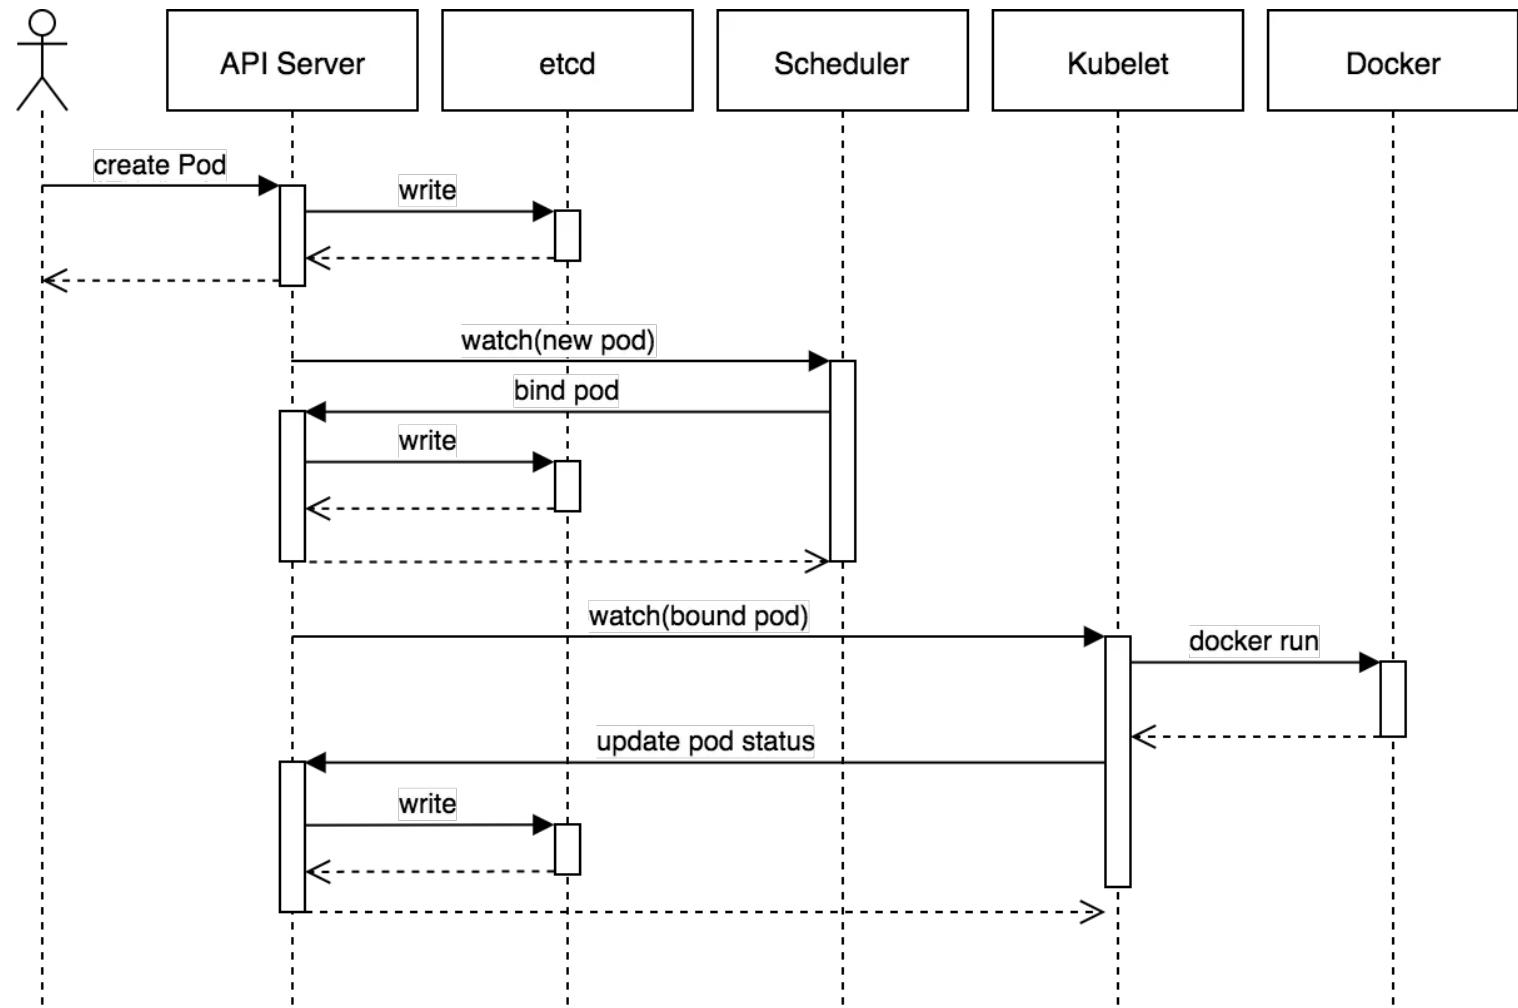
Node & Pod



State of Pod



Creation of Pod



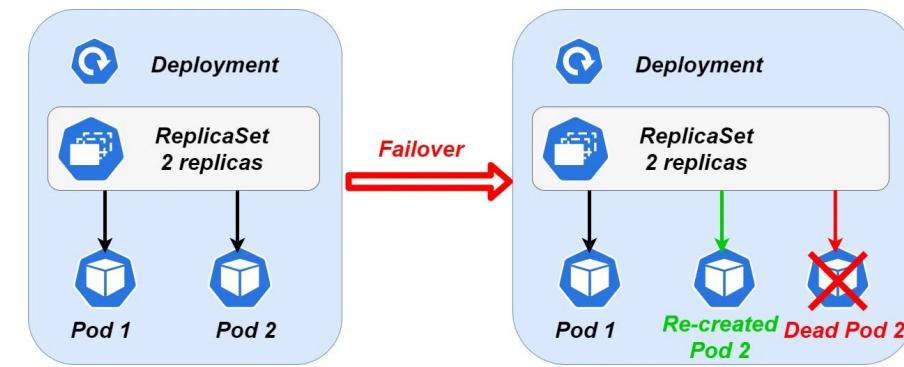
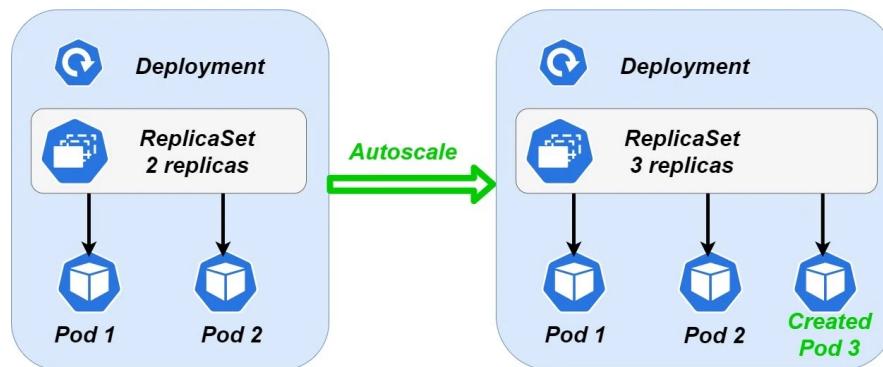


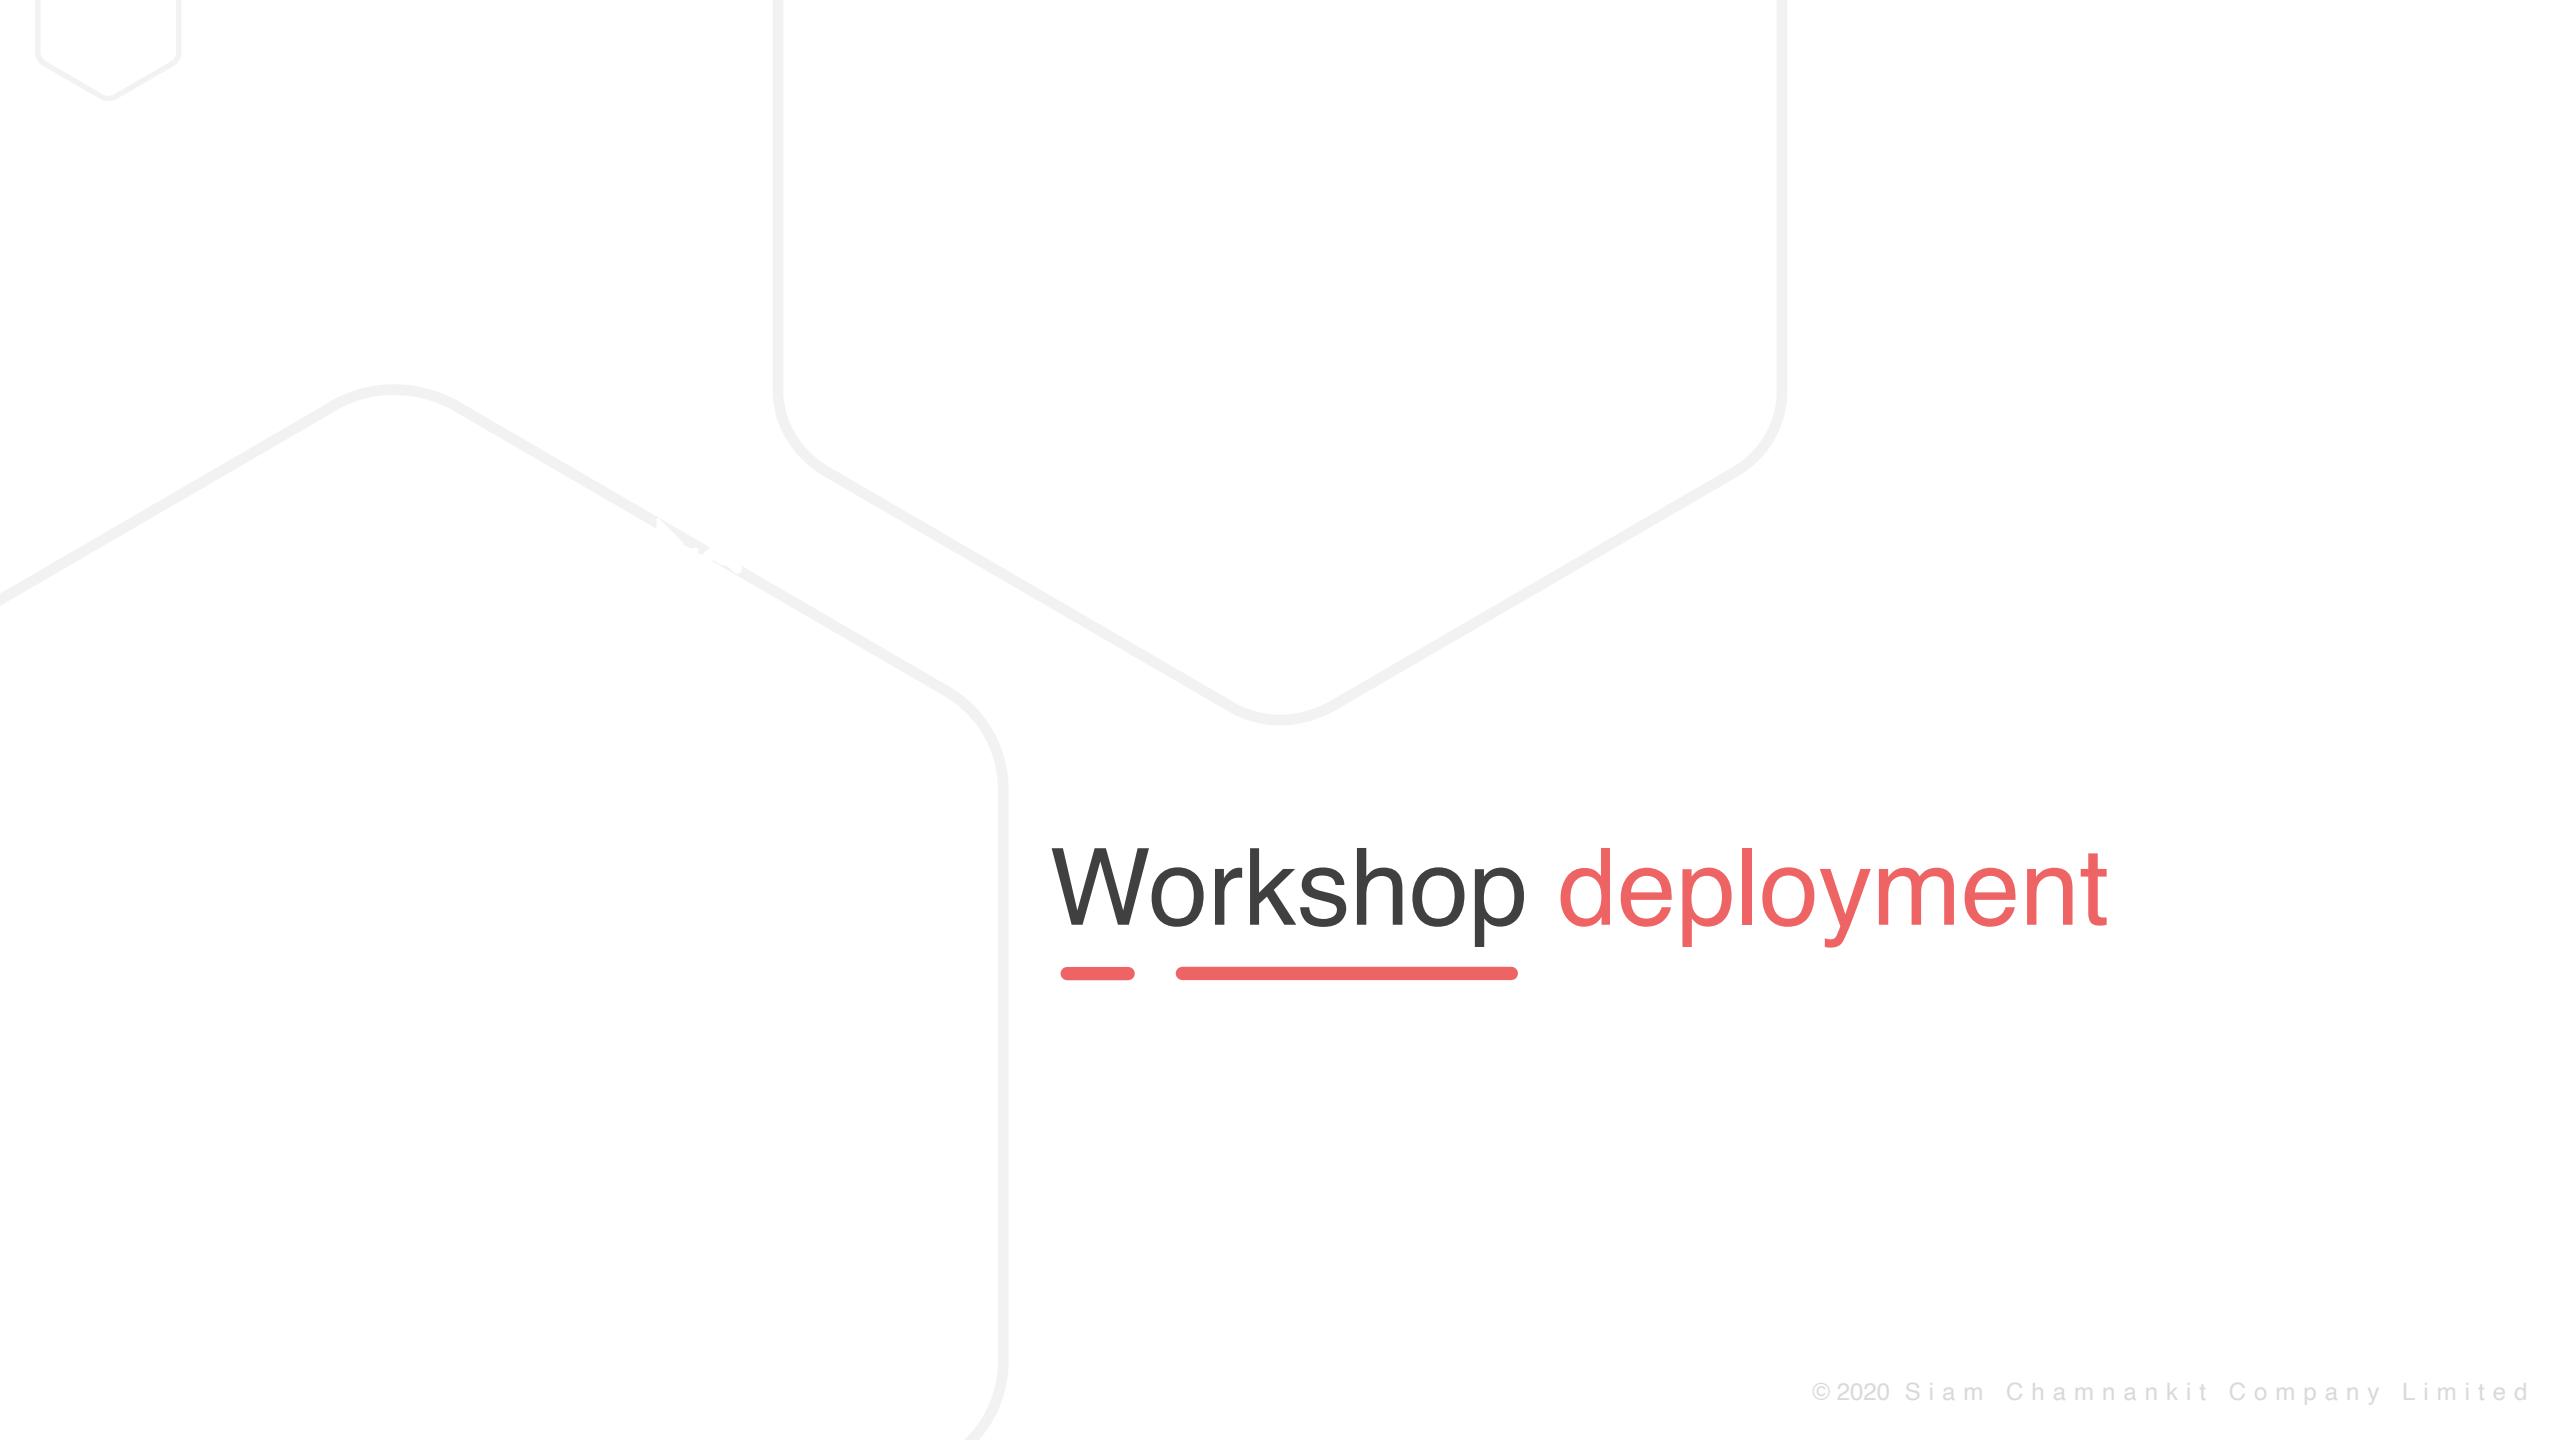
Workshop **pod**

Deployment

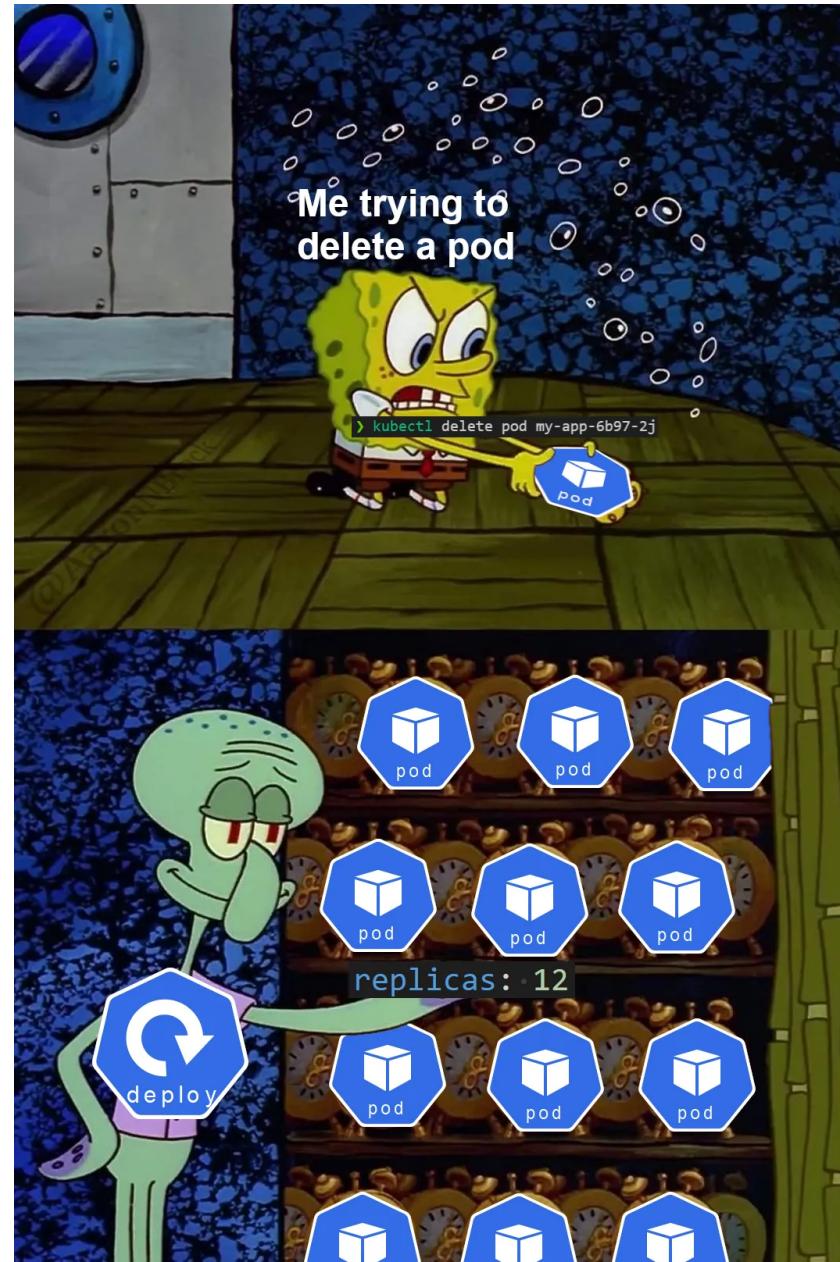
Deployment

- A Deployment provides declarative updates for Pods and ReplicaSets
- You describe a **desired state** in a Deployment, and the Deployment Controller **changes the actual state to the desired state** at a controlled rate
- Manage for rollout and rollback
- Deployment => ReplicaSet => pod

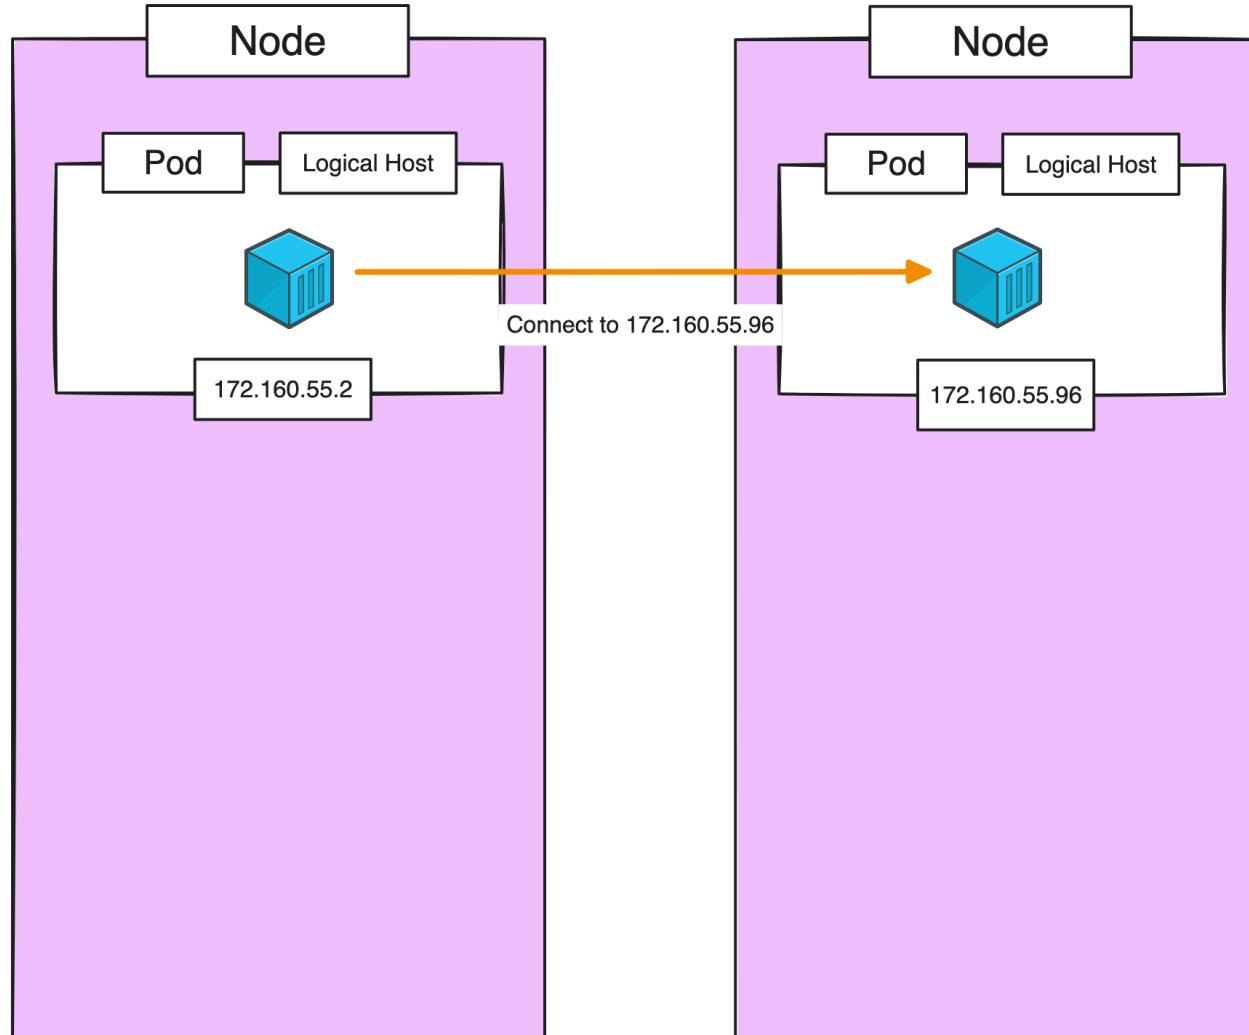




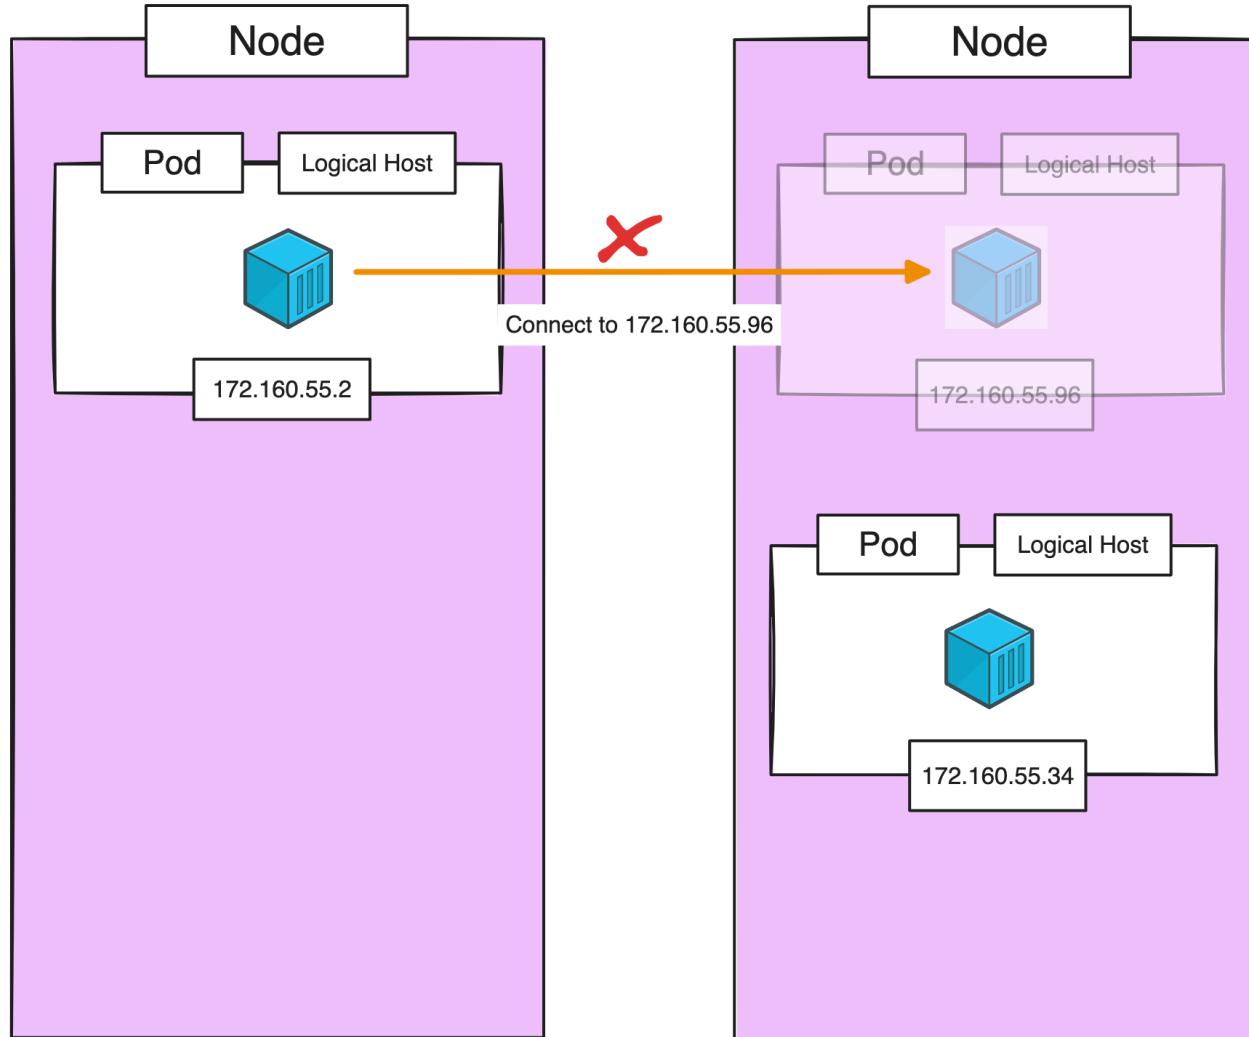
Workshop deployment



Problems



Problems



Service

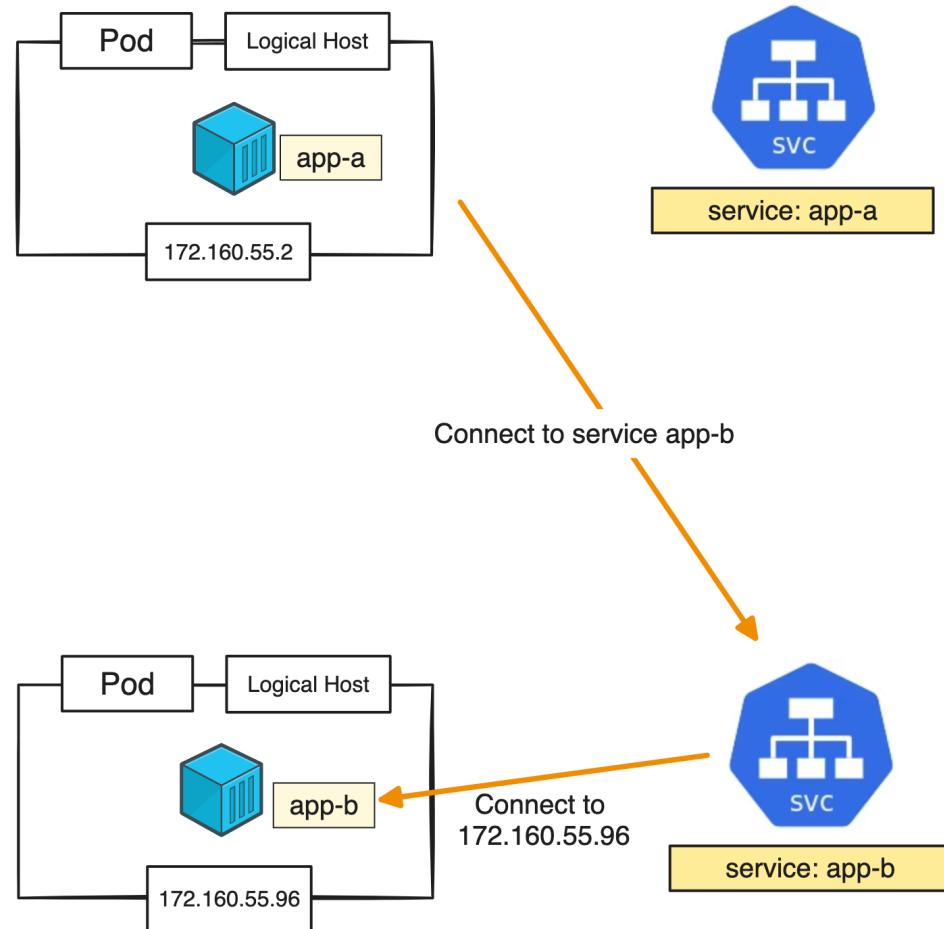


Service

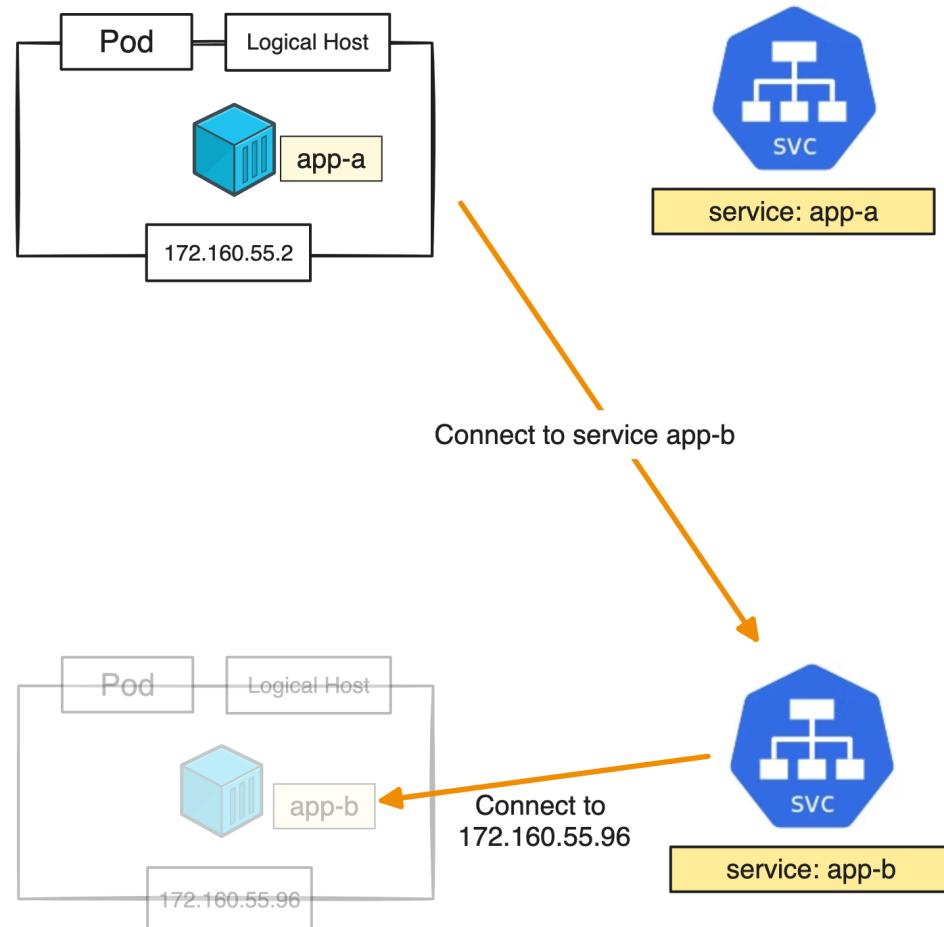
- Acts as a virtual endpoint for your application
- Providing a stable and discoverable way to access your pods without worrying about their individual IP addresses or network details
- Logical load balancer
- Communicating for outside and inside of cluster
- 4 types:
 - ClusterIP
 - NodePort
 - LoadBalancer
 - ExternalName



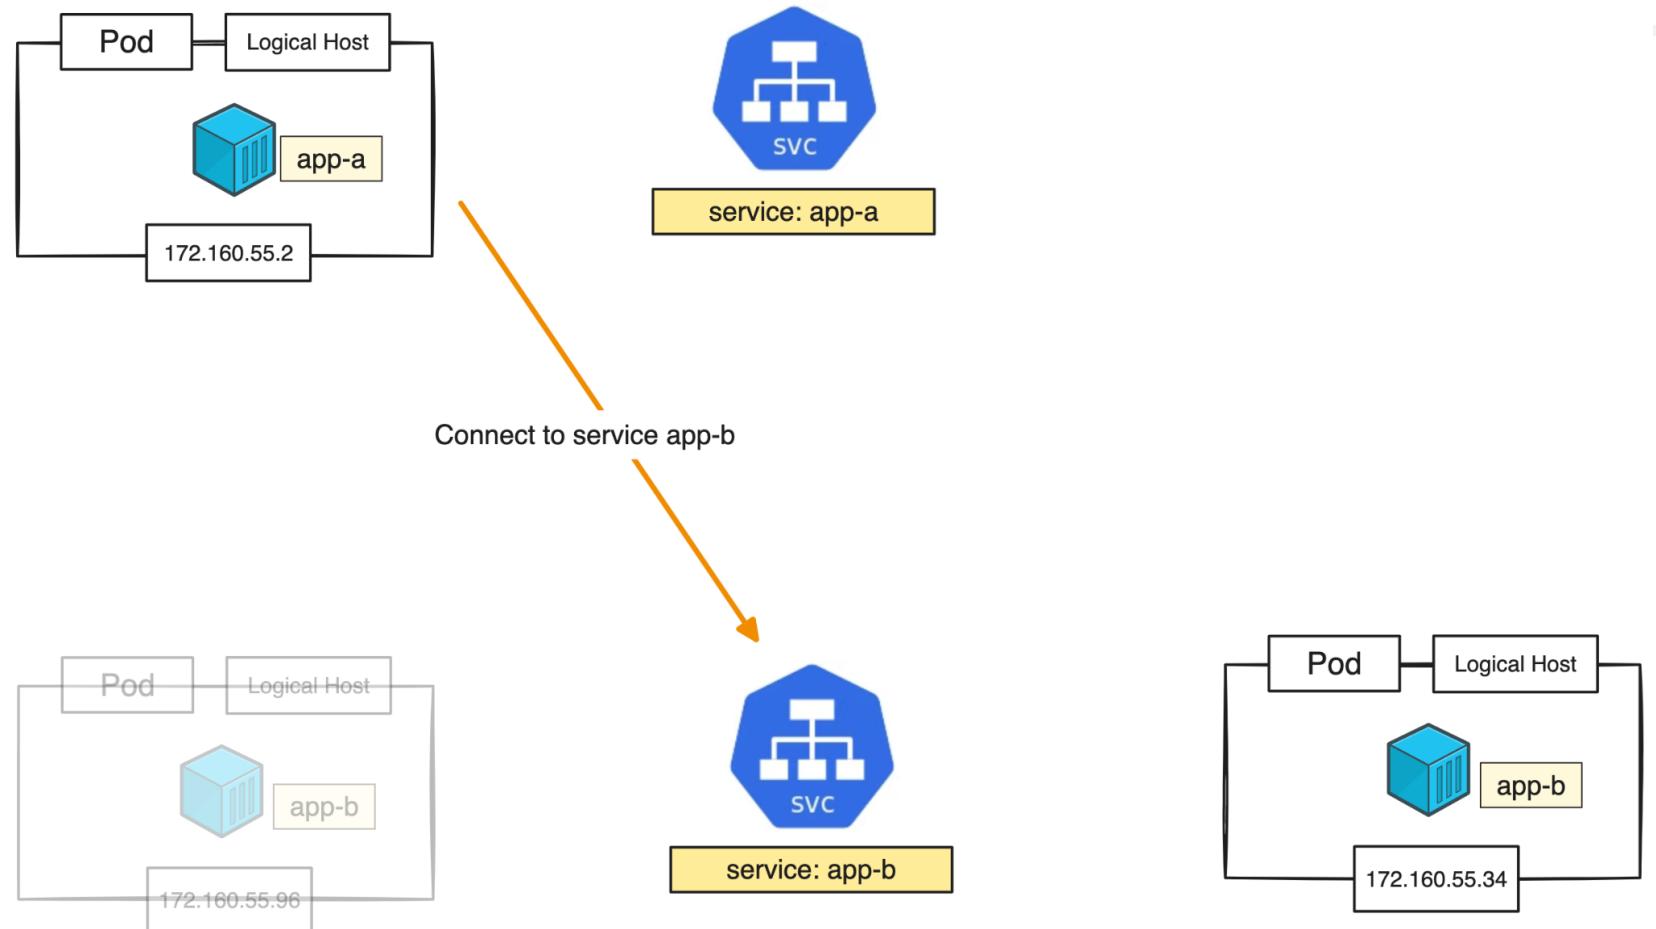
Service: Stable endpoint



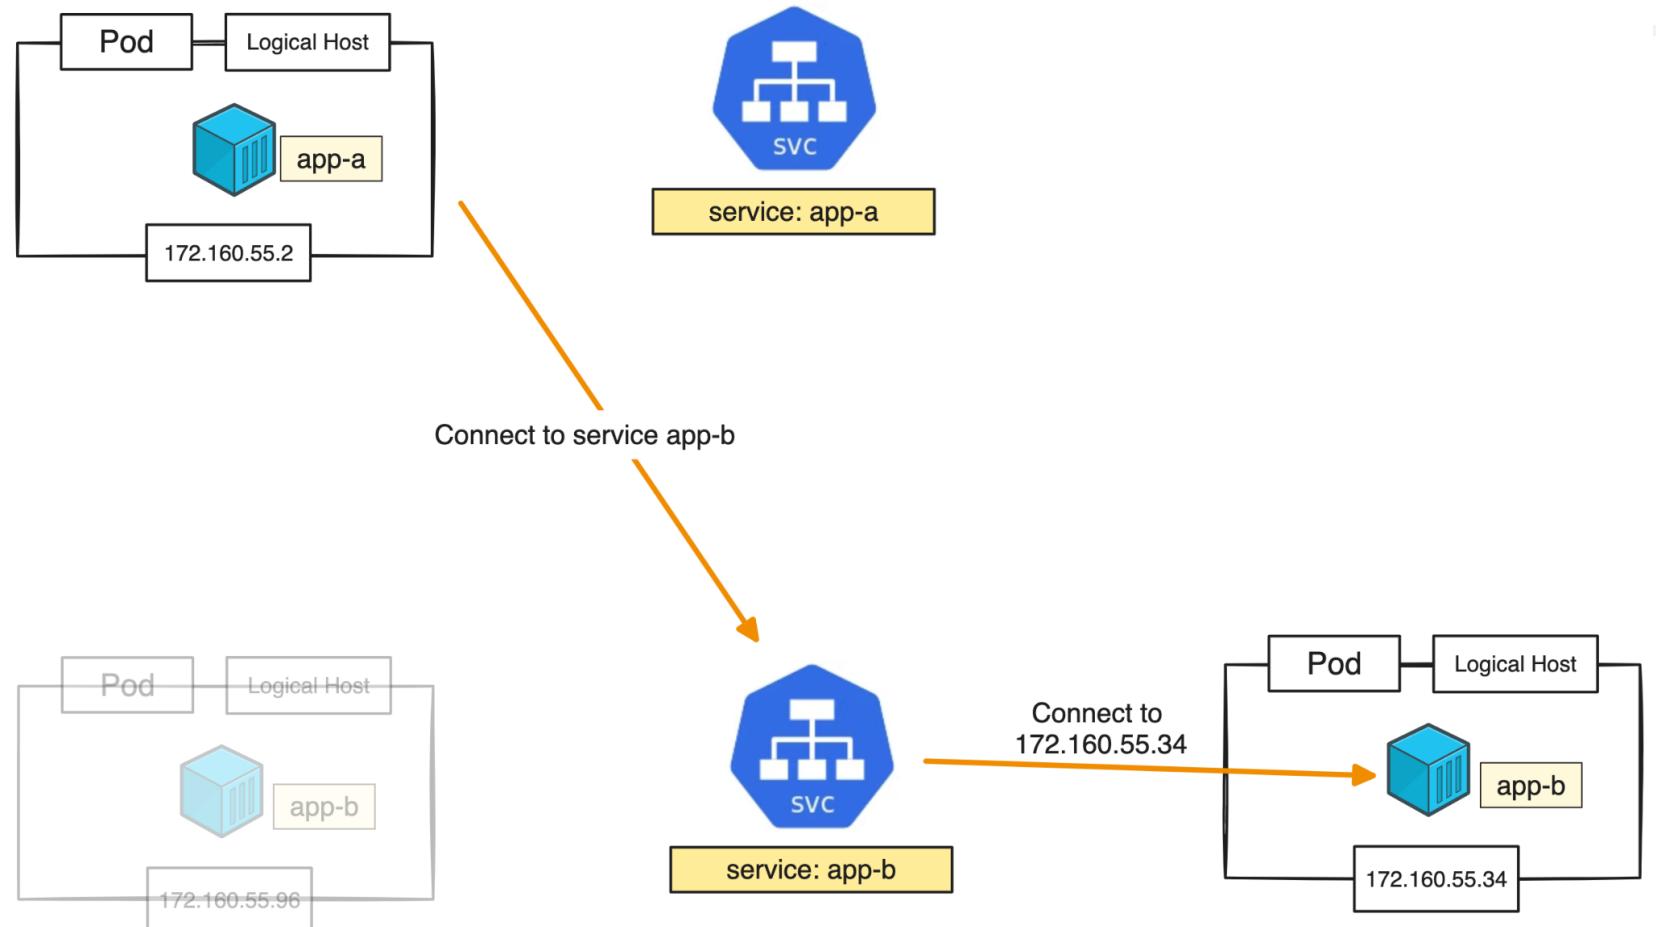
Service: Stable endpoint



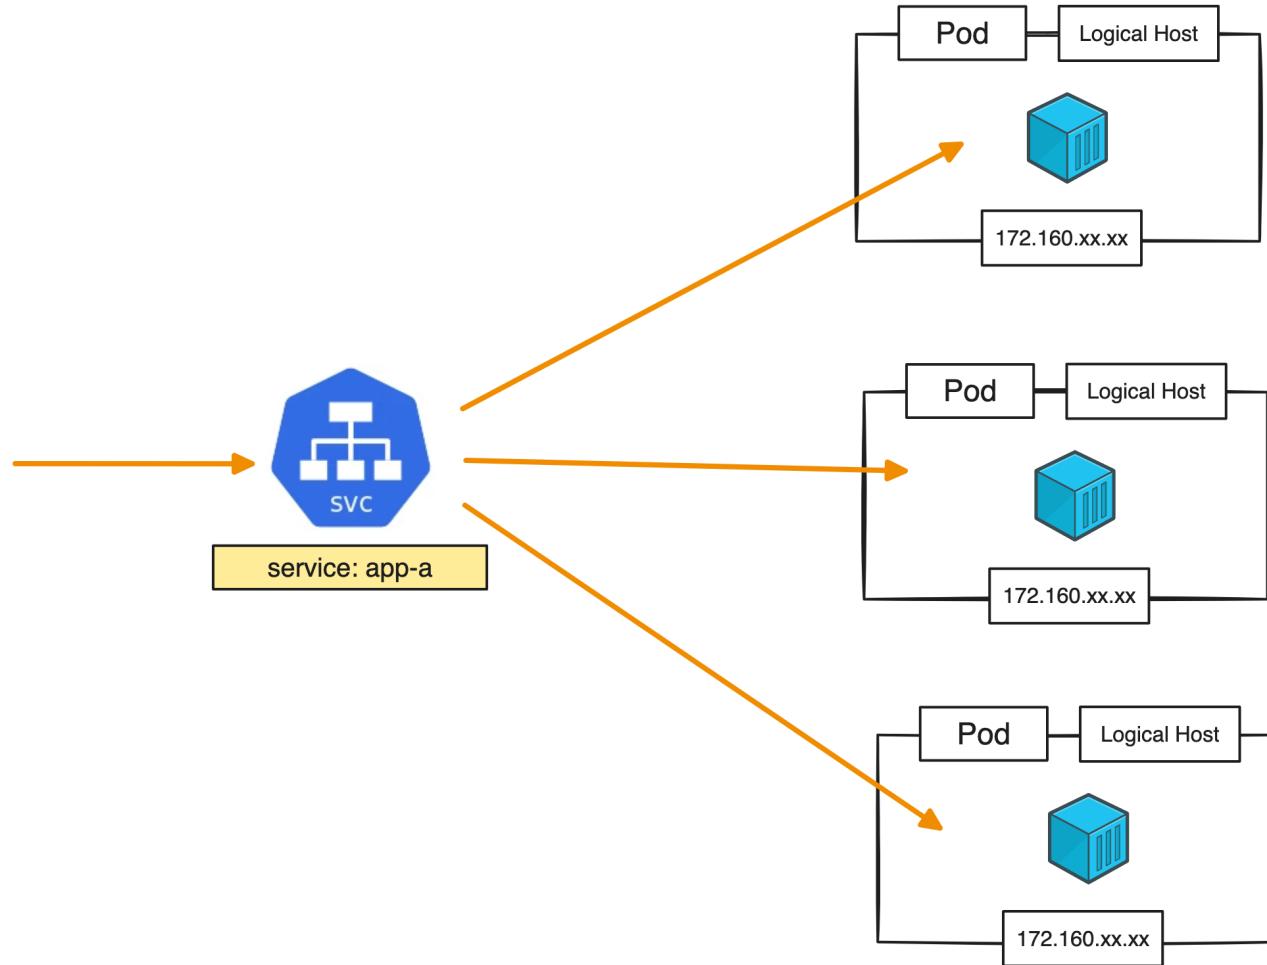
Service: Stable endpoint



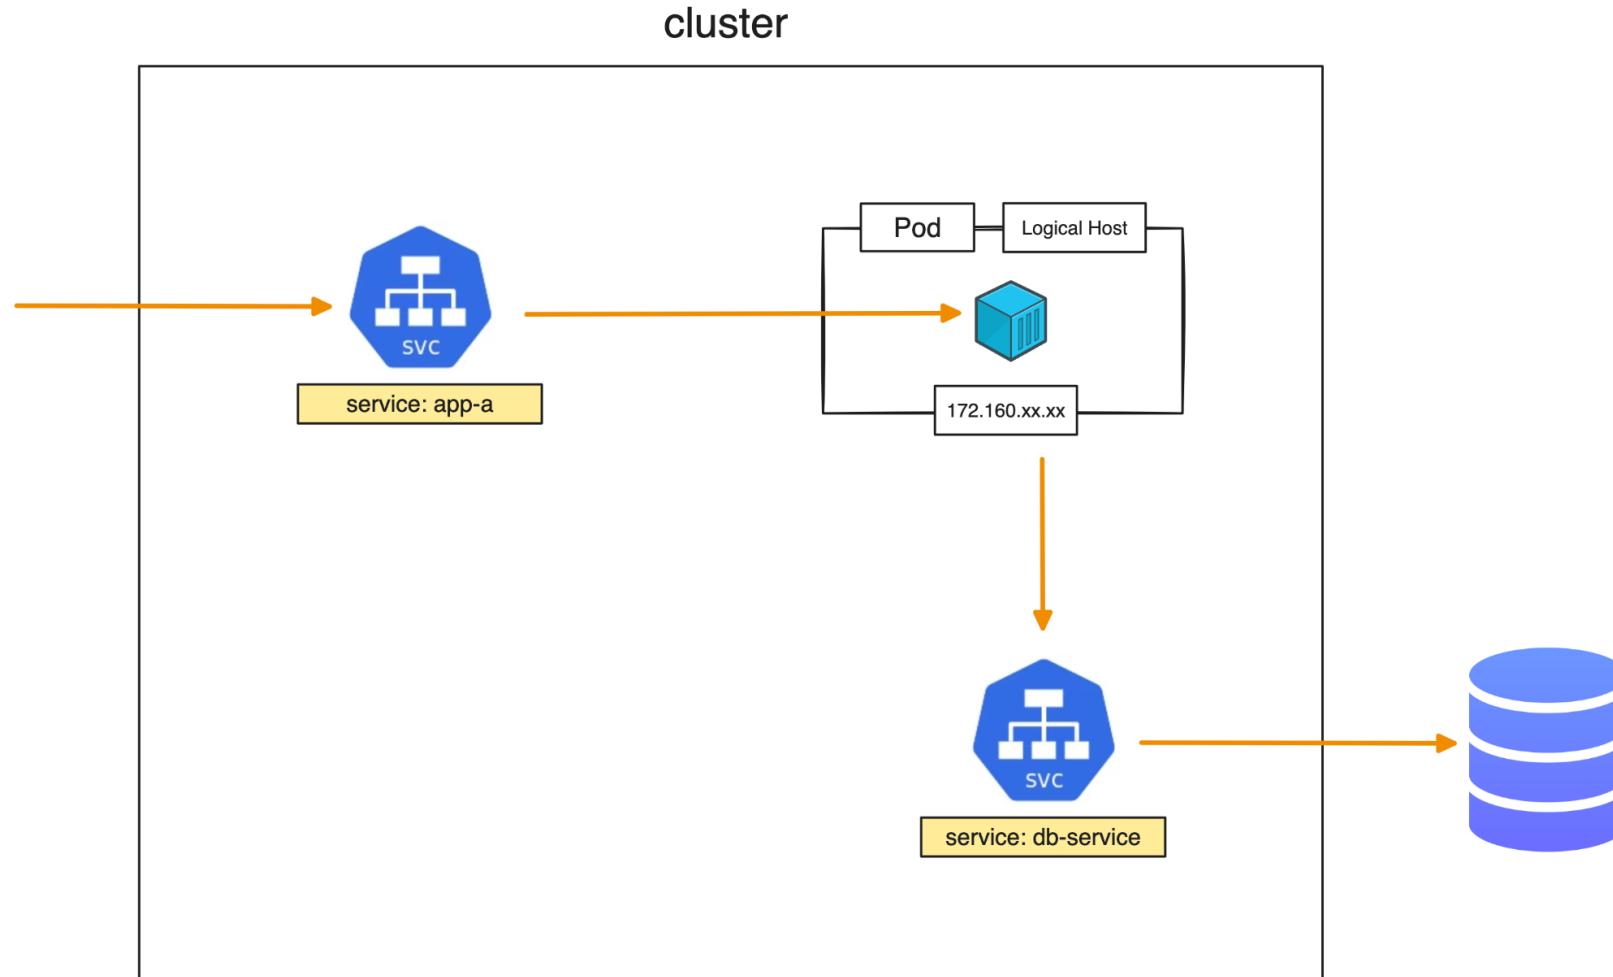
Service: Stable endpoint



Service: Load balancer



Service: Communicating inside & outside of cluster



Service type: ClusterIP

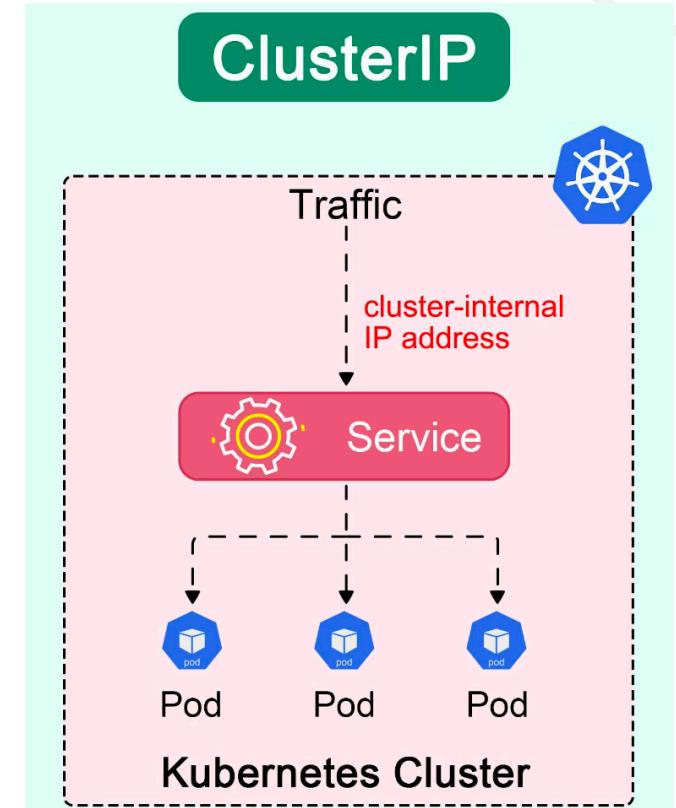
- ClusterIP is the default service type in Kubernetes
- It exposes the service on an internal IP address reachable **only within** the cluster

Advantages

- Suitable for inter-pod communication and internal services

Disdvantages

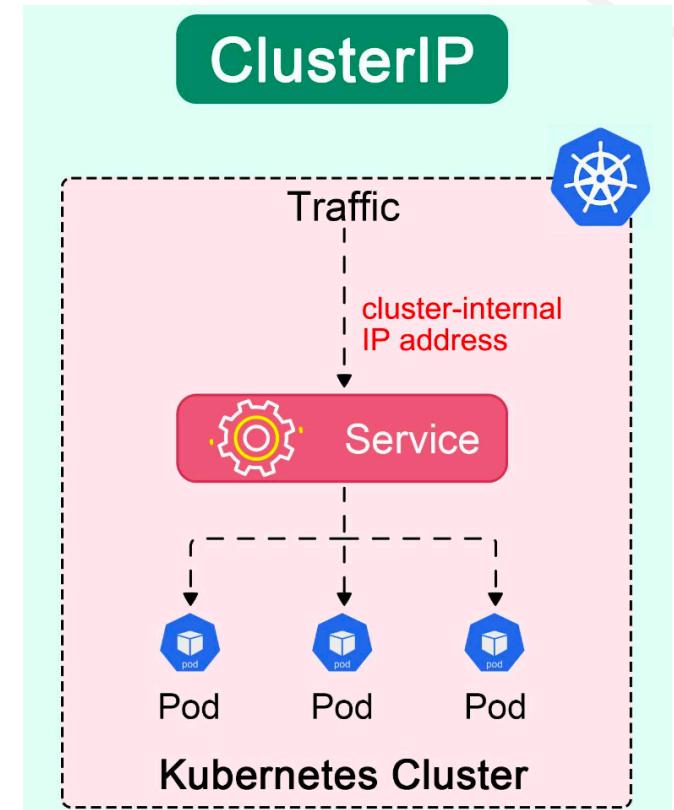
- Limited to internal cluster communication
- Not accessible externally



Service type: ClusterIP

Use Cases

- Inter-service communication within the cluster. For example, communication between the front-end and back-end components of your app





Workshop ClusterIP Service

Service type: NodePort

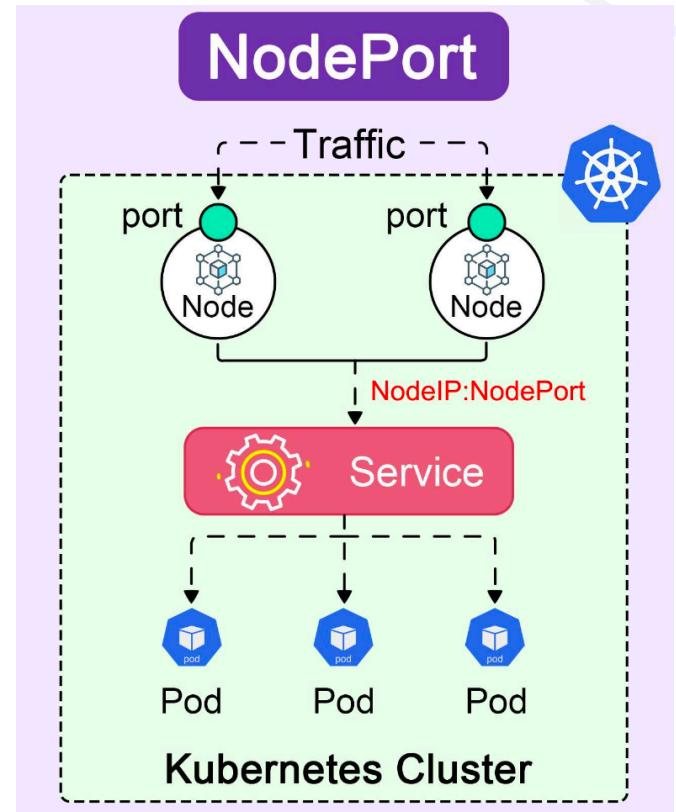
- It enables **external access** to the service using the cluster nodes' **IP address** and the assigned **NodePort**
- When a NodePort service is created, Kubernetes allocates a port from a predefined range (default: **30000-32767**)

Advantages

- Provides external access to the service
- Suitable for development and testing environments

Disadvantages

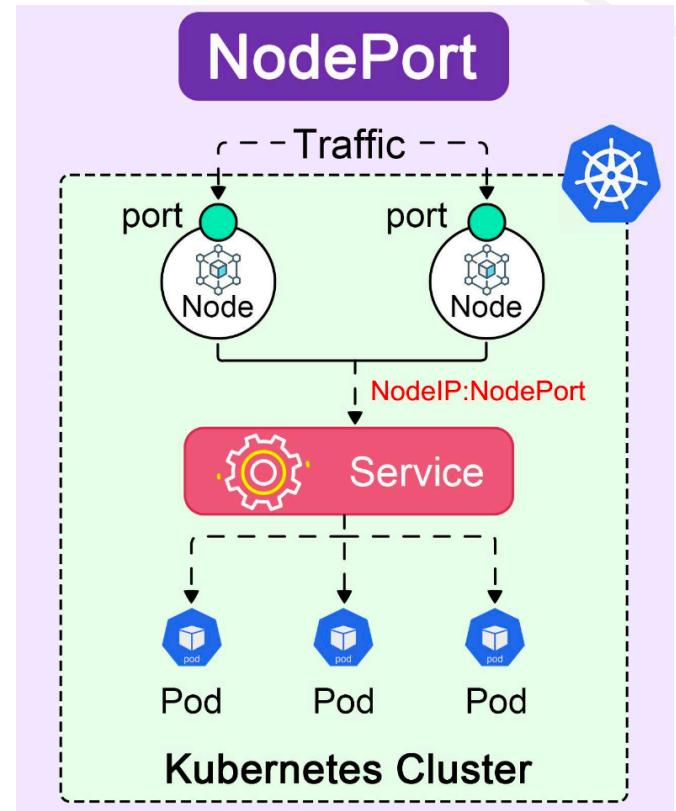
- Requires manual management of port allocations
- IP Node maybe have changed
- Should not use in production

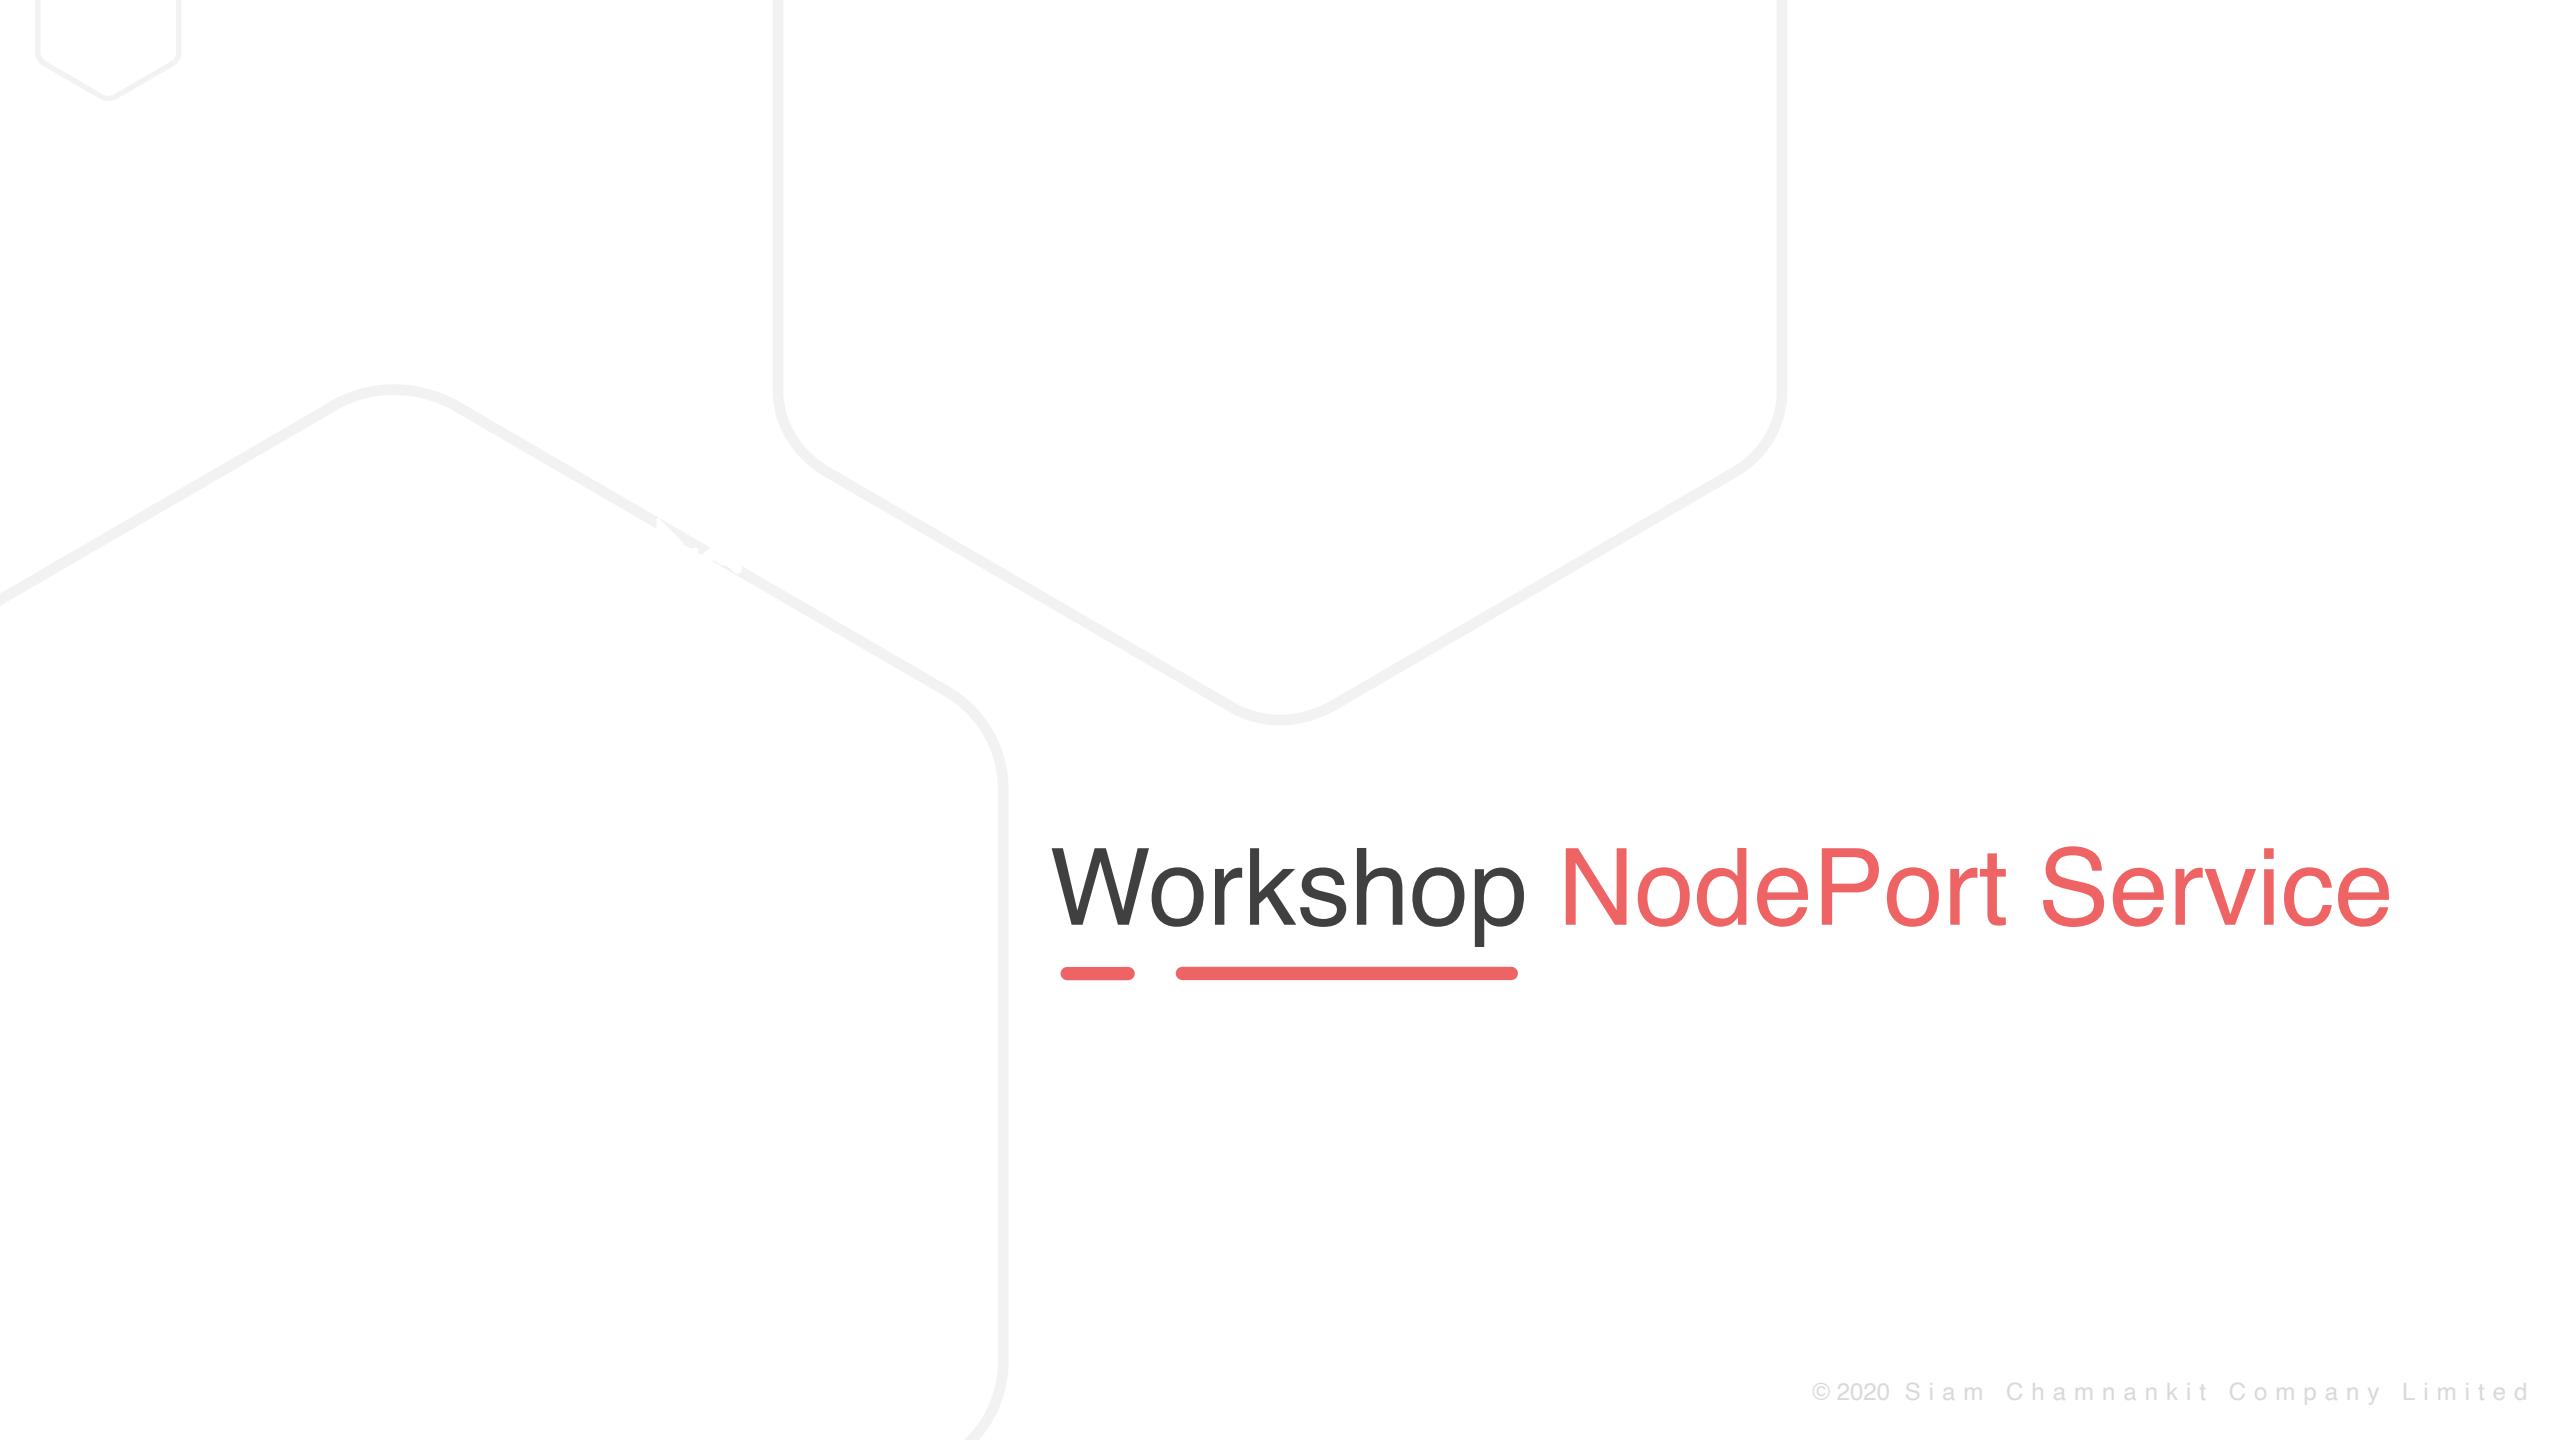


Service type: NodePort

Use case

- When you want to enable external connectivity to your service
- Best for testing access for a moment time





Workshop NodePort Service

Service type: LoadBalancer

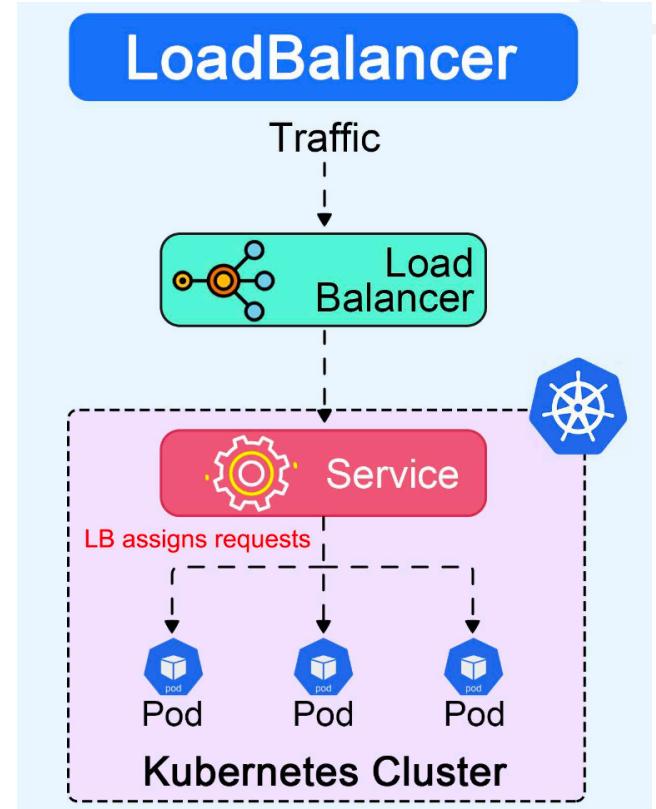
- LoadBalancer is a service type that **provisions** an **external load balancer** to distribute traffic across the pods
- Typically provided by a cloud provider's infrastructure
- When a LoadBalancer service is created, Kubernetes interacts with the cloud provider's API to an external load balancer

Advantages

- Provides external access to the service
- Automatically provisions an external load balancer for the service
- Supports scaling and high availability

Disadvantages

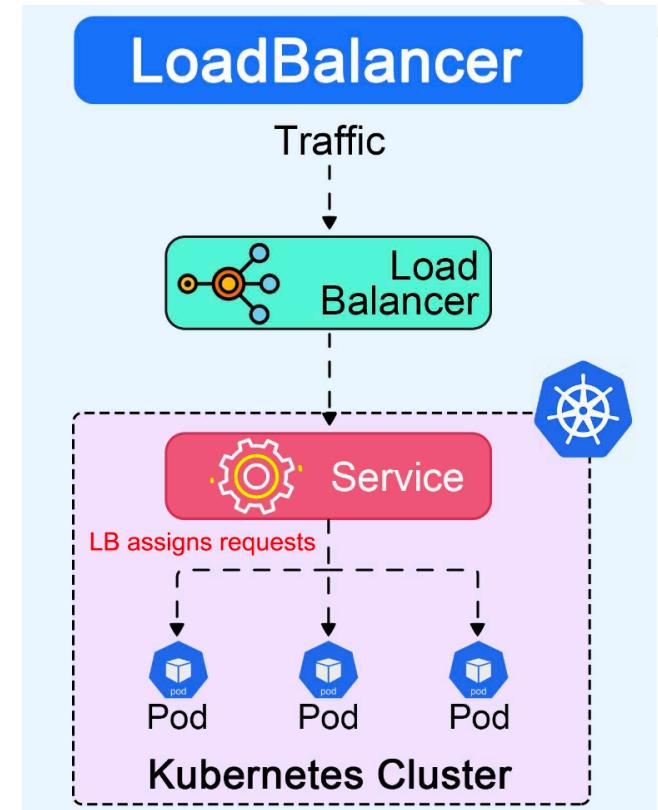
- Requires a cloud provider
- Additional costs for using the load balancer service on cloud provider



Service type: LoadBalancer

Use case

- When you want to enable external connectivity to your service
- When you are using a cloud provider to host your Kubernetes cluster





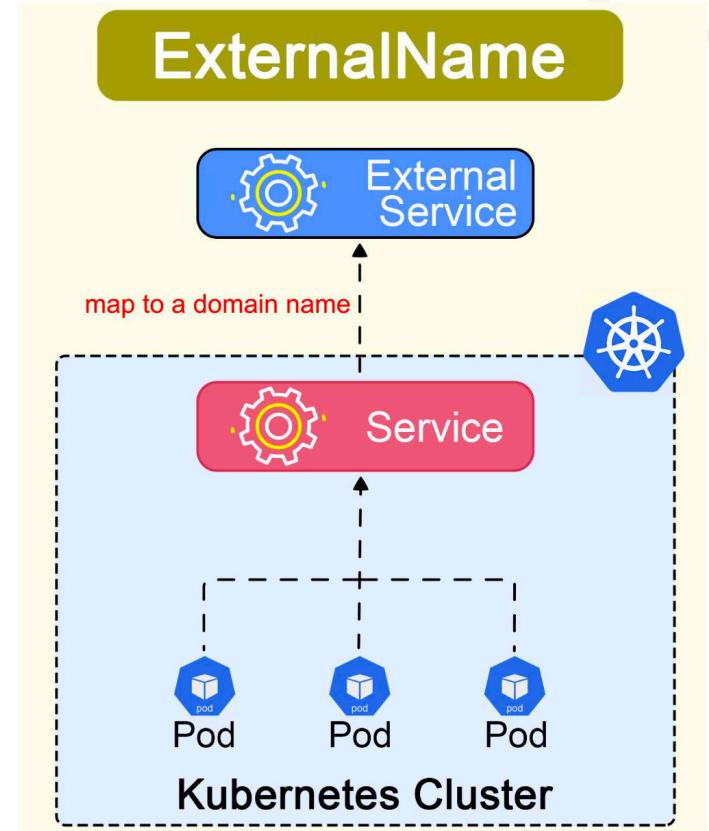
Workshop LoadBalancer Service

Service type: ExternalName

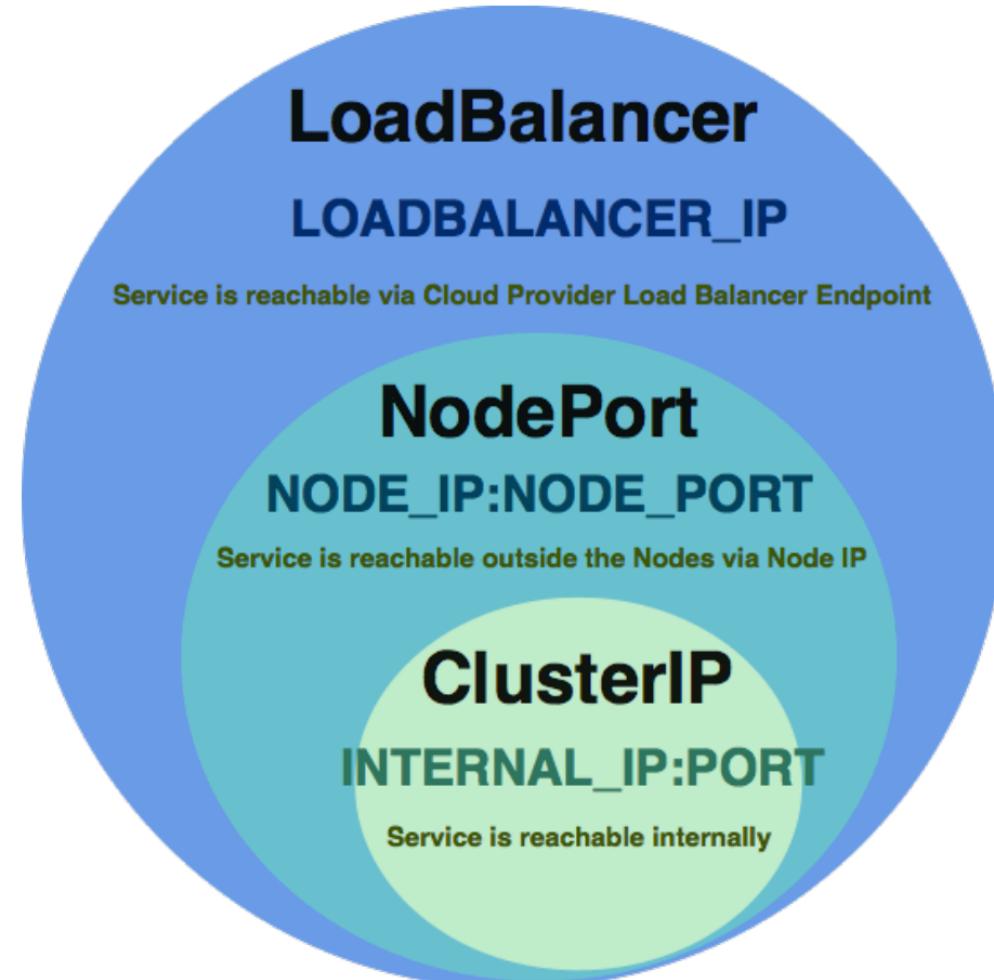
- An ExternalName Service is a special type of Kubernetes service that allows you to expose an existing external service (like a database or a third-party API) to your cluster without deploying any Pods

Use cases

- Accessing an external database to your Kubernetes cluster
- Connecting a third-party API from within your cluster



Service type

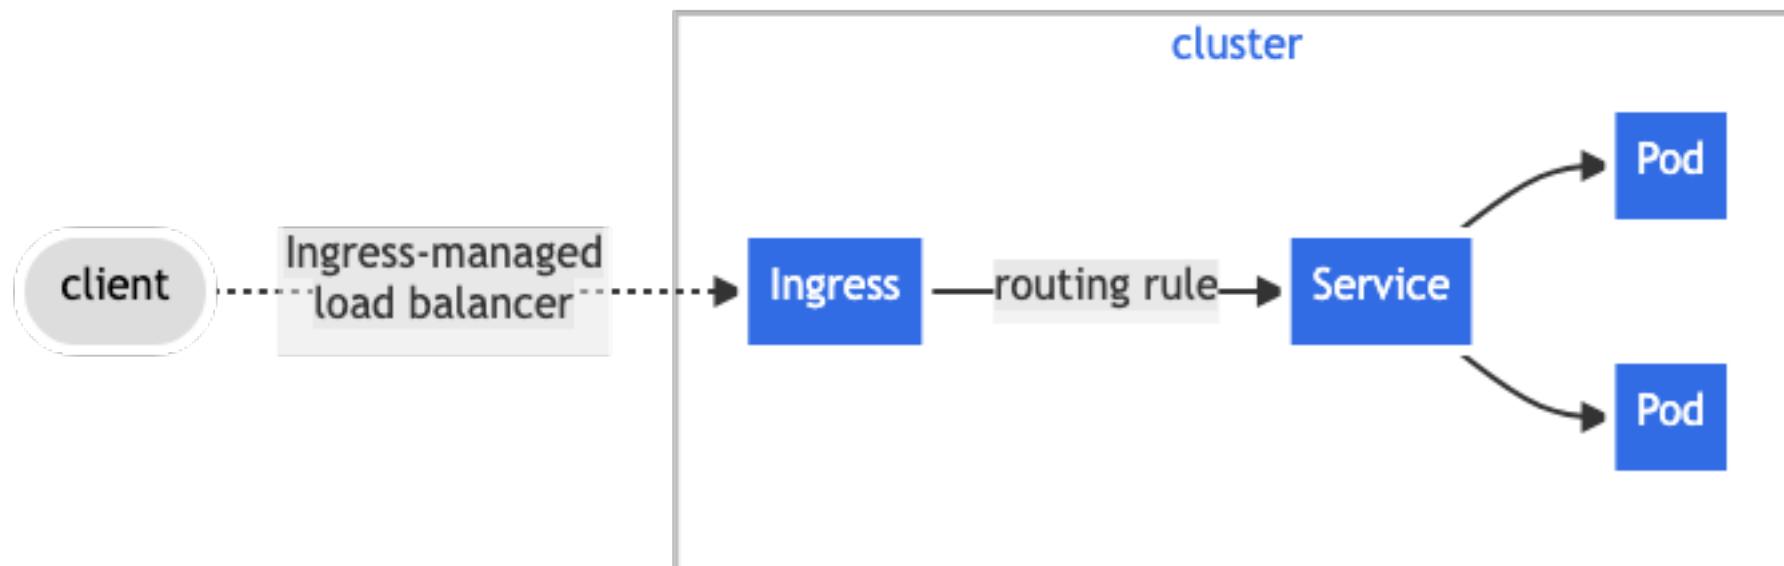


Ingress



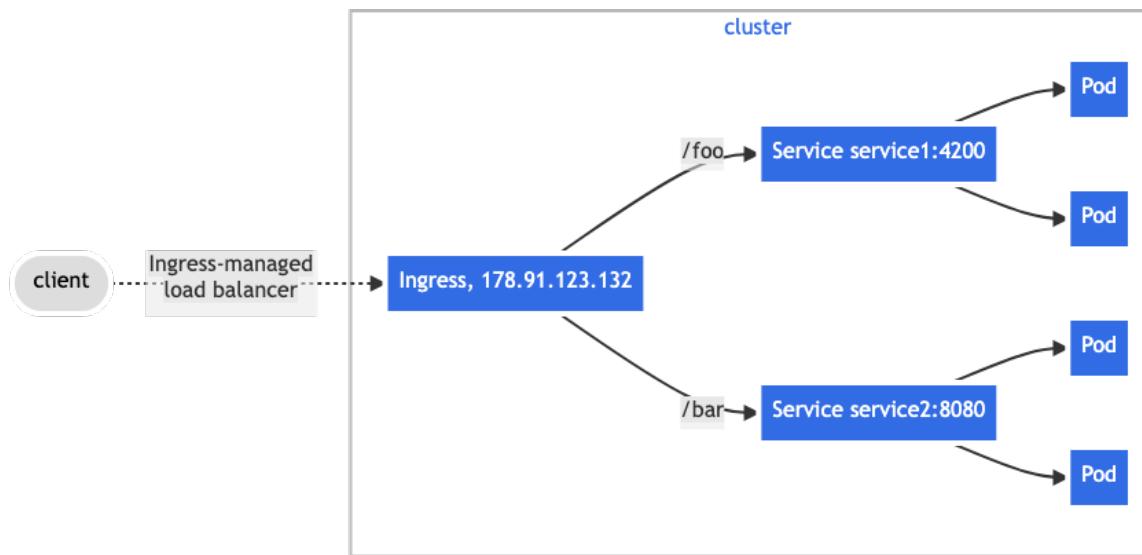
Ingress

- Routing from outside the cluster to services within the cluster
- Traffic routing is controlled by rules defined on the ingress resource
- An **Ingress controller** is responsible for fulfilling the Ingress



Ingress Controllers

- The Ingress Controller is an application hosted inside a Kubernetes cluster that actively **manages traffic** following Ingress Resources and pre-defined Ingress Rules
- Ingress controllers doesn't come with standard, have to be deployed separately
- Popular controllers include ingress-nginx, HAProxy, Traefik



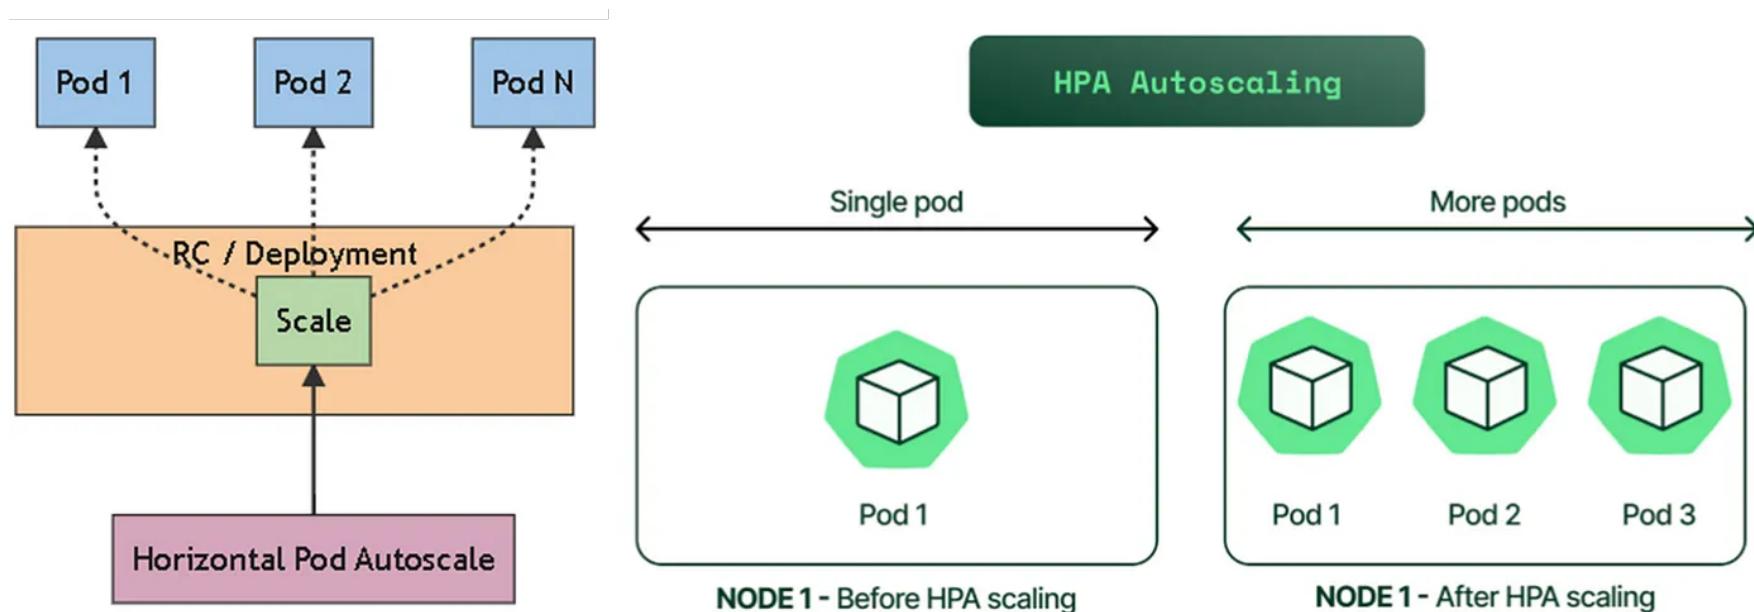


Workshop Ingress

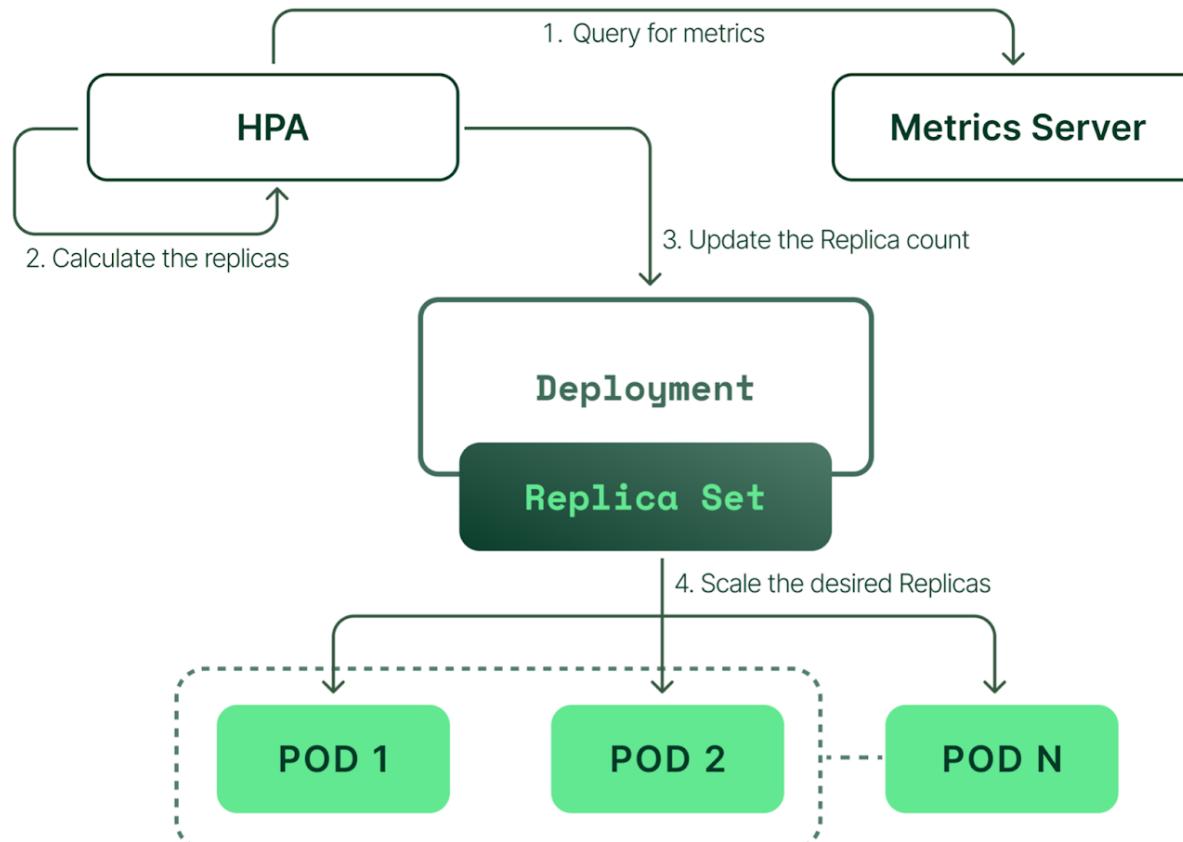
HPA (Horizontal Pod Autoscaler)

HPA

- A HorizontalPodAutoscaler (HPA for short) automatically updates a workload resource, with the aim of automatically scaling the workload to match demand



How does HPA work?





Workshop HPA

Namespace



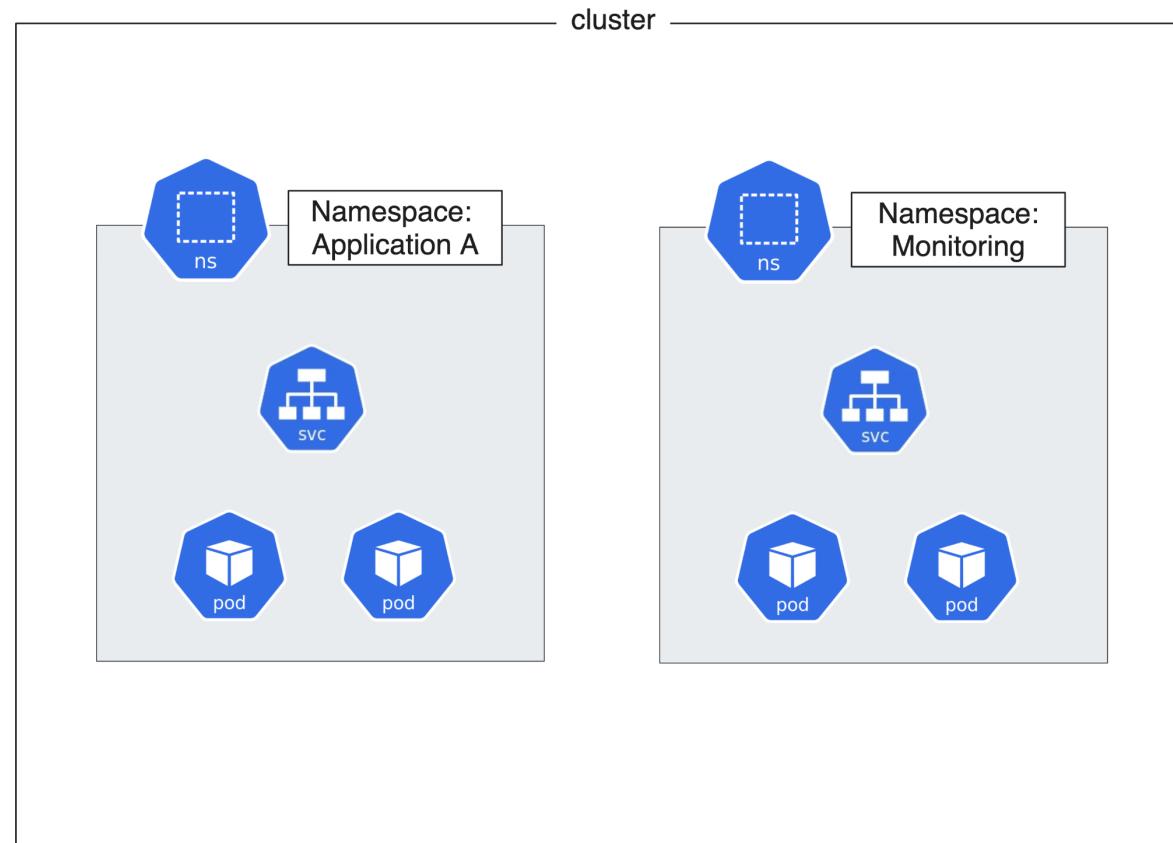
Namespace

- Partition a single kubernetes cluster into multiples clusters
- 4 defaults namespaces
 - **kube-system**: core system processes
 - **kube-public**: public accessible data, ex: ConfigMaps & Secrets
 - **kube-node-lease**: heartbeat of nodes, so that the control plane can detect node failure
 - **default**: your resources/objects



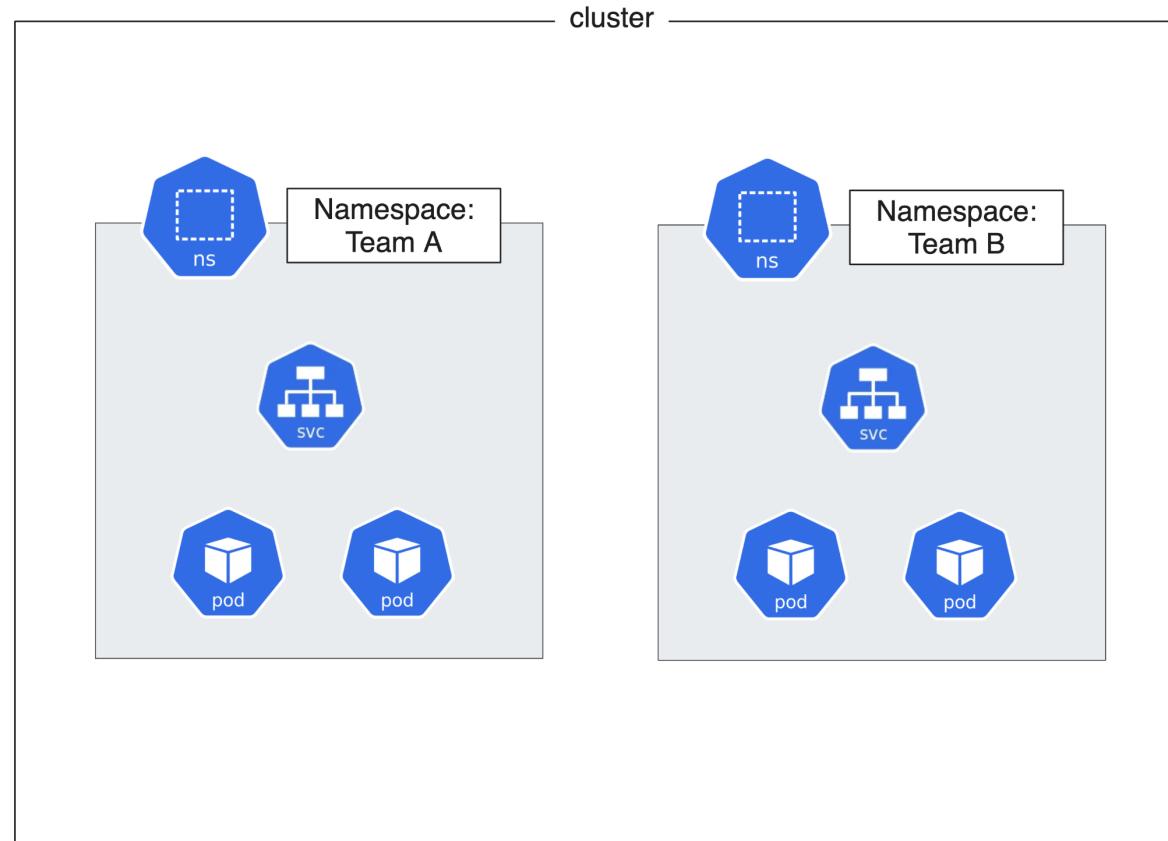
Namespace use cases

- Application domain



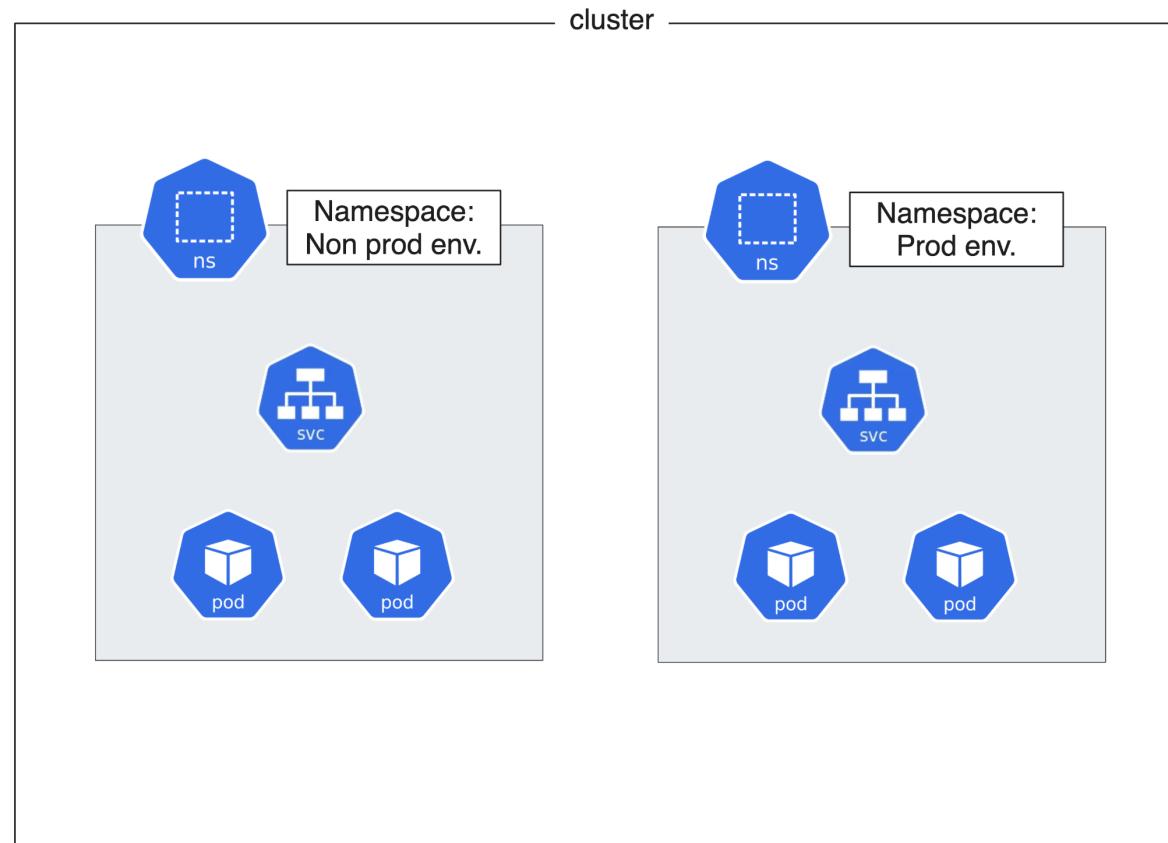
Namespace use cases

- Teams driven



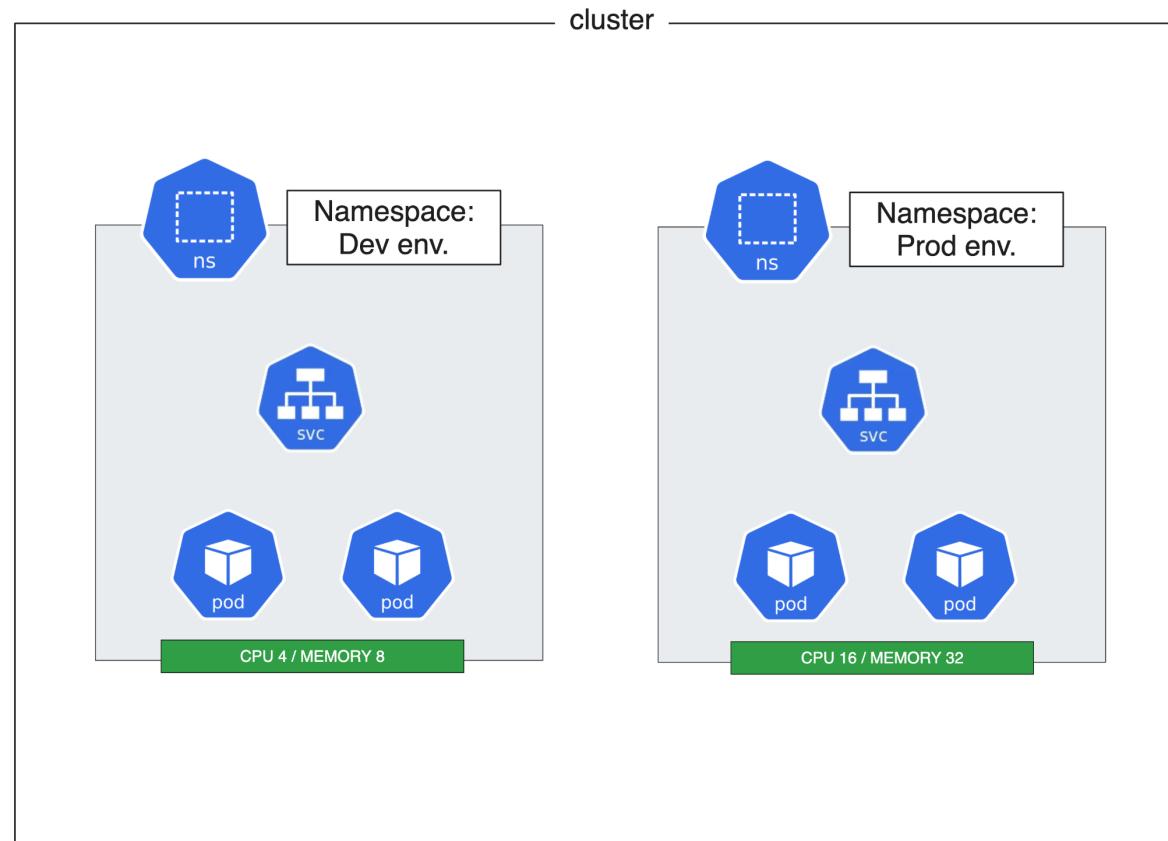
Namespace use cases

- Environments



Namespace use cases

- Resources quota limit



HELM



HELM

HELM is the package manager for Kubernetes



HELM

HELM is the package manager for Kubernetes

- Boosts productivity
- Reduces complexity
- Simple sharing
- Versioning and Upgrades



HELM Charts

Public HELM Chart: <https://artifacthub.io/>

The screenshot shows the Artifact Hub website with a teal header. The header includes the logo "Artifact HUB BETA", "SIGN UP", "SIGN IN", and a gear icon. Below the header, the text "Find, install and publish Kubernetes packages" is displayed. A search bar with the placeholder "Search packages" and a tip below it stating "Tip: Use or to combine multiple searches. Example: postgresql or mysql" is visible. Below the search bar, there are several buttons for filtering packages: "Packages from verified publishers", "Prometheus packages in official repositories", "Helm Charts provided by Bitnami", "Helm Charts in the storage category", and "OLM operators for databases". At the bottom, two large numbers are shown: "2086 PACKAGES" and "38119 RELEASES". The footer contains the text "Artifact Hub is an Open Source project" and links to GitHub, Slack, and Twitter.

SIGN UP SIGN IN

Find, install and publish
Kubernetes packages

Search packages

Tip: Use **or** to combine multiple searches. Example: [postgresql](#) [or](#) [mysql](#)

You can also [browse all packages](#) - or - try one of the sample queries:

Packages from verified publishers | Prometheus packages in official repositories | Helm Charts provided by Bitnami
Helm Charts in the storage category | OLM operators for databases

2086 | 38119

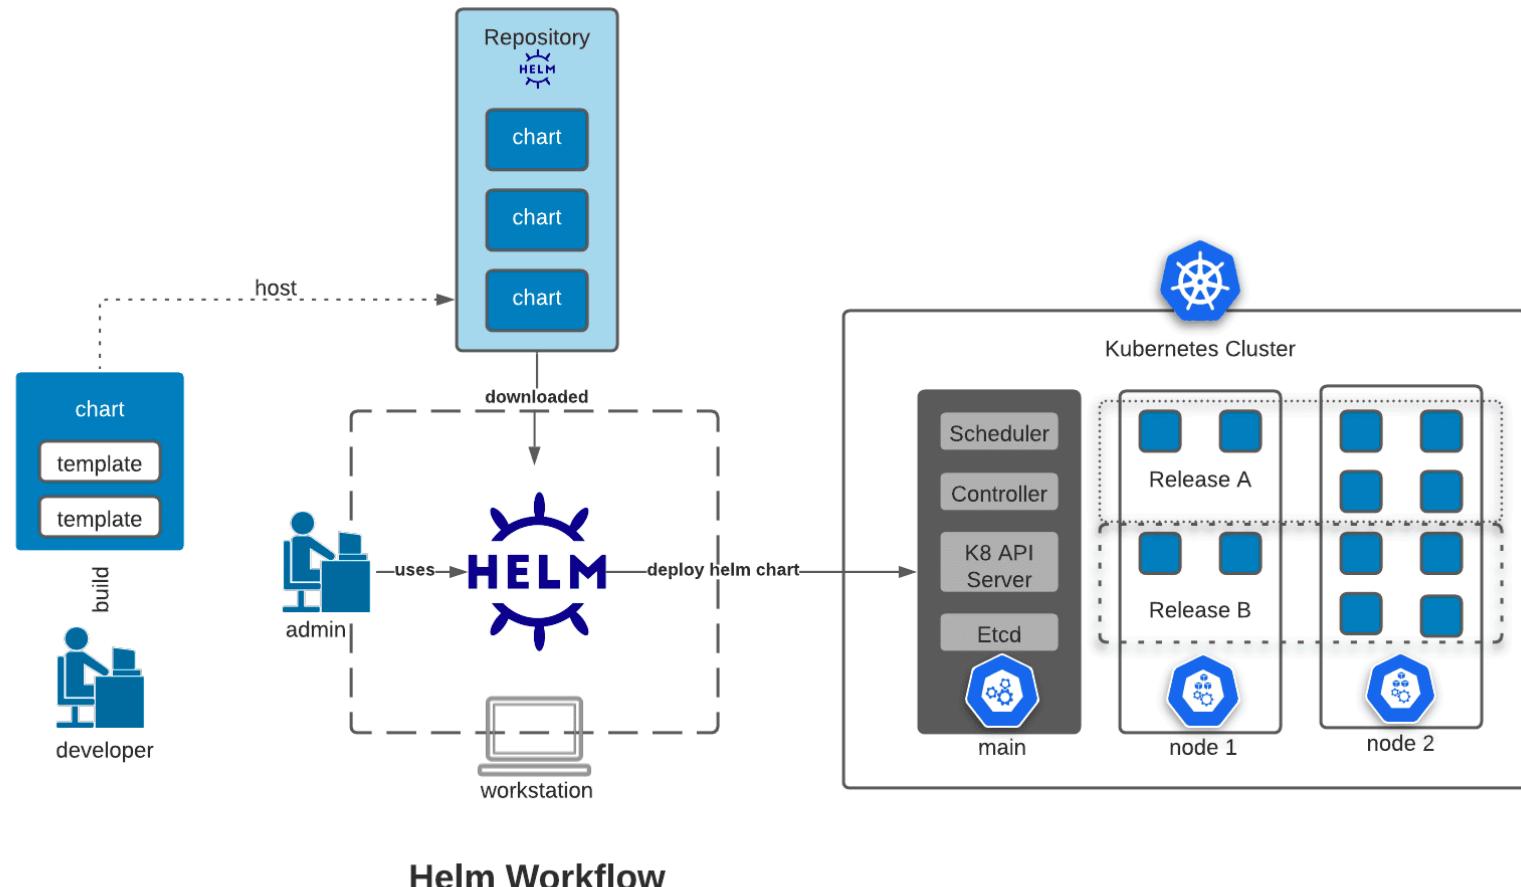
PACKAGES | RELEASES

Artifact Hub is an [Open Source](#) project

GitHub Slack Twitter



HELM Charts





Workshop HELM

