

# Kubernetes

---



# Agenda

- Introduction
- Kubernetes Architecture
- Cluster
- Resources / Objects
- HPA
- Namespace
- Helm



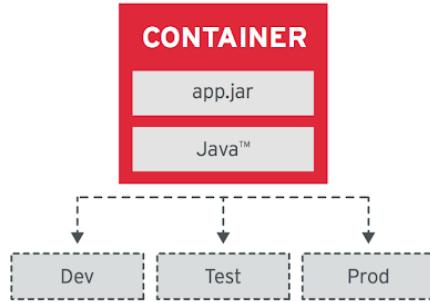
# Introduction

---



# Principles of Container-based

Image Immutability Principle



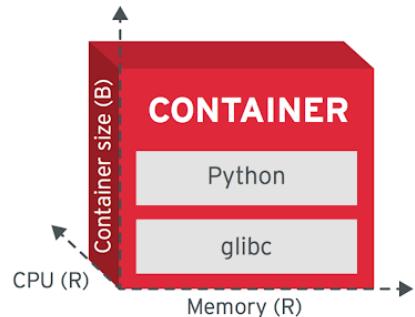
High Observability Principle



Process Disposability Principle



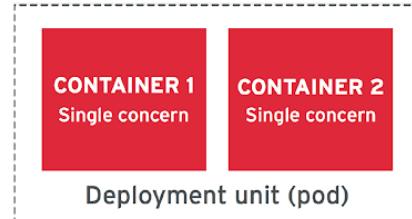
Runtime Confinement Principle



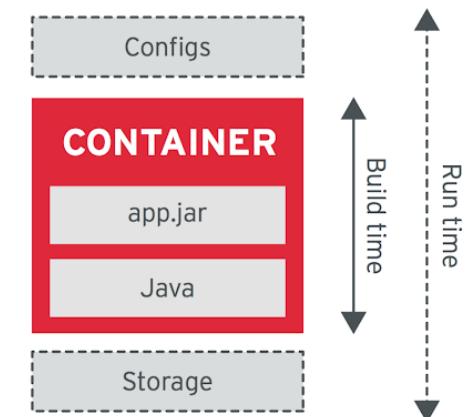
Lifecycle Conformance Principle



Single Concern Principle



Self-Containment Principle



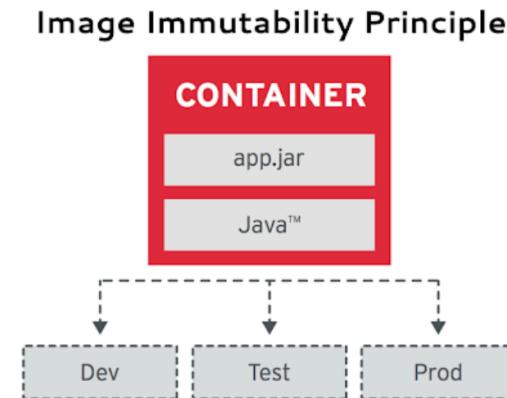
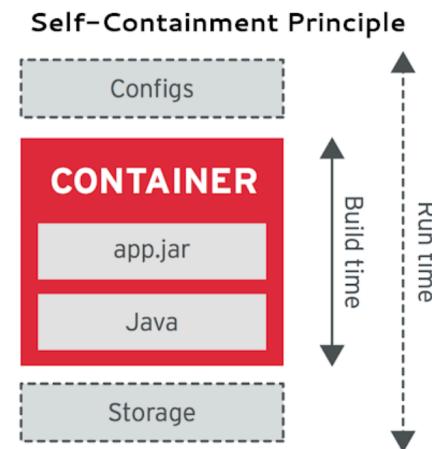
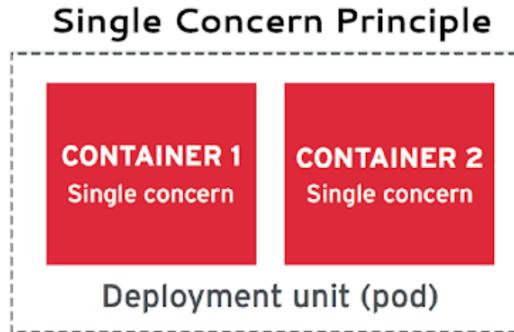
# Principles of Container-based

## Build Time

**Single Concern:** Each container addresses a single concern and does it well.

**Self-Containment:** A container relies only on the presence of the Linux kernel. Additional libraries are added when the container is built.

**Image Immutability:** Containerized applications are meant to be immutable, and once built are not expected to change between different environments.



# Principles of Container-based

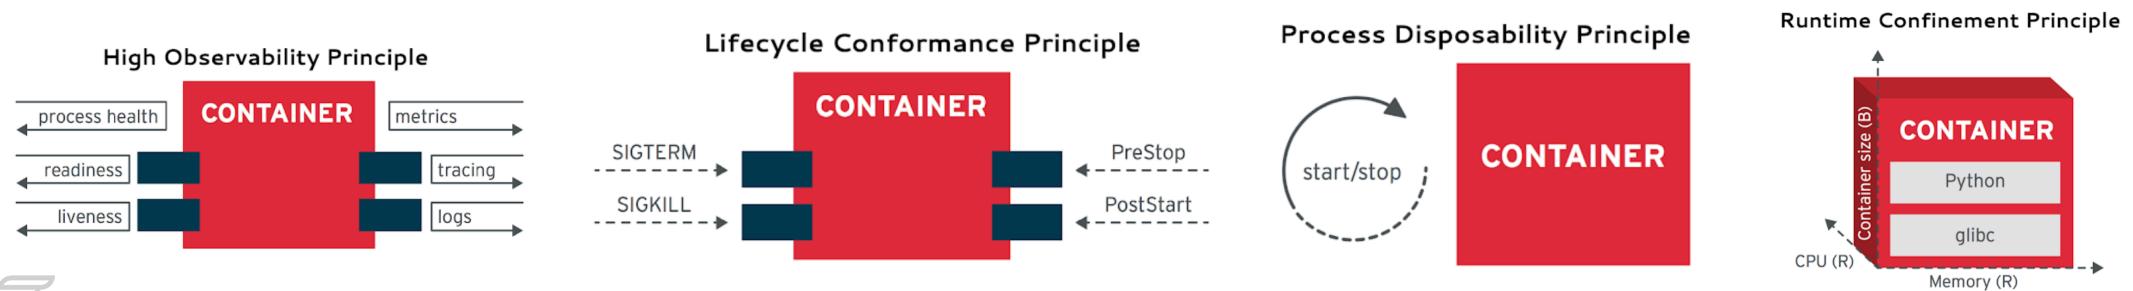
## Run Time

**High Observability:** Every container must implement all necessary APIs to help the platform observe and manage the application in the best way possible.

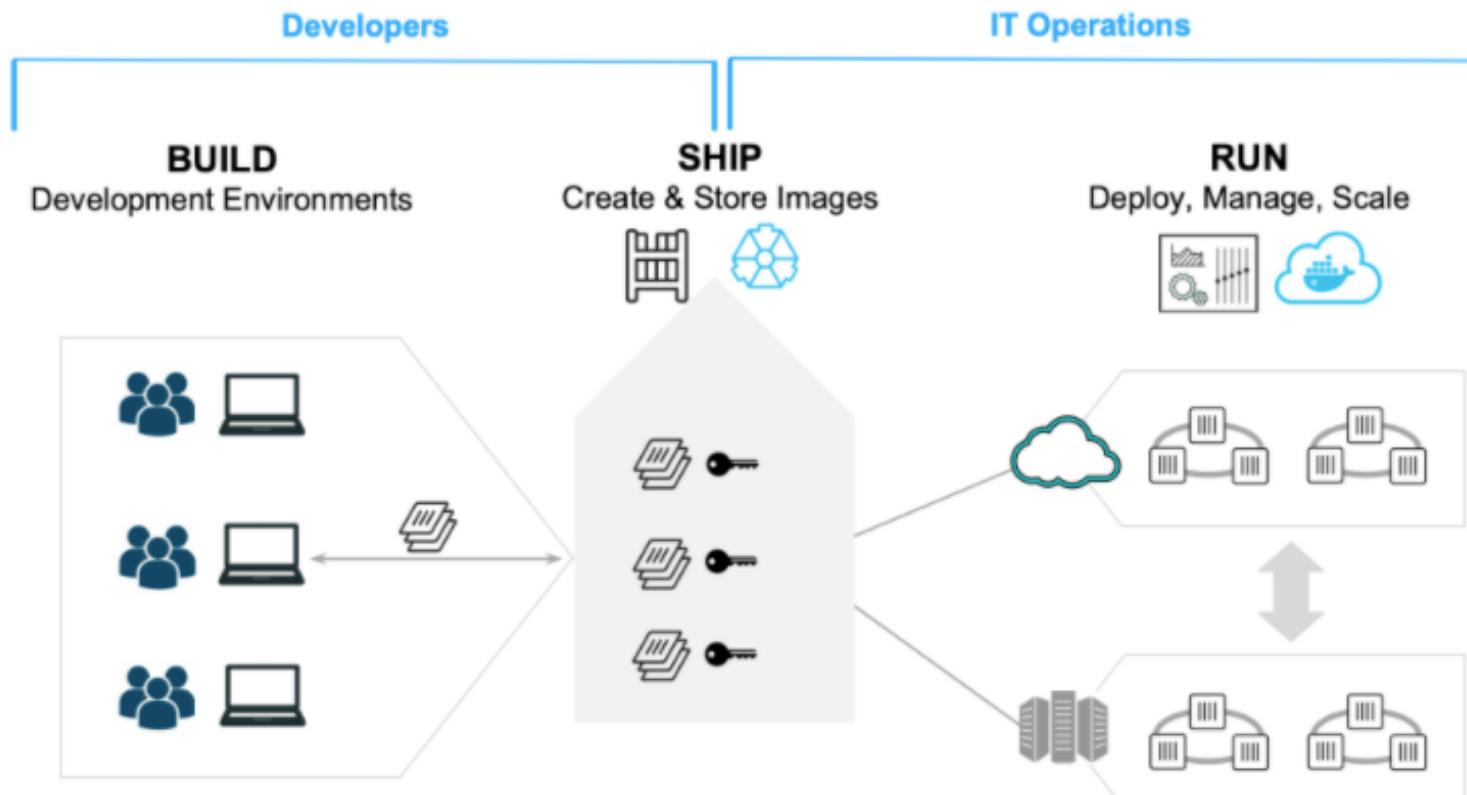
**Lifecycle Conformance:** A container must have a way to read events coming from the platform and conform by reacting to those events.

**Process Disposability:** Containerized applications must be as ephemeral as possible and ready to be replaced by another container instance at any point in time.

**Runtime Confinement:** Every container must declare its resource requirements and restrict resource use to the requirements indicated.

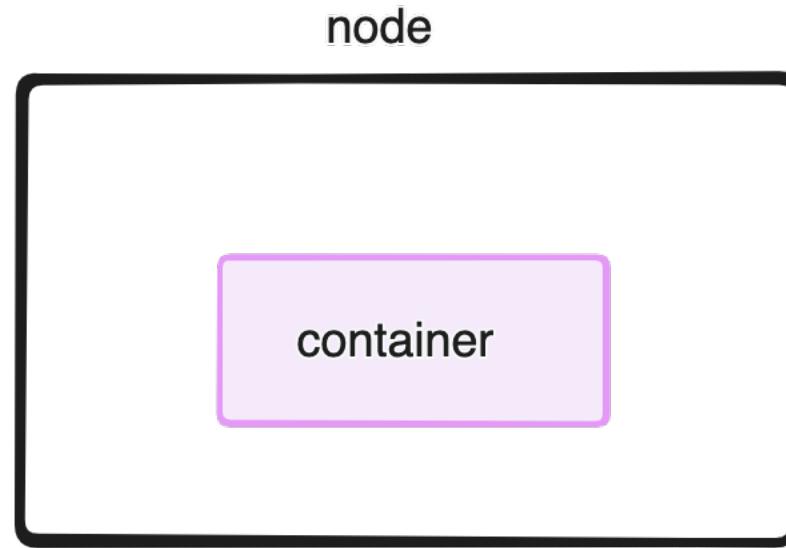


# Using Docker: Build, Ship, Run



# Docker standalone server

---



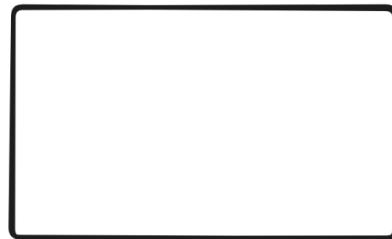
# Docker multi server

---

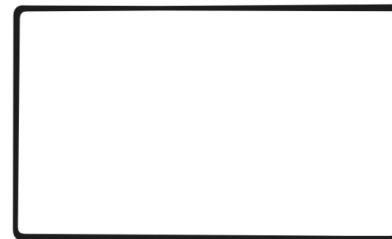
container

???

node



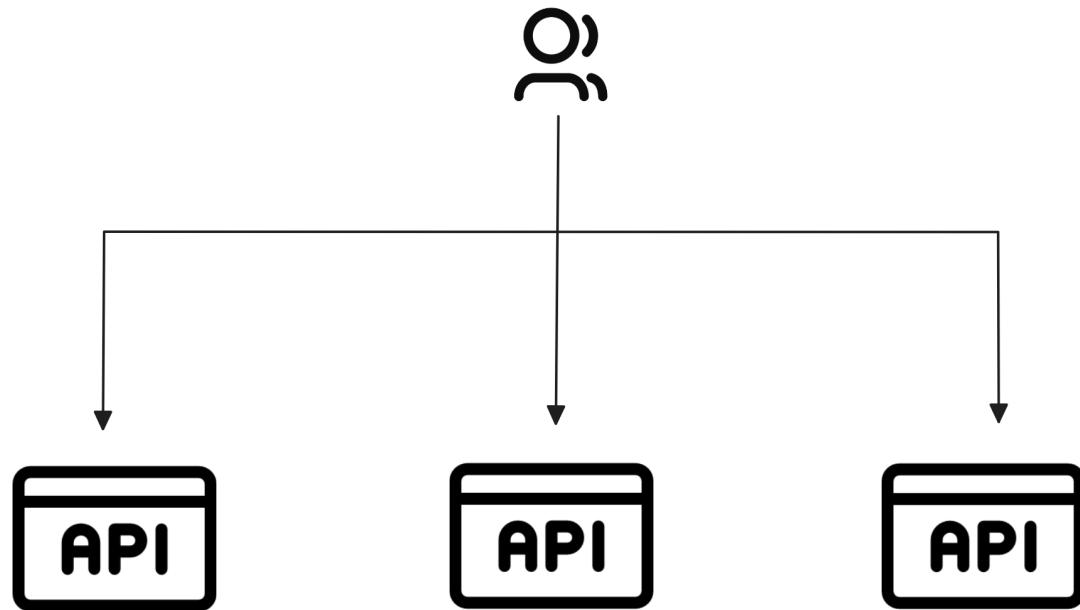
node



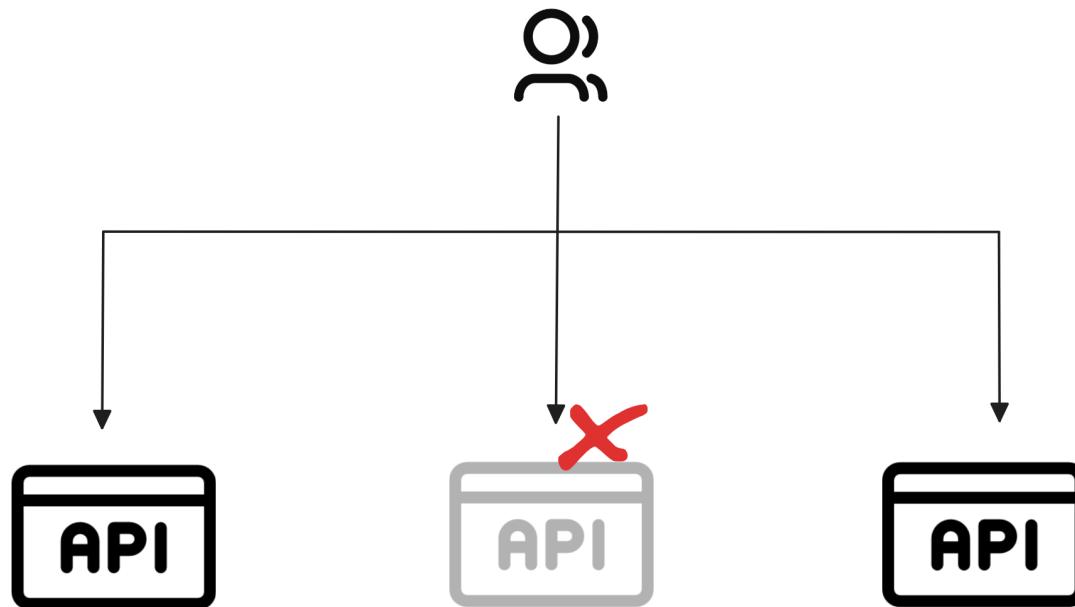
node



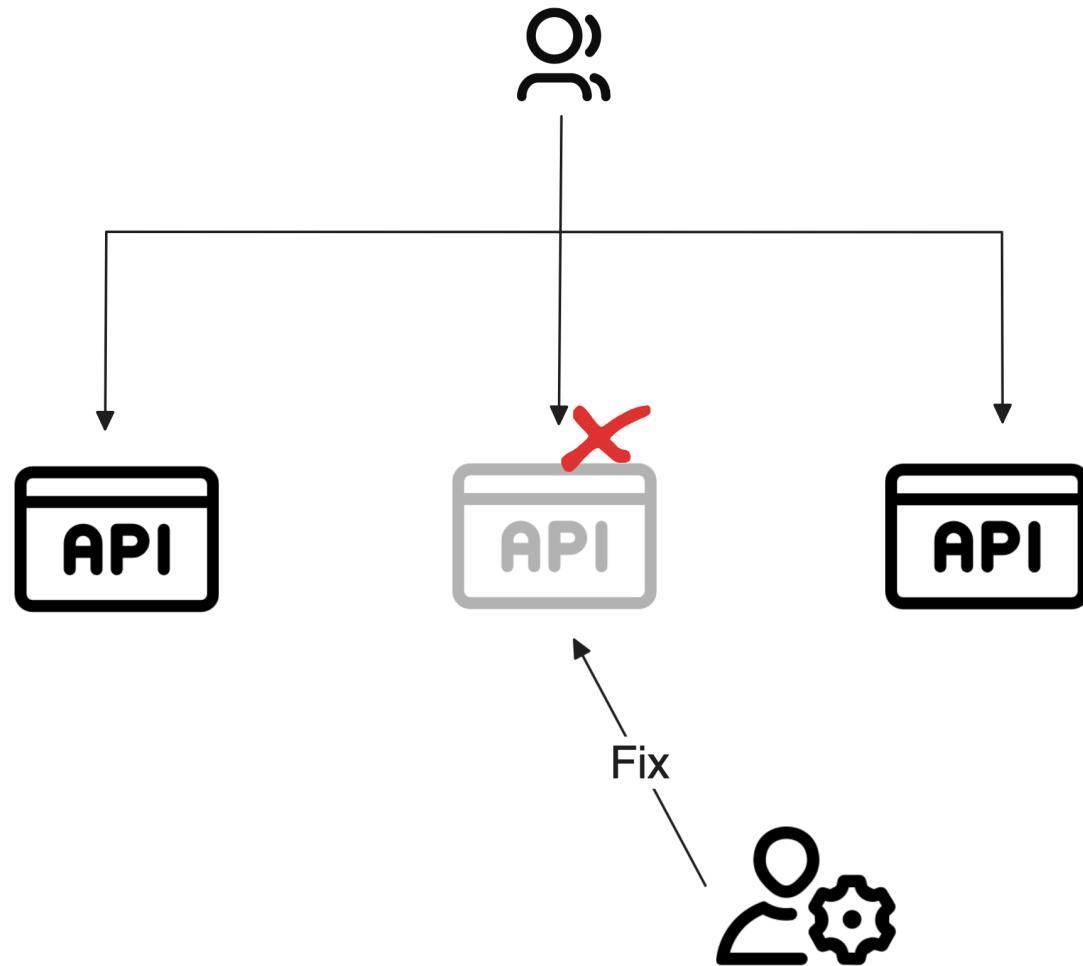
# Applications problem



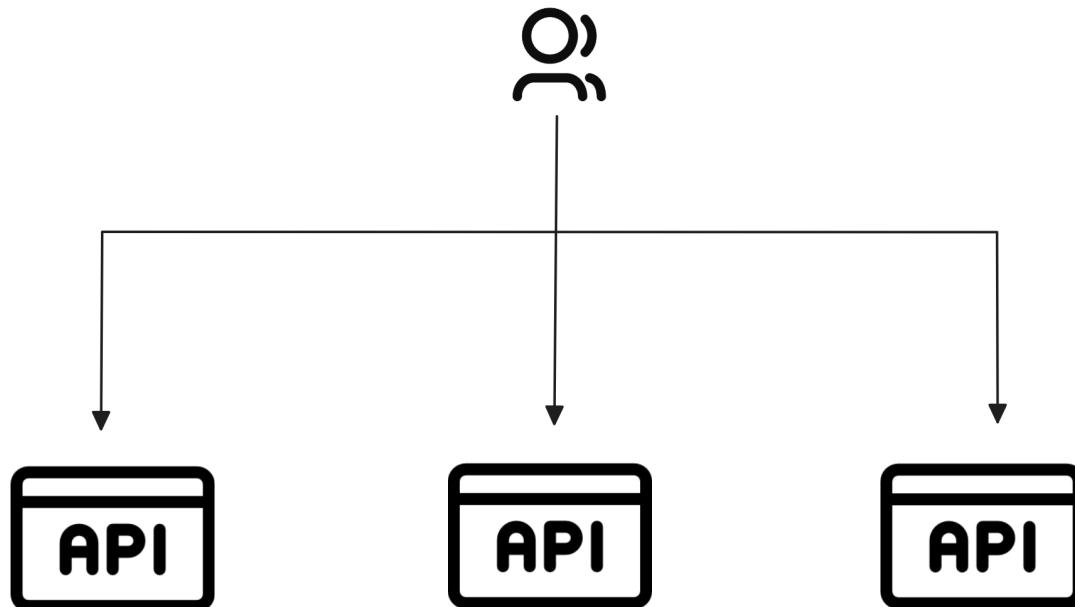
# Applications problem



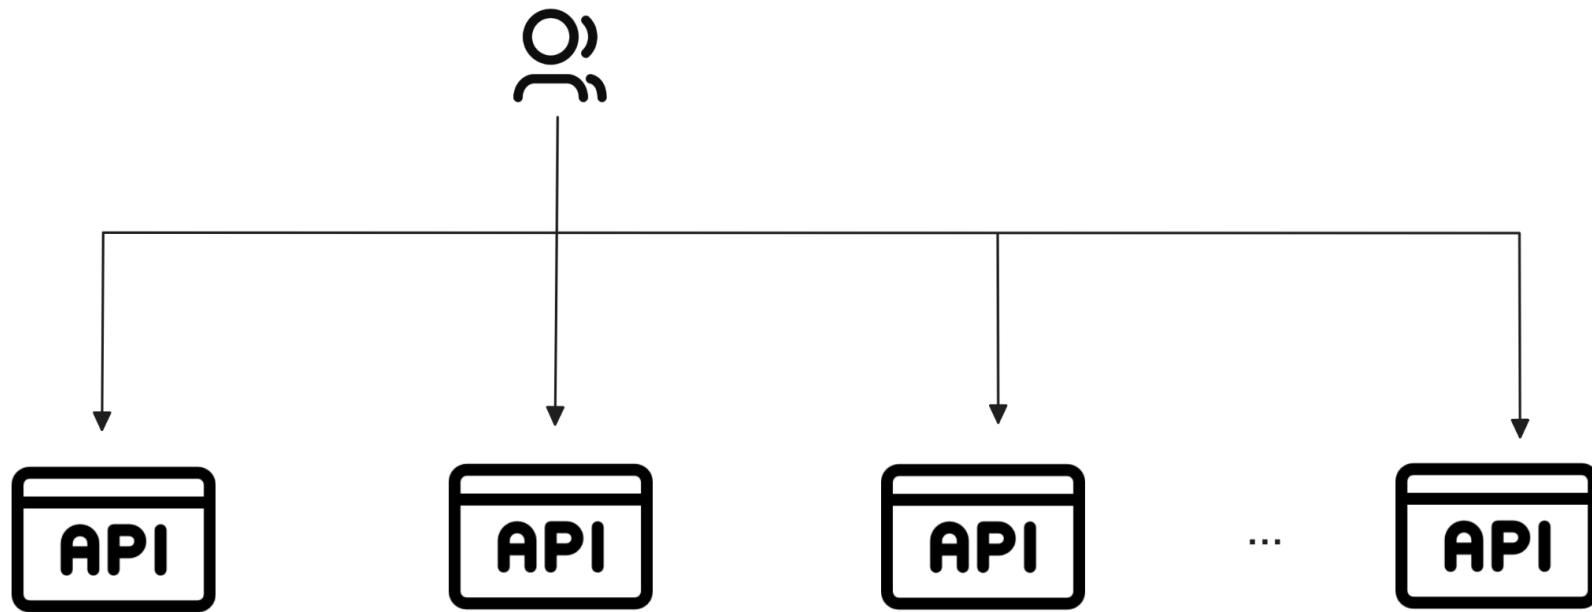
# Applications problem



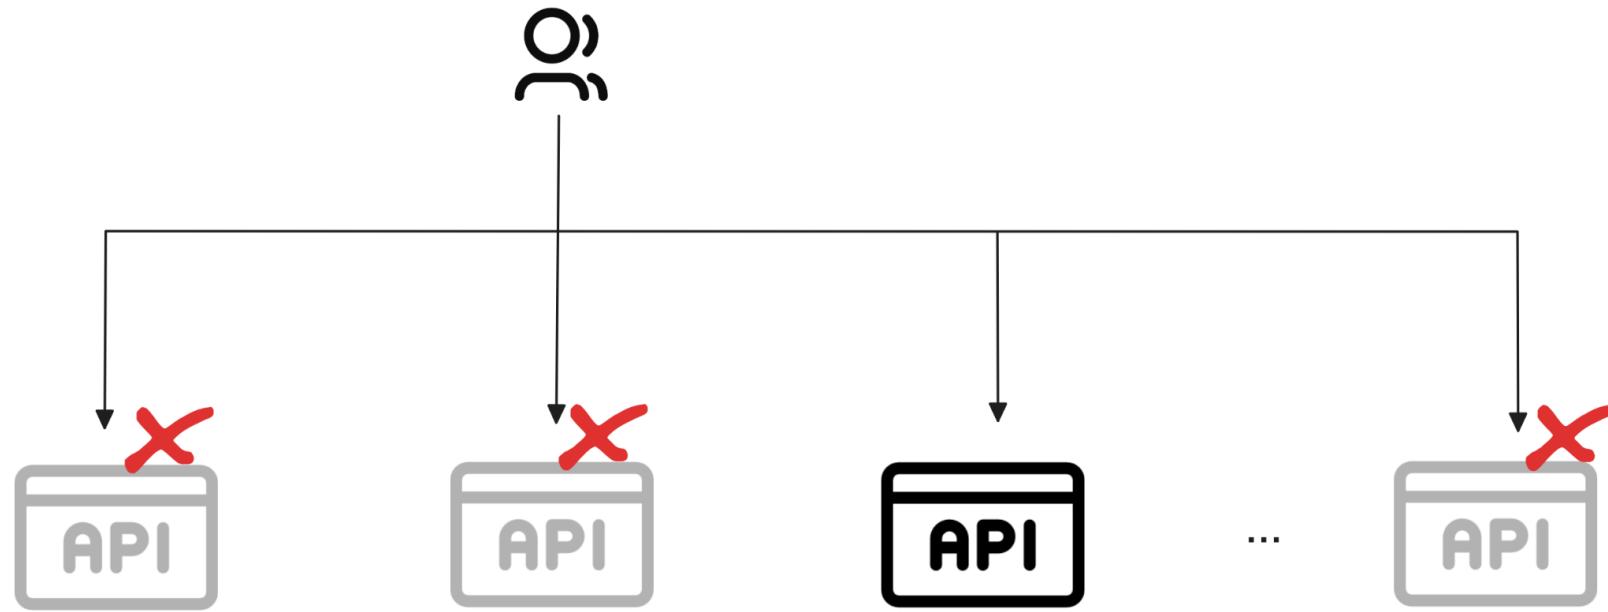
# Applications problem



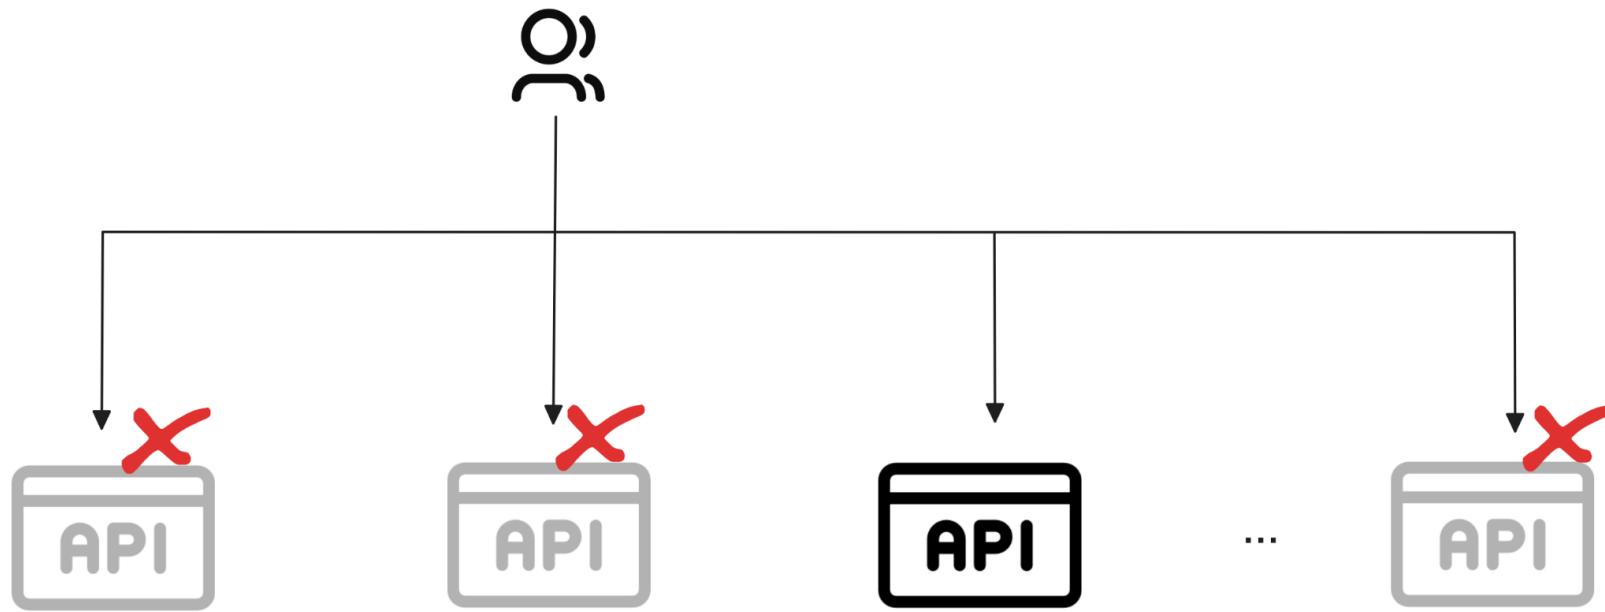
# Applications problem



# Applications problem



# Applications problem

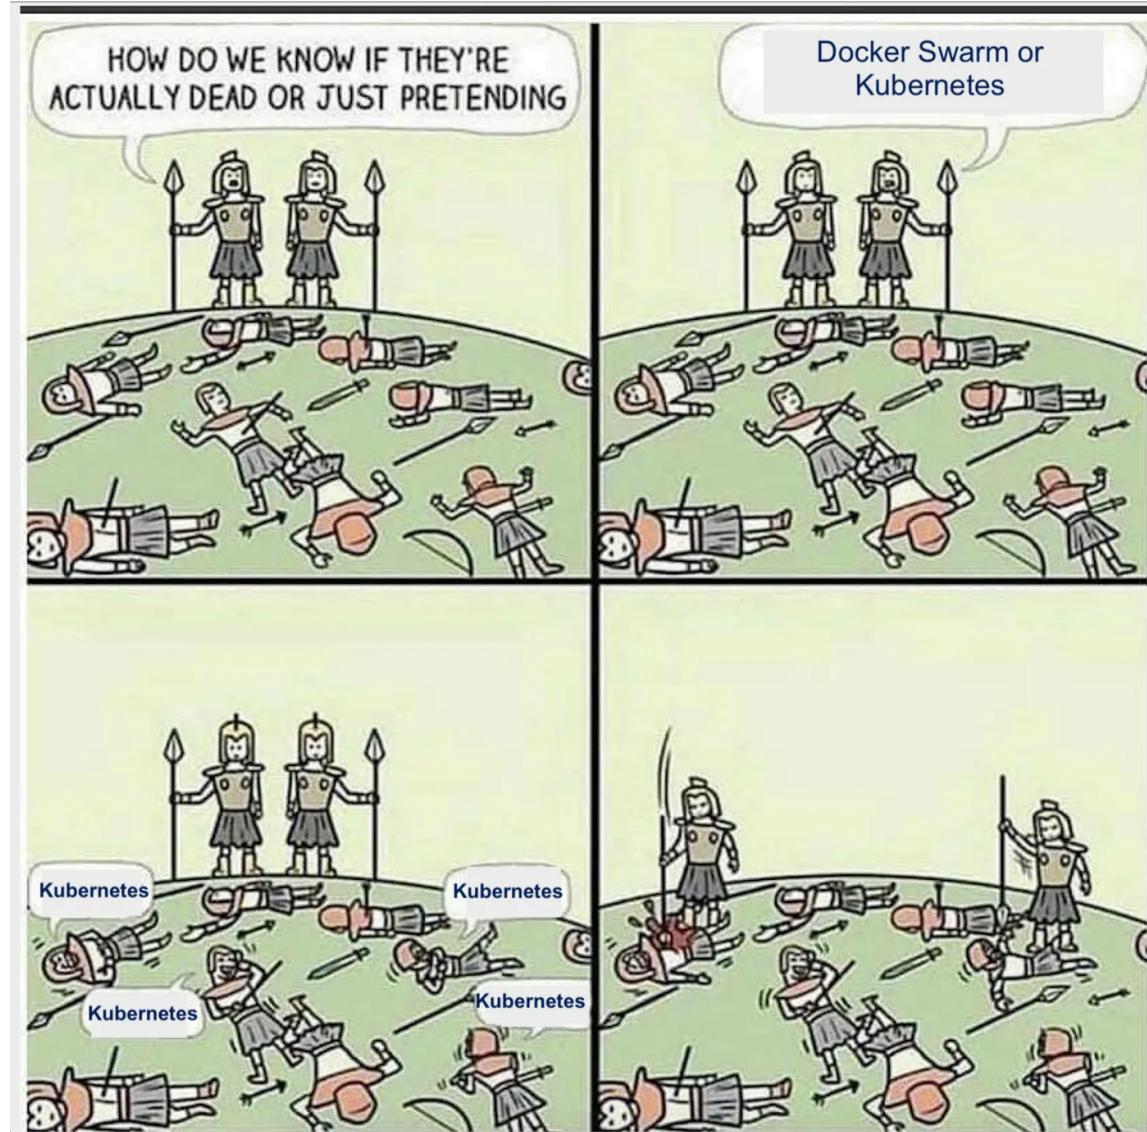


# What is Kubernetes ?

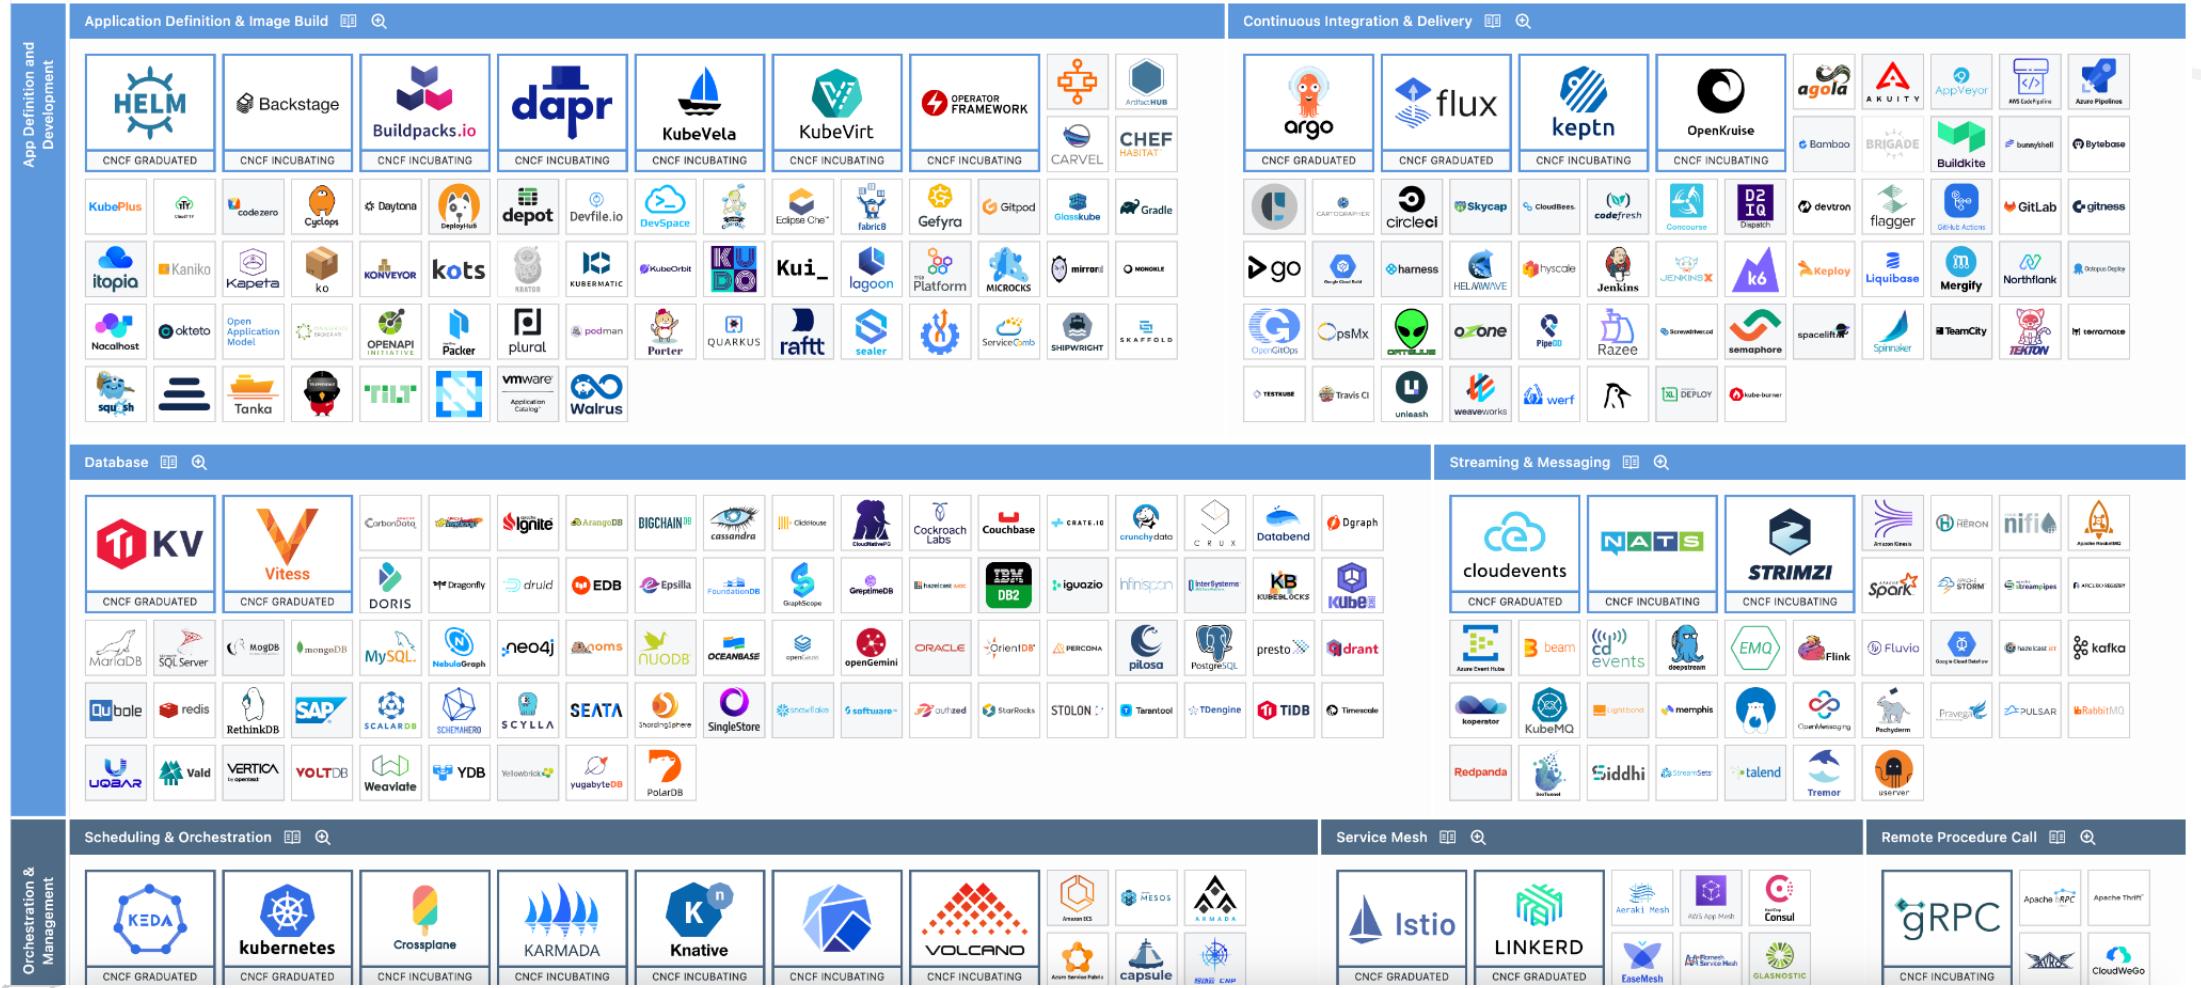
---

- Kubernetes in Greek means the Helmsman
- Also known as k8s
- Inspired and informed by Google's experiences and internal systems
- 100% opensource, written in GO
- Container orchestrator, runs and manages containers
- Kubernetes v1.0 released on July 2015, join CNCF on March 2016
- Container choose Docker, Container Orchestrator choose Kubernetes





# CNCF Landscape



# The features of Kubernetes

---

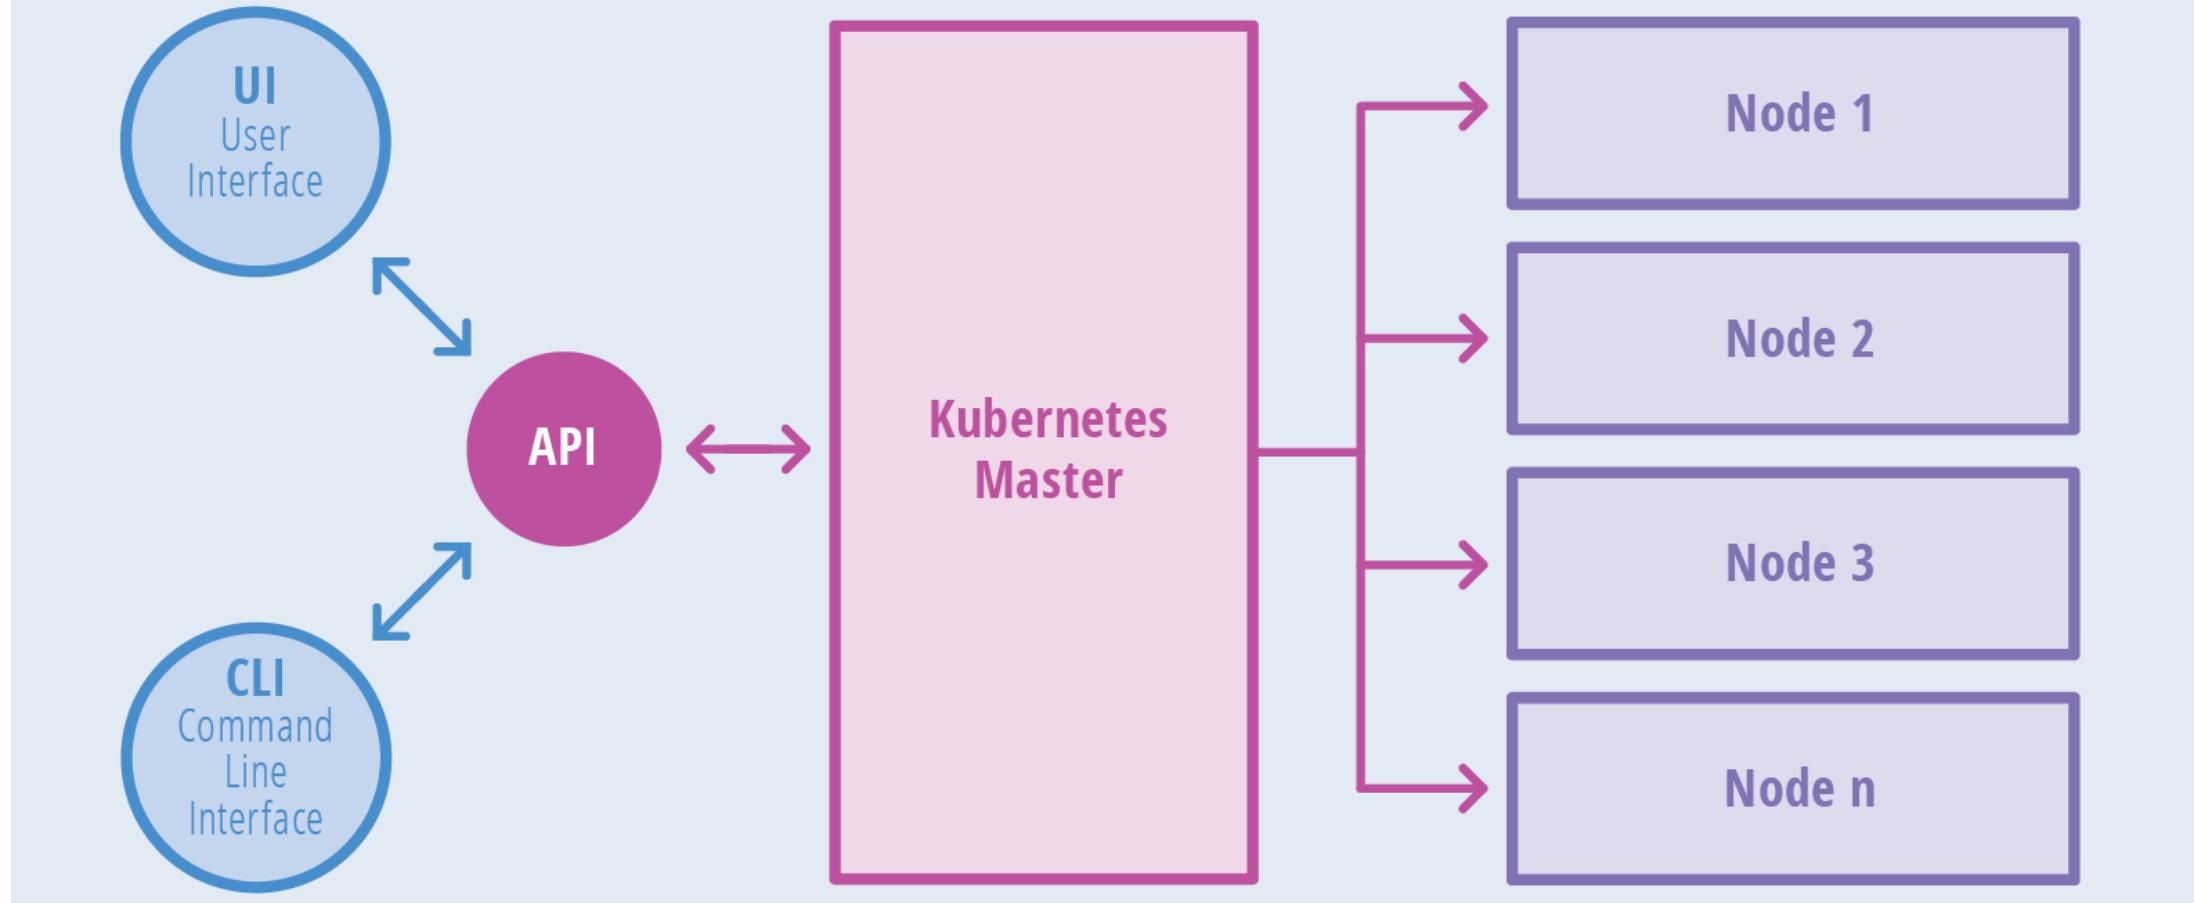
- Automatic bin packing
- Self healing
- Automated rollout & rollback
- Service Discovery & load Balancing
- Horizontal scaling
- Declarative Configuration



# Kubernetes Architecture

---

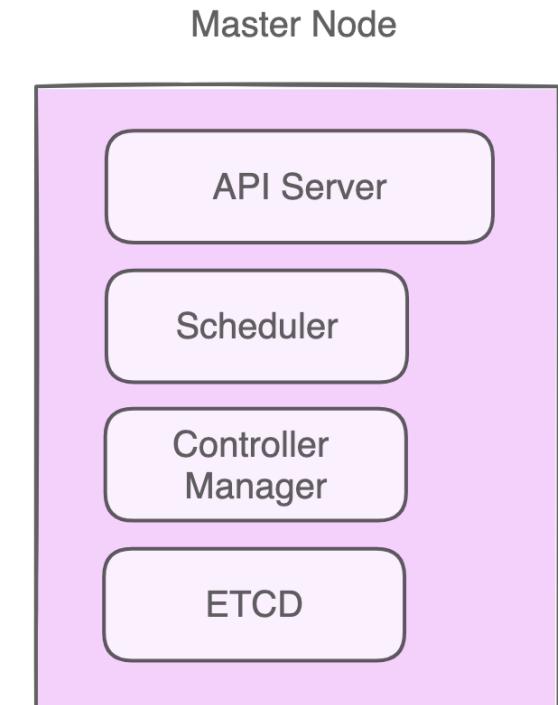
# Kubernetes Architecture



# Kubernetes Architecture

## Master node

- Responsible for managing cluster
- Access master node via CLI, UI or API
- For fault tolerance, can be more than 1 master in cluster
- Include 4 components: API Server, Scheduler, Controller manager and ETCD

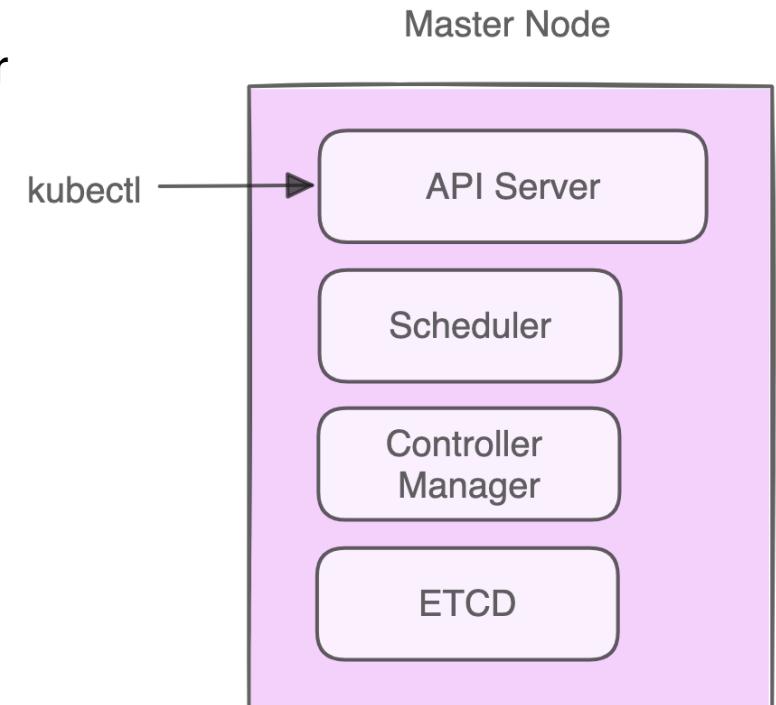


# Kubernetes Architecture

## Master node

### API Server

- Centralized component where all the cluster components are communicated
- Execute and validate from user command

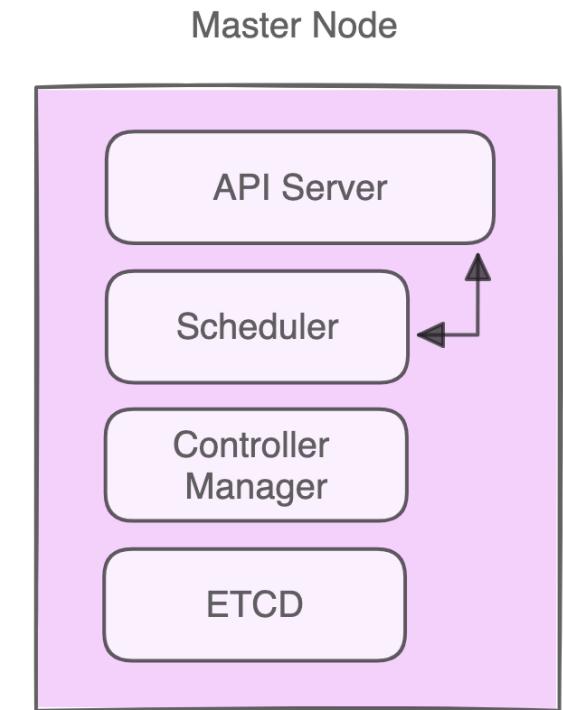


# Kubernetes Architecture

## Master node

### Scheduler

- Assigning the application to worker node
- Automatically detect which pod to place on which node based on all factors

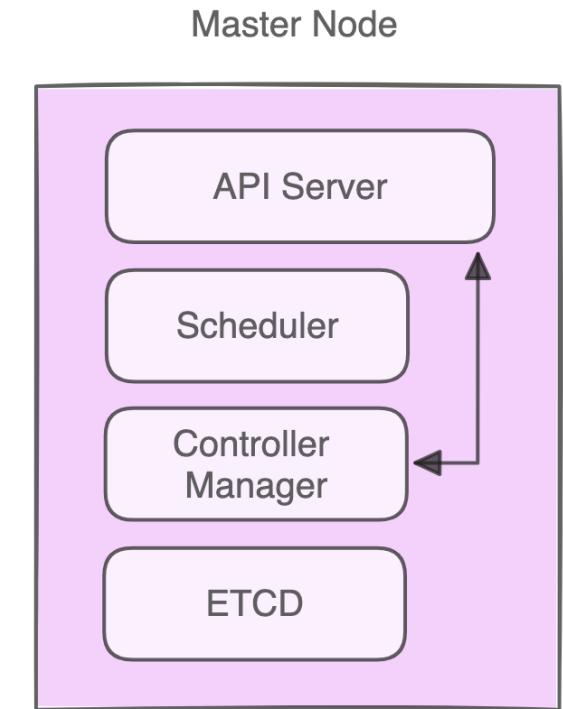


# Kubernetes Architecture

## Master node

### Controller Manager

- The Controller Manager maintains the cluster
- Handles node failures, replicating components, maintaining the correct number of pods, etc.

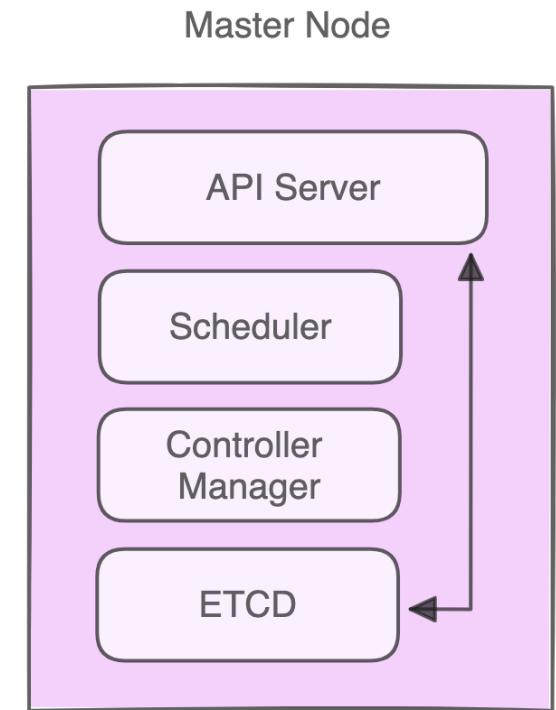


# Kubernetes Architecture

## Master node

### ETCD

- ETCD is a distributed reliable key-value store used by Kubernetes to store all data used to manage the cluster
- When you have multi master node in cluster, ETCD stores all information in a distributed manner
- Application data doesn't store in ETCD

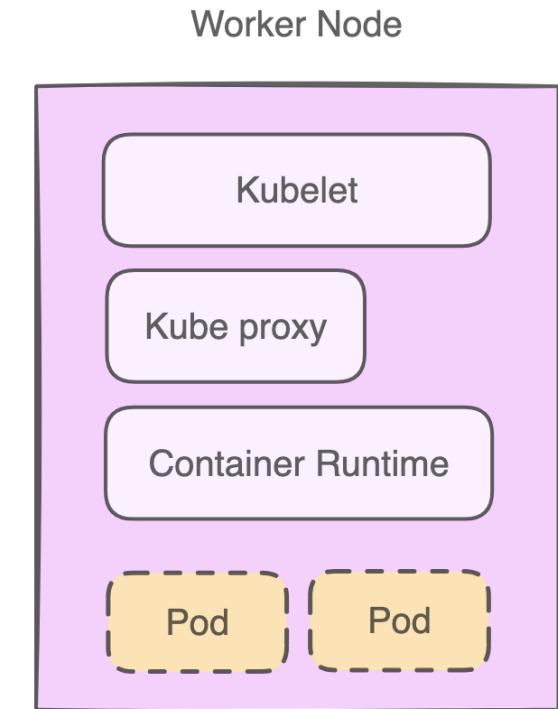


# Kubernetes Architecture

## Worker node

- Responsible about applications
- Include 3 components: **Container Runtime, Kubelet and Kube**

## Proxy

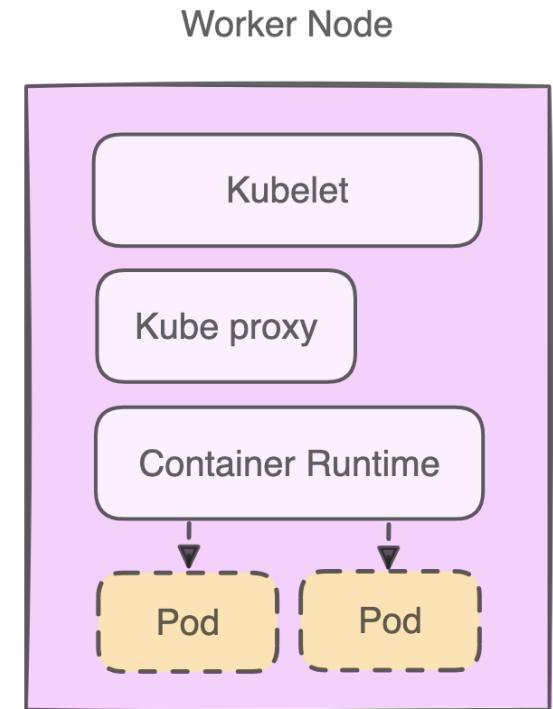


# Kubernetes Architecture

## Worker node

### Container Runtime

- Container runtime which runs the containers like Docker, crio or containers
- Pulling the images and run the containers

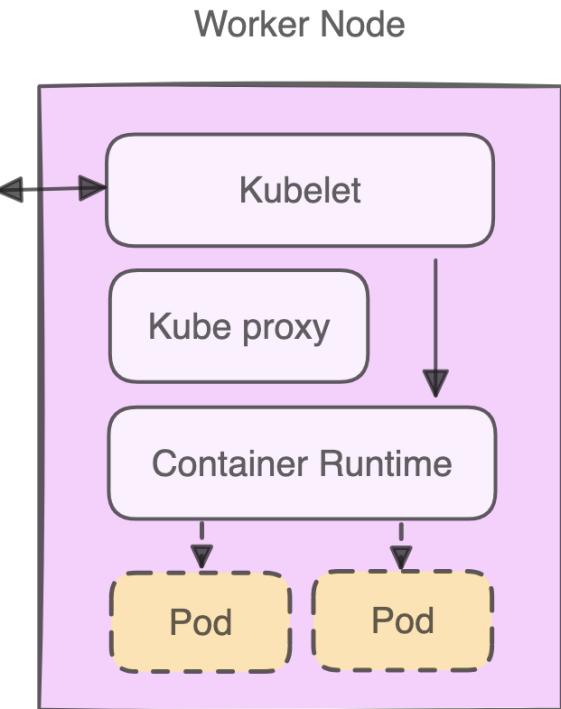


# Kubernetes Architecture

## Worker node

### Kubelet

- Interacting with the master node
- Contact the container runtime for create pod depend on Po
- Provide health information of the worker node

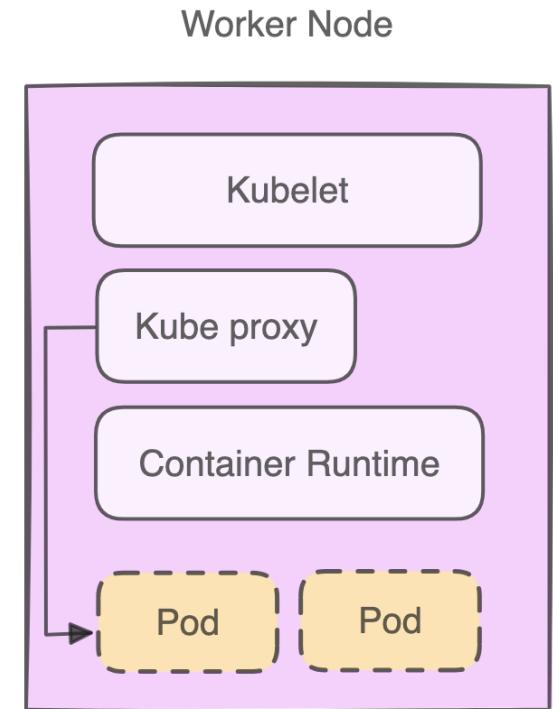


# Kubernetes Architecture

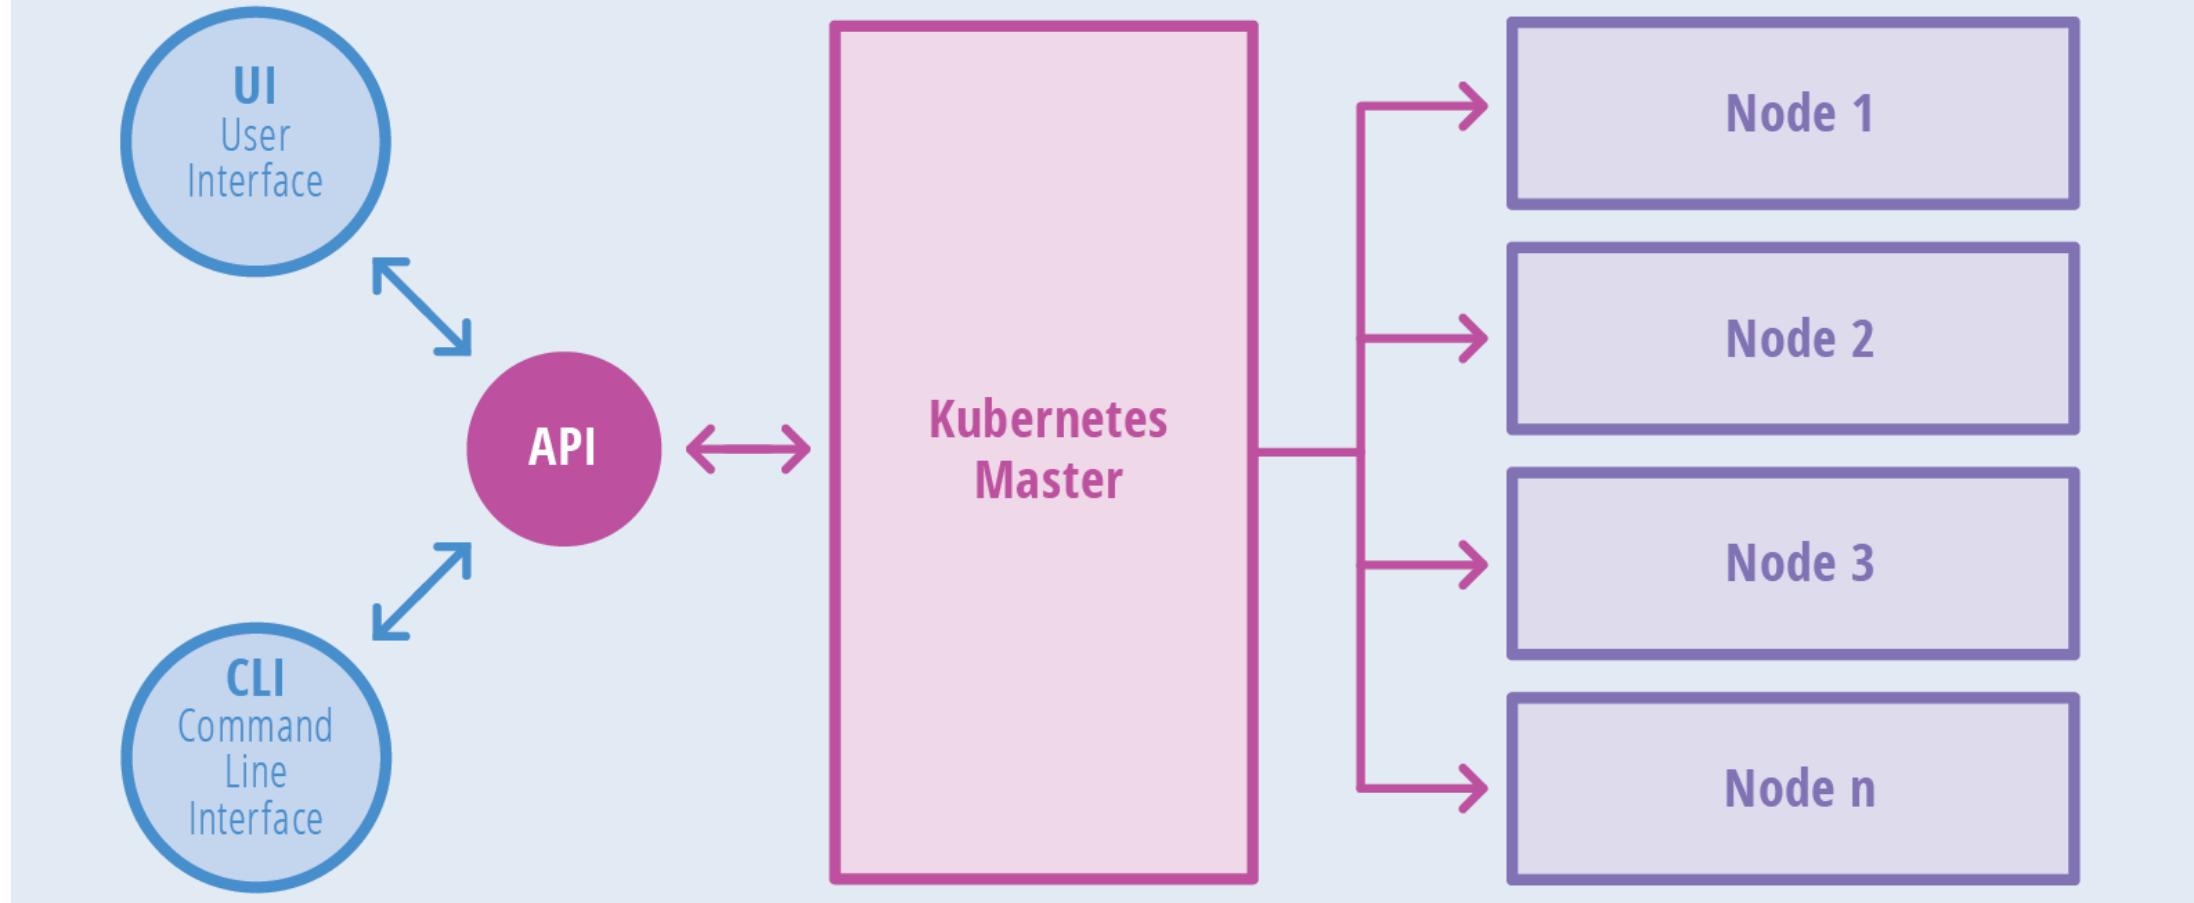
## Worker node

### Kube Proxy

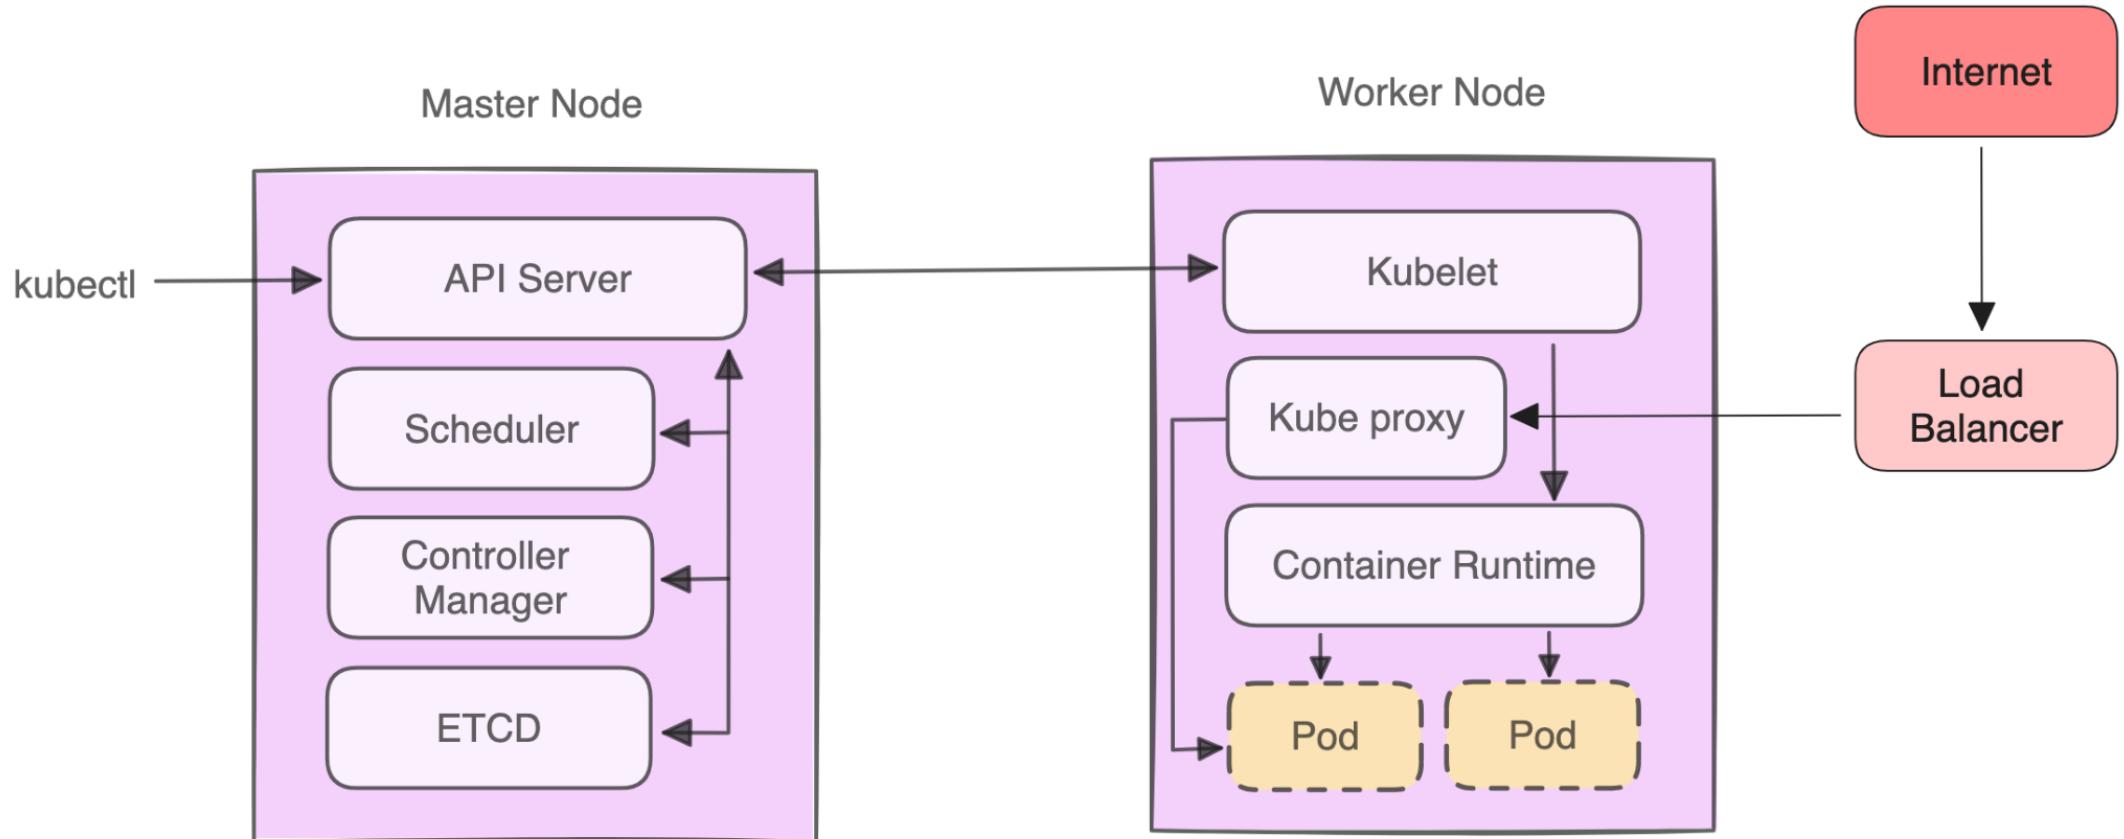
- Responsible for ensuring network traffic is routed properly to internal and external services
- Manage route for reduce latency time for example pod will communicate with another pod on same node first



# Kubernetes Architecture



# Kubernetes Architecture



# Kubectl

---



# Kubectl

---

- Command line
- Control the Kubernetes cluster manager
- CRUD Kubernetes resources



# Kubectl example

---

- \$ kubectl apply
- \$ kubectl create
- \$ kubectl get
- \$ kubectl describe
- \$ kubectl delete



# Cluster

---



# Cluster

---

- On Cloud
- On Premise
- On local development



# Cluster on cloud

---



Amazon **EKS**



Google Kubernetes Engine



Azure Kubernetes Service (**AKS**)

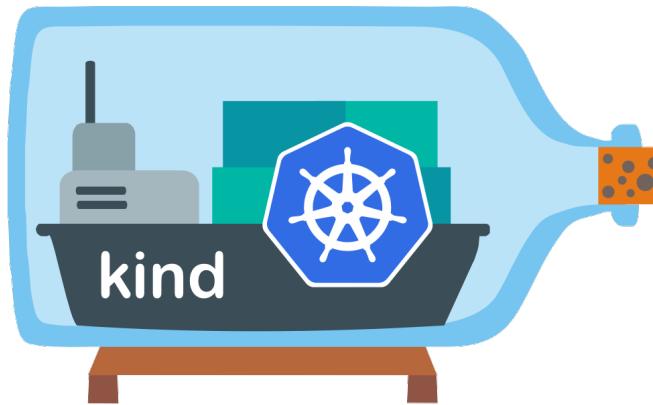
 Alibaba Cloud



# Cluster on premise



# Cluster on local development



# Cluster comparing

Feature	Kubernetes on Cloud	Kubernetes On-Premise
Setup & Maintenance	<ul style="list-style-type: none"><li>- Easy to set up with automated updates and patches</li><li>- Minimal operational overhead</li></ul>	<ul style="list-style-type: none"><li>- Requires manual setup</li><li>- Ongoing maintenance and upgrades are the user's responsibility</li></ul>
Cost	<ul style="list-style-type: none"><li>- High depend cloud provider</li><li>- Costs tied to usage</li><li>- No upfront infrastructure investment</li></ul>	<ul style="list-style-type: none"><li>- Lower costs</li><li>- High initial investment for hardware and infrastructure</li></ul>
Scalability	<ul style="list-style-type: none"><li>- Highly scalable with near-infinite resources.</li><li>- Automatic scaling options available</li></ul>	<ul style="list-style-type: none"><li>- Limited by on-premise hardware resources.</li><li>- Scaling may require additional hardware and setup.</li></ul>
Disaster Recovery	<ul style="list-style-type: none"><li>- Built-in disaster recovery and backups.</li></ul>	<ul style="list-style-type: none"><li>- Must design and implement your own</li></ul>
Customization	<ul style="list-style-type: none"><li>- Limited by cloud provider's offerings and configurations.</li></ul>	<ul style="list-style-type: none"><li>- Full control over configurations and infrastructure.</li></ul>



# Cluster comparing

Feature	Kubernetes on Cloud	Kubernetes On-Premise
Security	- Security managed by cloud providers	- Full control over security configurations
Networking	- Advanced networking features available (e.g., VPCs, load balancers). - Simplified setup with integration to other cloud services	- Full control over network configurations. - Requires in-house expertise for complex networking setups
Upgrades & New Features	- Automatic access to the latest features and updates	- Full control over when and how to upgrade - Delayed to the latest features unless manually updated

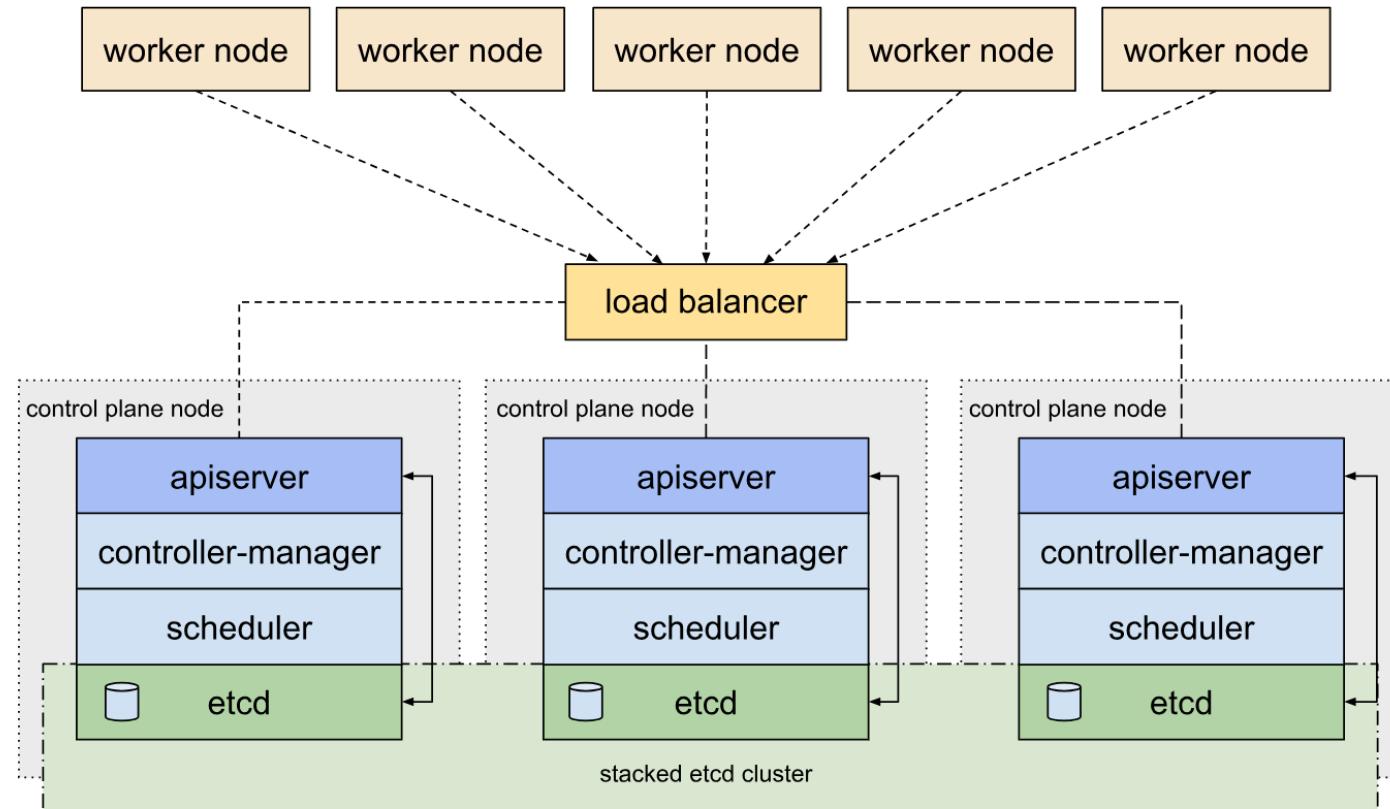


# Options for Highly Available

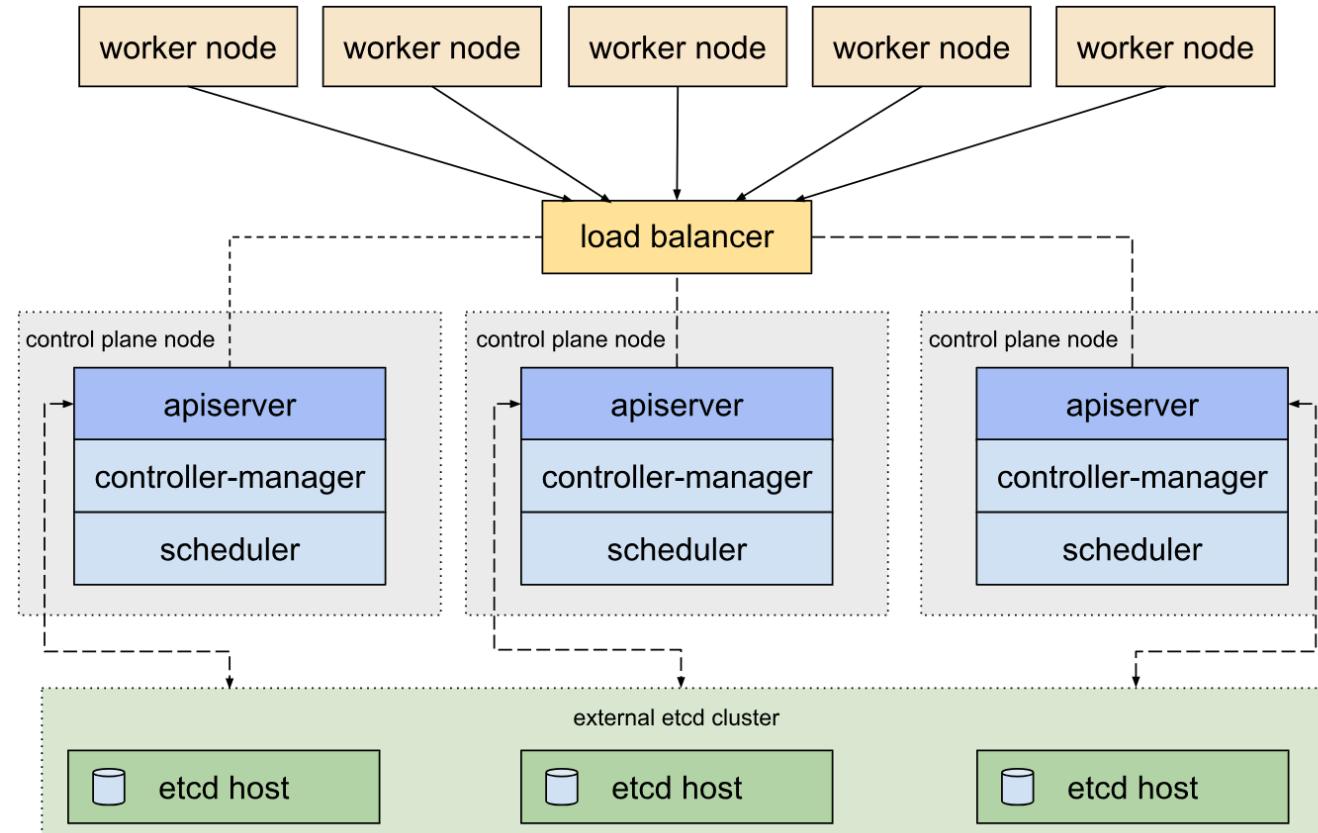
- Stacked etcd topology
- External etcd topology



# Stacked etcd topology



# External etcd topology



# Fault Tolerance Table

Cluster Size	Majority	Failure Tolerance
1	1	0
2	2	0
3	2	<b>1</b>
4	3	1
5	3	<b>2</b>
6	4	2
7	4	<b>3</b>
8	5	3
9	5	<b>4</b>



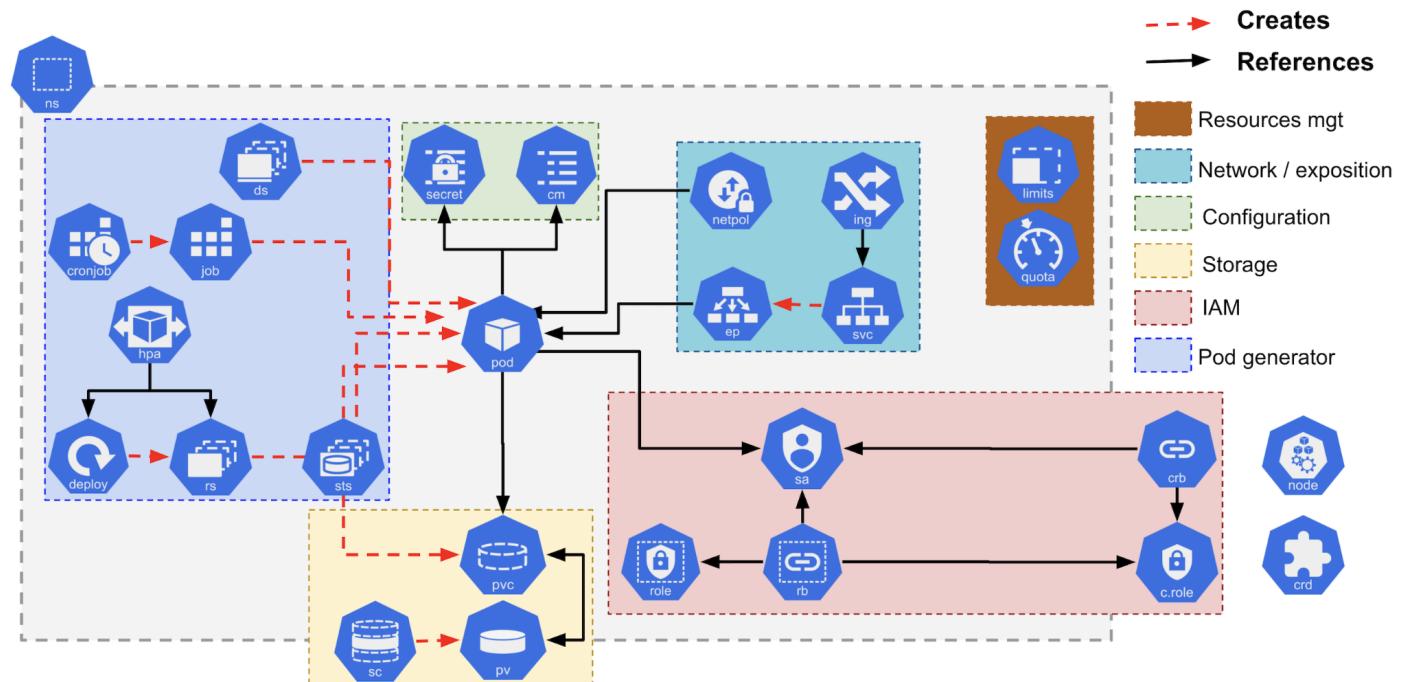
# Resources / Objects

---

# Resources / Objects

- Pod
- Deployment
- Service
- Ingress
- Config map
- Secret
- PV & PVC

## Kubernetes Resources Map



# Pod

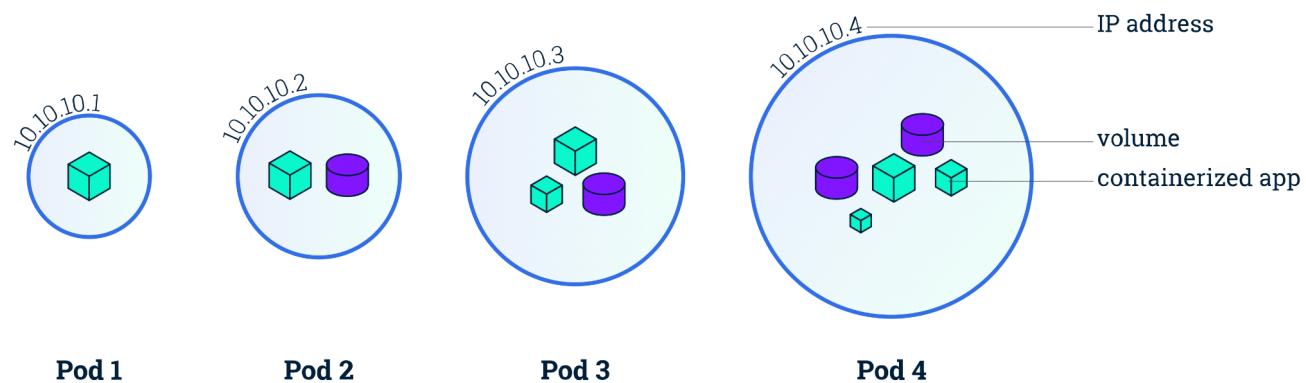
---



# Pod

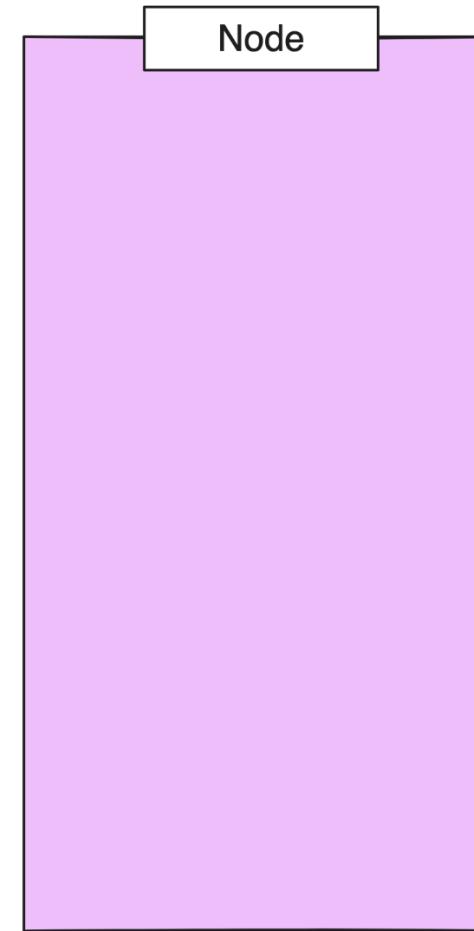
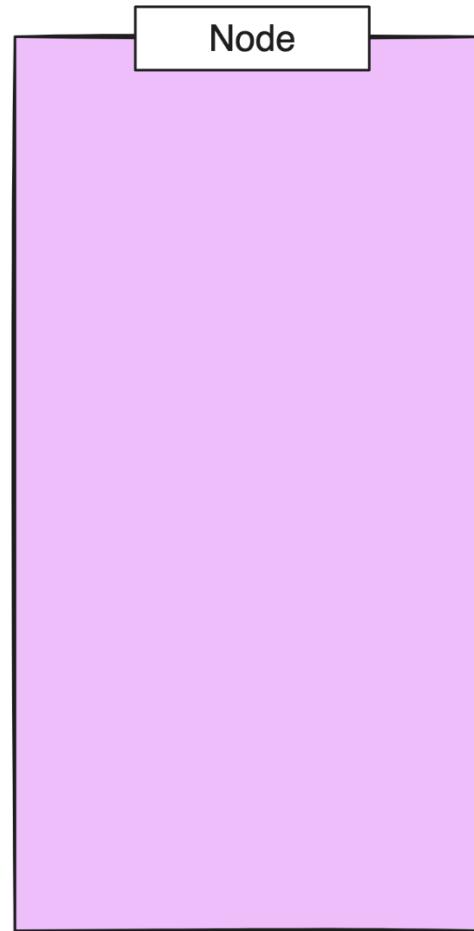
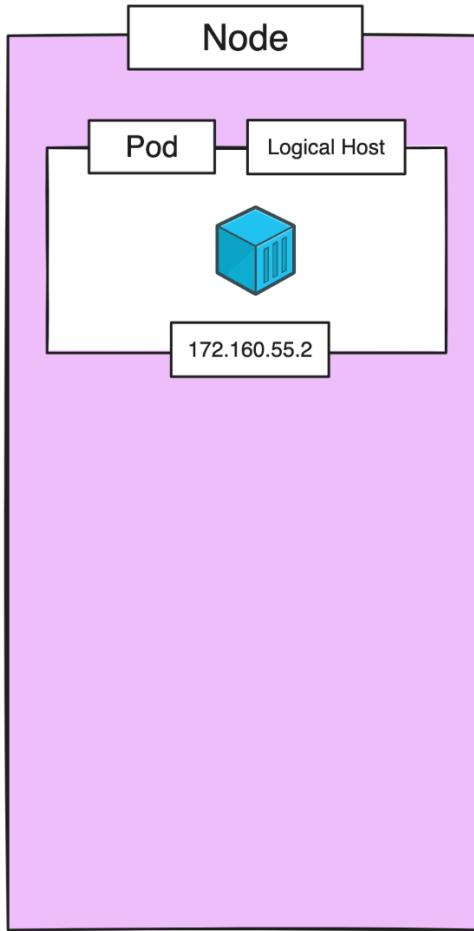
---

- Smallest deployable units of computing that you can create and manage in Kubernetes
- Logical Application
- One or more container in Pod (recommend 1 container per pod)
- Any containers in same pod will share storage volumes and network resources
- Each pod is assigned a unique ip



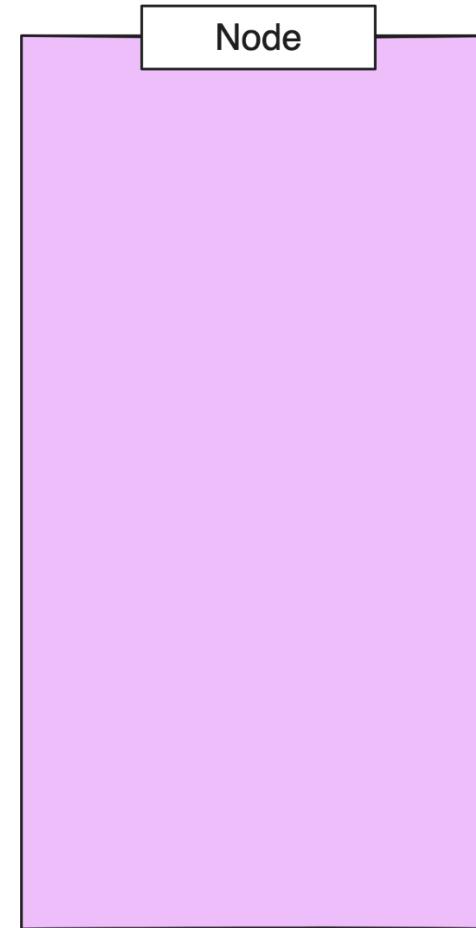
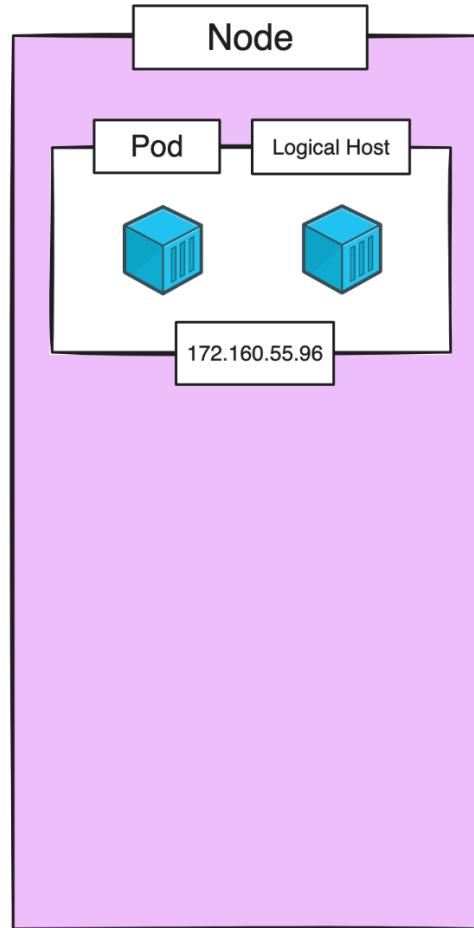
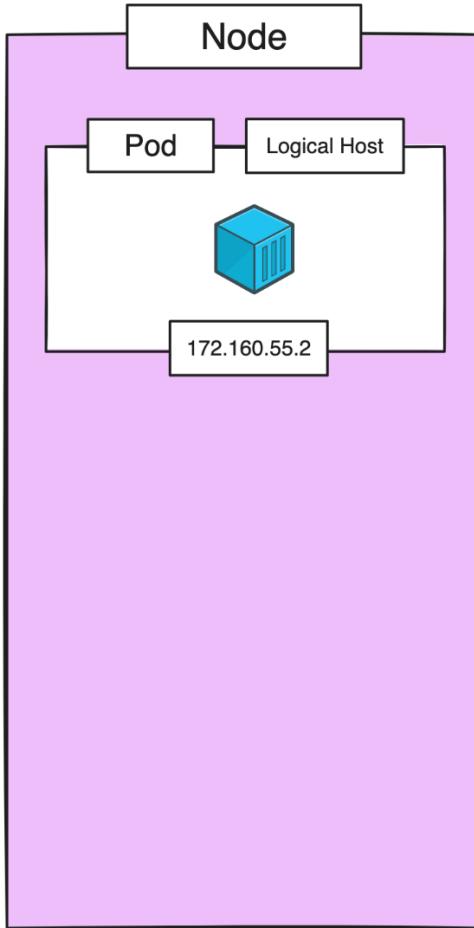
# Node & Pod

---



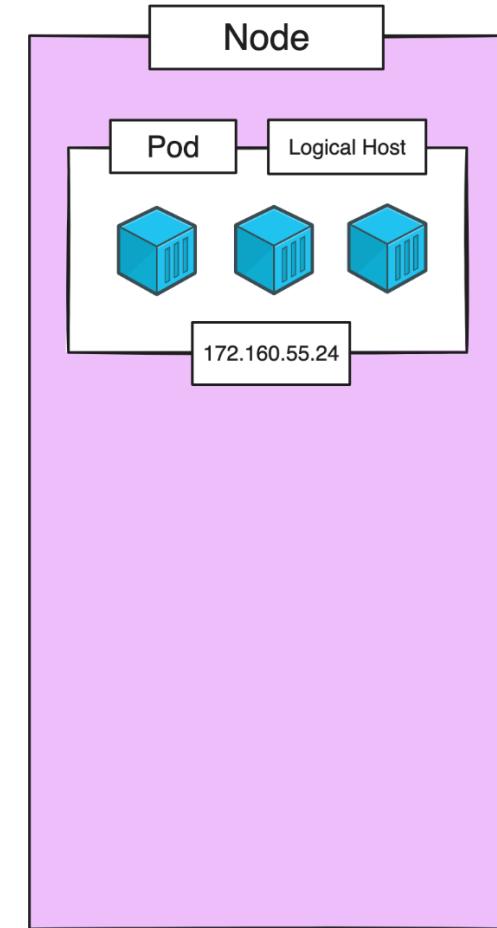
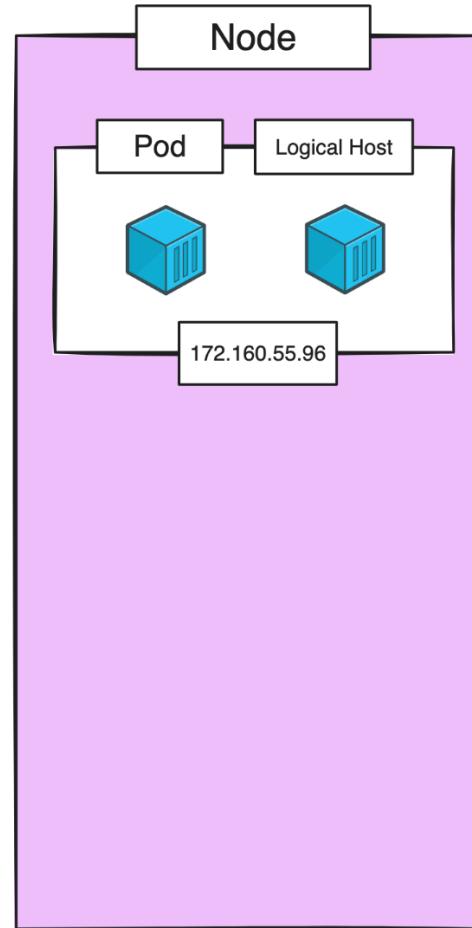
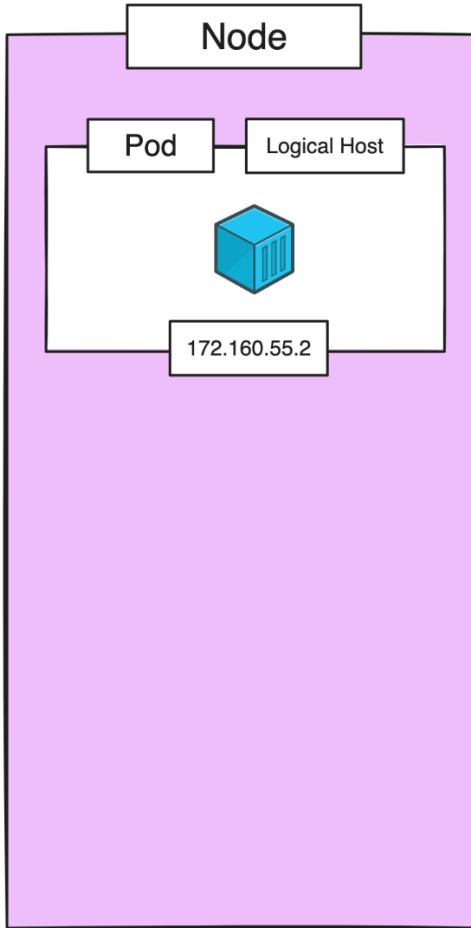
# Node & Pod

---



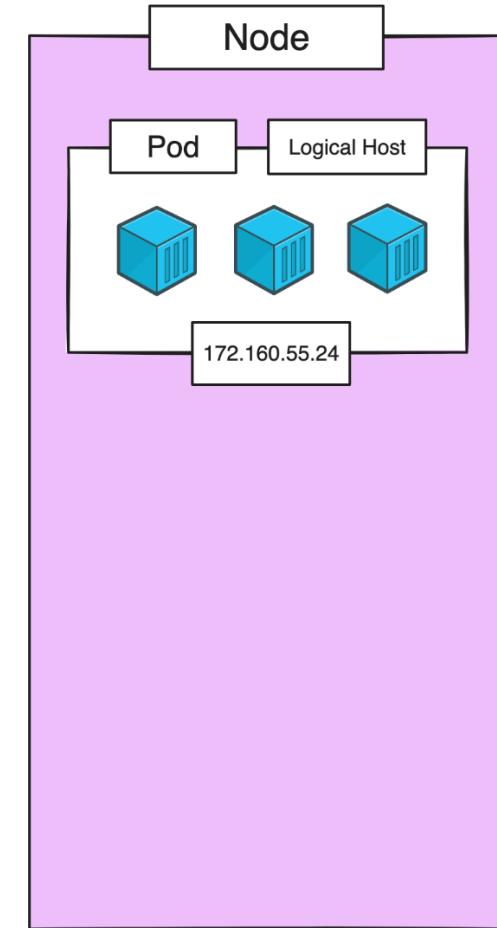
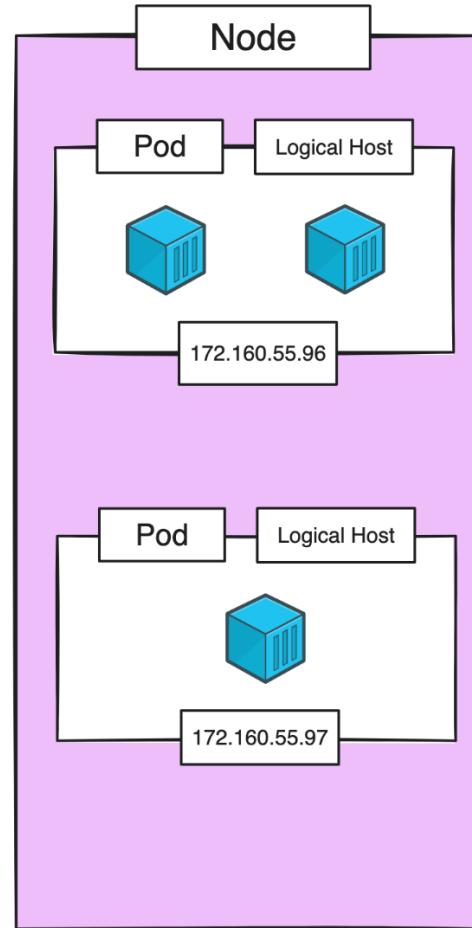
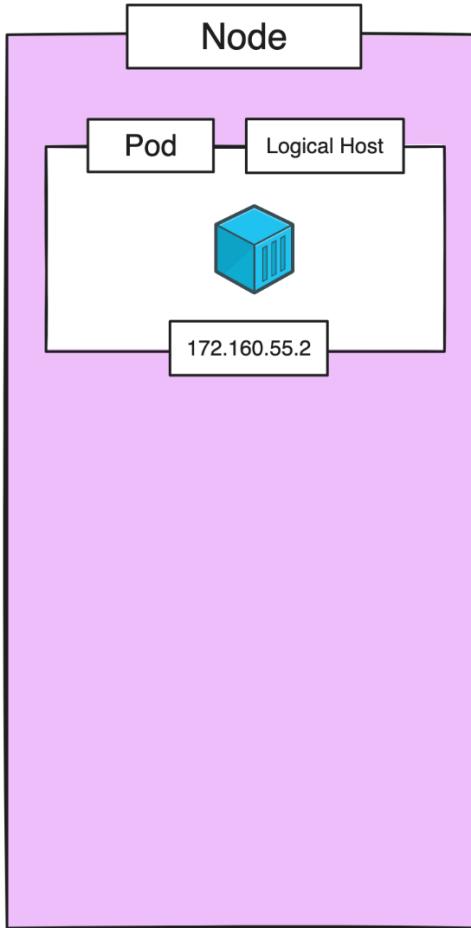
# Node & Pod

---

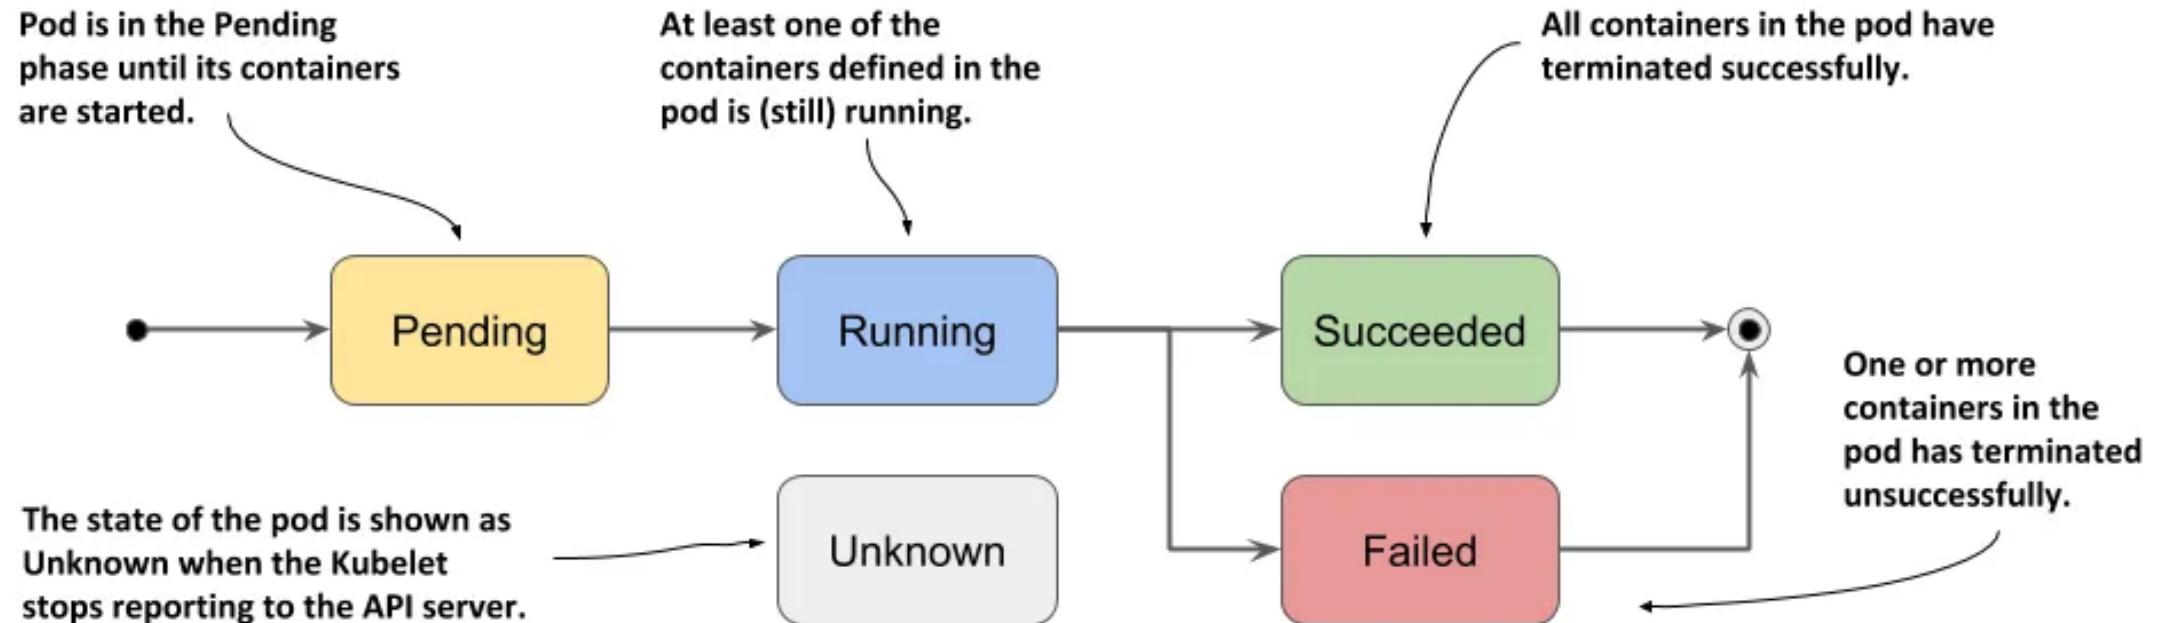


# Node & Pod

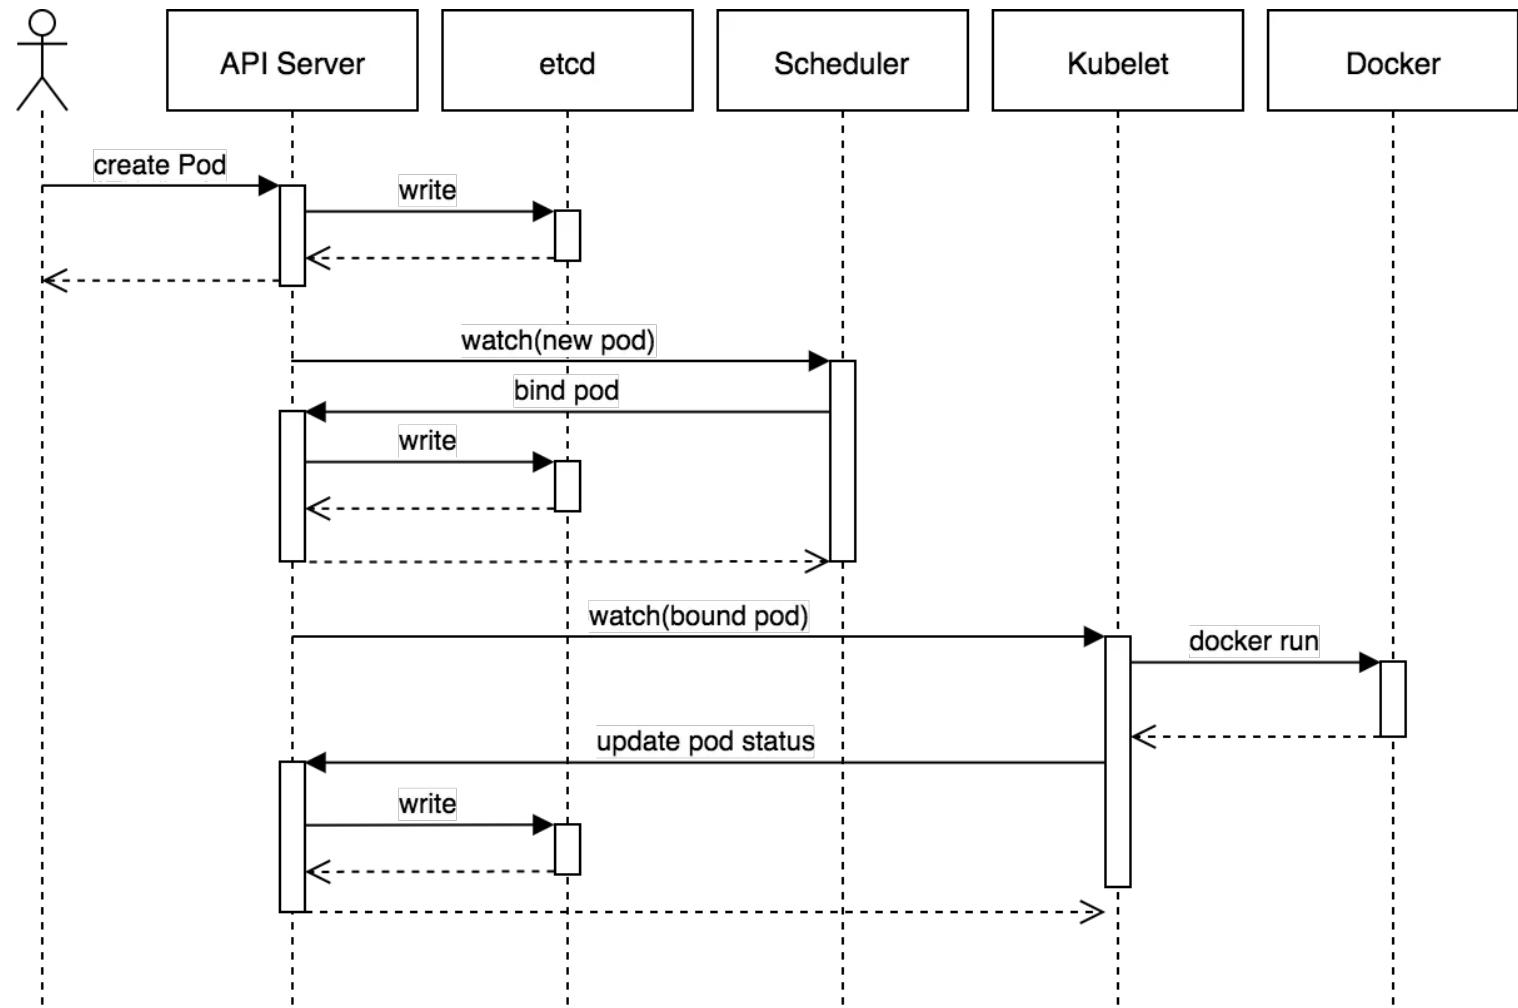
---



# State of Pod



# Creation of Pod





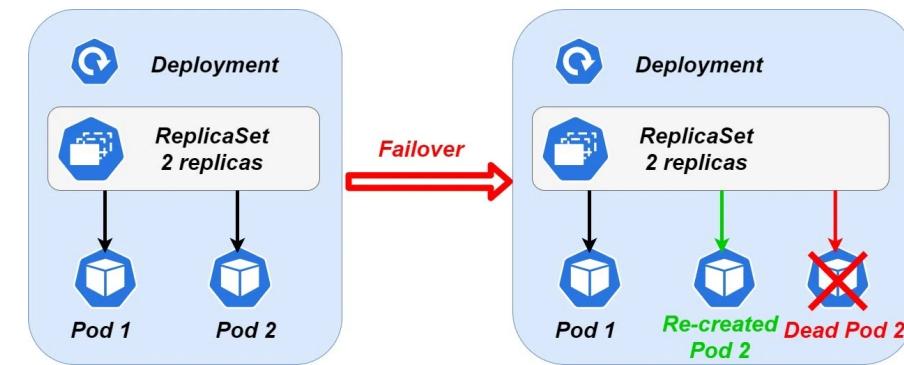
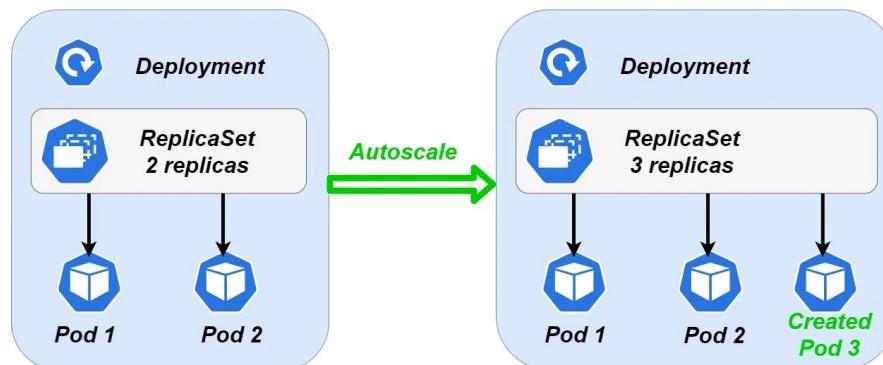
**Workshop** **pod**

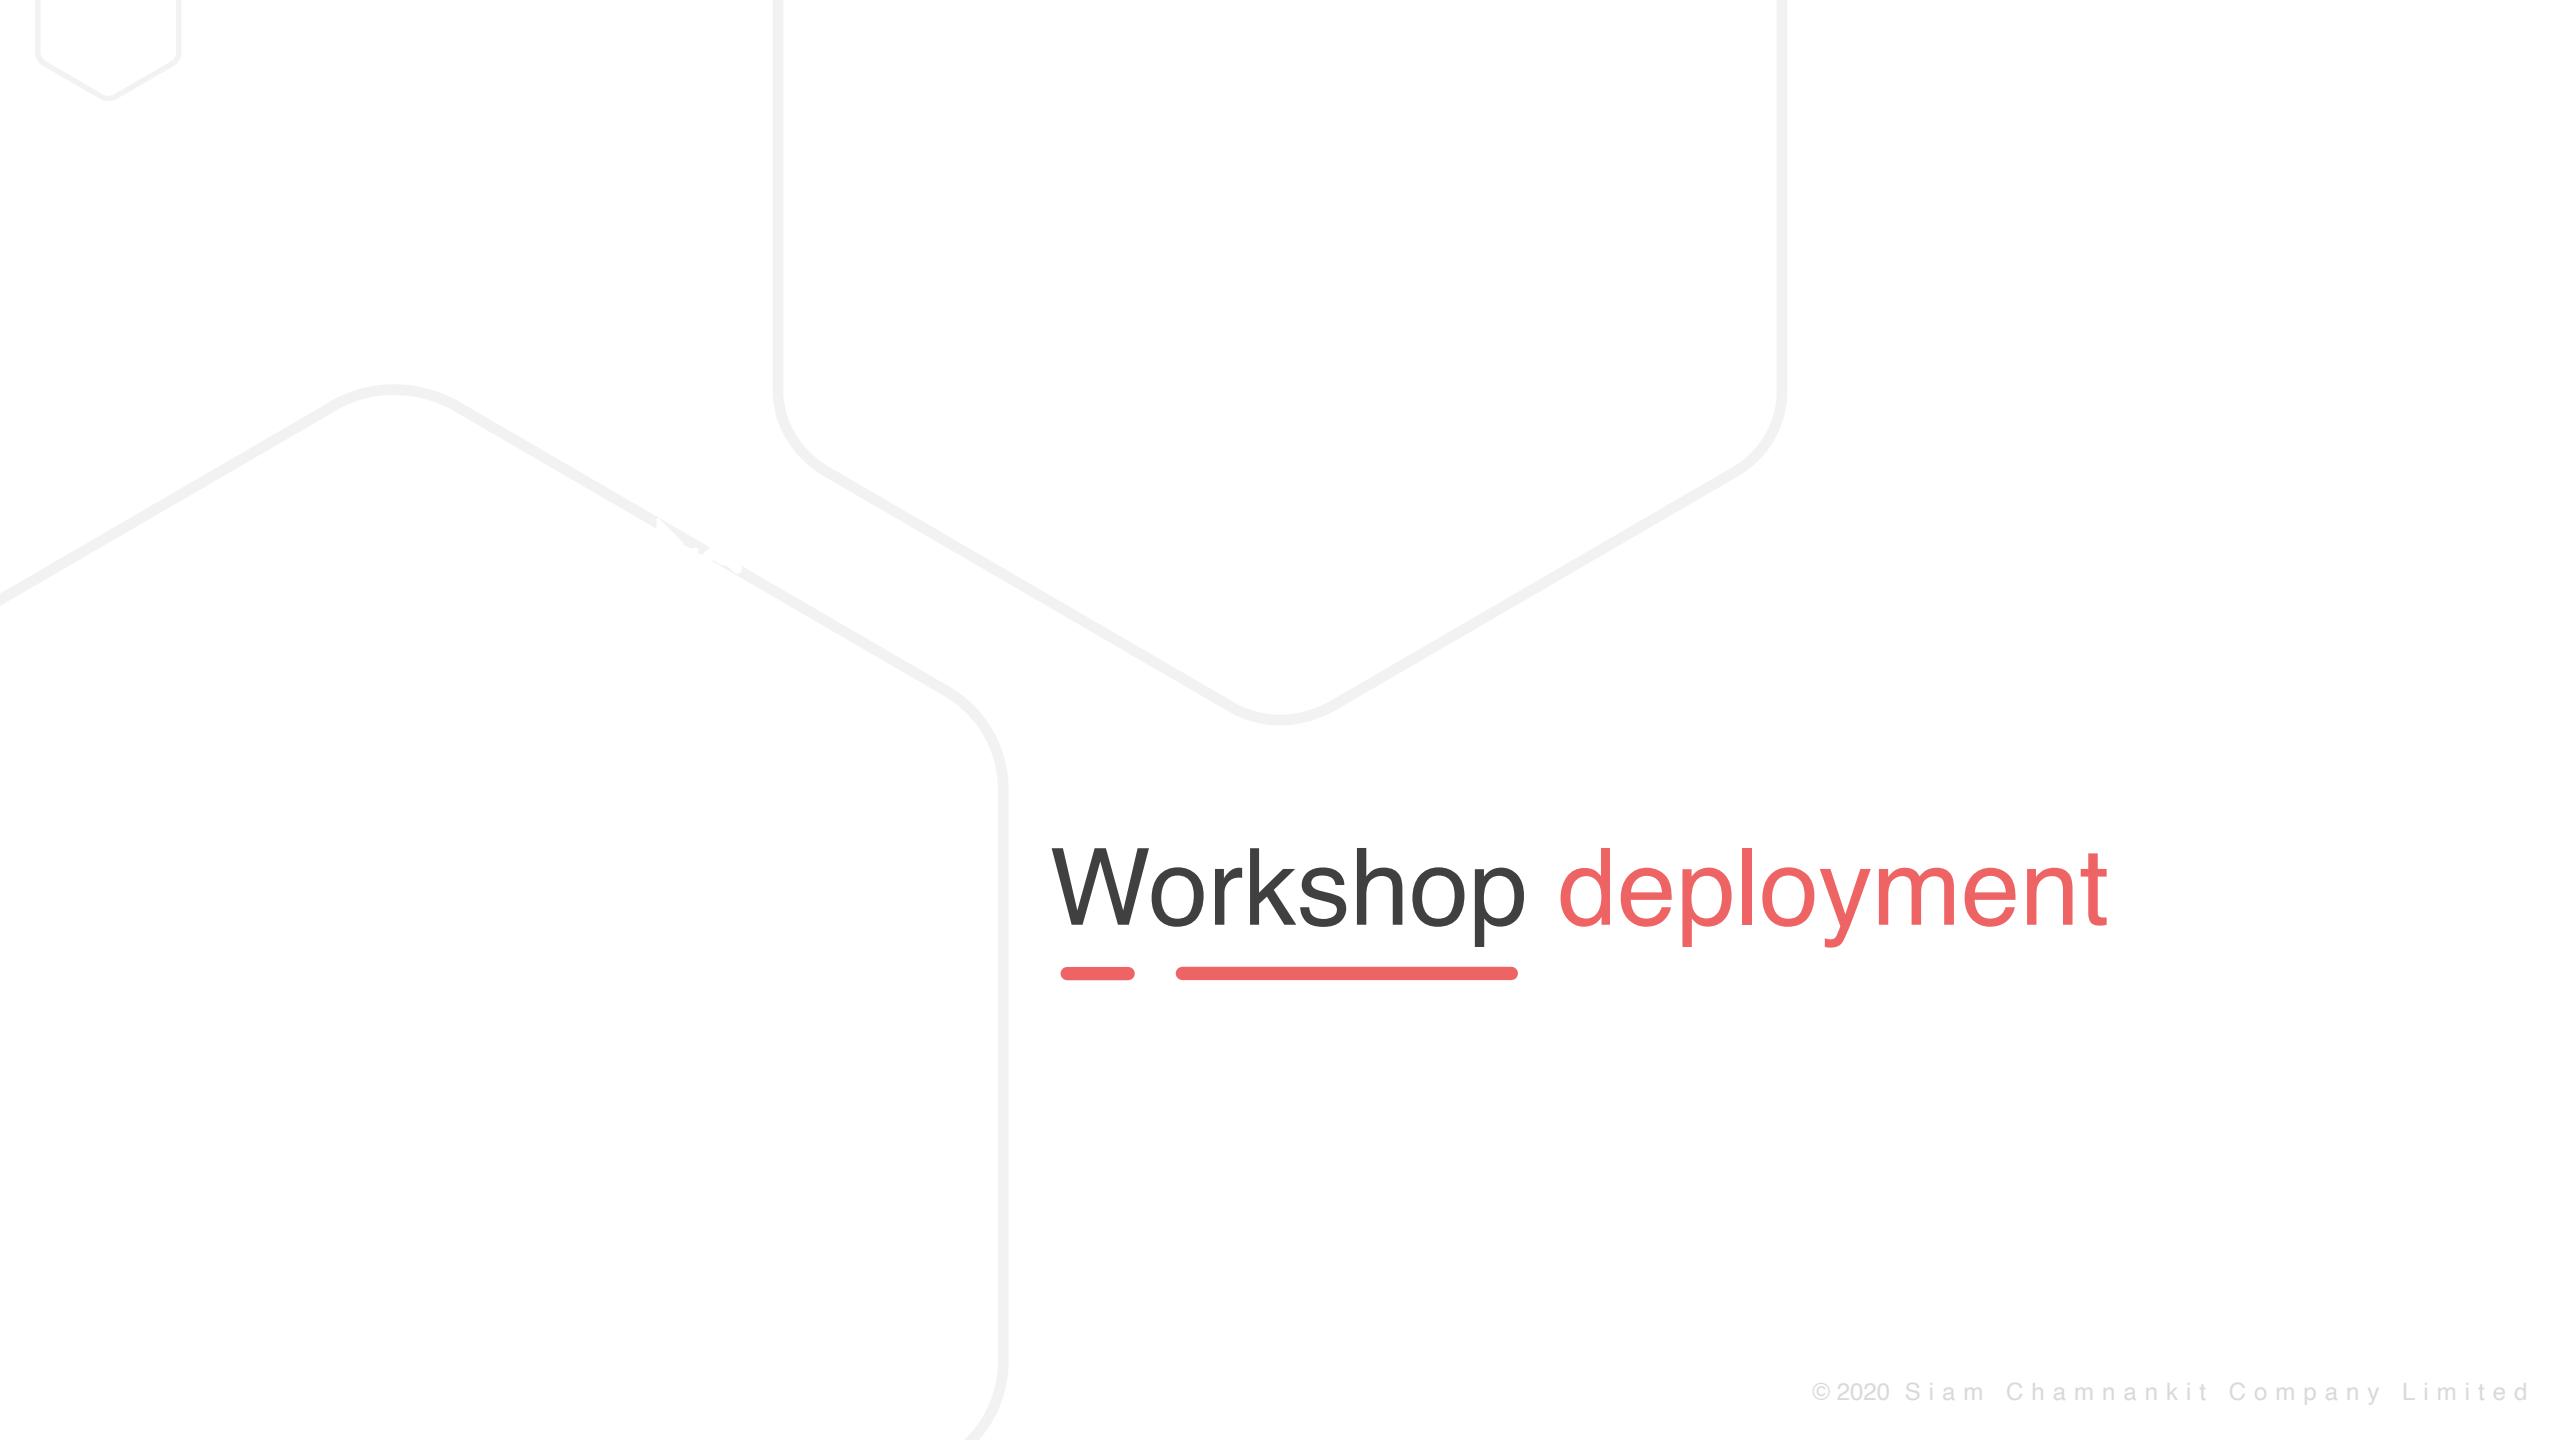
---

# Deployment

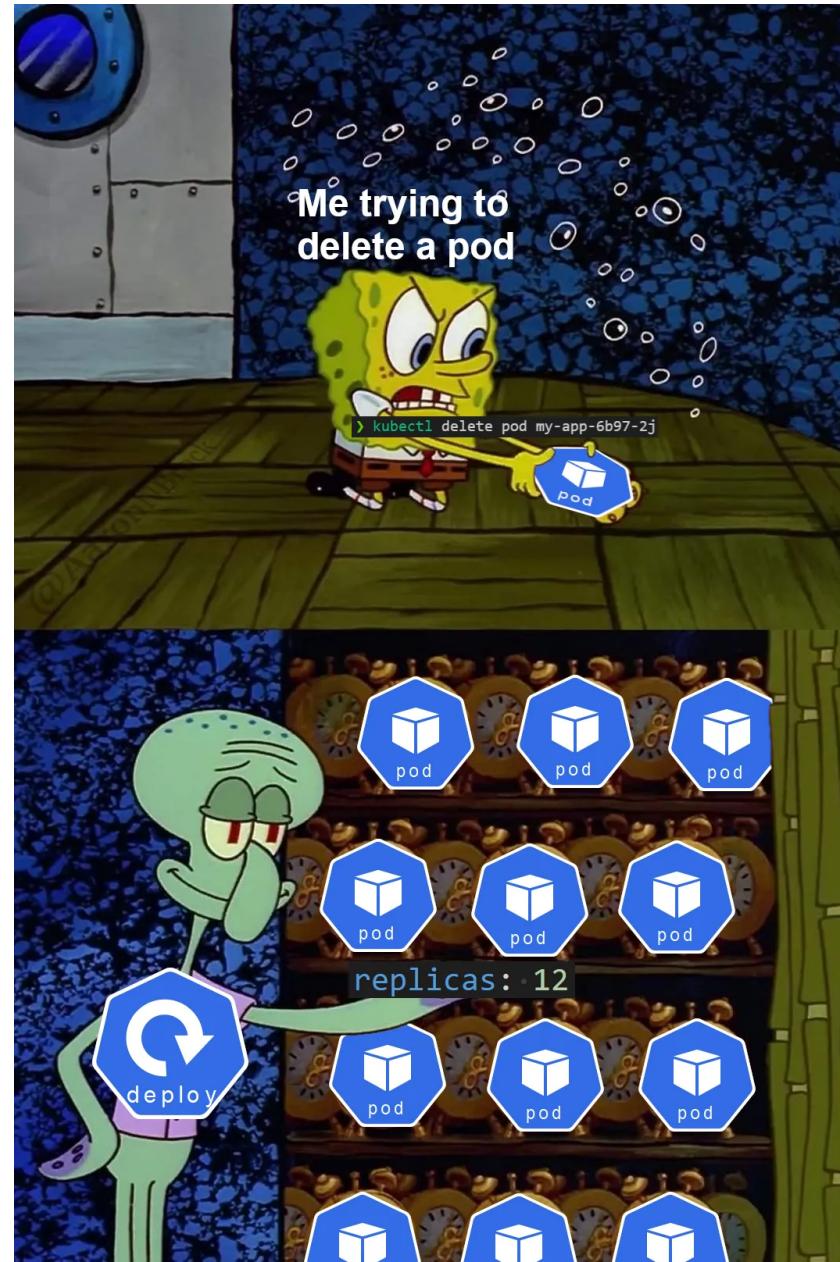
# Deployment

- A Deployment provides declarative updates for Pods and ReplicaSets
- You describe a **desired state** in a Deployment, and the Deployment Controller **changes the actual state to the desired state** at a controlled rate
- Manage for rollout and rollback
- Deployment => ReplicaSet => pod



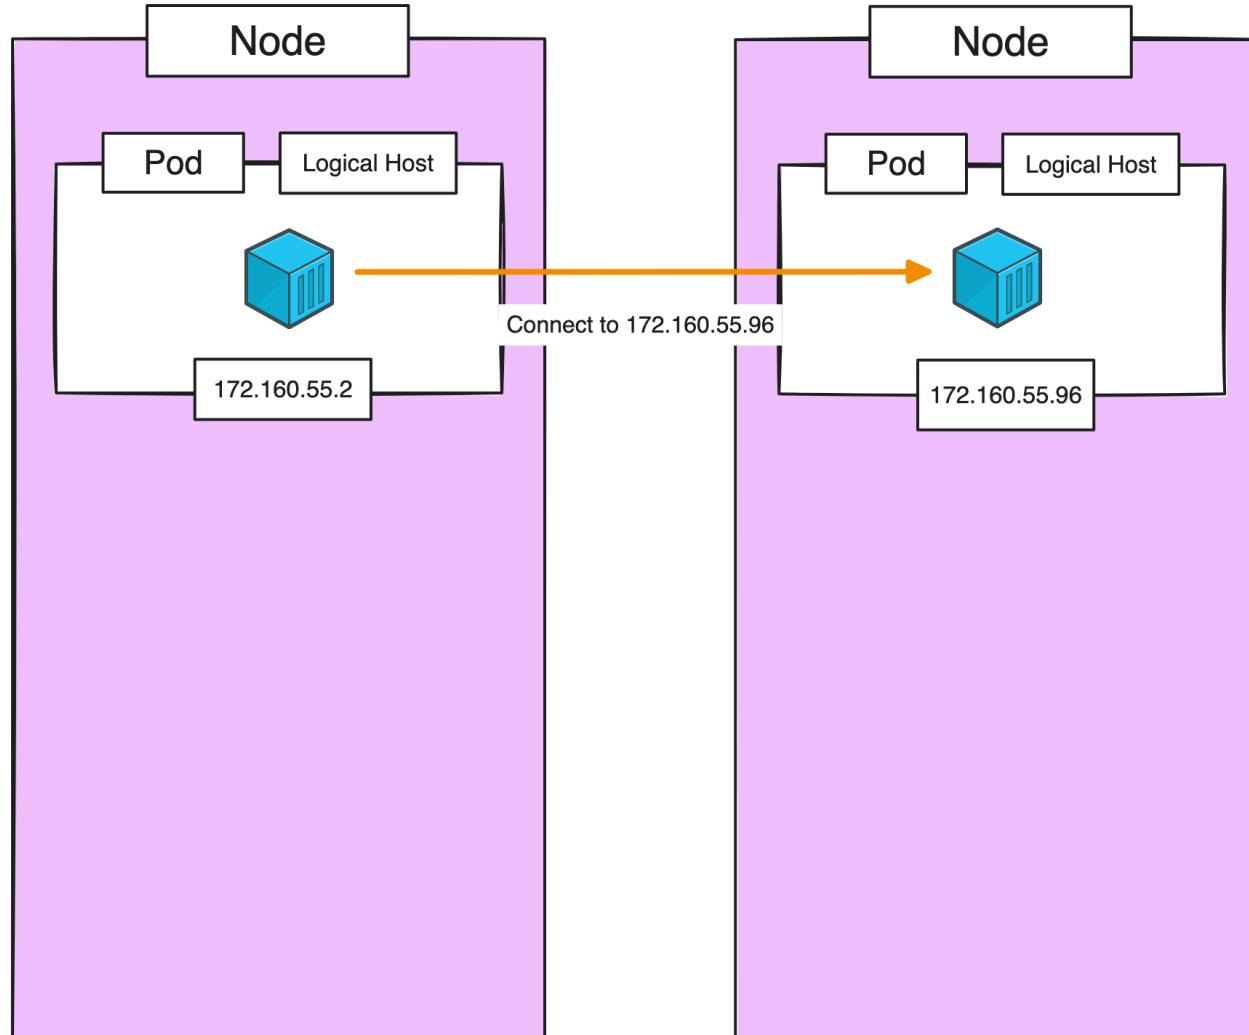


# Workshop deployment



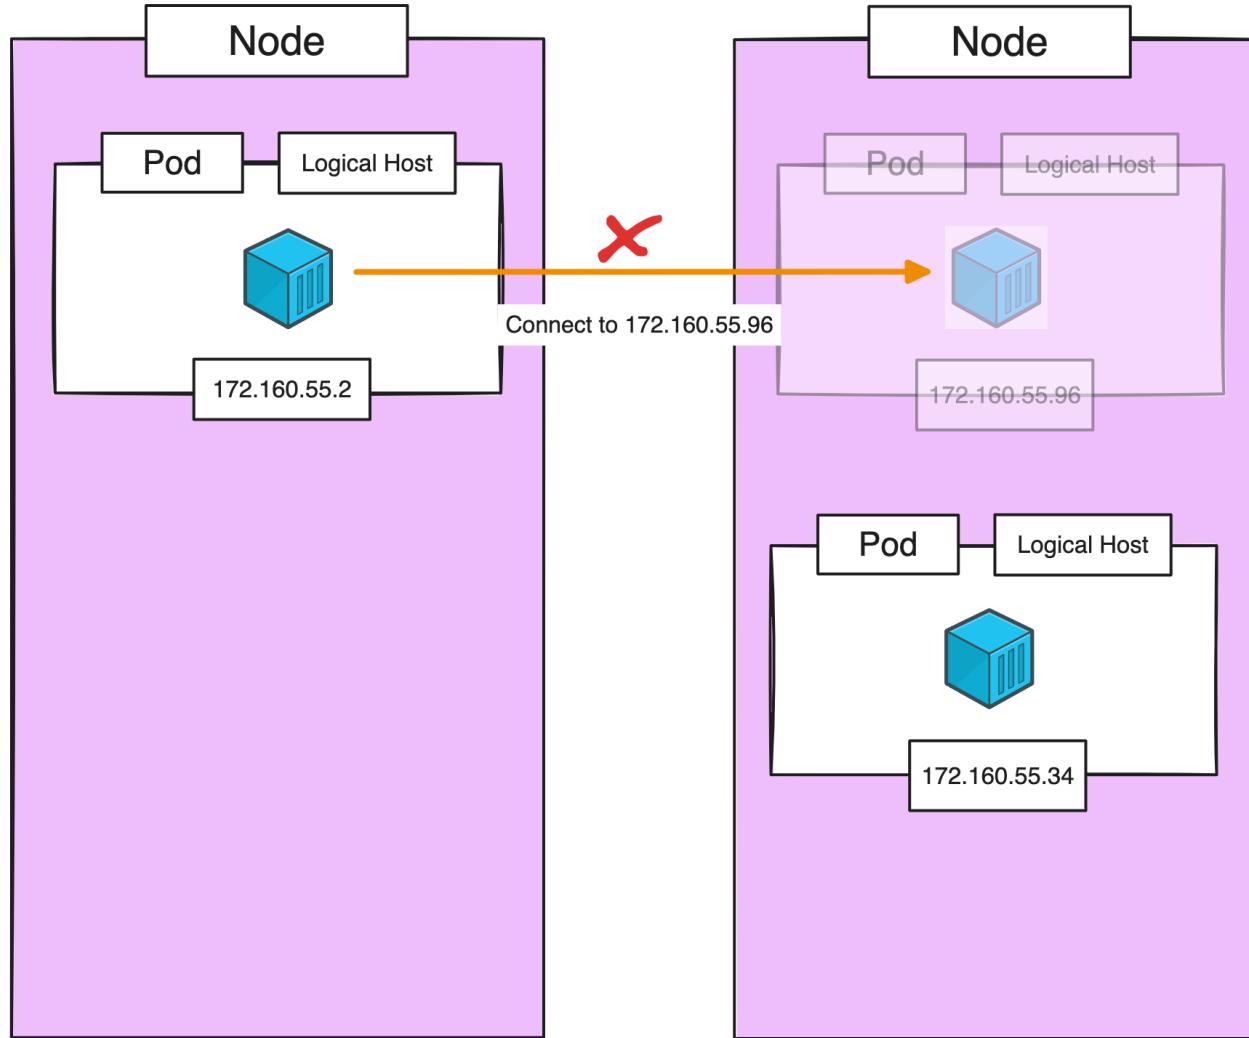
# Problems

---



# Problems

---



# Service

---



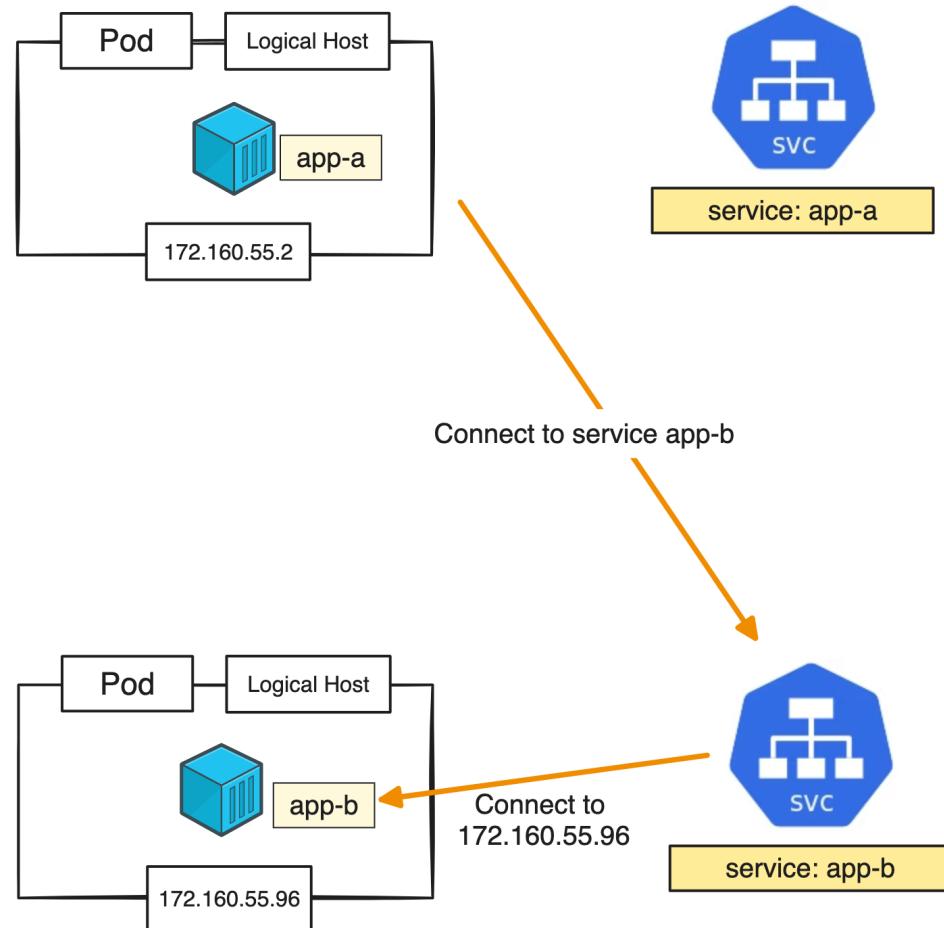
# Service

---

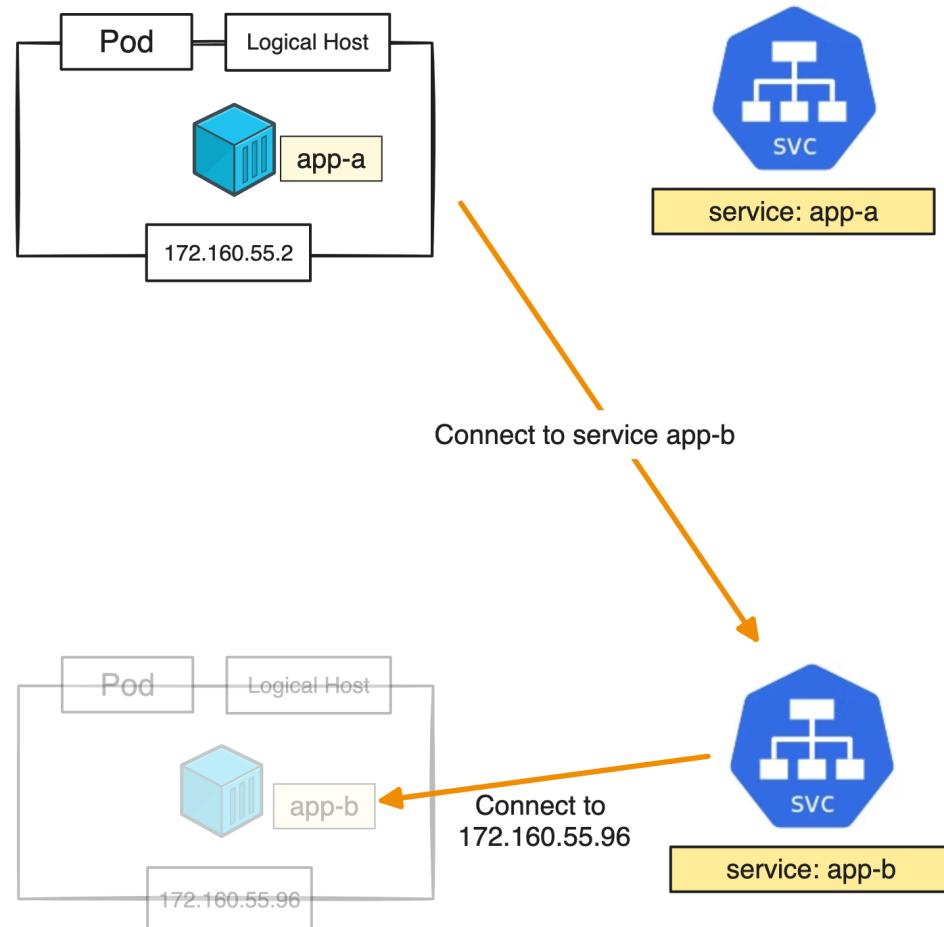
- Acts as a virtual endpoint for your application
- Providing a stable and discoverable way to access your pods without worrying about their individual IP addresses or network details
- Logical load balancer
- Communicating for outside and inside of cluster
- 4 types:
  - ClusterIP
  - NodePort
  - LoadBalancer
  - ExternalName



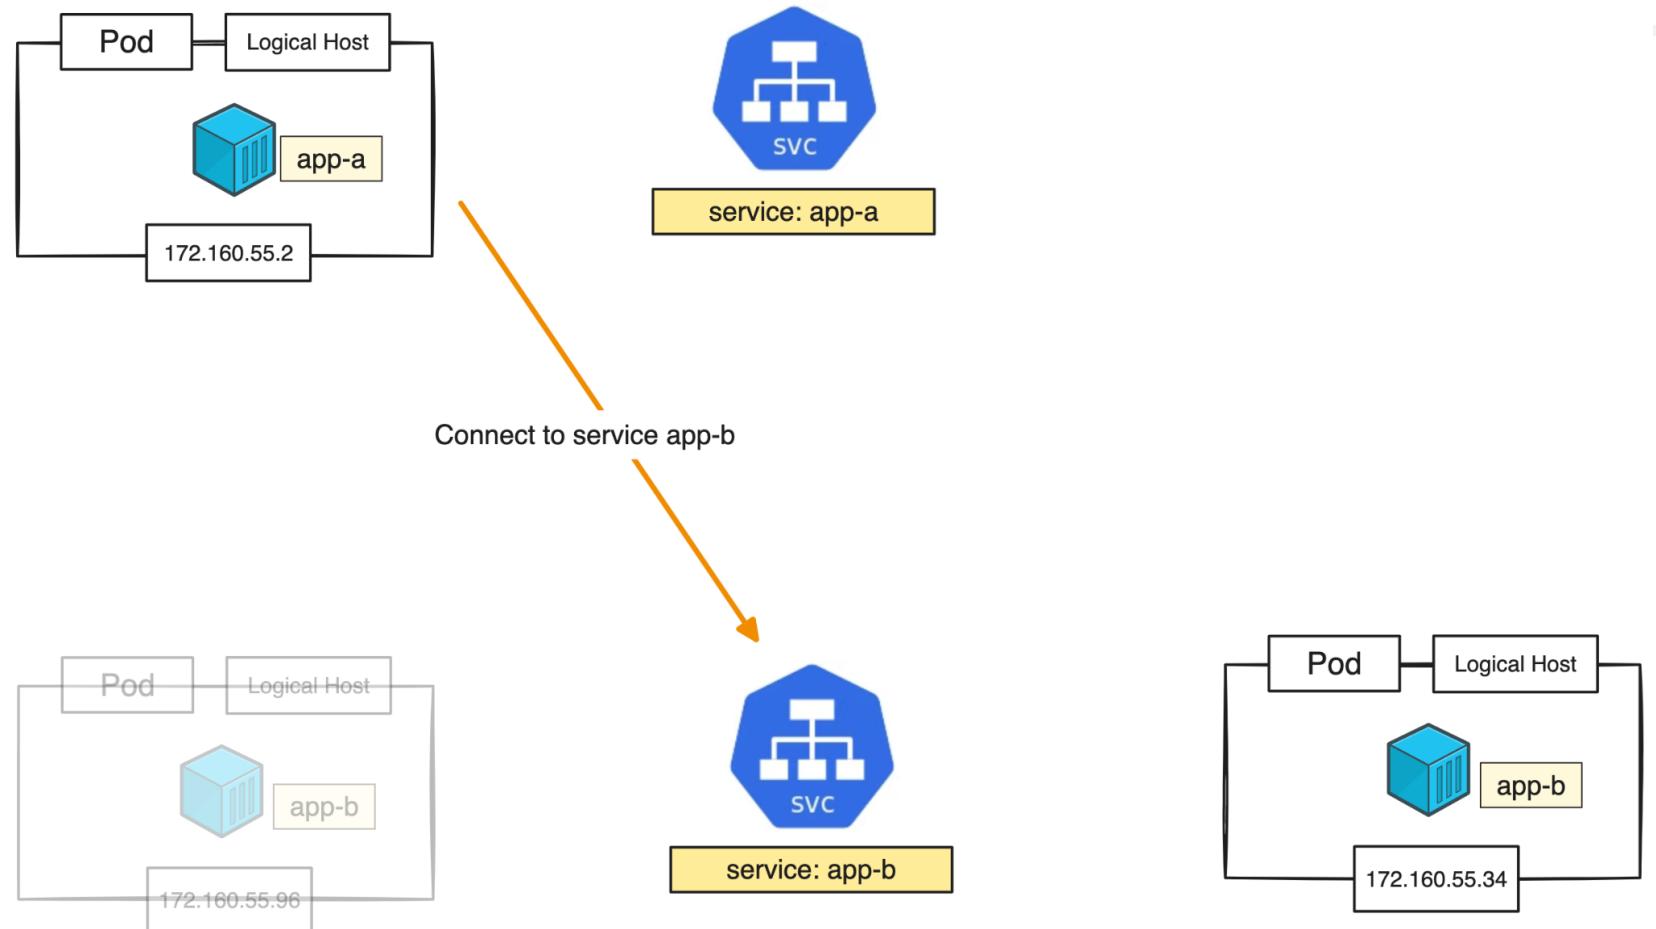
# Service: Stable endpoint



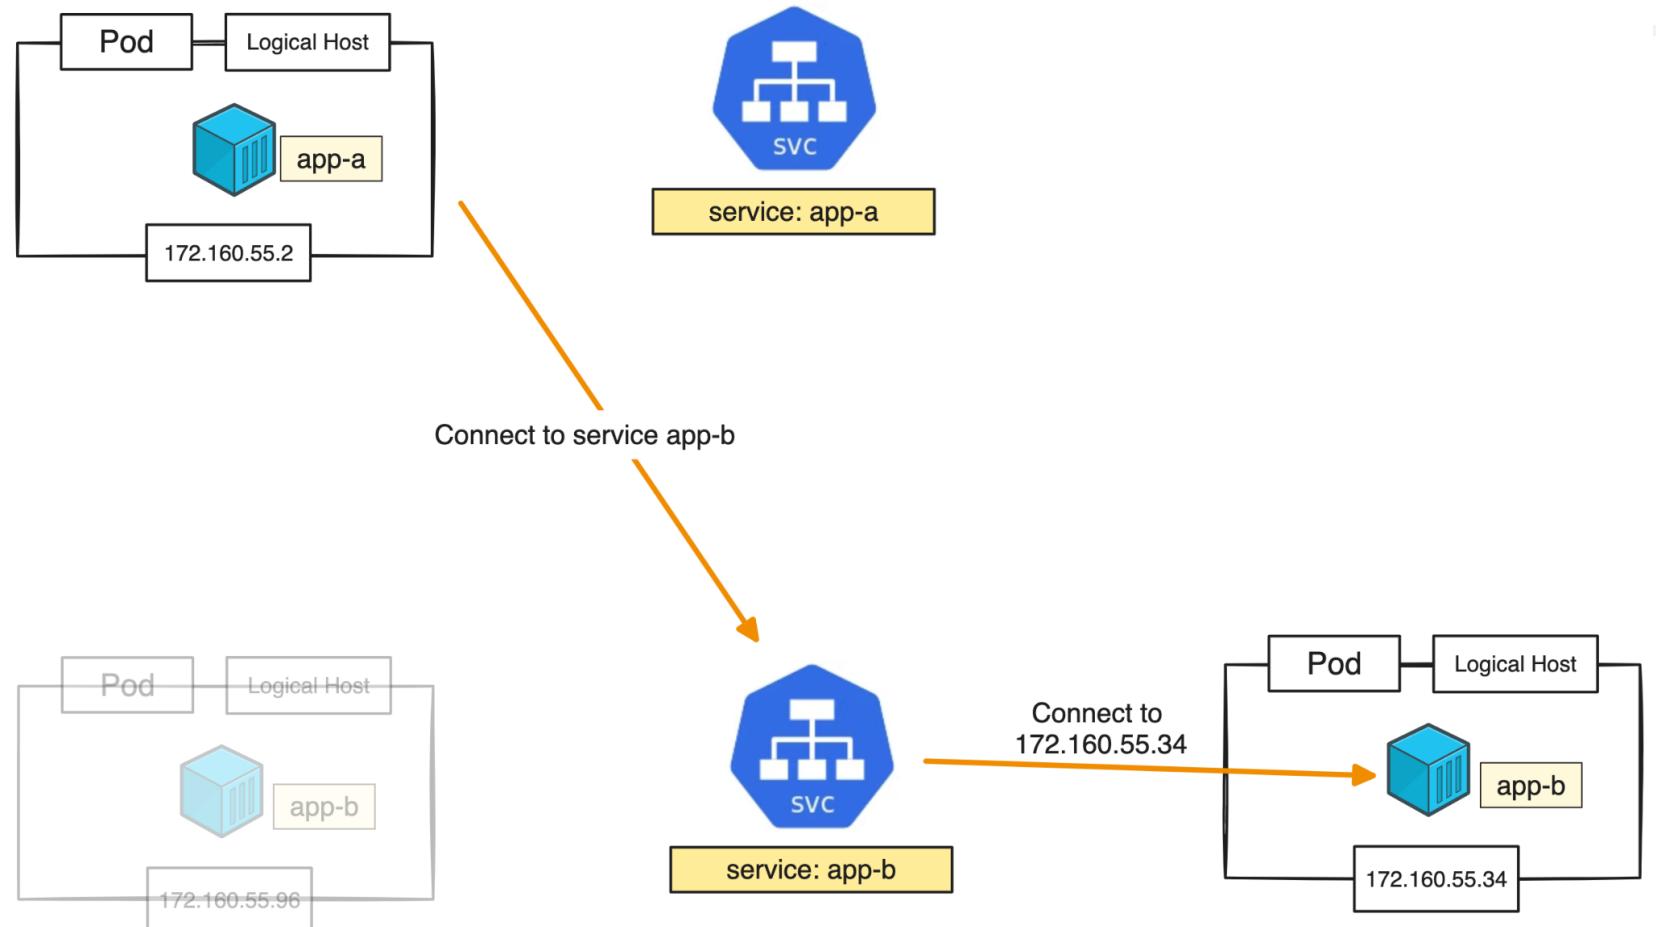
# Service: Stable endpoint



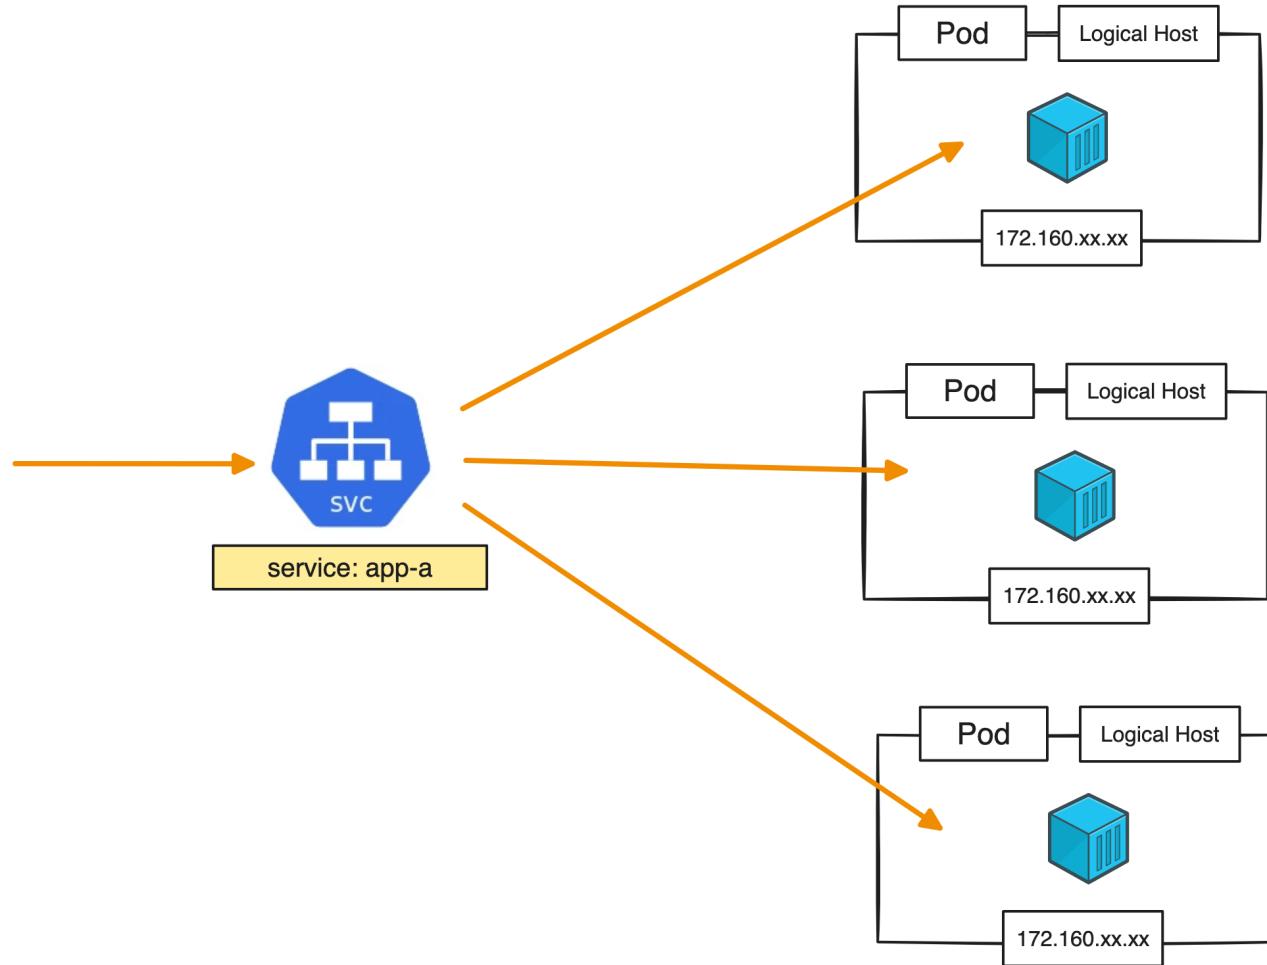
# Service: Stable endpoint



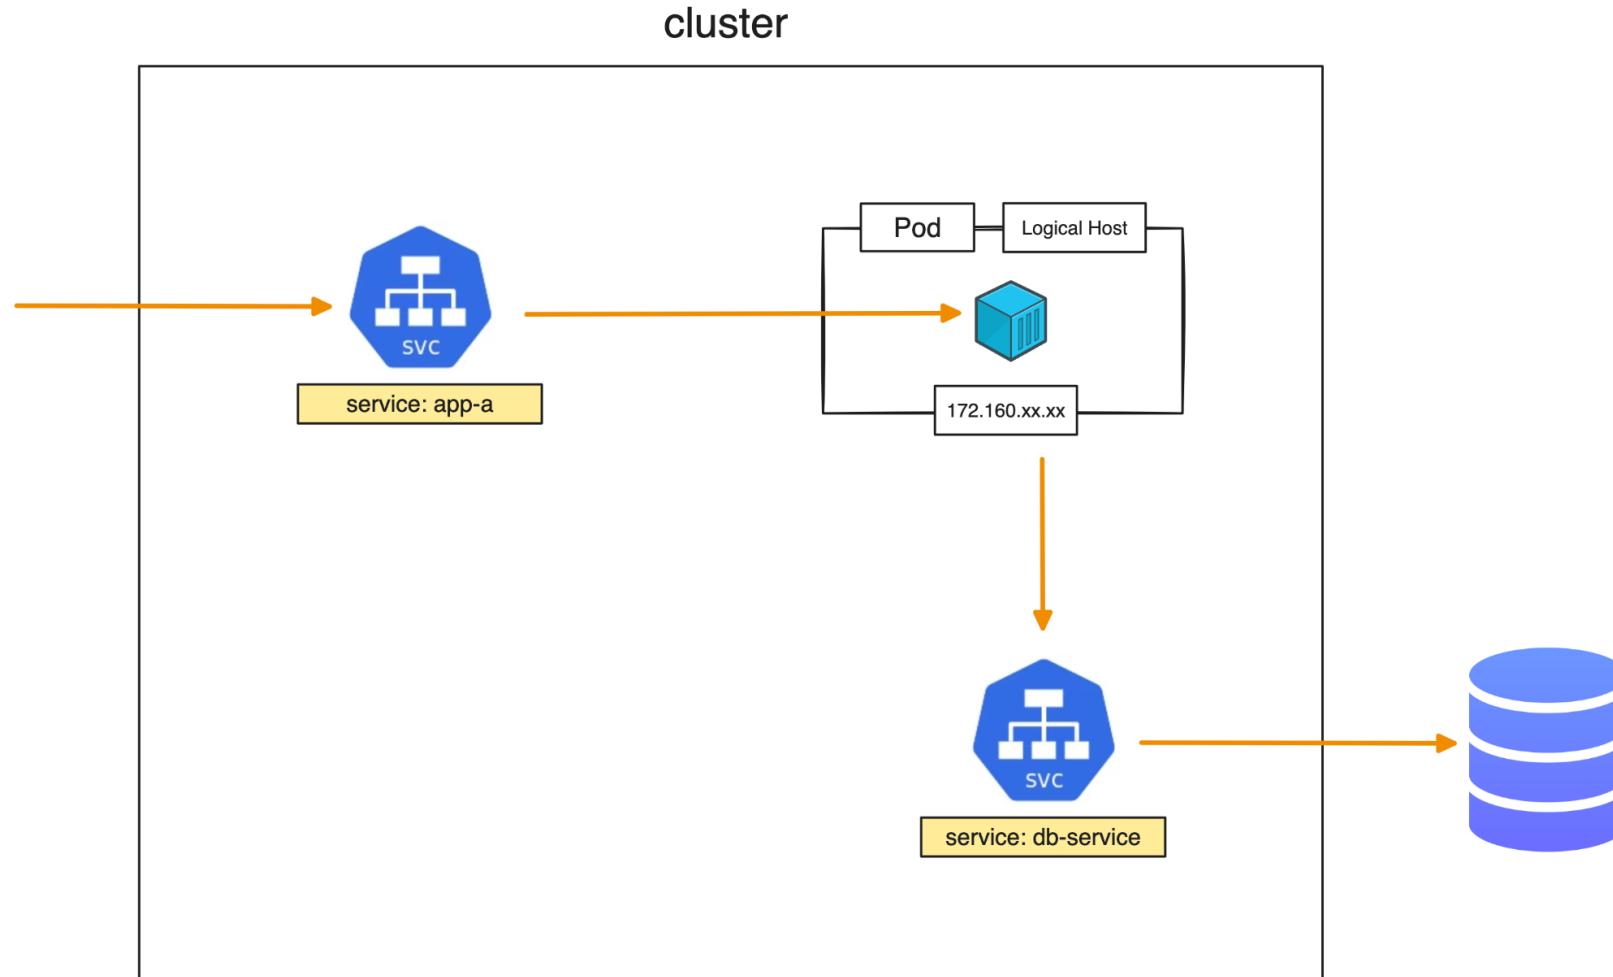
# Service: Stable endpoint



# Service: Load balancer



# Service: Communicating inside & outside of cluster



# Service type: ClusterIP

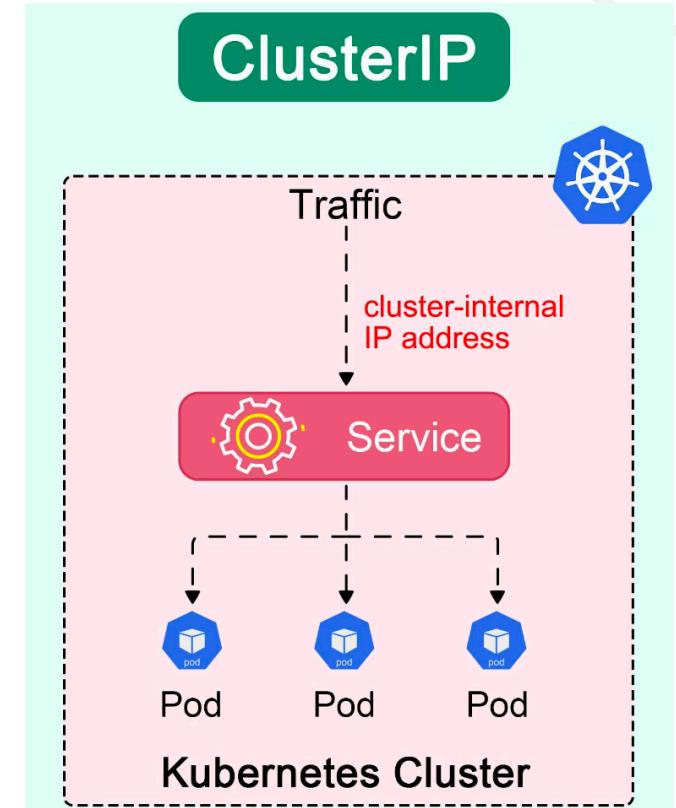
- ClusterIP is the default service type in Kubernetes
- It exposes the service on an internal IP address reachable **only within** the cluster

## Advantages

- Suitable for inter-pod communication and internal services

## Disadvantages

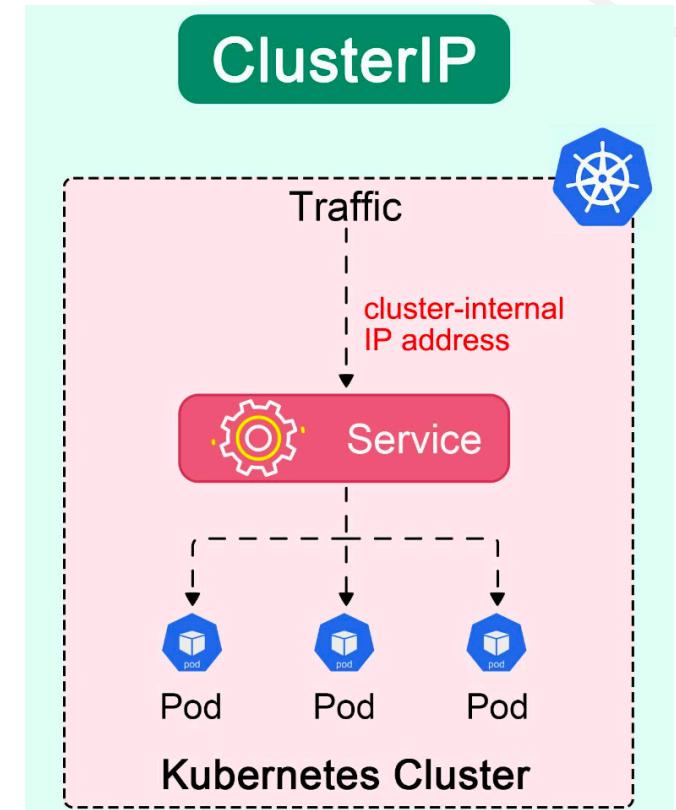
- Limited to internal cluster communication
- Not accessible externally



# Service type: ClusterIP

## Use Cases

- Inter-service communication within the cluster. For example, communication between the front-end and back-end components of your app





# Workshop ClusterIP Service

# Service type: NodePort

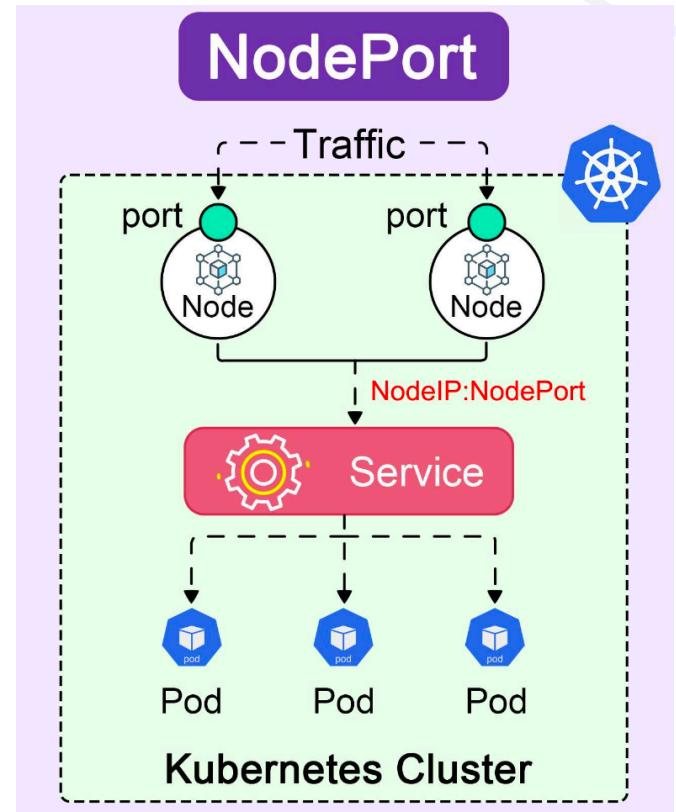
- It enables **external access** to the service using the cluster nodes' **IP address** and the assigned **NodePort**
- When a NodePort service is created, Kubernetes allocates a port from a predefined range (default: **30000-32767**)

## Advantages

- Provides external access to the service
- Suitable for development and testing environments

## Disadvantages

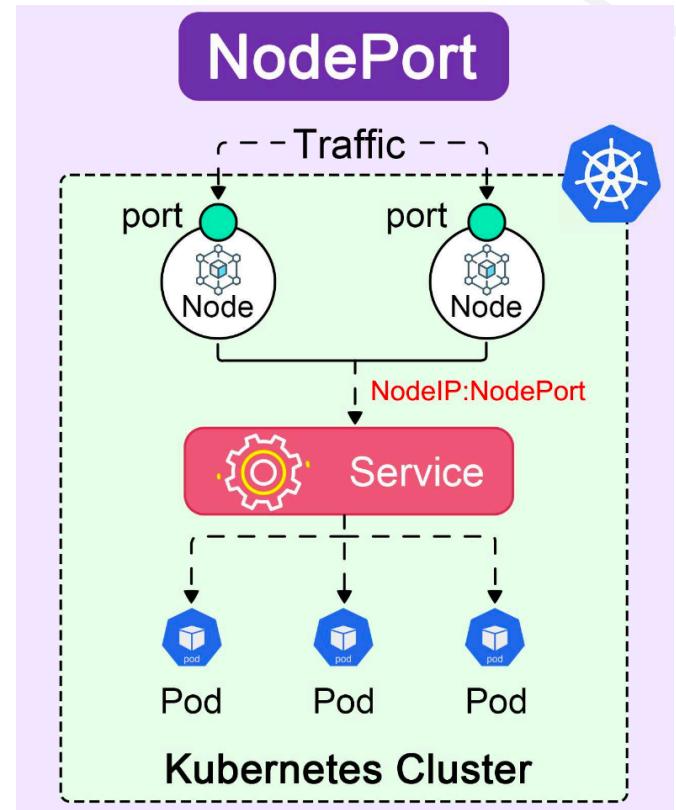
- Requires manual management of port allocations
- IP Node maybe have changed
- Should not use in production

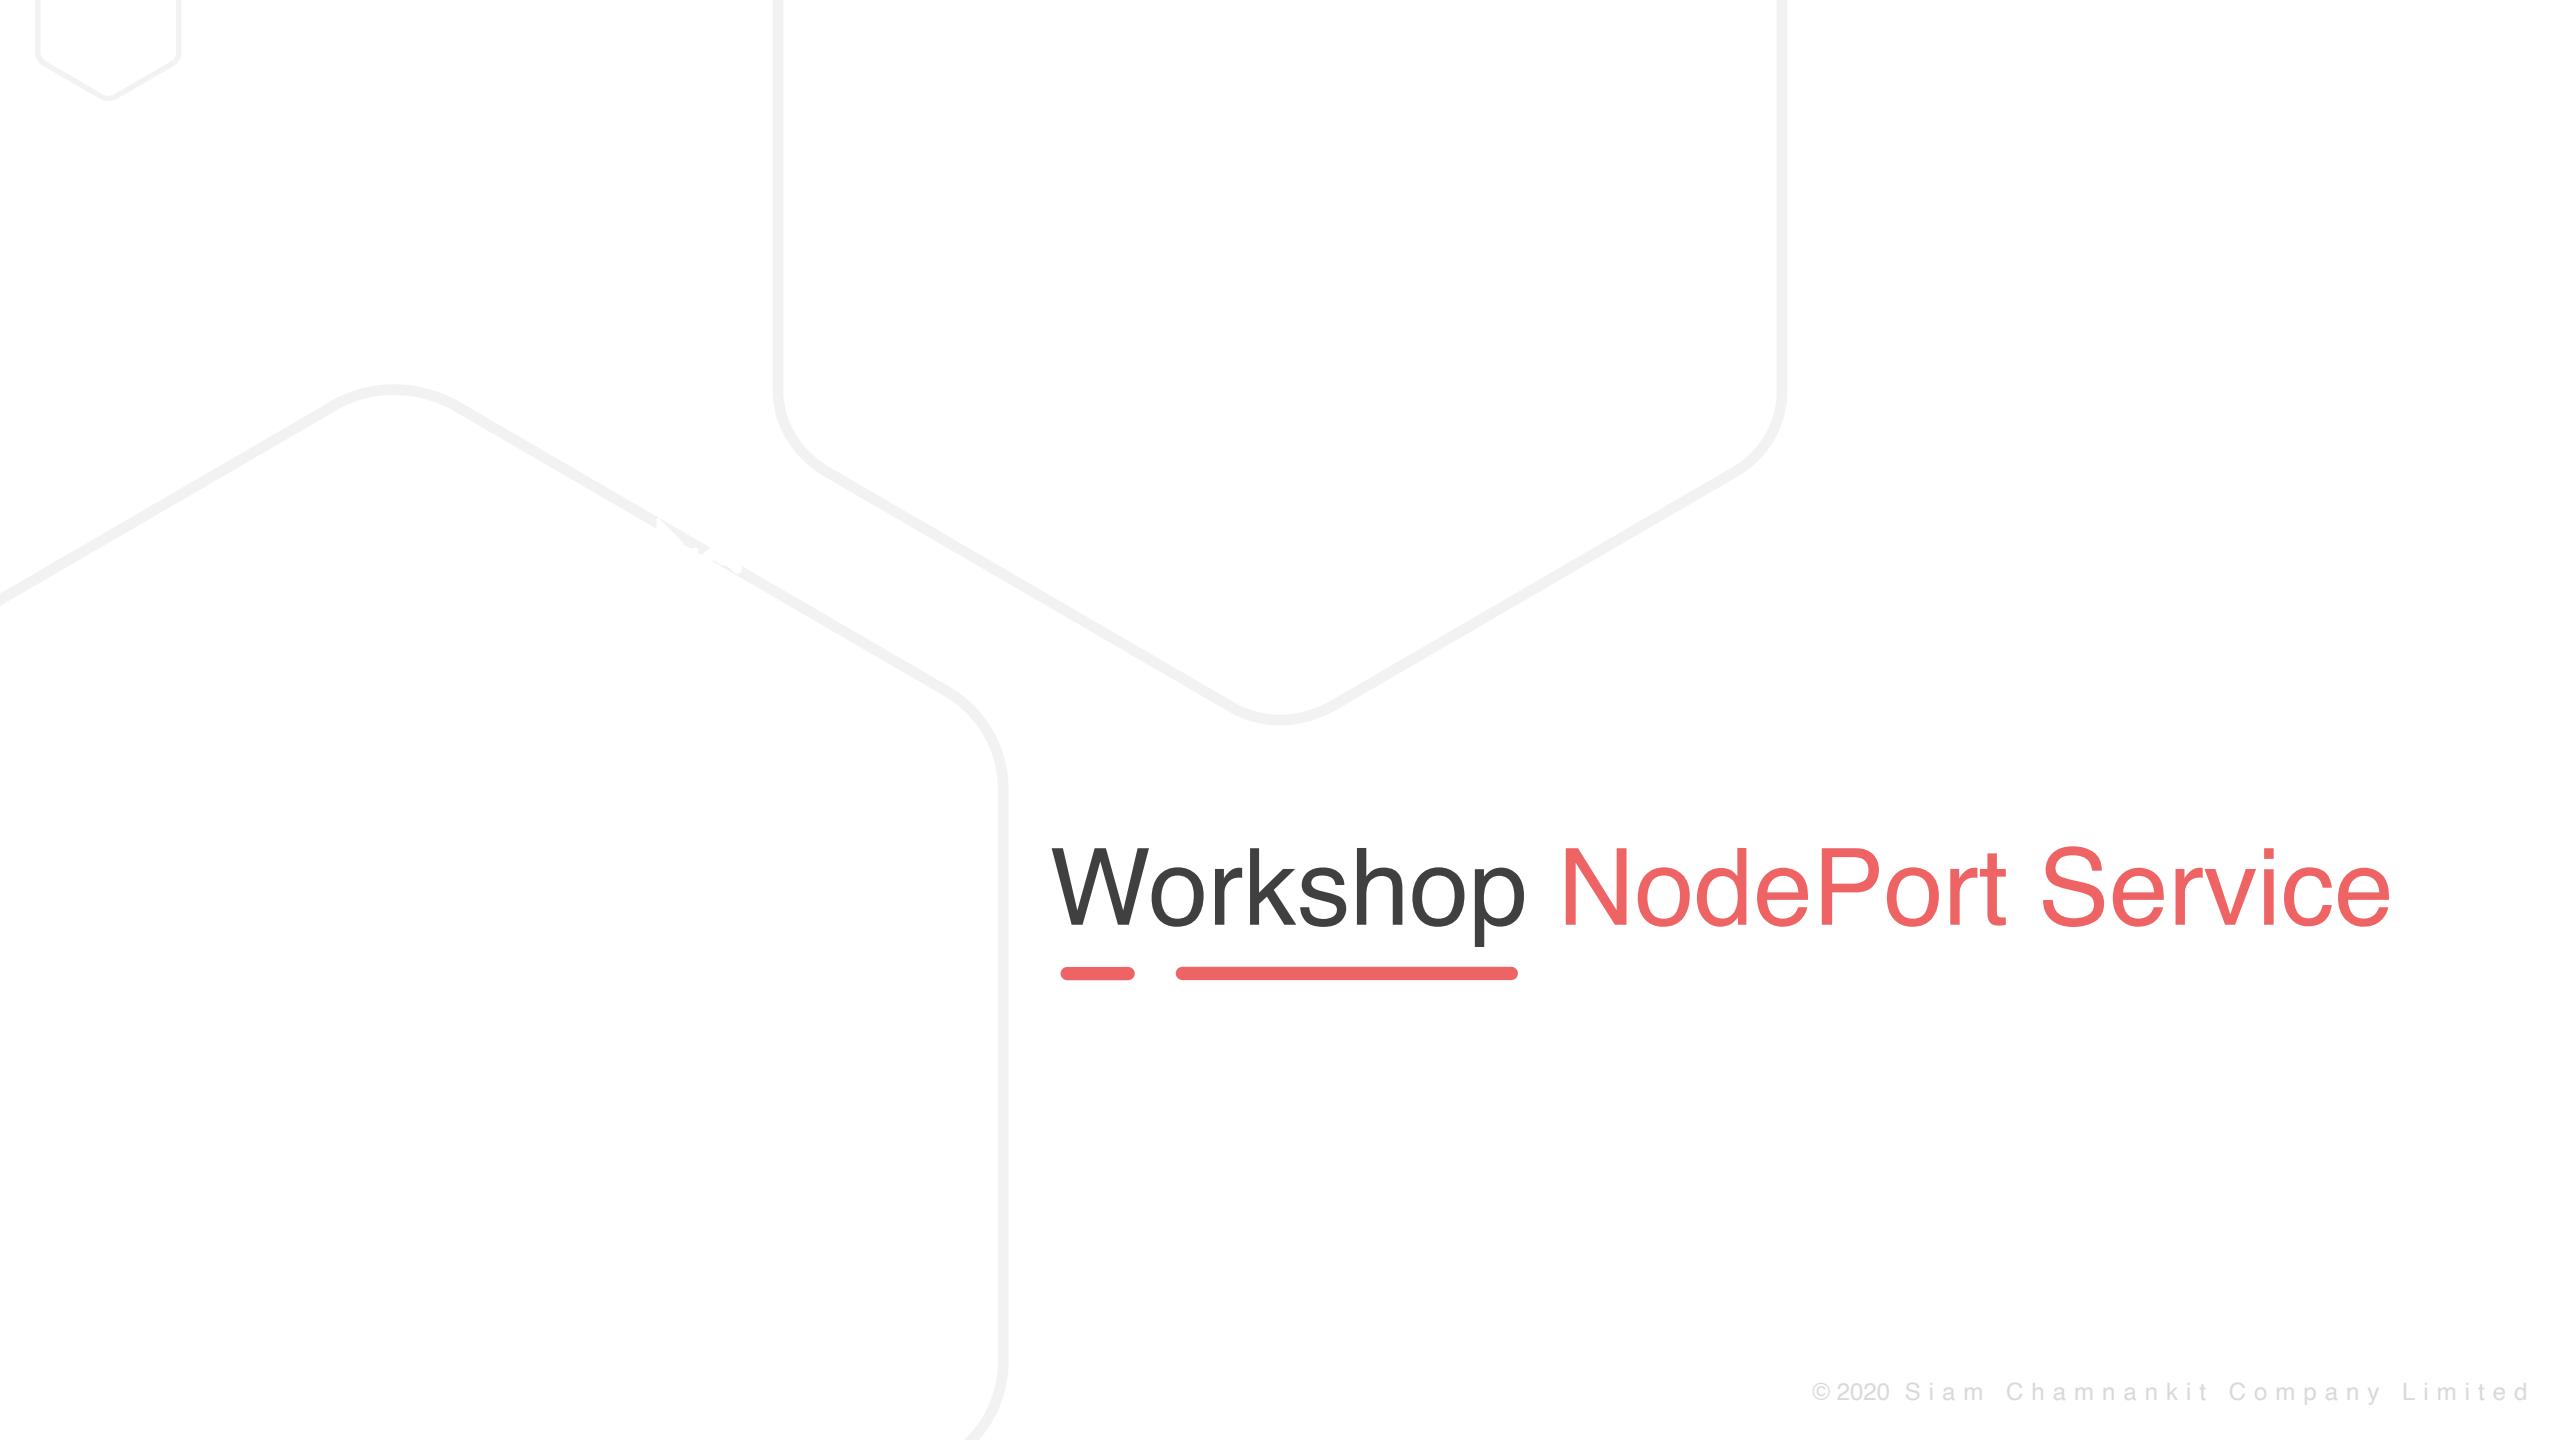


# Service type: NodePort

## Use case

- When you want to enable external connectivity to your service
- Best for testing access for a moment time





# Workshop NodePort Service

# Service type: LoadBalancer

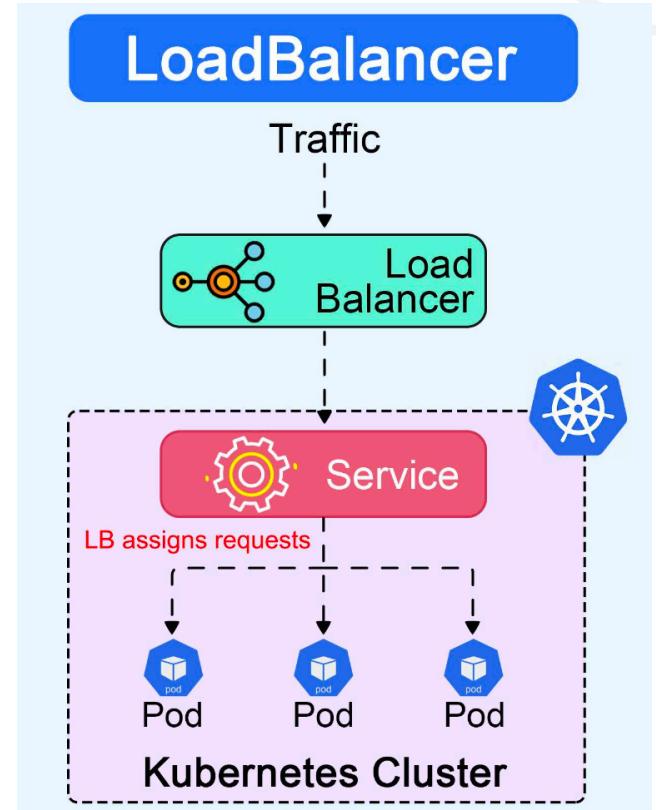
- LoadBalancer is a service type that **provisions** an **external load balancer** to distribute traffic across the pods
- Typically provided by a cloud provider's infrastructure
- When a LoadBalancer service is created, Kubernetes interacts with the cloud provider's API to an external load balancer

## Advantages

- Provides external access to the service
- Automatically provisions an external load balancer for the service
- Supports scaling and high availability

## Disadvantages

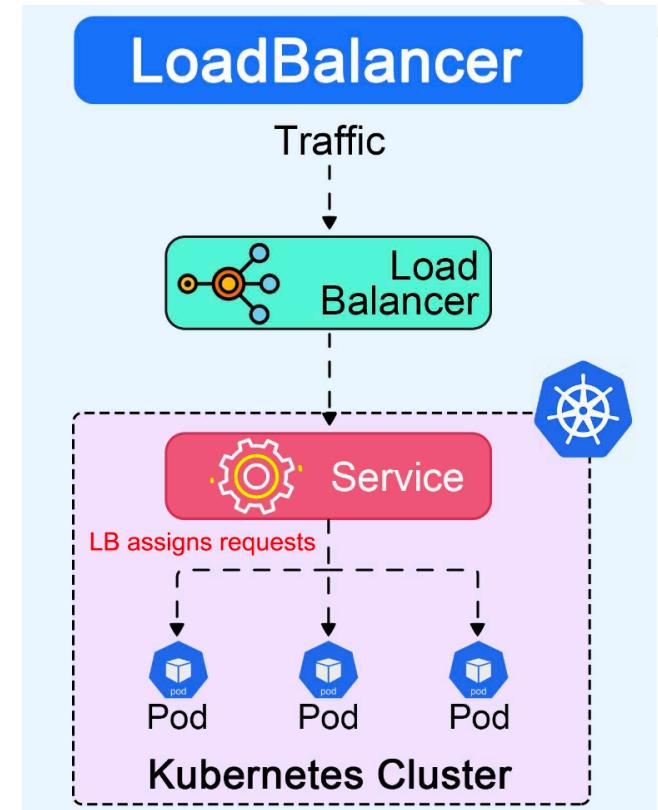
- Requires a cloud provider
- Additional costs for using the load balancer service on cloud provider



# Service type: LoadBalancer

## Use case

- When you want to enable external connectivity to your service
- When you are using a cloud provider to host your Kubernetes cluster





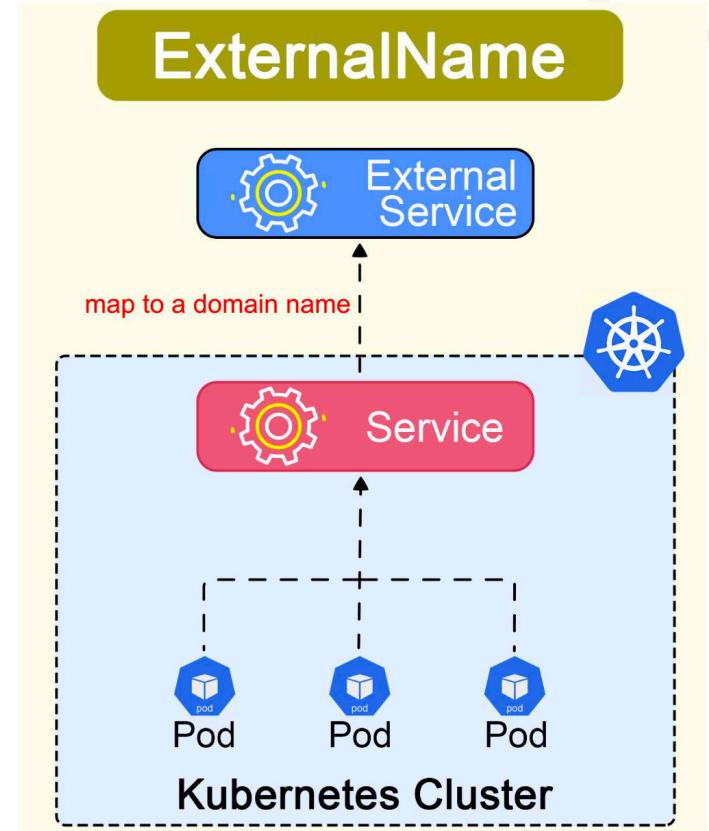
# Workshop LoadBalancer Service

# Service type: ExternalName

- An ExternalName Service is a special type of Kubernetes service that allows you to expose an existing external service (like a database or a third-party API) to your cluster without deploying any Pods

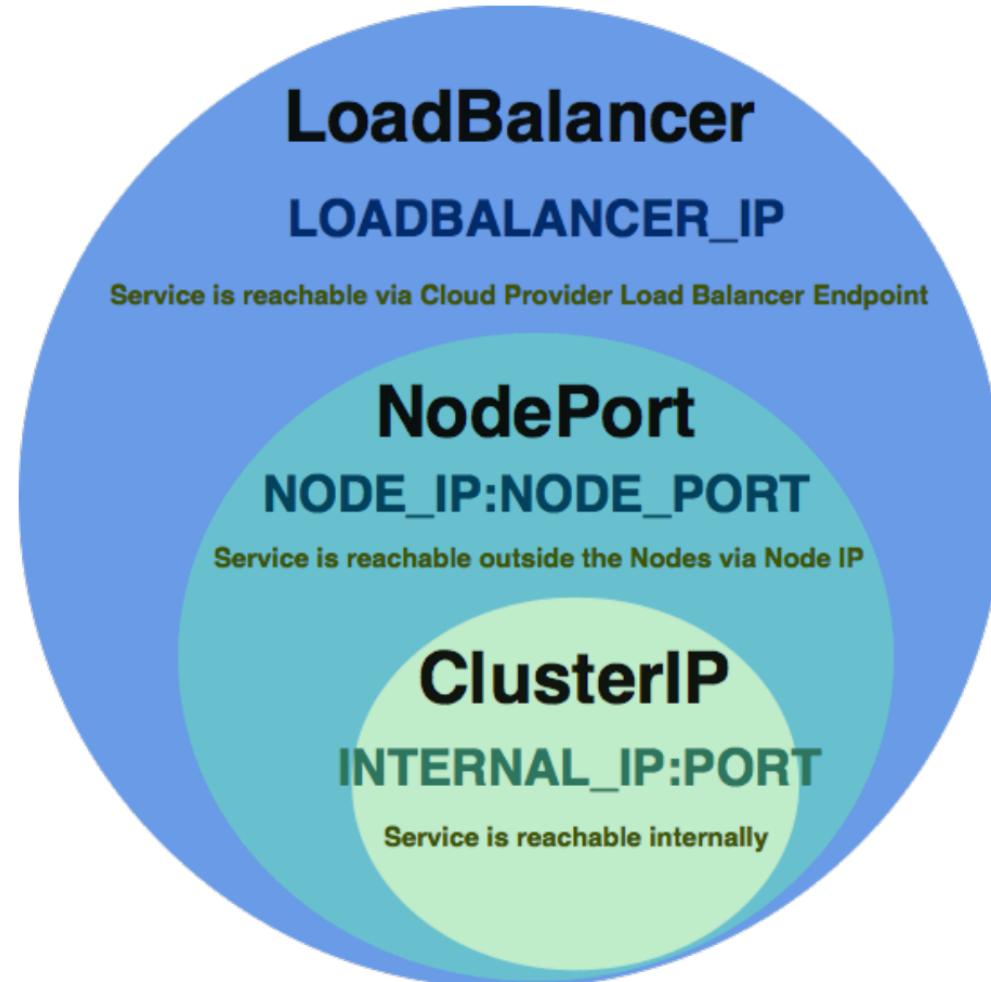
## Use cases

- Accessing an external database to your Kubernetes cluster
- Connecting a third-party API from within your cluster



# Service type

---



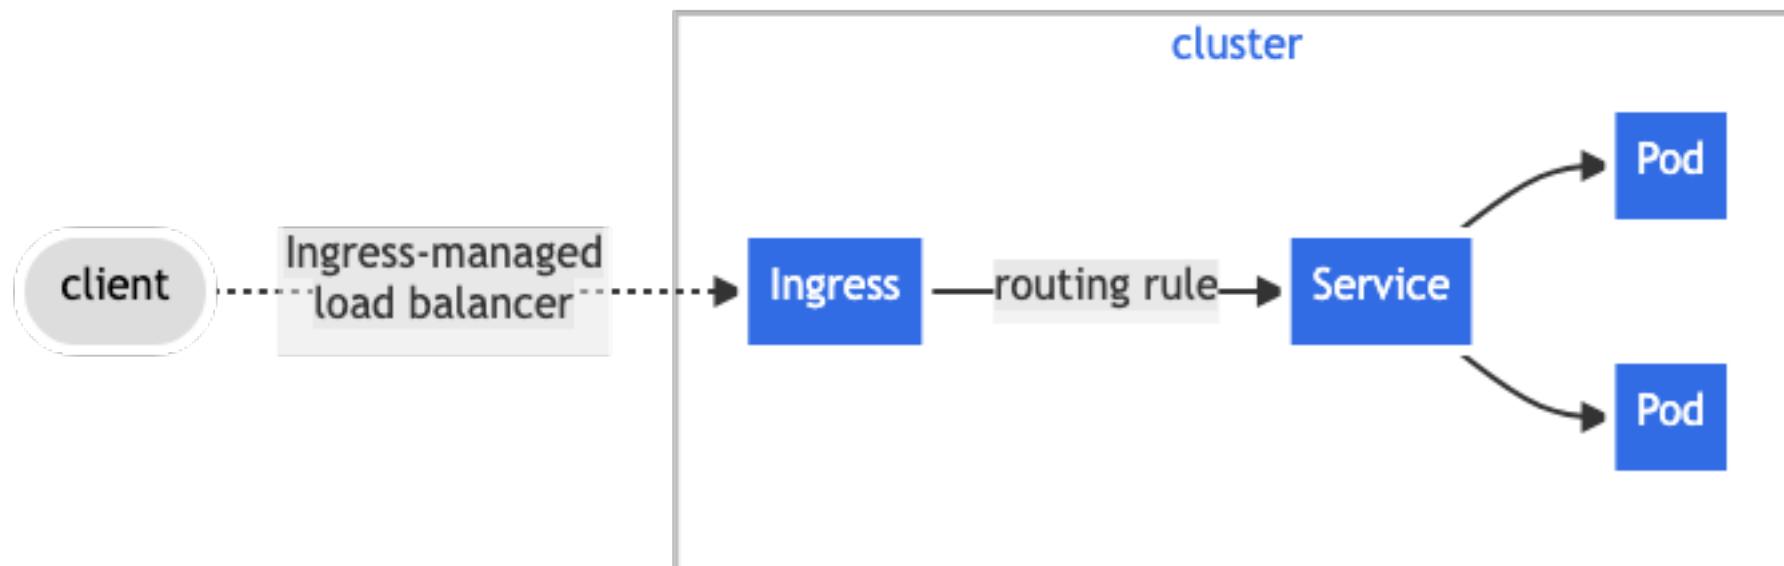
# Ingress

---



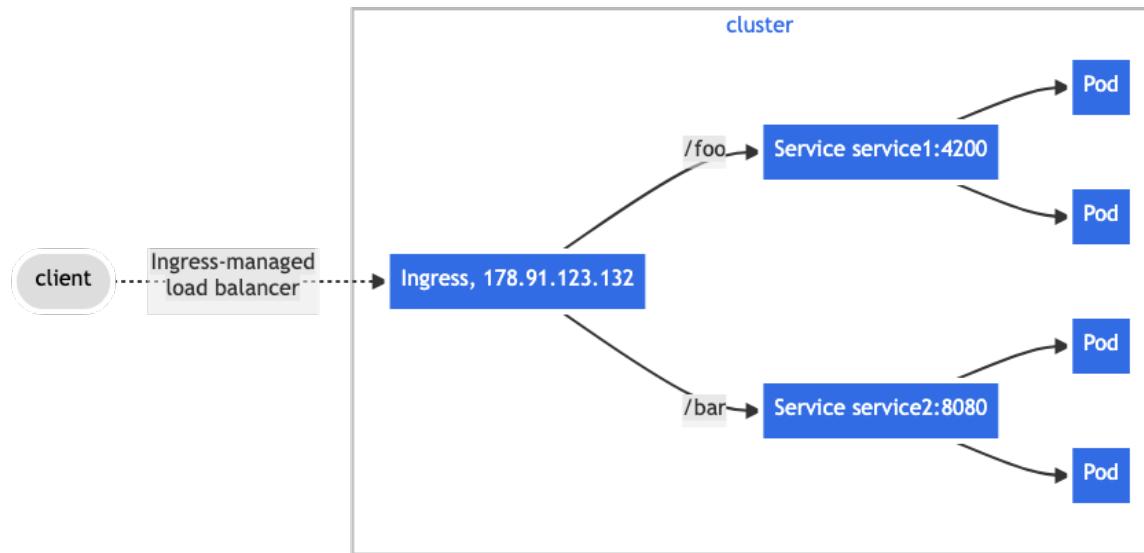
# Ingress

- Routing from outside the cluster to services within the cluster
- Traffic routing is controlled by rules defined on the ingress resource
- An **Ingress controller** is responsible for fulfilling the Ingress



# Ingress Controllers

- The Ingress Controller is an application hosted inside a Kubernetes cluster that actively **manages traffic** following Ingress Resources and pre-defined Ingress Rules
- Ingress controllers doesn't come with standard, have to be deployed separately
- Popular controllers include ingress-nginx, HAProxy, Traefik





# Workshop Ingress

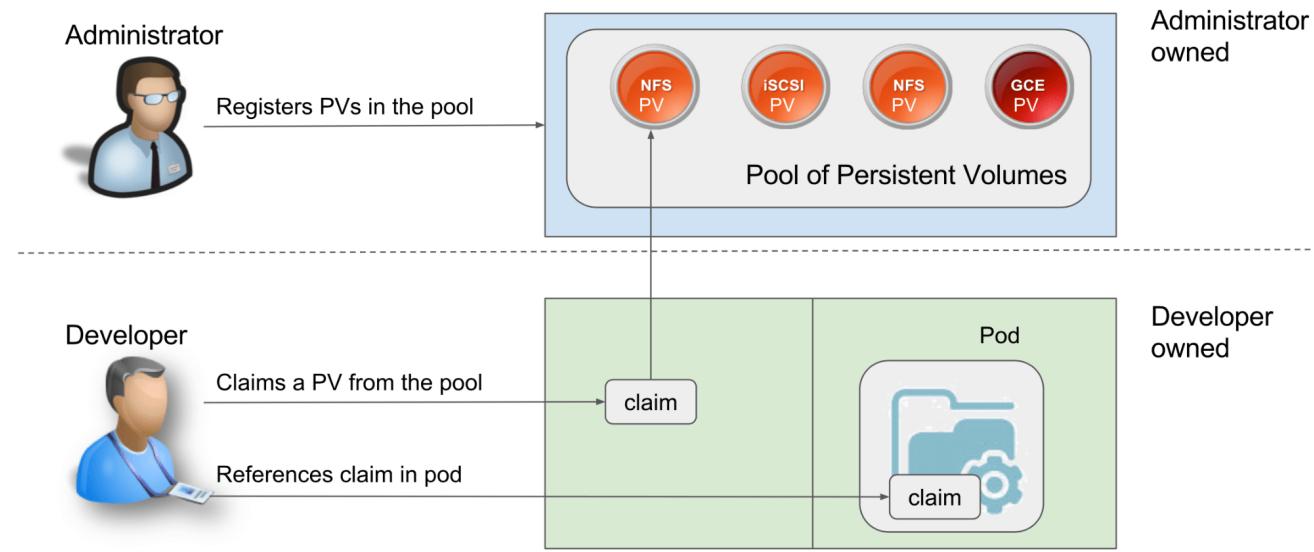
# Persistent Volume

---

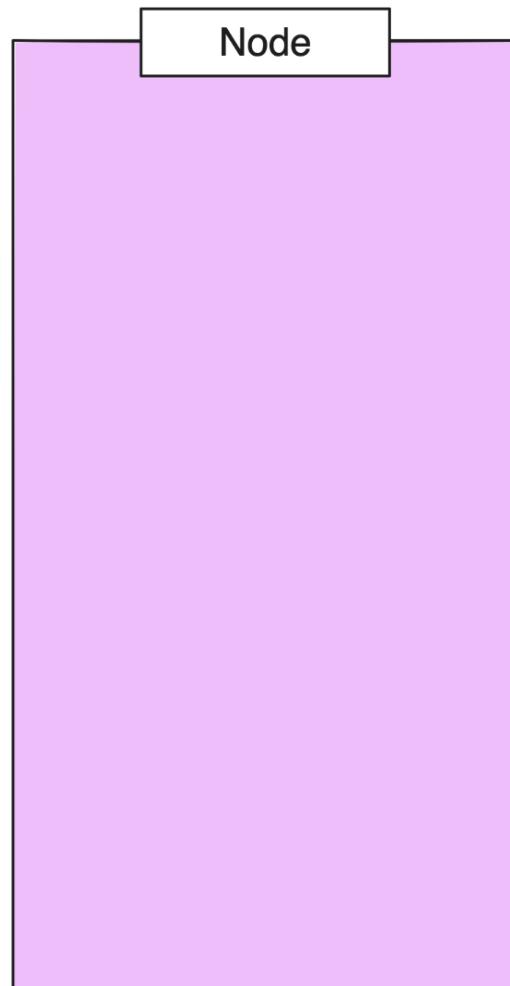
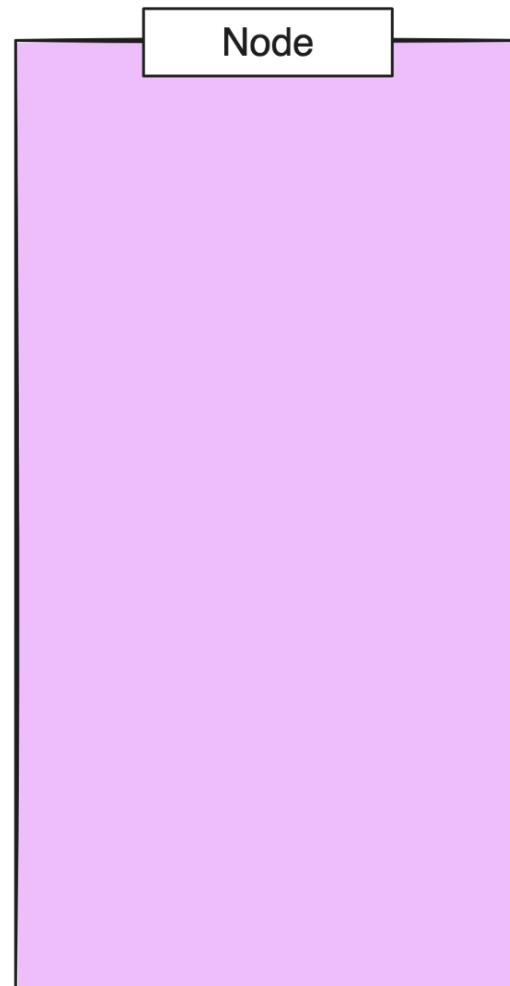
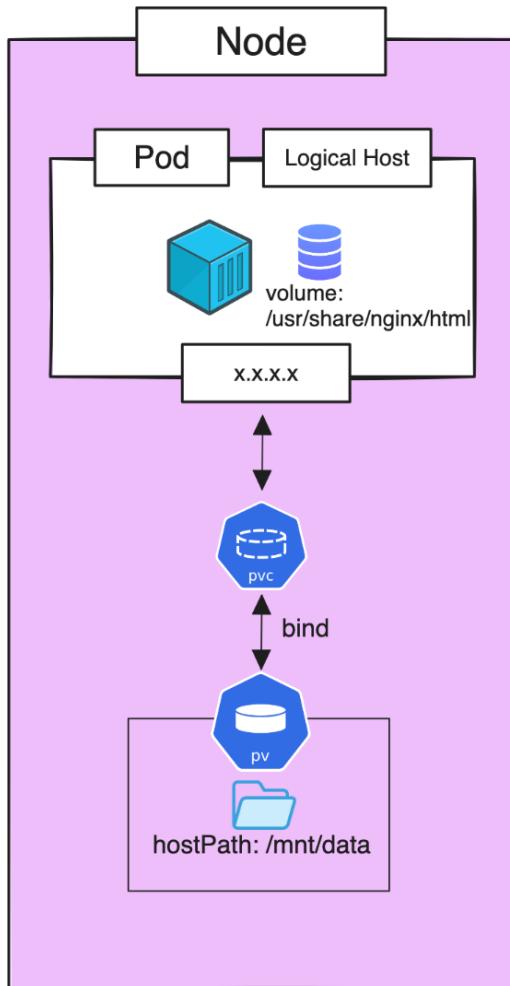


# Persistent Volume

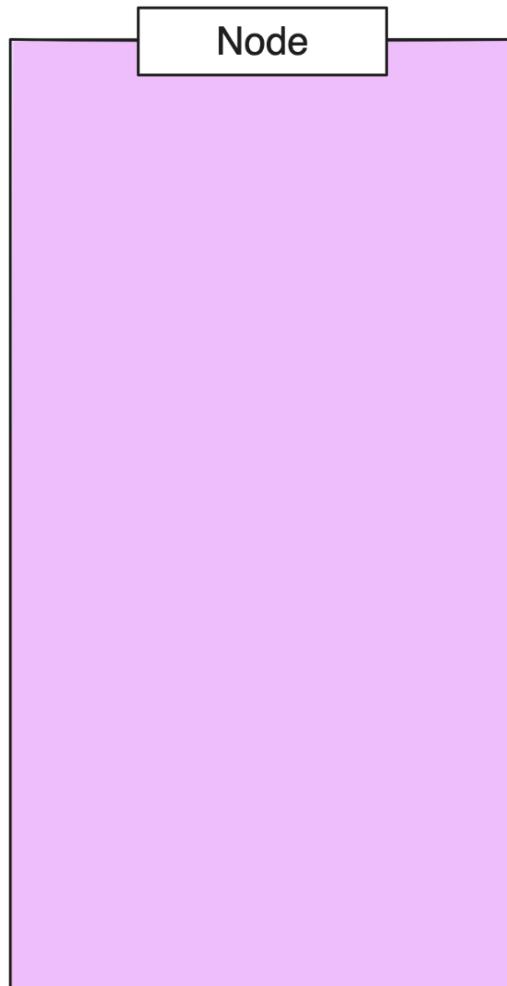
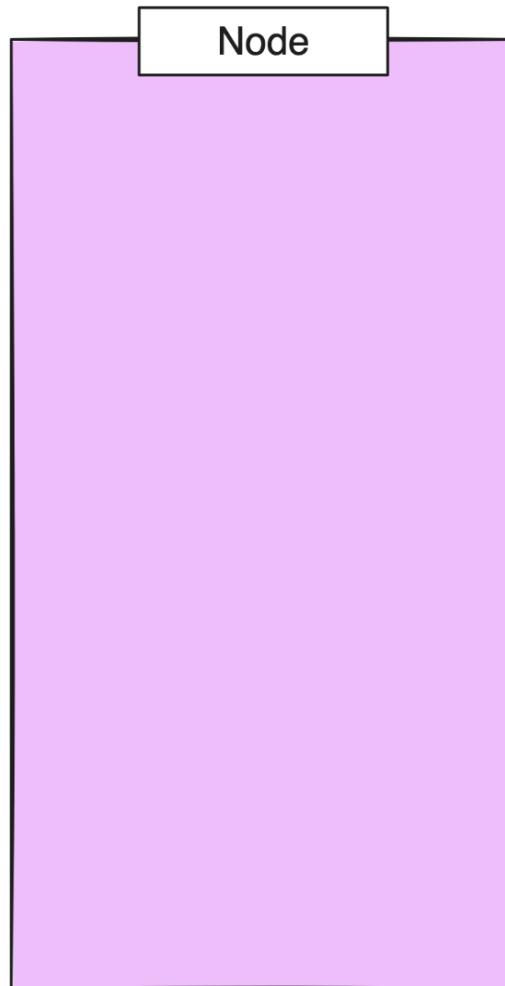
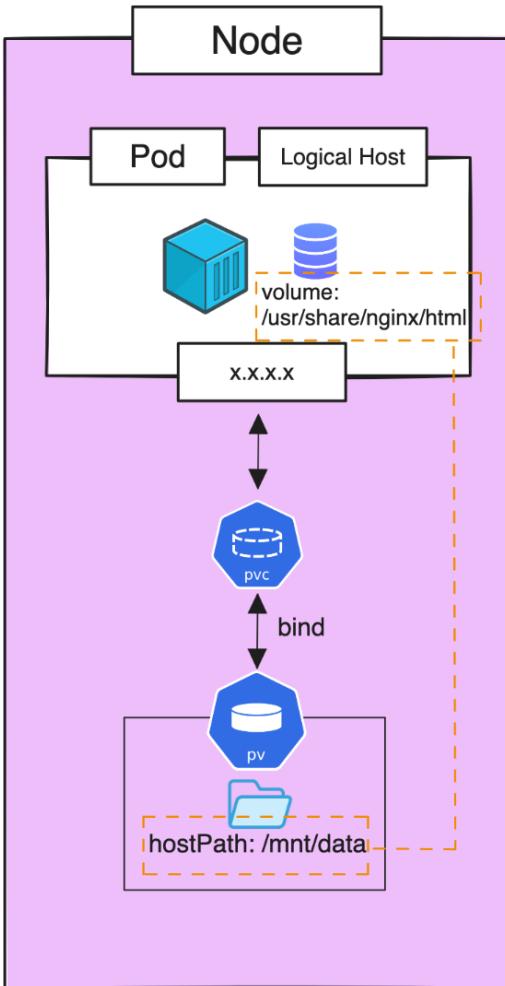
- Persistent Volume (PV) is a piece of storage in the cluster
- Can be provisioned from various sources like host local storage, cloud providers storage, or NFS.
- Persistent Volume Claims (PVC) allow a user to consume abstract storage resources
- A PVC to PV binding is a one-to-one mapping



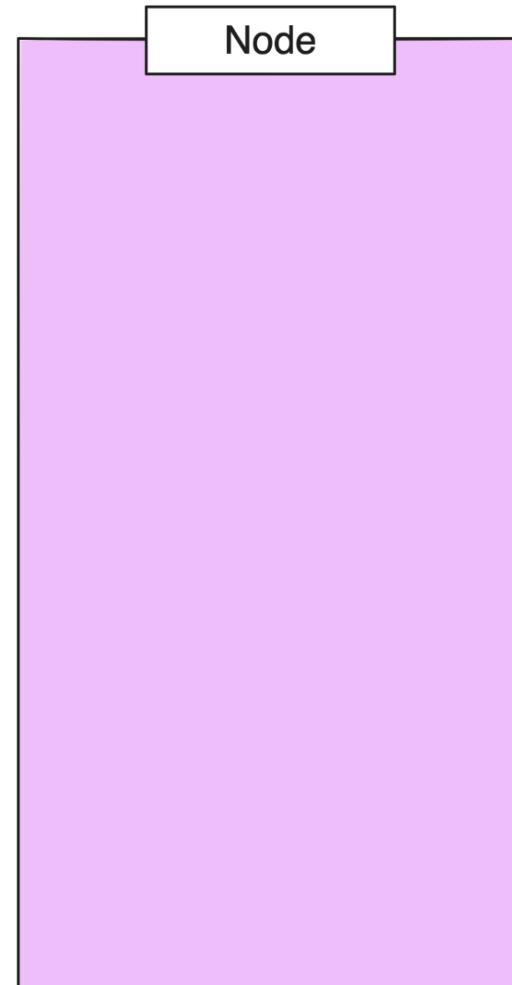
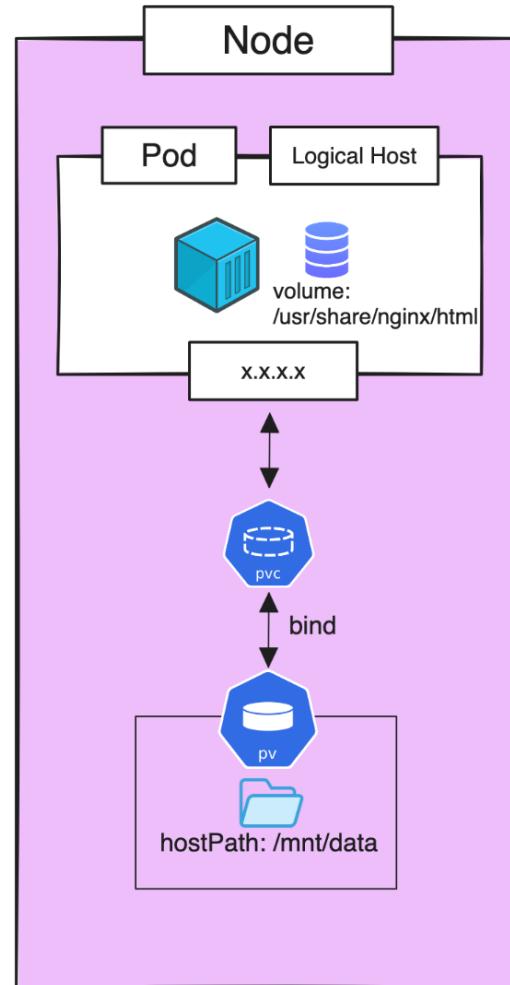
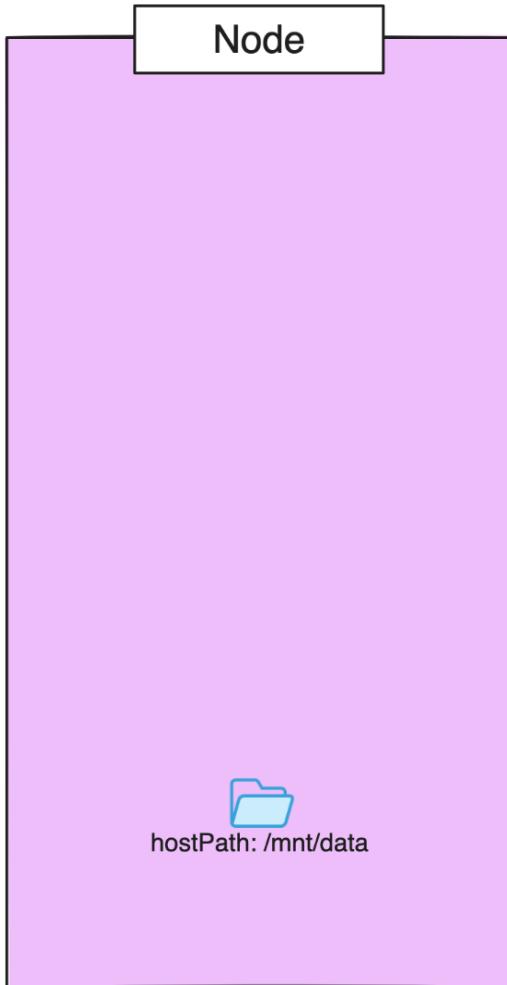
# PV & PVC



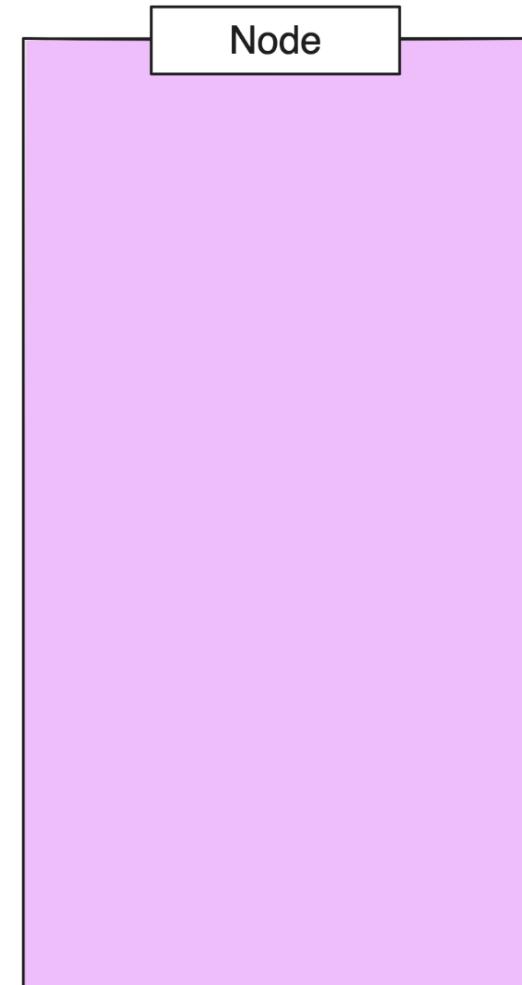
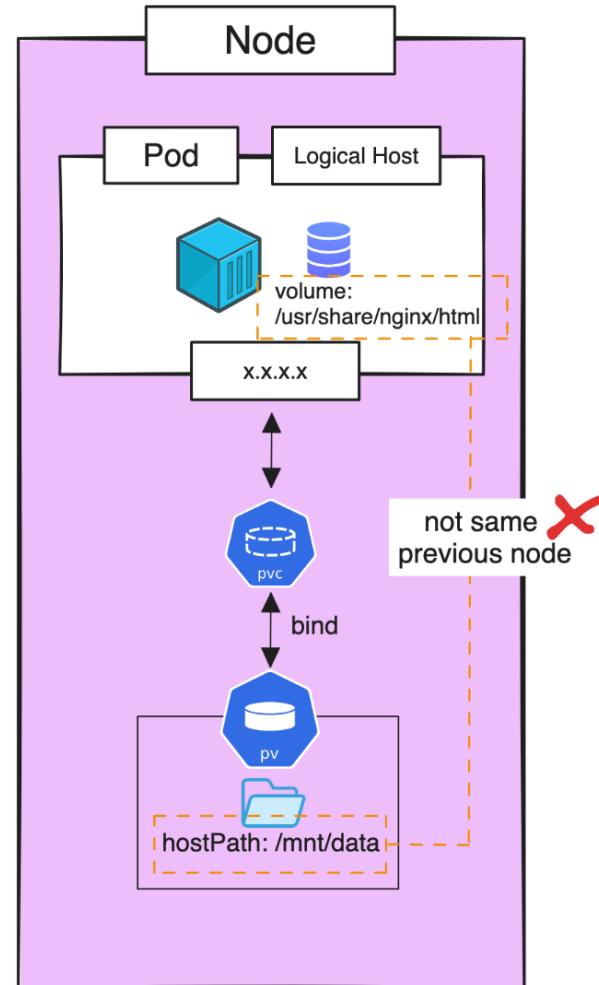
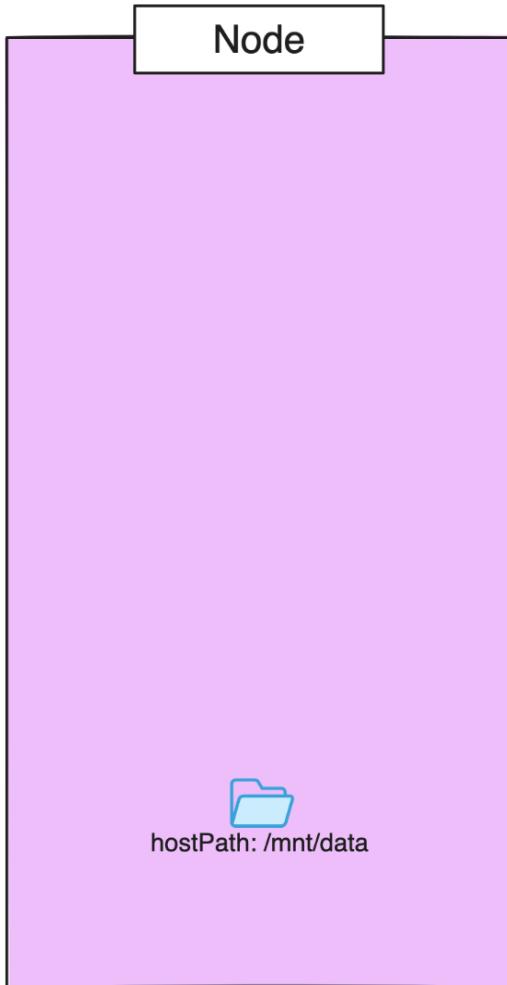
# PV & PVC



# PV & PVC



# PV & PVC



# PV & PVC

---

## Solutions

- Using selector node or nodeAffinity for ensure the target node is same
- Using shared storages like NFS or cloud storages

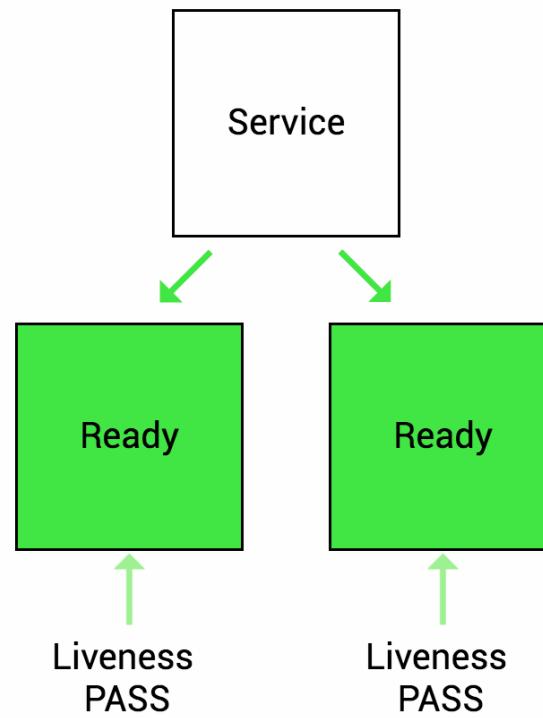


# **Pod liveness & readiness**

---

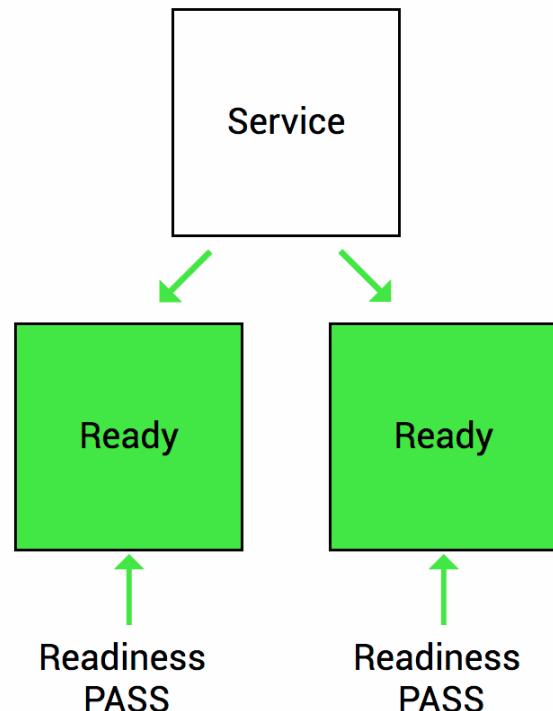
# Pod liveness & readiness

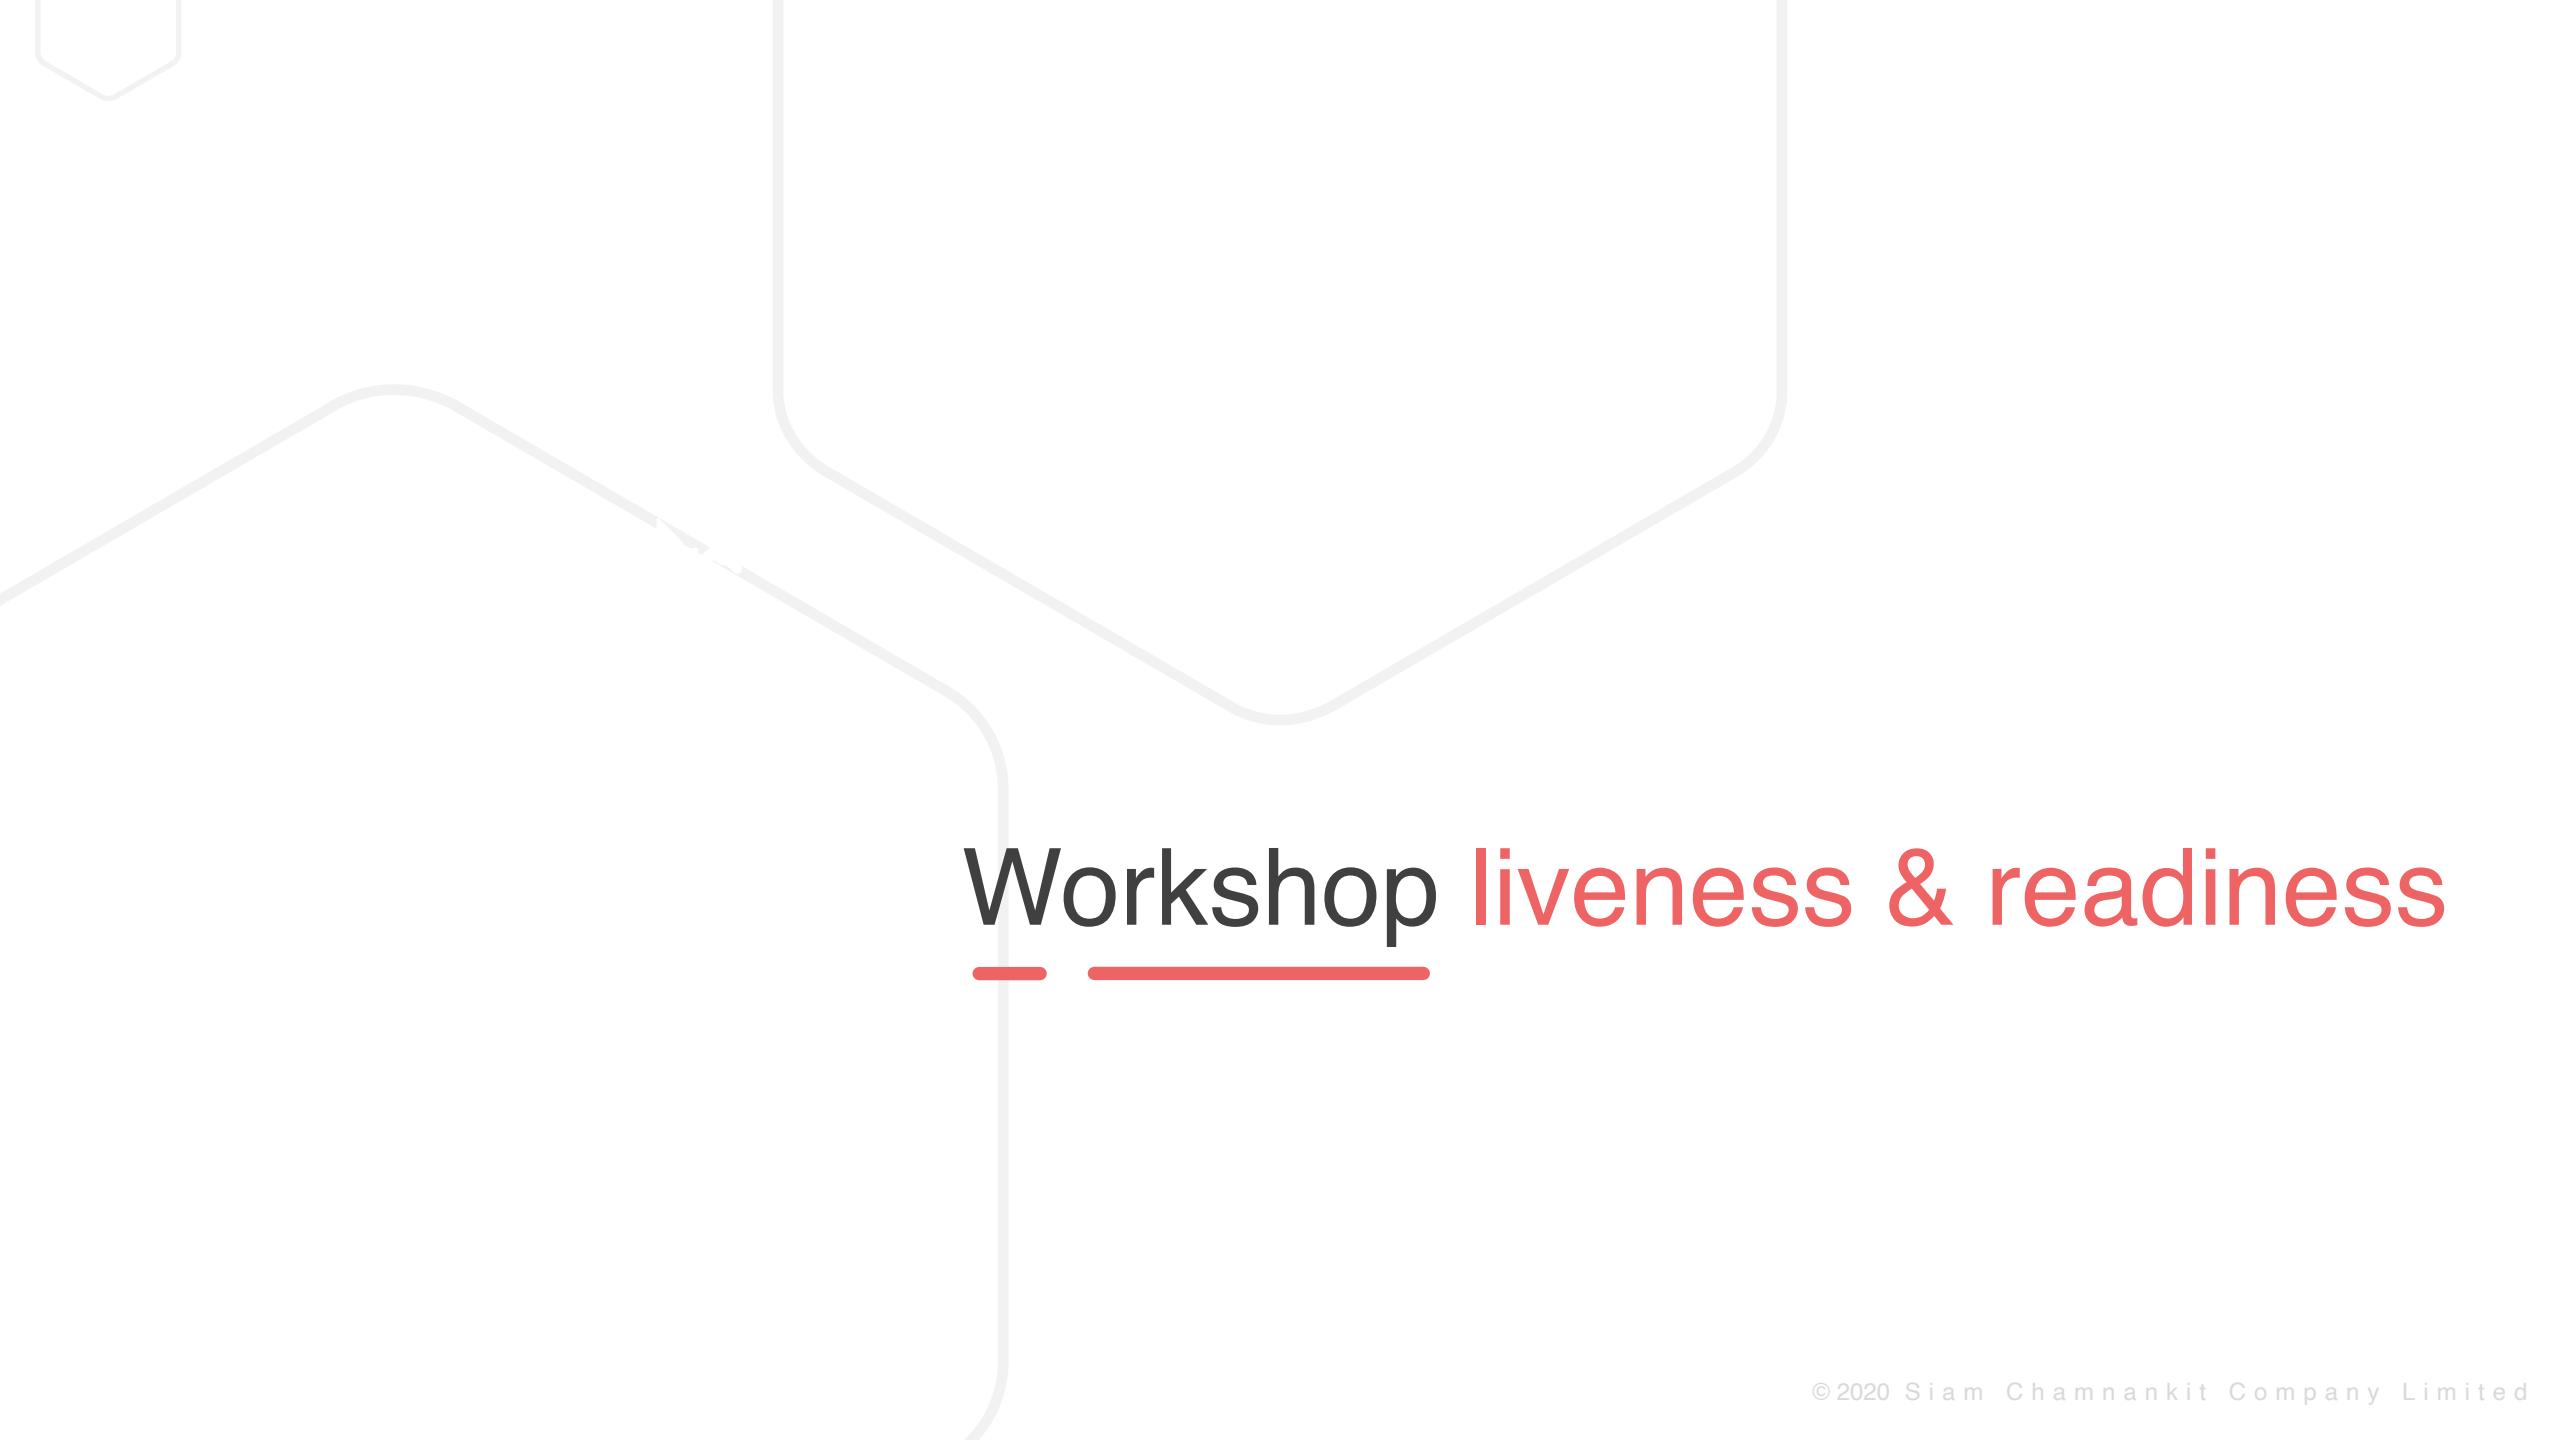
**Liveness:** Indicates whether the Container is running. and if for some reason the liveness probe fails, it restarts the container



# Pod liveness & readiness

**Readiness:** Indicates whether the Container is ready to service requests. If the condition inside readiness probe passes, only then our application can serve traffic.





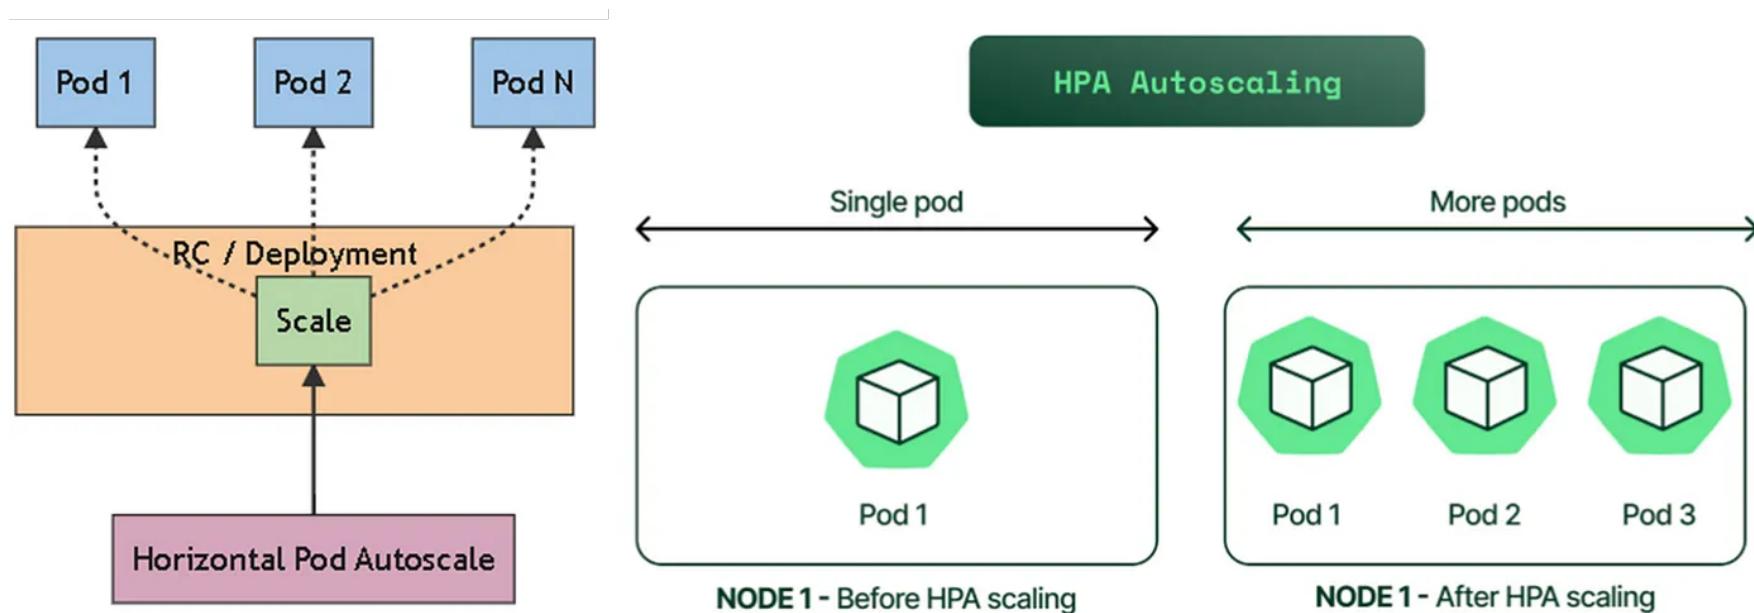
# Workshop liveness & readiness

# HPA (Horizontal Pod Autoscaler)

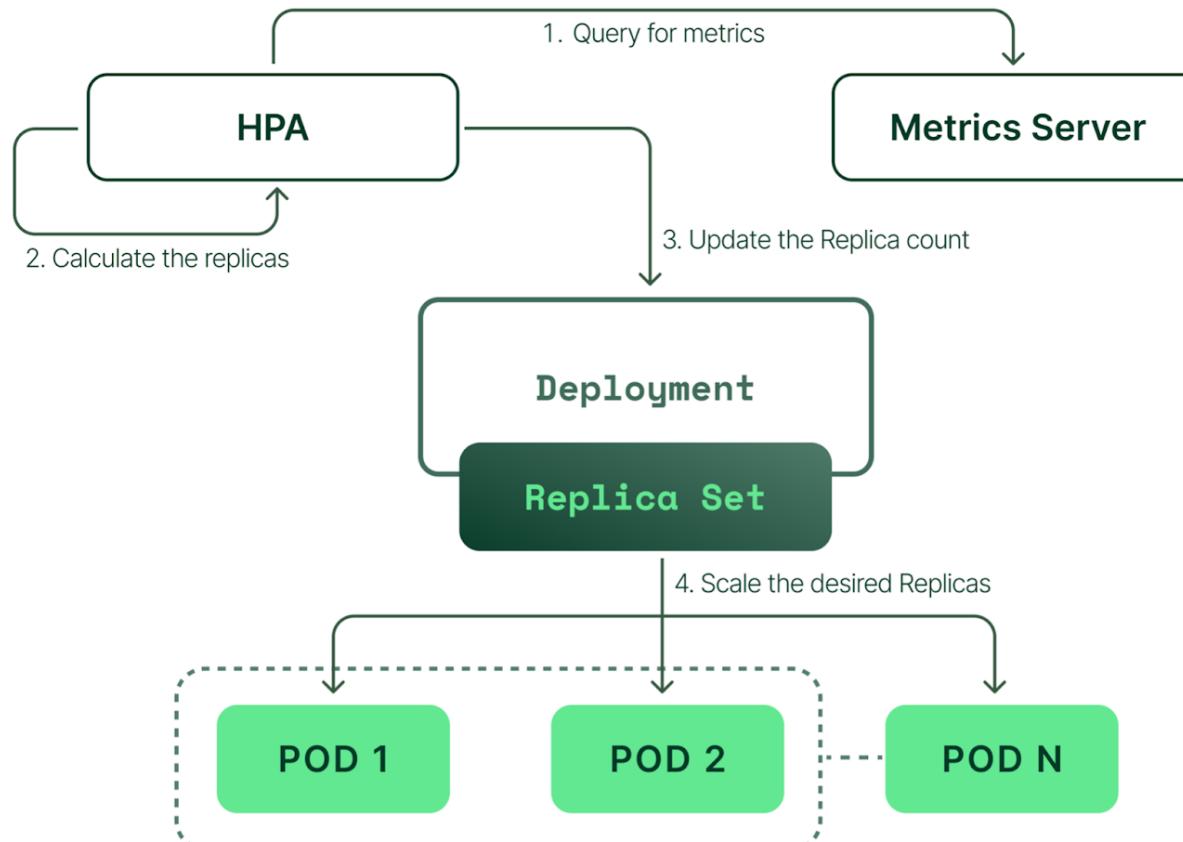
# HPA

---

- A HorizontalPodAutoscaler (HPA for short) automatically updates a workload resource, with the aim of automatically scaling the workload to match demand



# How does HPA work?





# Workshop HPA

# Namespace

---



# Namespace

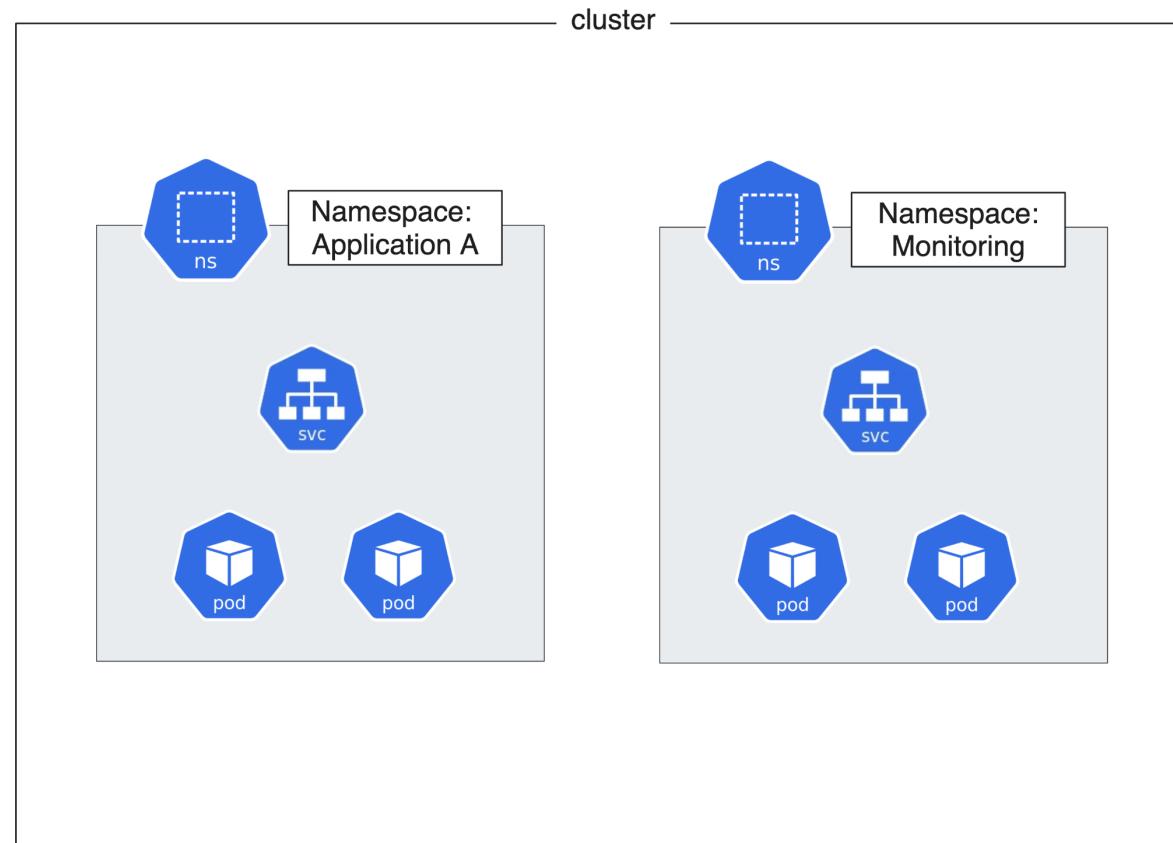
---

- Partition a single kubernetes cluster into multiples clusters
- 4 defaults namespaces
  - **kube-system**: core system processes
  - **kube-public**: public accessible data, ex: ConfigMaps & Secrets
  - **kube-node-lease**: heartbeat of nodes, so that the control plane can detect node failure
  - **default**: your resources/objects



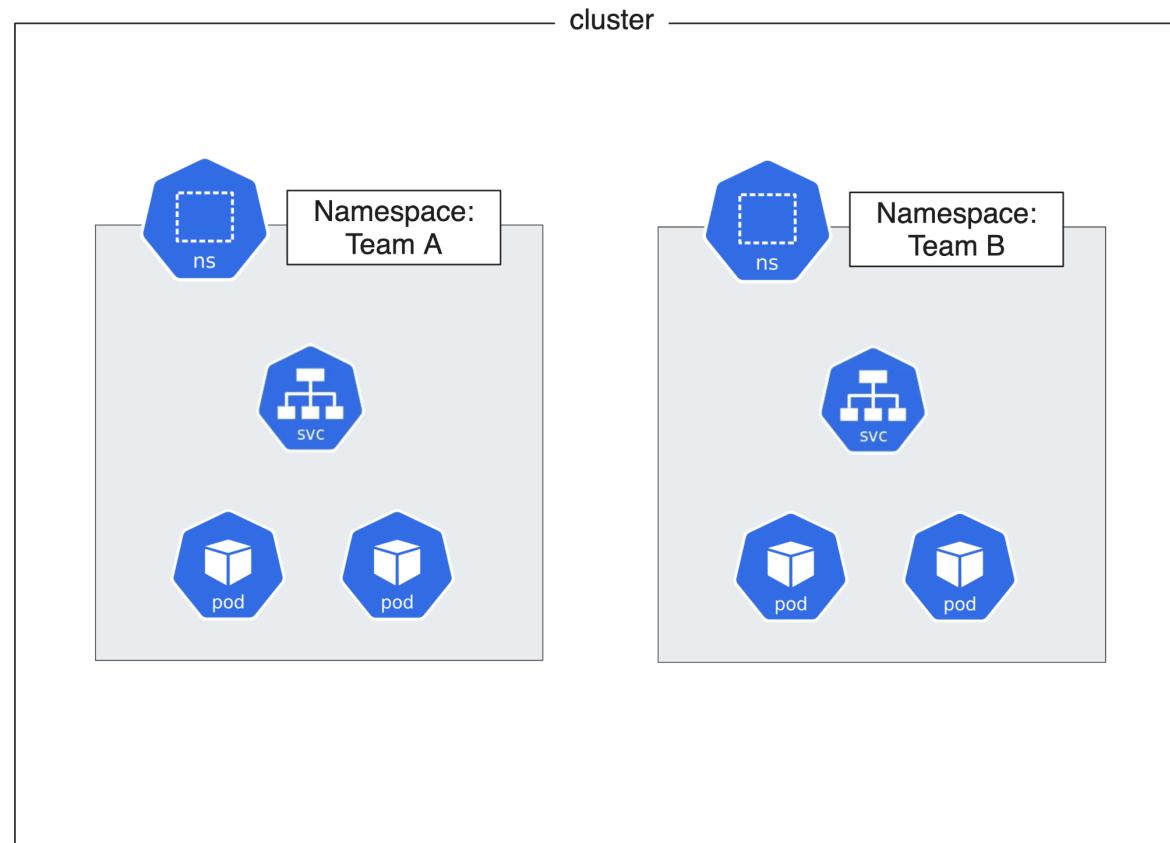
# Namespace use cases

- Application domain



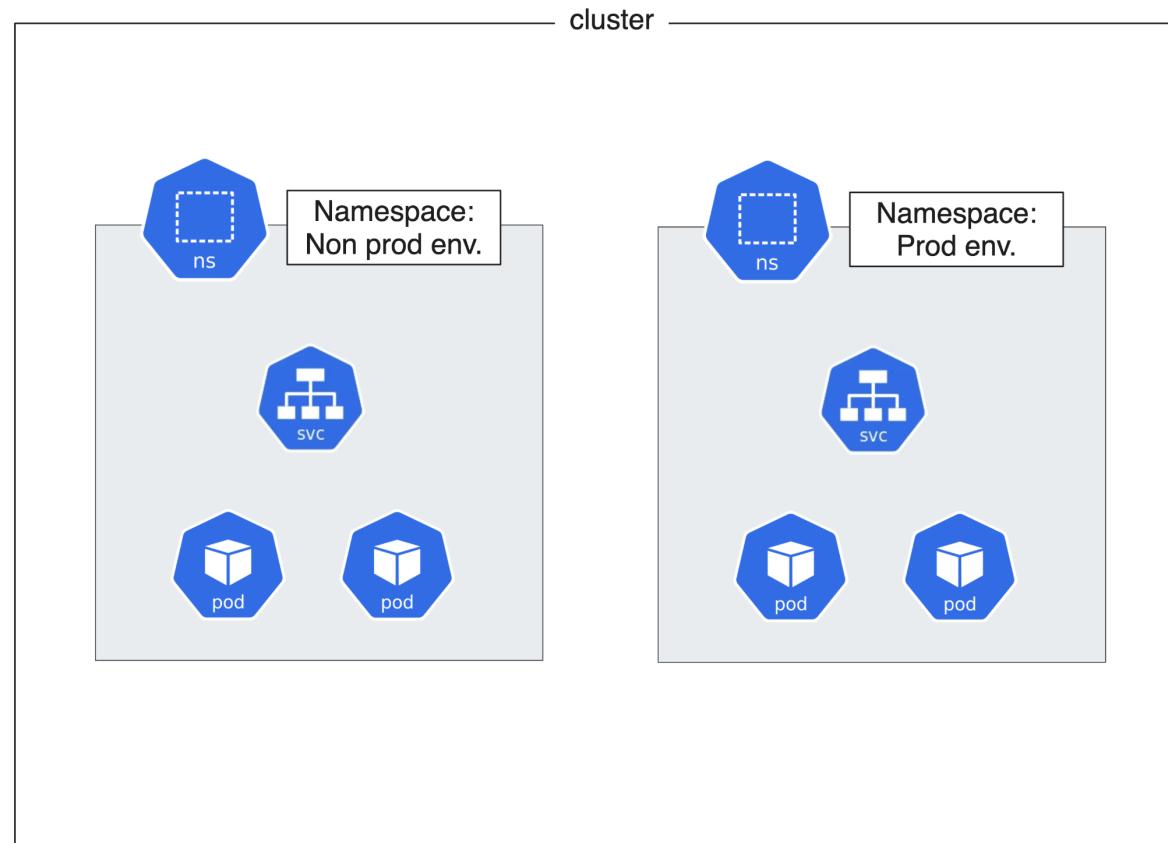
# Namespace use cases

- Teams driven



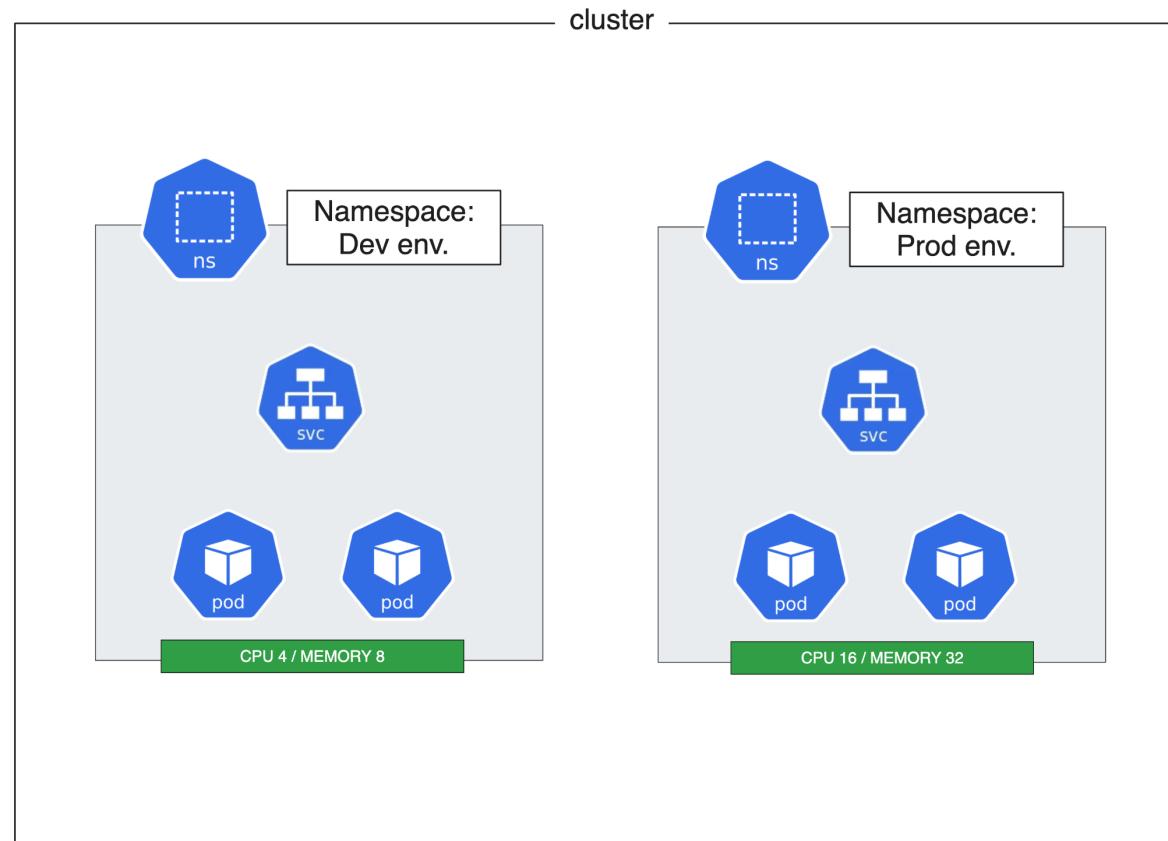
# Namespace use cases

- Environments



# Namespace use cases

- Resources quota limit



# HELM

---



# HELM

---

HELM is the package manager for Kubernetes



# HELM

---

HELM is the package manager for Kubernetes

- Boosts productivity
- Reduces complexity
- Simple sharing
- Versioning and Upgrades



# HELM Charts

---

Public HELM Chart: <https://artifacthub.io/>

The screenshot shows the Artifact Hub website with a teal header. The header includes the logo "Artifact HUB BETA", "SIGN UP", "SIGN IN", and a gear icon. Below the header, the text "Find, install and publish Kubernetes packages" is displayed. A search bar with the placeholder "Search packages" and a tip below it stating "Tip: Use or to combine multiple searches. Example: postgresql or mysql" is visible. Below the search bar, there are several buttons for filtering packages: "Packages from verified publishers", "Prometheus packages in official repositories", "Helm Charts provided by Bitnami", "Helm Charts in the storage category", and "OLM operators for databases". At the bottom, two large numbers are shown: "2086 PACKAGES" and "38119 RELEASES". The footer contains the text "Artifact Hub is an Open Source project" and links to GitHub, Slack, and Twitter.

SIGN UP SIGN IN

Find, install and publish  
Kubernetes packages

Search packages

Tip: Use **or** to combine multiple searches. Example: [postgresql](#) [or](#) [mysql](#)

You can also [browse all packages](#) - or - try one of the sample queries:

Packages from verified publishers | Prometheus packages in official repositories | Helm Charts provided by Bitnami  
Helm Charts in the storage category | OLM operators for databases

2086 | 38119

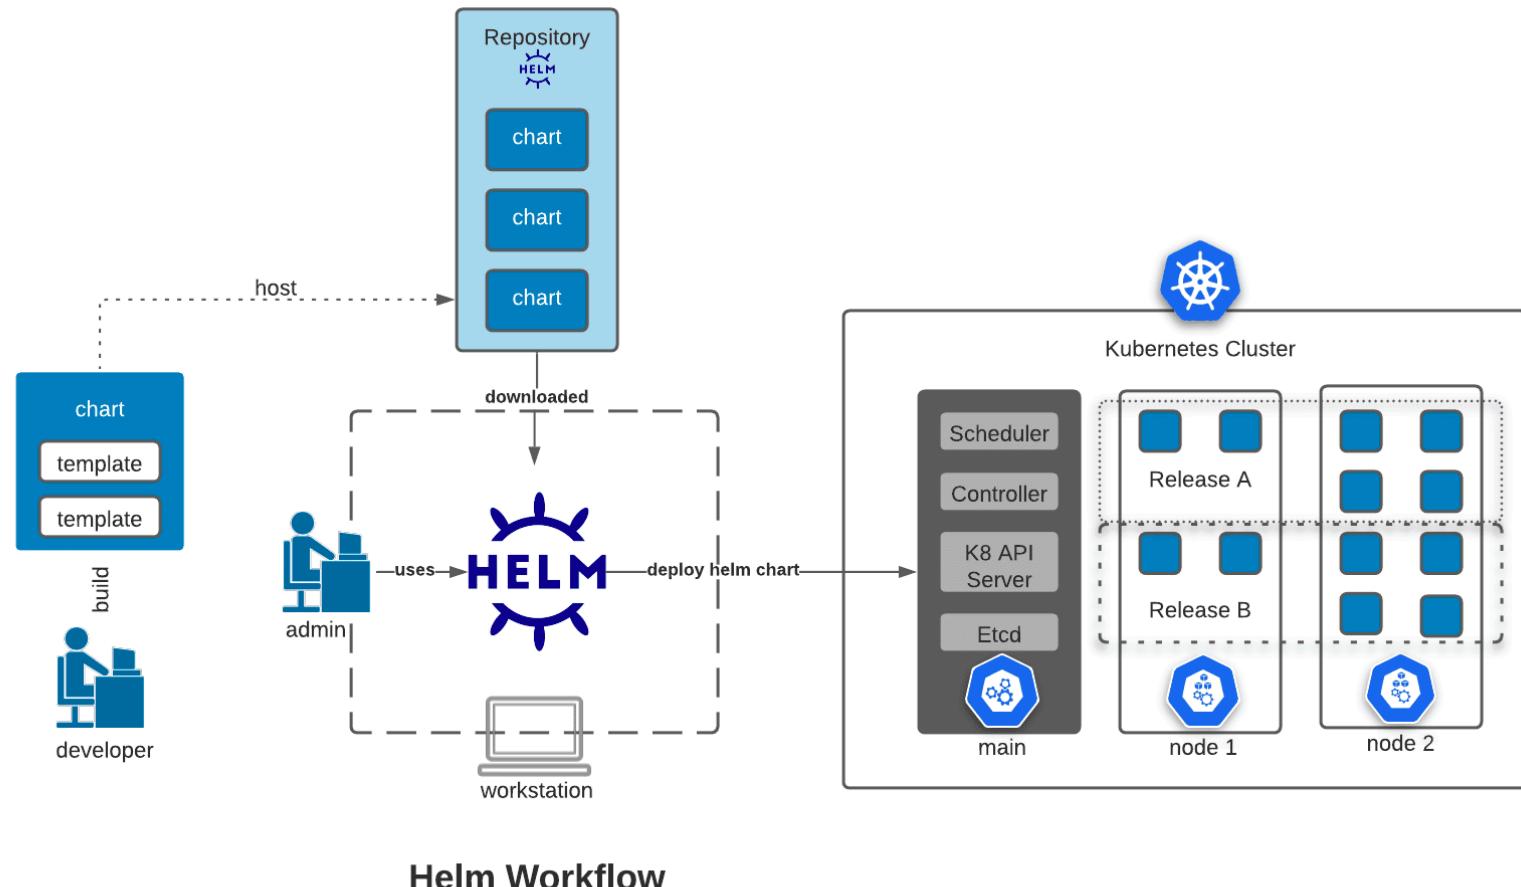
PACKAGES | RELEASES

Artifact Hub is an [Open Source](#) project

GitHub Slack Twitter

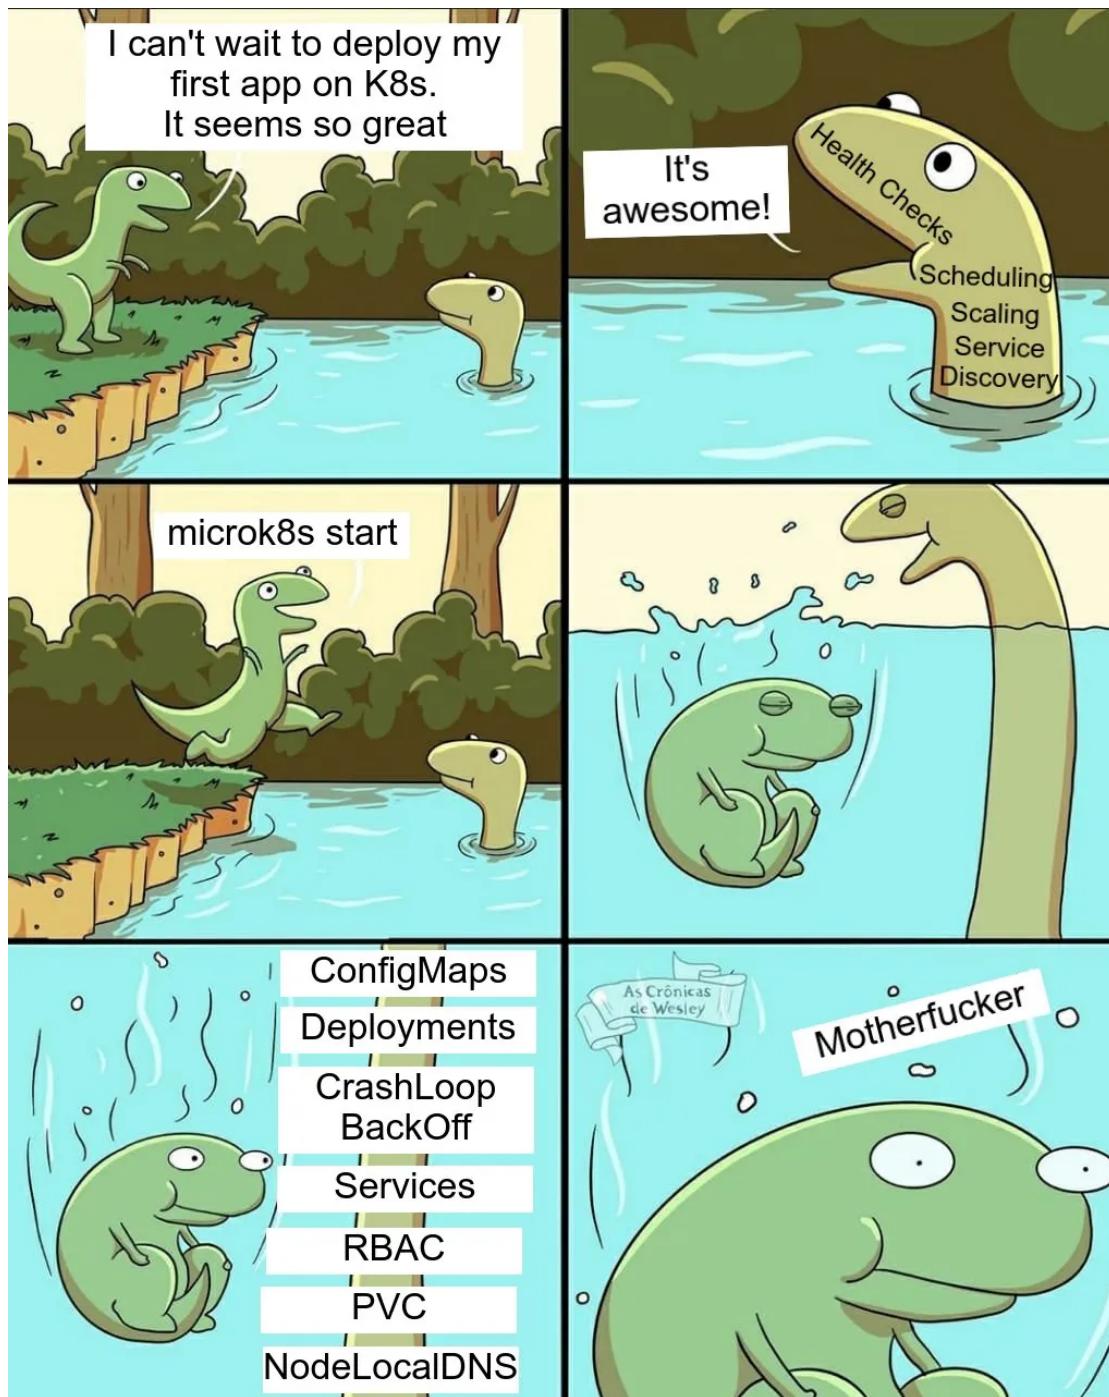


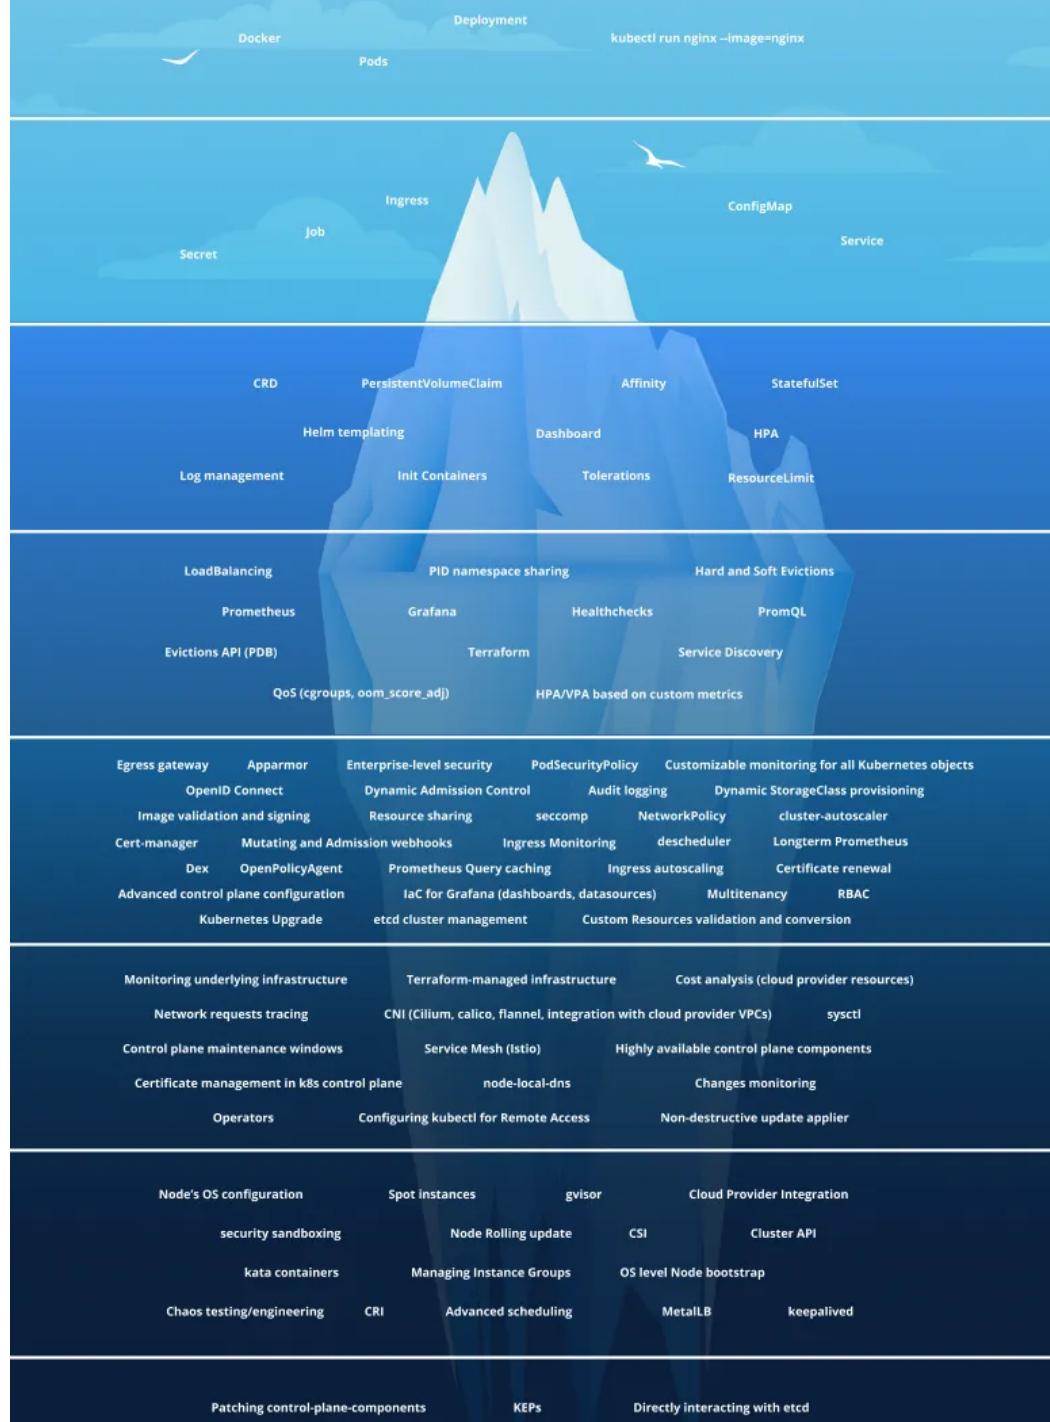
# HELM Charts



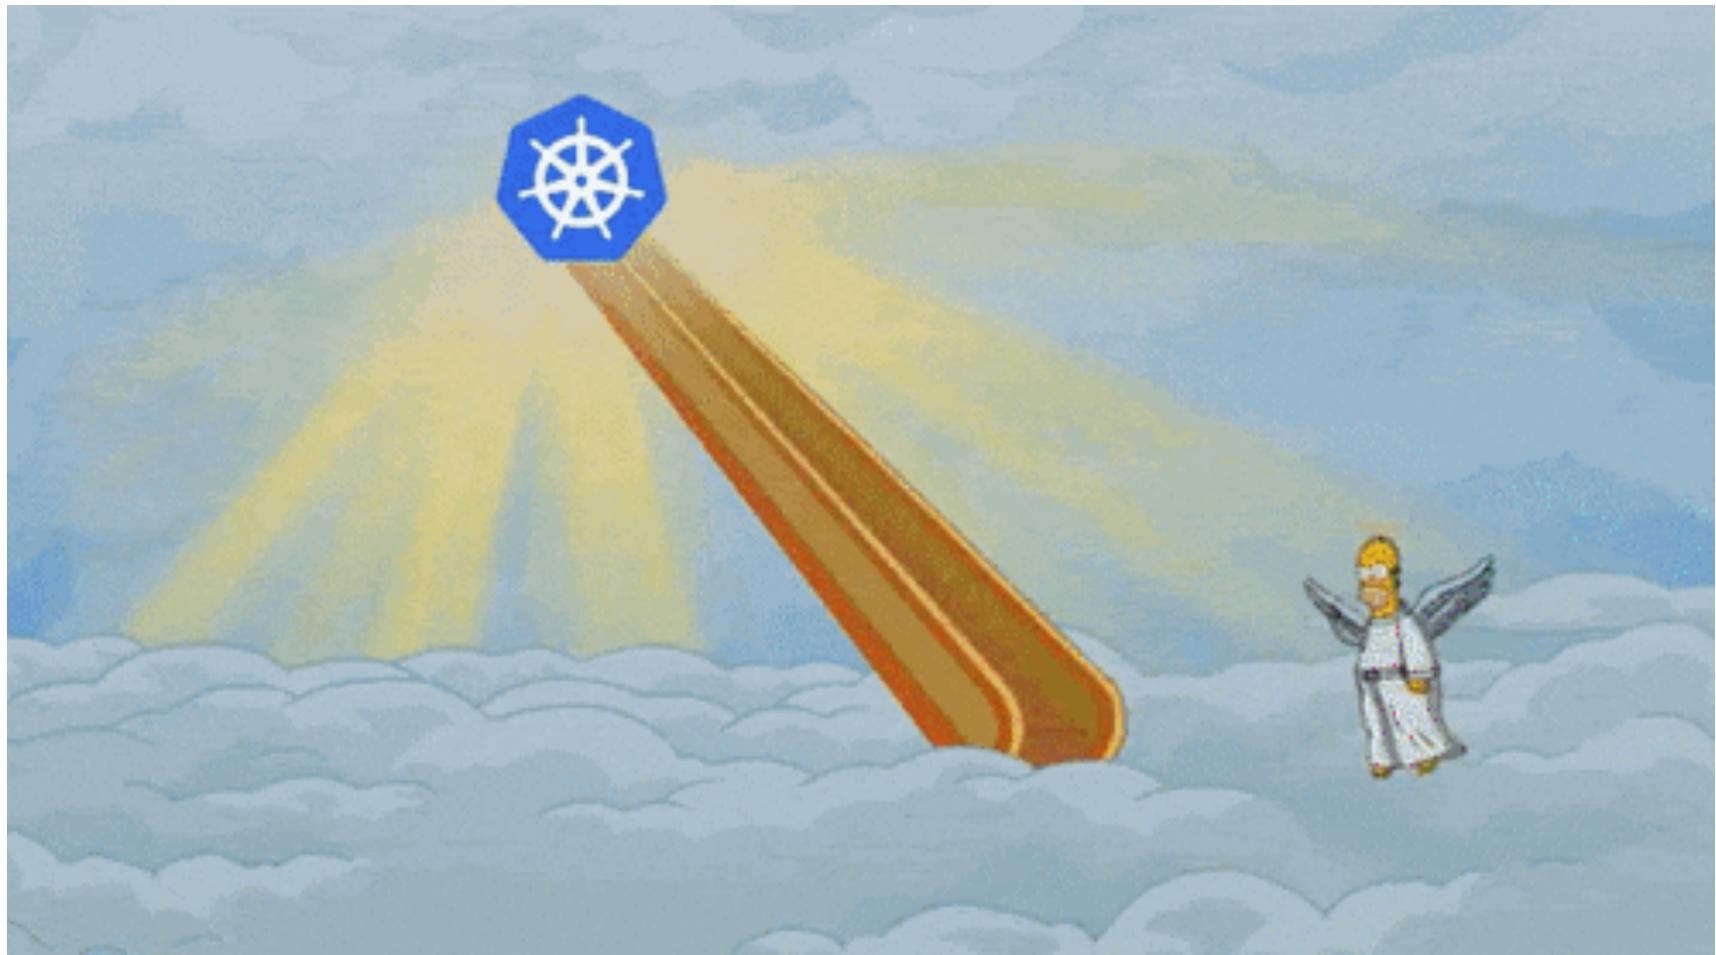


# Workshop HELM









# DevOps



# Agenda

- Introduction
- What is DevOps ?
- How to start DevOps ?
- CI/CD
  - Continuous integration
  - Continuous deployment
- Deployment strategy
- Observability



# Introduction

---



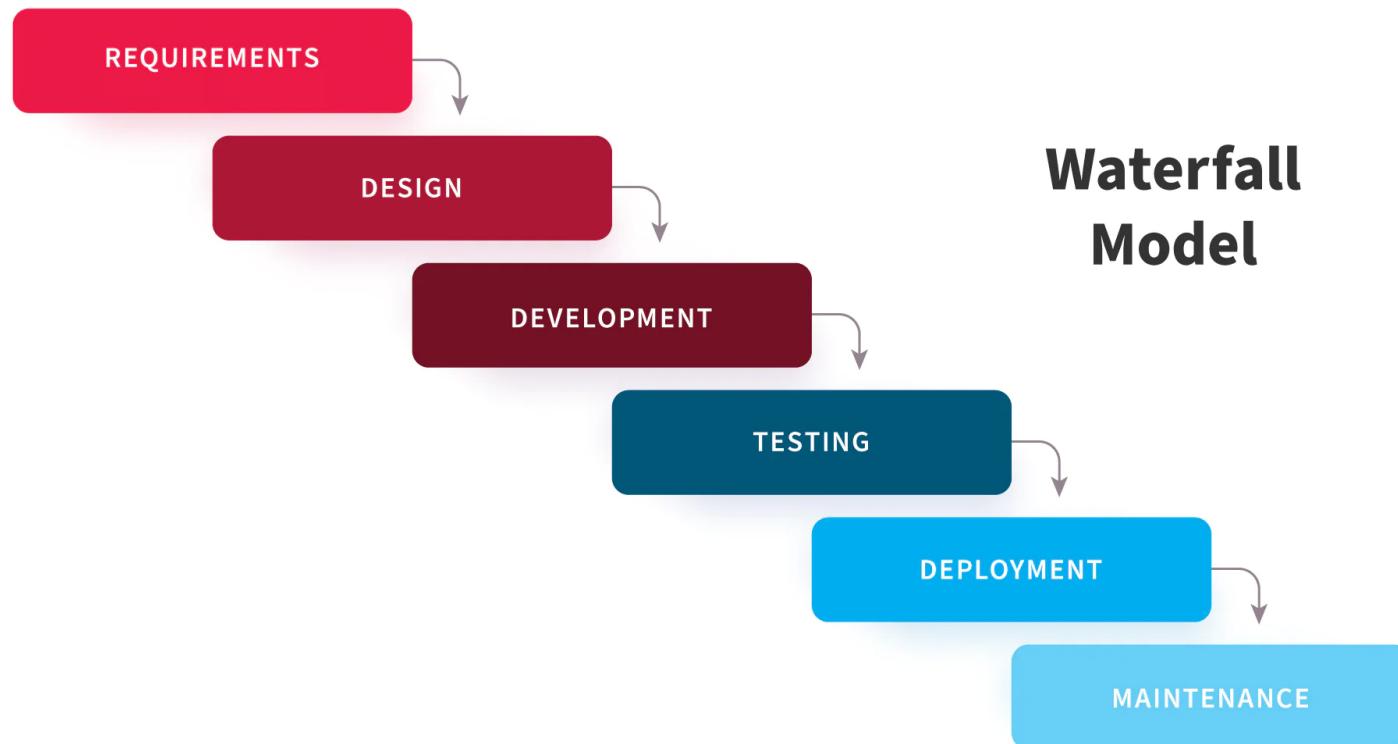
# Before DevOps

---

- In the early days, software was developed and maintained by a small group of specialists
- The development process known as “Waterfall”
- Spending months or even years building and testing software before it was ready to be deployed
- Breaks project activities into phases, each phase must be completed before the next phase can begin



# Waterfall Model



**Waterfall  
Model**



# Waterfall Model

---

- Lack of flexibility
- No room for errors
- Difficulty in dealing with changes
- No customer involvement



# The rise of the Agile methodology

- Focusing on the clean delivery of individual pieces or parts, not the entire application
- Get new features into the hands of the users quickly
- Focusing on rapid iteration
- Emphasizing collaboration, flexibility and continuous improvement



# Agile



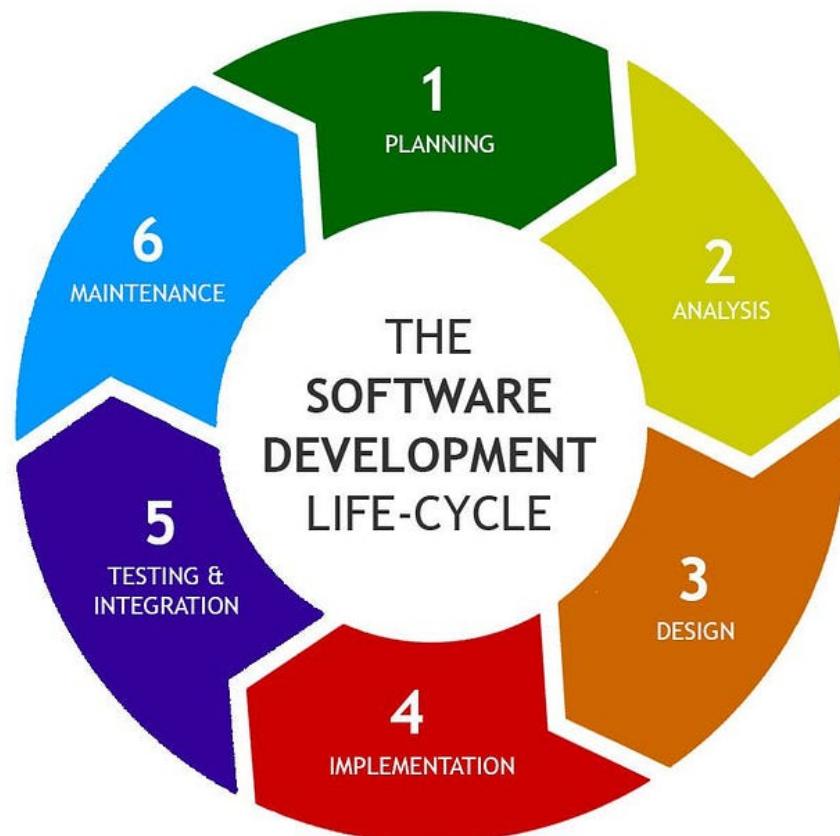
# Agile

## Advantages

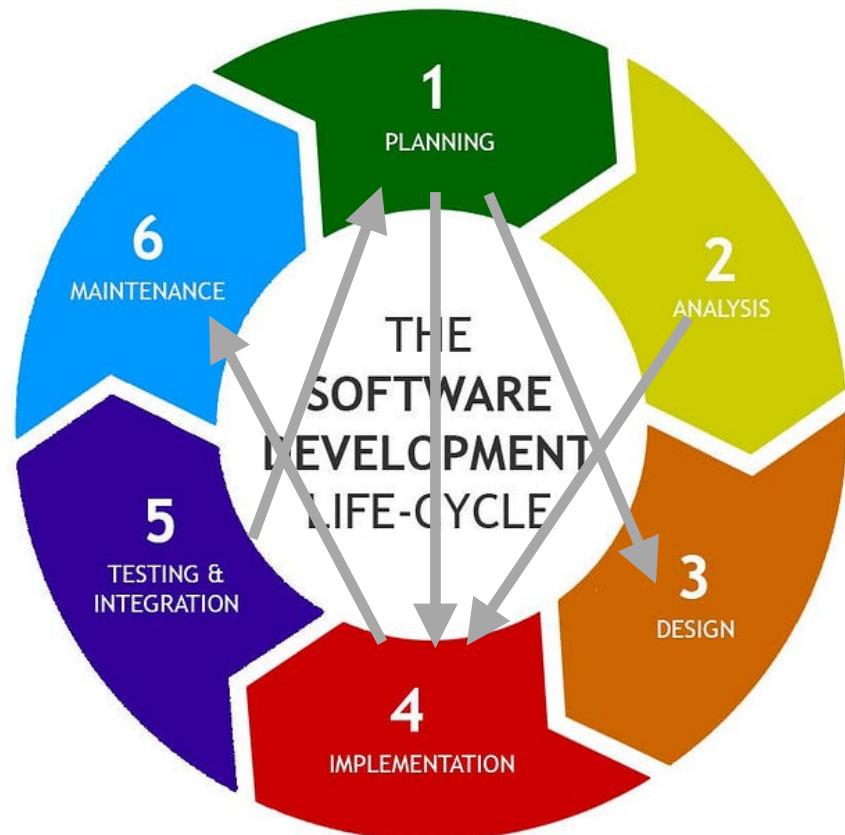
- Faster time-to-market
- Testing and superior quality
- Flexible priorities
- Lower risk
- Project visibility and transparency, etc.,



# SDLC (Software Development Life Cycle)



# SDLC (Software Development Life Cycle)



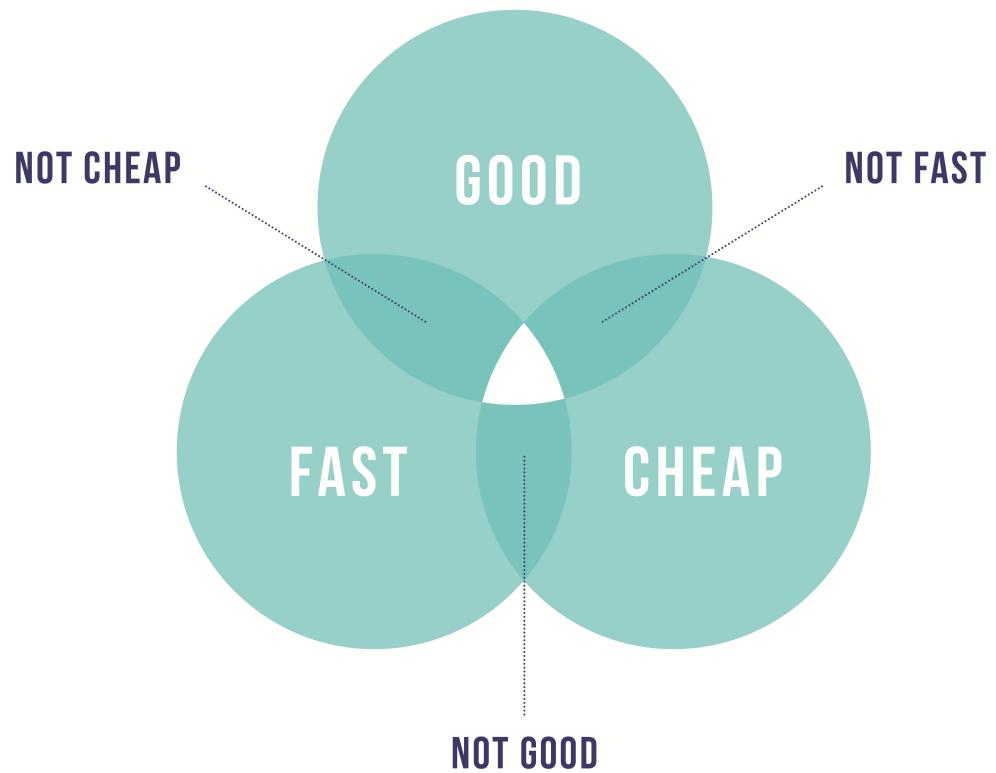
# Cost vs Time vs Quality

---



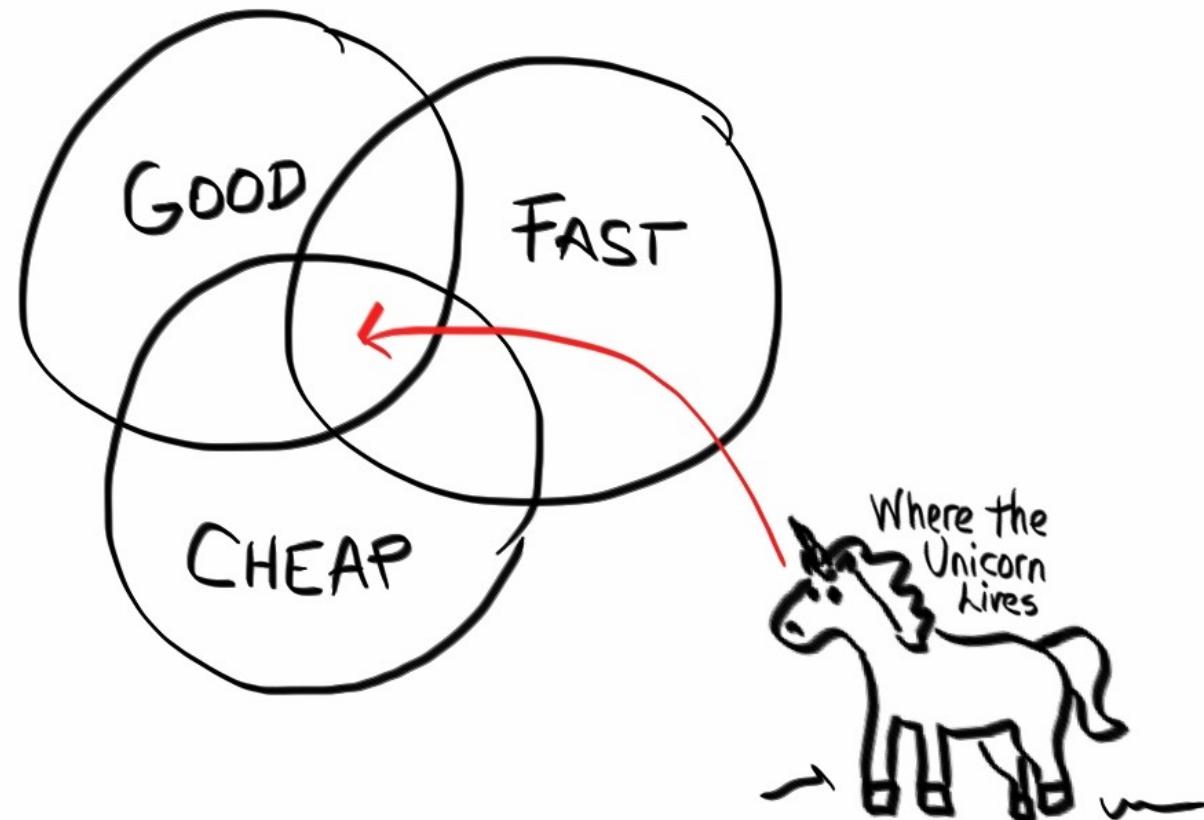
# Cost vs Time vs Quality

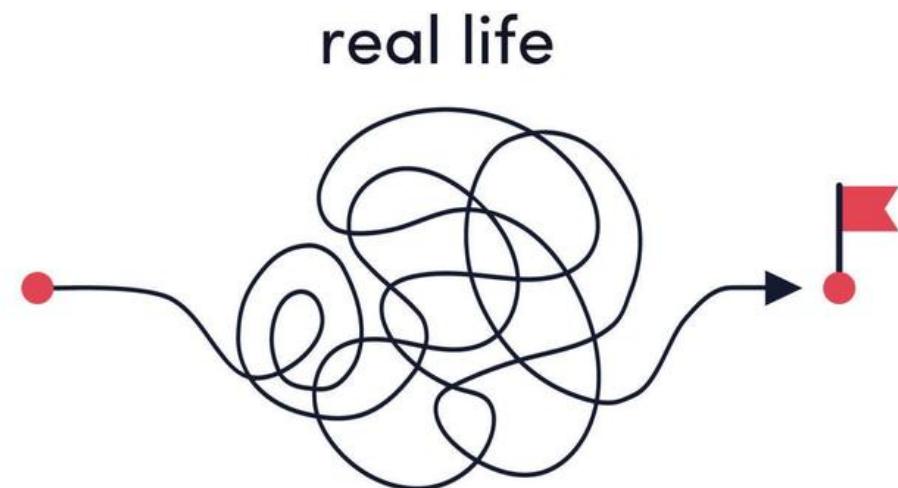
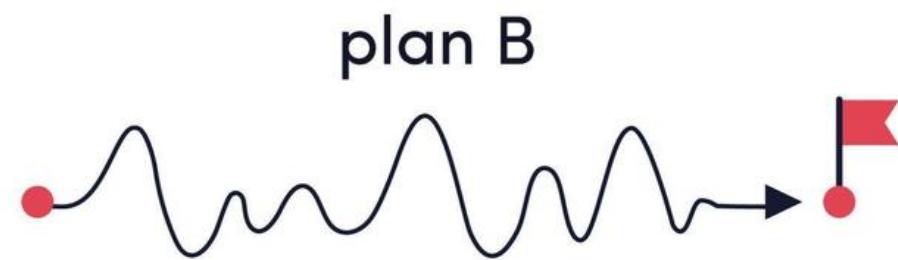
---

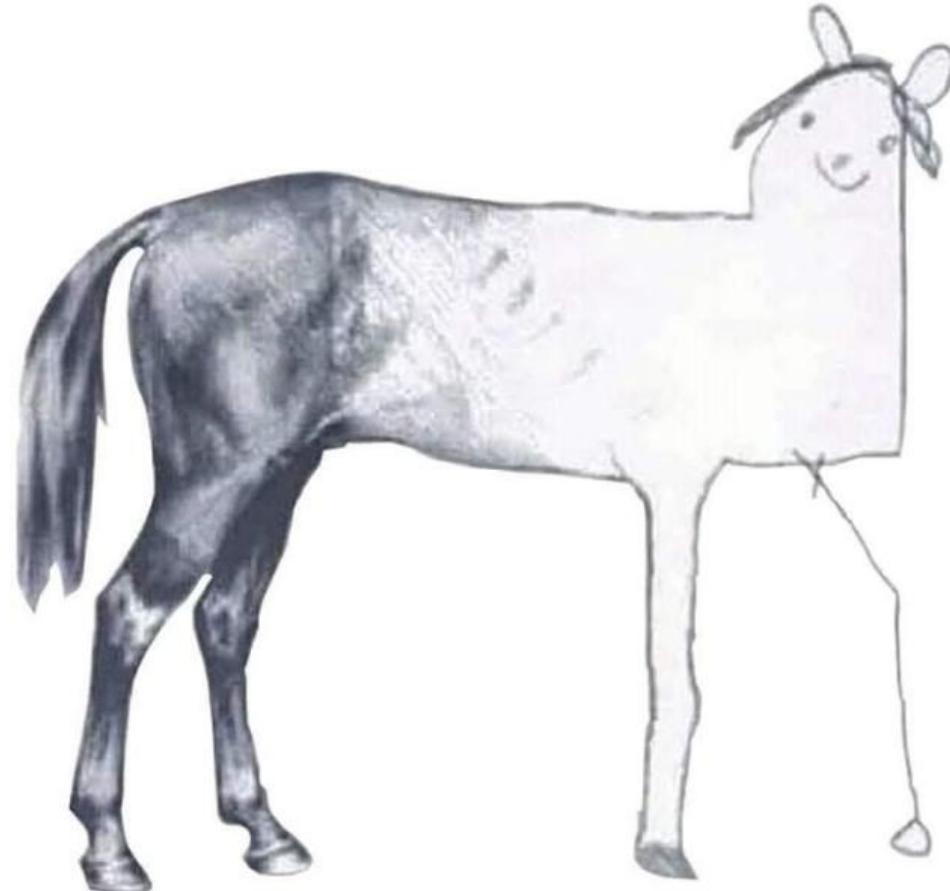


# Cost vs Time vs Quality

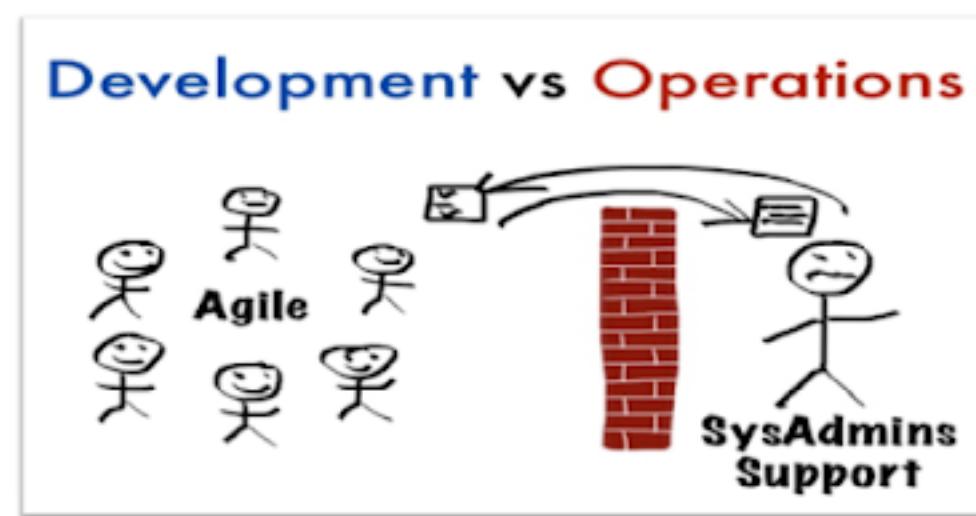
---







# Development vs Operations



# Development vs Operations



# 3 C's of Software Quality

- **Correctness:** perform its intended functions correctly and without errors or bugs
- **Completeness:** fulfills all its specified requirements. Ensures meets the needs of users
- **Consistency:** Ensures that the software is reliable and can be used in different situations.

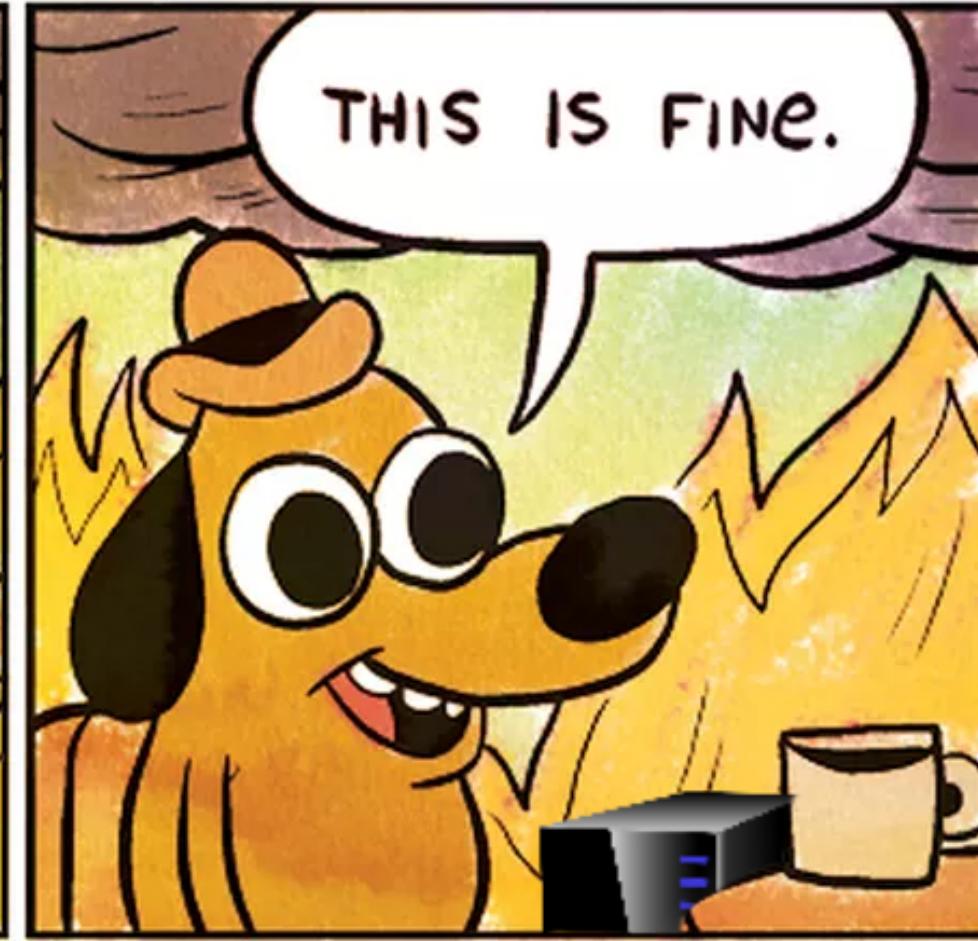


# Problems

---

- Cannot deliver on time
- Deliver software with low quality
- Low customer satisfaction
- Collaboration in team have problem
- Automate not common and/or cannot do repeat



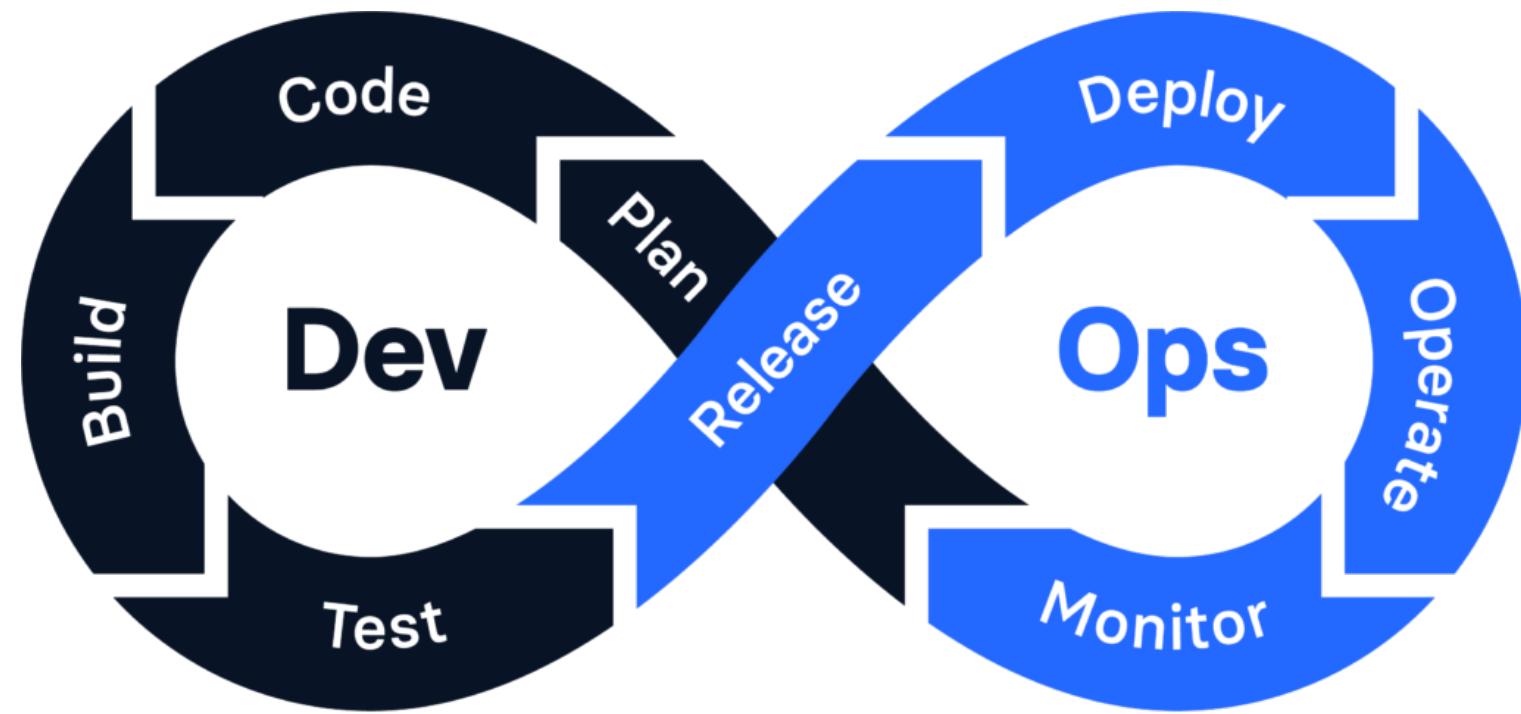




# What's DevOps ?

# DevOps

---



# What's DevOps

---

The Microsoft-approved definition of DevOps

- " **the union of people**, processes and technology to **continually provide value** to customers. "



# What's DevOps

---

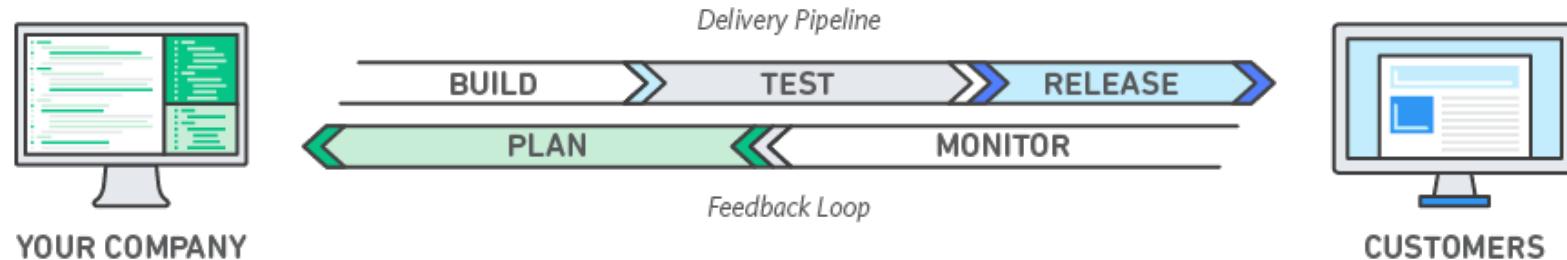
Patrick Debois (The Father of DevOps)

- " It's a movement of **people** who think it's changing in the IT Industry time to stop wasting money, time to start delivering **great software**, and building systems that **scale** and last **DEV** Integration **Ops** Communication Collaboration. "



# What's DevOps

- DevOps emphasizes **people** (and culture), and seeks to improve collaboration between operations and development teams. DevOps implementations utilize technology—especially **automation** tools for **more rapid** and **more quality**



# DevOps

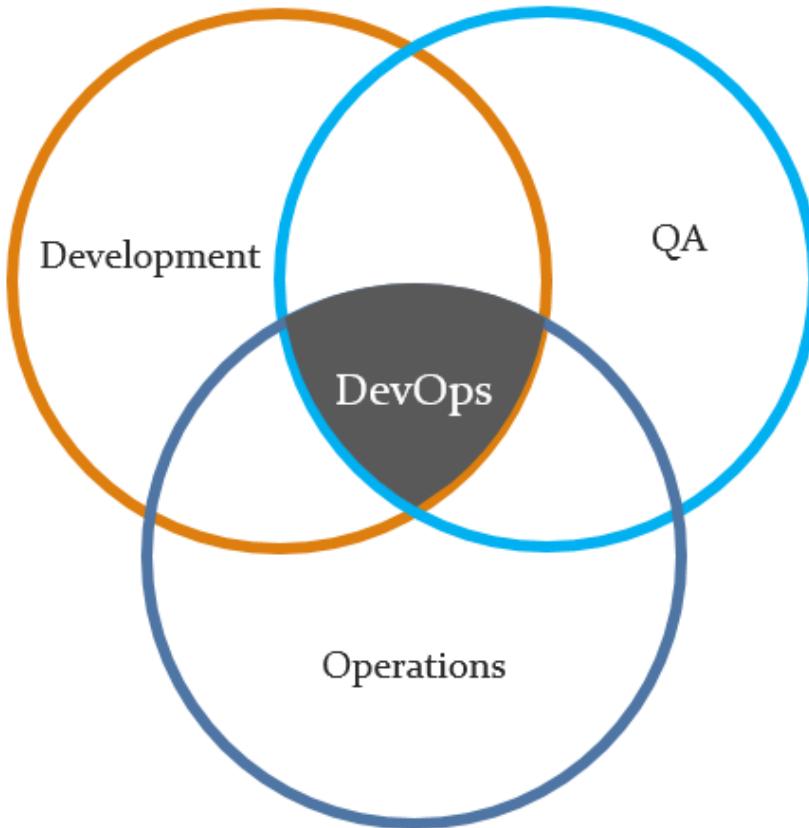
---

- It's not a product
- It's not a job title
- It's not all about tools

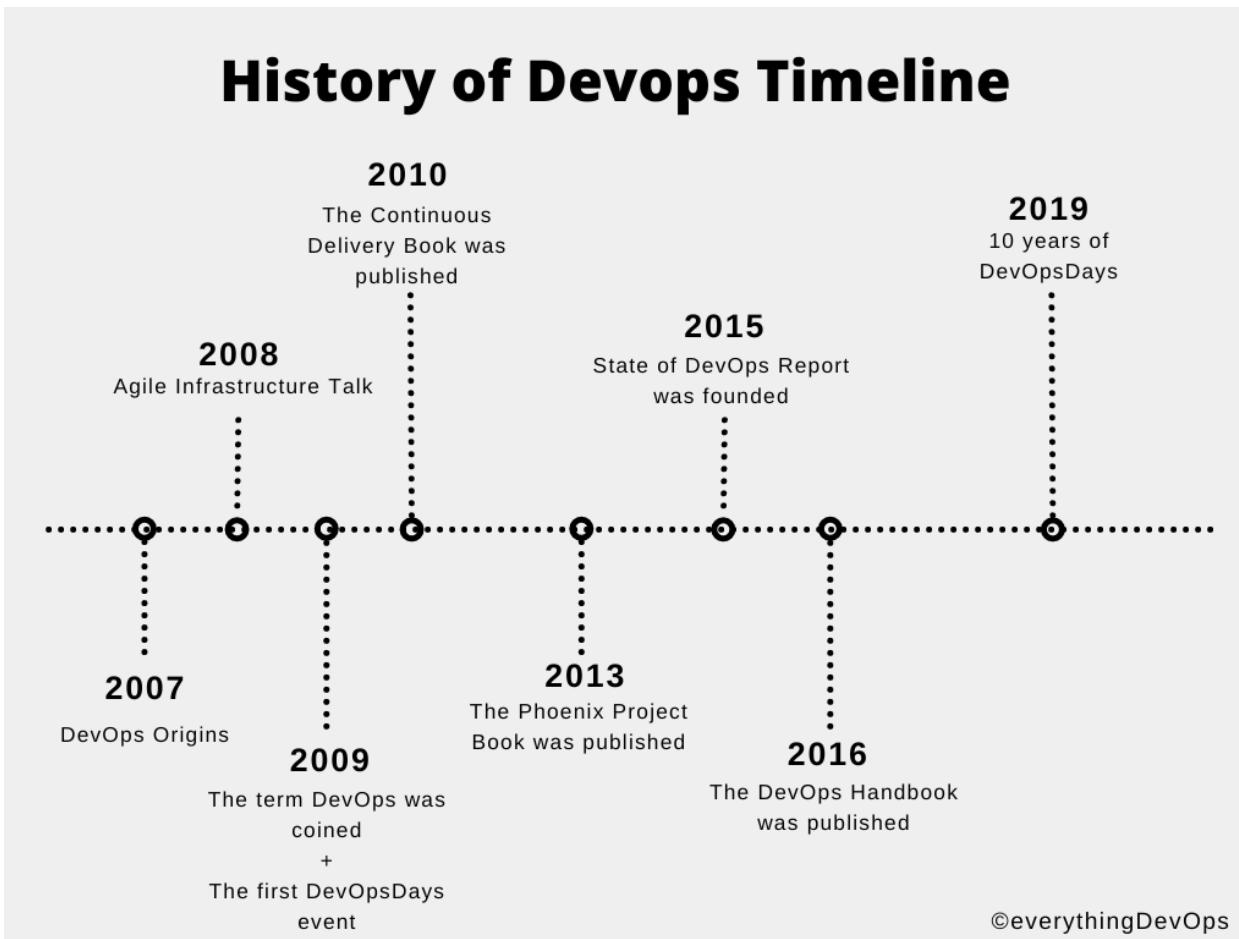


# DevOps

---



# History of DevOps



# Benefits

---

- Faster Time-to-Market
- Improved Collaboration
- Increased Quality, Reliability
- Iteration & Continuous Feedback
- Better Customer Satisfaction



# How to start DevOps ?

# CALMS framework

---



**KEEP  
CALMS  
AND  
DEVOPS**



# CALMS framework

---

- Culture
- Automation
- Lean
- Measurement
- Sharing



# Culture

---

- Open **communication** with inter and Internal team
- **Collaboration:** shared ownership and responsibility
- **Trust:** Teams and individuals
- **Empathy:** Understanding the challenges faced



# Automation

---

- Continuous Integration (CI)
- Continuous Delivery (CD)
- Infrastructure as Code (IaC)



# Lean

---

- Small Batches
- Continuous Improvement
- Eliminating repeated tasks



# Measurement

---

- Lead time for changes
- Deployment frequencies
- Meantime to restore
- Failure rate



# DORA Metrics

Software delivery performance metric	Elite	High	Medium	Low
<input checked="" type="checkbox"/> <b>Deployment frequency</b> For the primary application or service you work on, how often does your organization deploy code to production or release it to end users?	On-demand (multiple deploys per day)	Between once per week and once per month	Between once per month and once every 6 months	Fewer than once per six months
<input checked="" type="checkbox"/> <b>Lead time for changes</b> For the primary application or service you work on, what is your lead time for changes (i.e., how long does it take to go from code committed to code successfully running in production)?	Less than one hour	Between one day and one week	Between one month and six months	More than six months
<input checked="" type="checkbox"/> <b>Time to restore service</b> For the primary application or service you work on, how long does it generally take to restore service when a service incident or a defect that impacts users occurs (e.g., unplanned outage or service impairment)?	Less than one hour	Less than one day	Between one day and one week	More than six months
<input checked="" type="checkbox"/> <b>Change failure rate</b> For the primary application or service you work on, what percentage of changes to production or released to users result in degraded service (e.g., lead to service impairment or service outage) and subsequently require remediation (e.g., require a hotfix, rollback, fix forward, patch)?	0%-15%	16%-30%	16%-30%	16%-30%



# Sharing

---

- Knowledge Exchange
- Documentation



# CI/CD

---



# What is a CI/CD pipelines ?

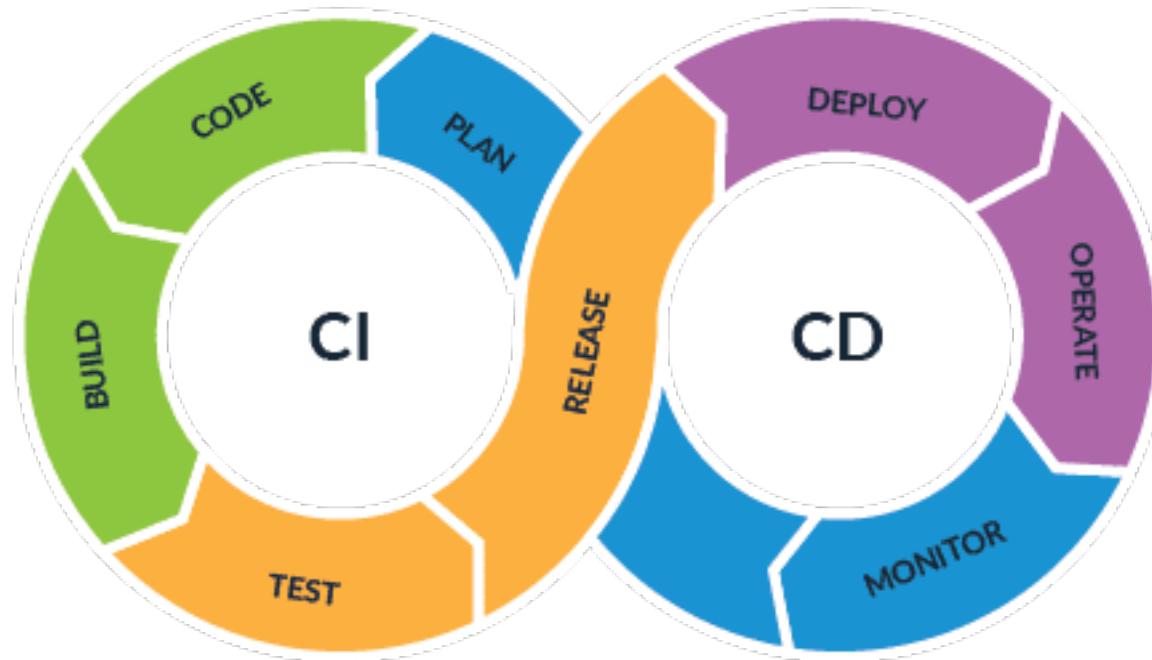
---

- Continuous Integration and Continuous Deployment (CI/CD) pipeline represents an **automatic workflow** that continuously integrates code developed by software developers and **deploys** them into a target environment with less human intervention.



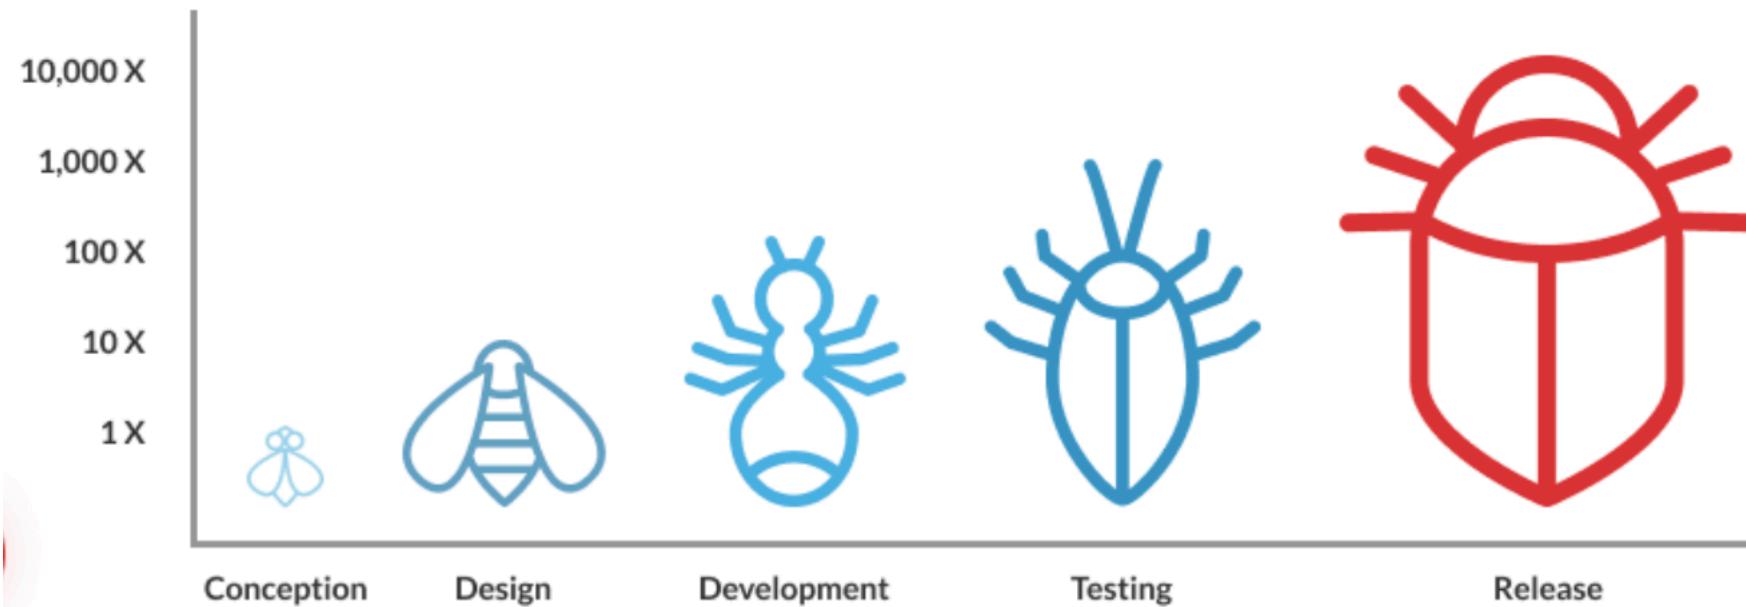
# What is a CI/CD pipelines ?

---



# Cost of integration

Resolving bugs early and often reduces associated costs

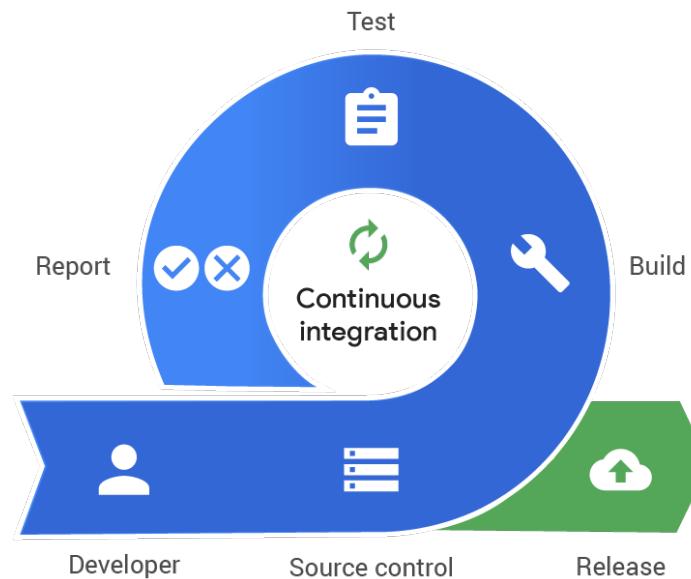


# Continuous Integration (CI)

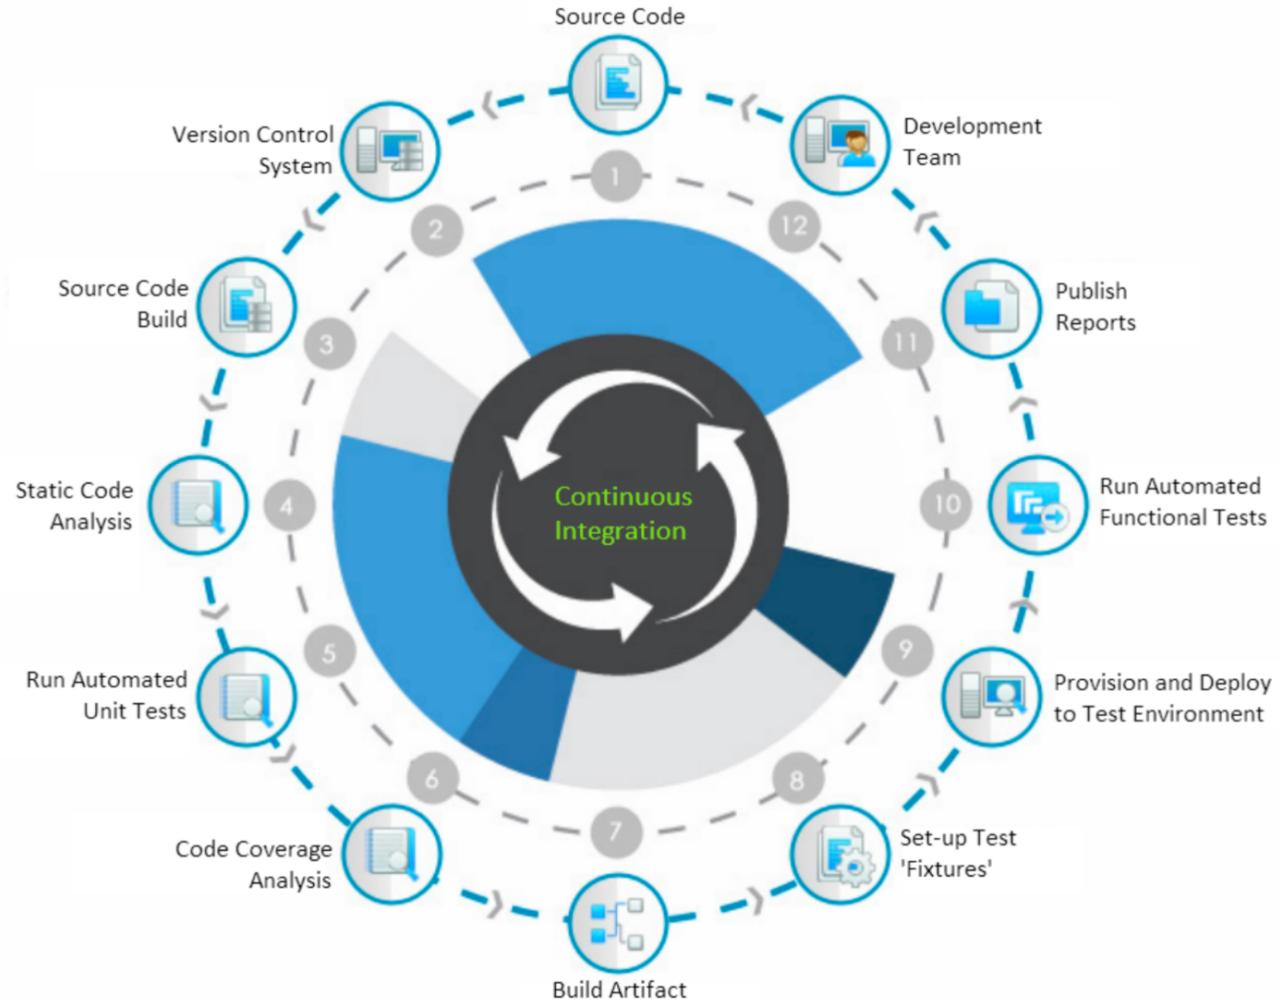
# Continuous Integration (CI)

CI is part of the development process. **Automatically** build the software whenever new software is **committed** by developers

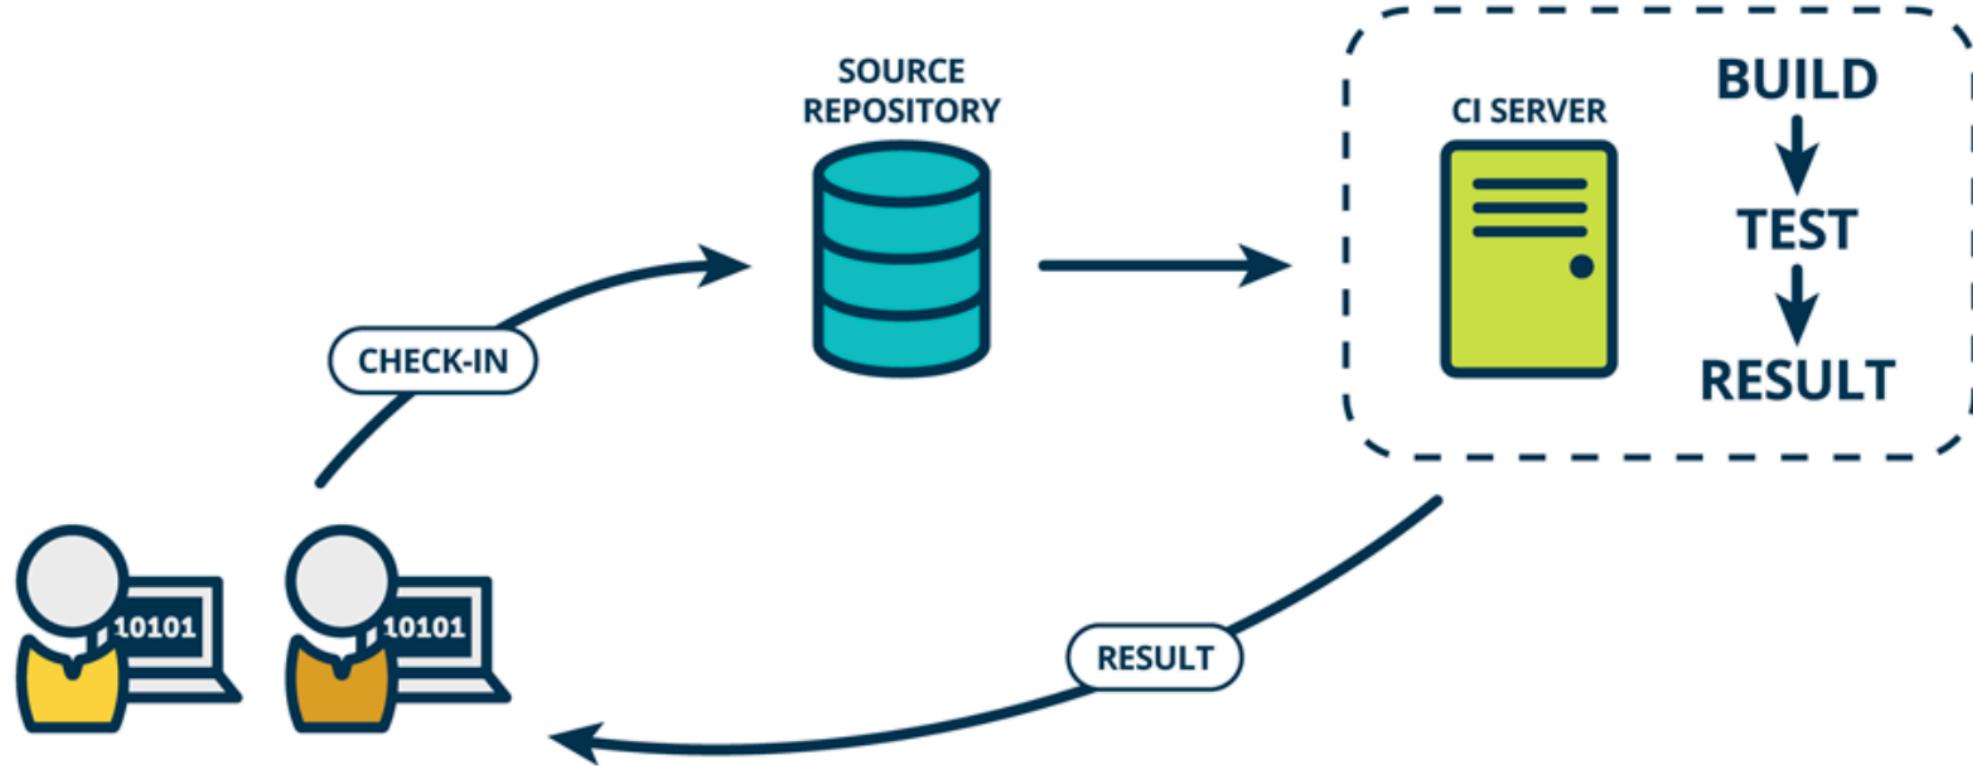
- Focuses on **frequently** merging code changes into a central repository
- Each code change, **trigger automated** builds and tests
- **Early** detection
- This ensures that the code compiles correctly and **passes all** the essential **tests**



# Continuous Integration (CI)

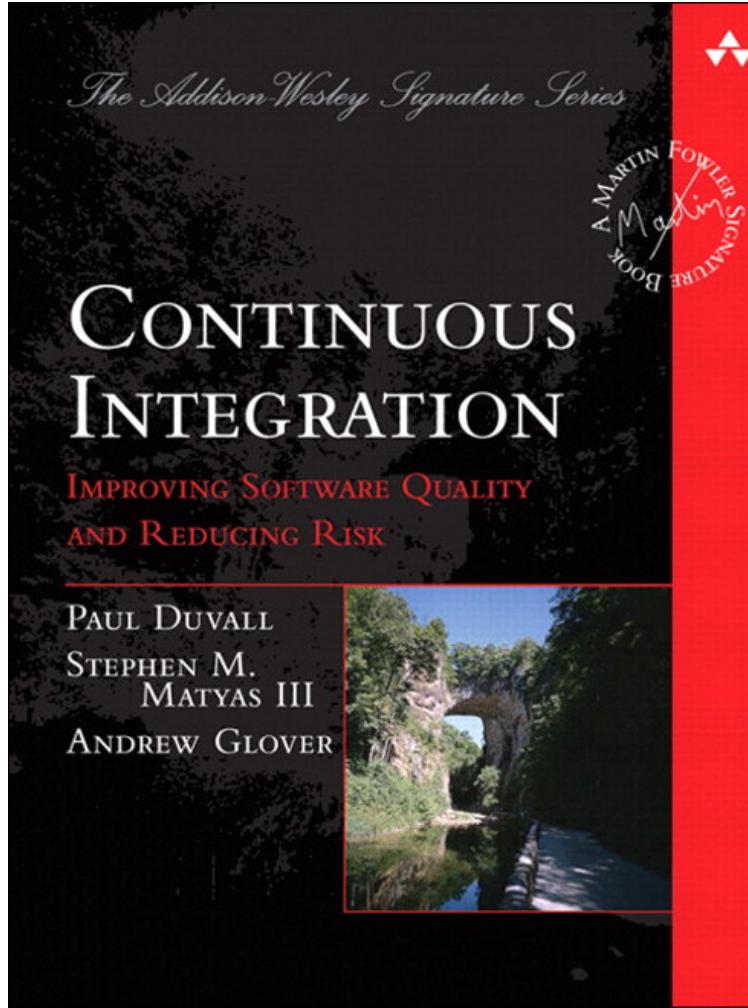


# Continuous Integration (CI)



# CI's practices

---



# CI's practices

## 1. Maintain a **single** source repository

The key take away here is to have all artifacts needed to build the project in one repository. Not require additional dependencies or manual steps

### **Caution**

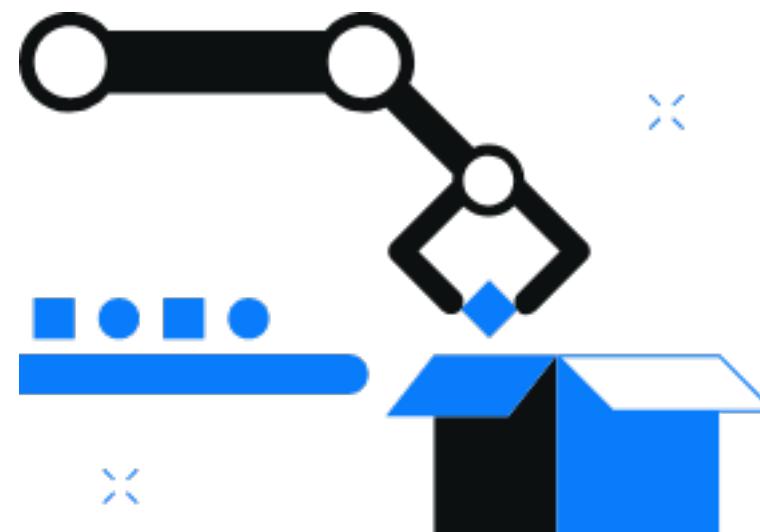
Recommends that **only a single branch** of the project, **NOT** to use **branching**. However depend on your organize with team



# CI's practices

## 2. Automate the build

Automation of the build should include steps such as compiling the code, executing unit tests and integration tests



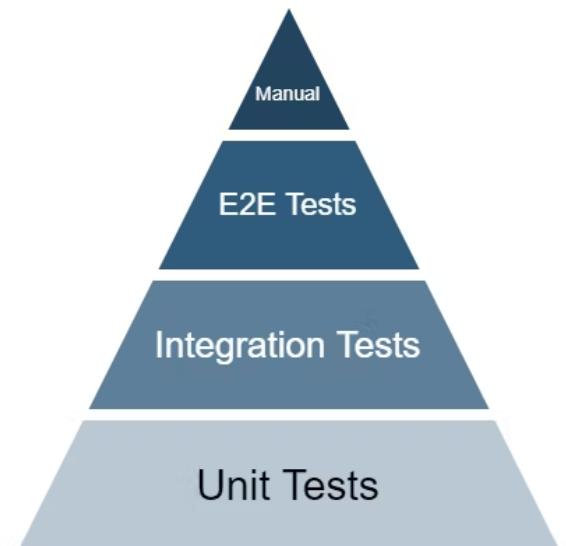
# CI's practices

## 3. Make the build self-testing

All tests should run to confirm that it behaves as the developers expect it to behave, and always test on your machine first

### **Caution**

Unit tests should test behavior, not implementation details



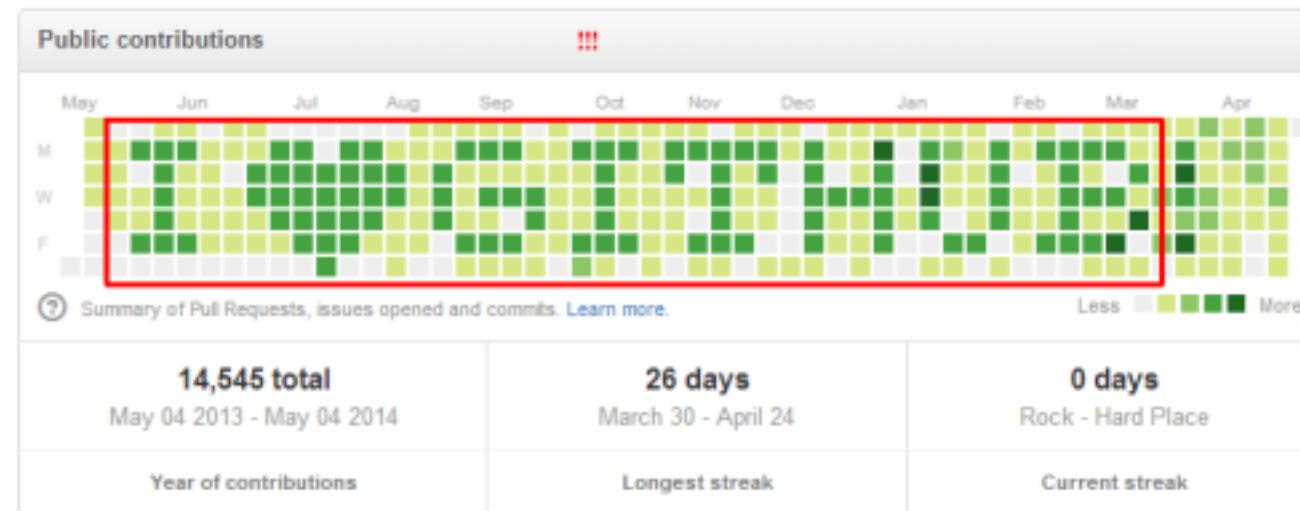
# CI's practices

4. Everyone commits to the baseline every day, (frequency and small)

Every committer can reduce the number of conflicting changes.

## Caution

Do not commits source code that do not work



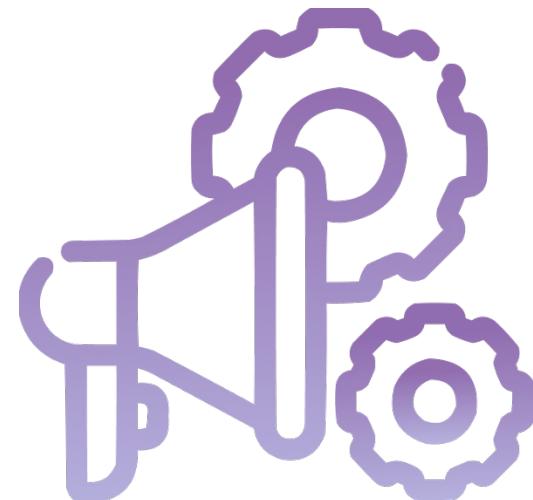
# CI's practices

## 5. Every commit (to baseline) should trigger a build

You should separate the CI workflows. This includes the steps from compilation right up to packaging and testing.

### **Caution**

There should be no commented out tests in the mainline branch



# CI's practices

## 6. Fix Broken Builds Immediately

Nobody has a higher priority task than fixing the build



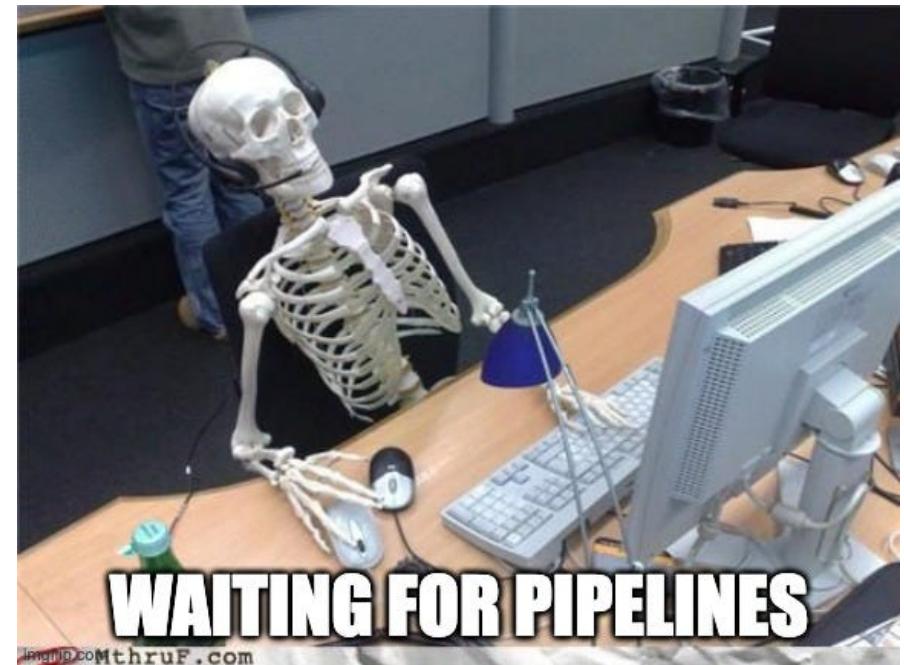
# CI's practices

## 7. Keep the build fast

You should aim to have faster-executing tests than slow tests. This would mean that you need to have more unit tests than other types of tests

### **Caution**

Balancing your workflows with build fast



**WAITING FOR PIPELINES**

imgur.com/Mthruf.com



# CI's practices

## 8. Test in a clone of the production environment

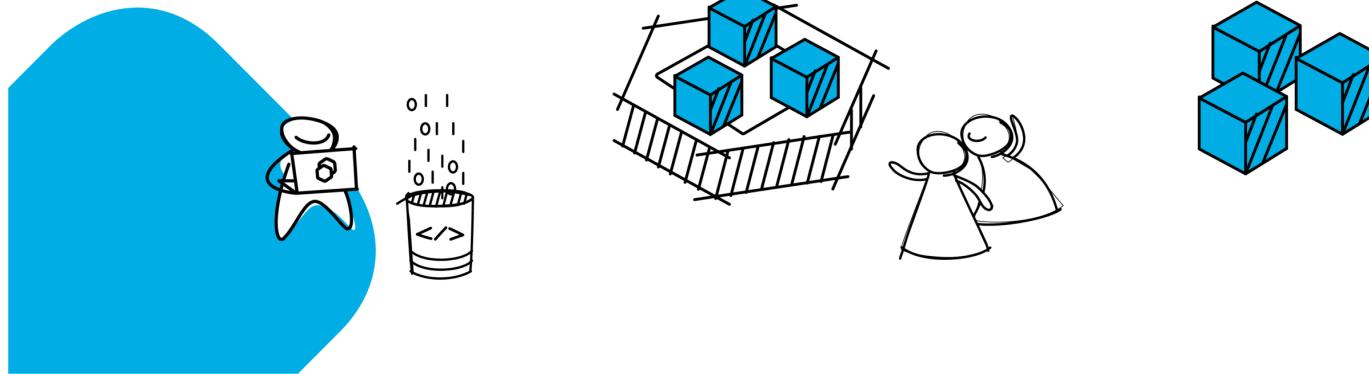
This is the hardest principle to put into practice in real world development. This needs the build automation system to create and deploy the packages into a staging environment that reflects the real production environment



# CI's practices

## 9. Make it easy to get the latest deliverables (artifacts)

Make sure well known place where people can find



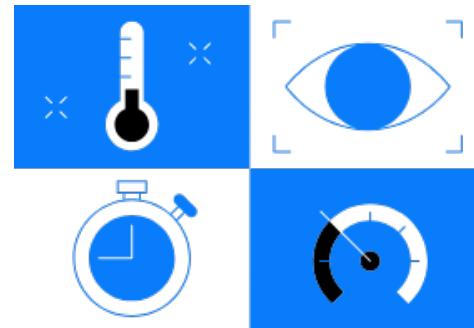
# CI's practices

## 10. Everyone can see the results of the latest build

It should be easy to find out whether the build breaks and, if so, who made the relevant change. It is recommended that the emails are sent to the whole team when the build fails so that it can be fixed as soon as possible.

### **Caution**

CI is not only for developers. The various metrics that can be derived using CI can help improve the quality of software, not just the quality of the software.



# CI's practices

## 11. Automate deployment

Most CI systems allow the running of scripts after a build finishes. In most situations, it is possible to write a script to deploy the application to a live test server that everyone can look at



# CI's benefits

---

- Early Bug Detection
- Improved Code Quality
- Reduced Integration Issues
- Automated Build and Testing
- Improve Developer Productivity
- Consistency in Environments





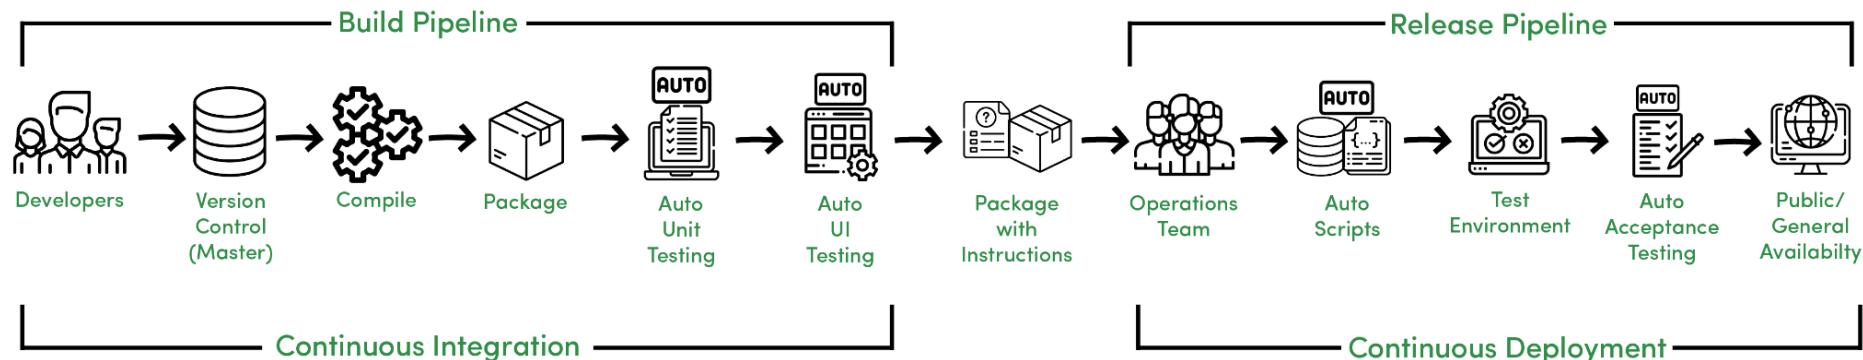
# Workshop Continuous Integration

# Continuous Deployment (CD)

# Continuous Deployment (CD)

Continuous deployment stage refers to pulling the artifact from the repository and deploying it frequently and safely. Automate the deployment into the each environments with less human intervention

- Once the code passes CI checks, CD takes over
- Some approaches call Continuous Delivery: Makes the built code readily available for deployment, allowing for manual review before pushing it live



# Continuous Delivery vs Continuous Deployment

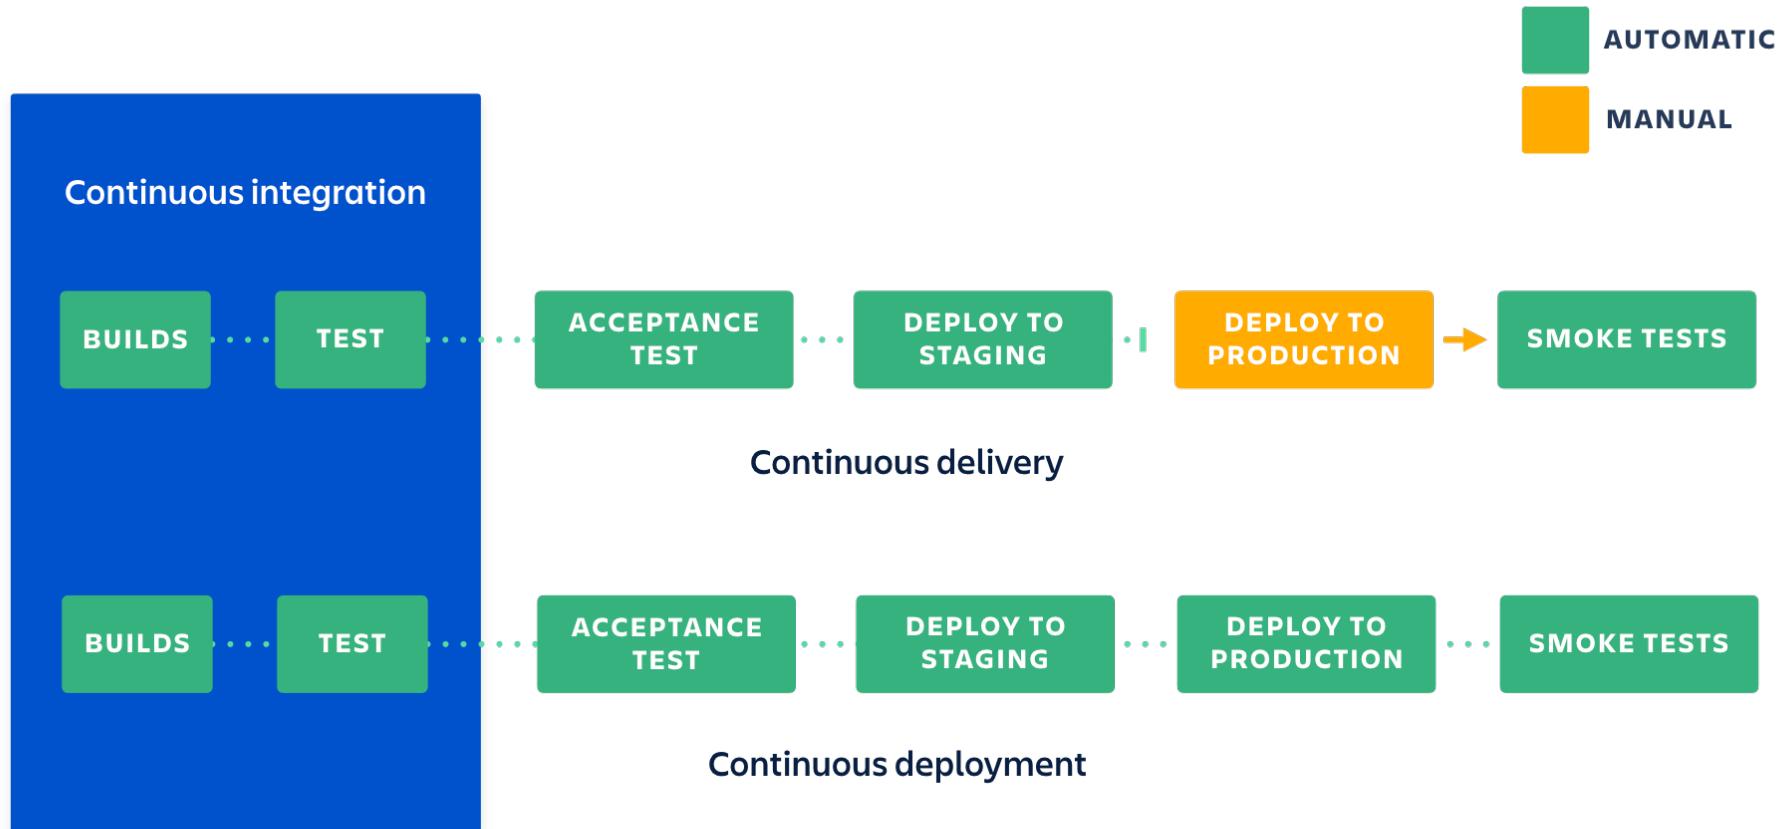
---

Both is an extension of continuous integration where code changes are automatically prepared for a release to production

But continuous delivery is the presence of a **manual approval** to update to production. However continuous deployment, production happens **automatically** without explicit approval

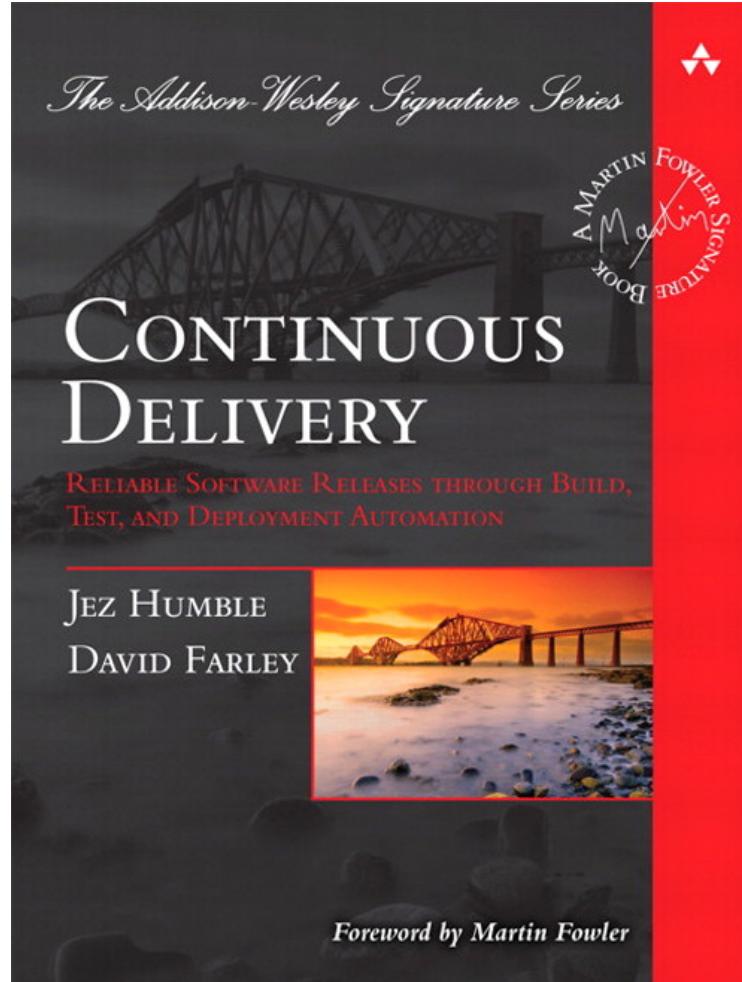


# Continuous Delivery vs Continuous Deployment



# CD's practices

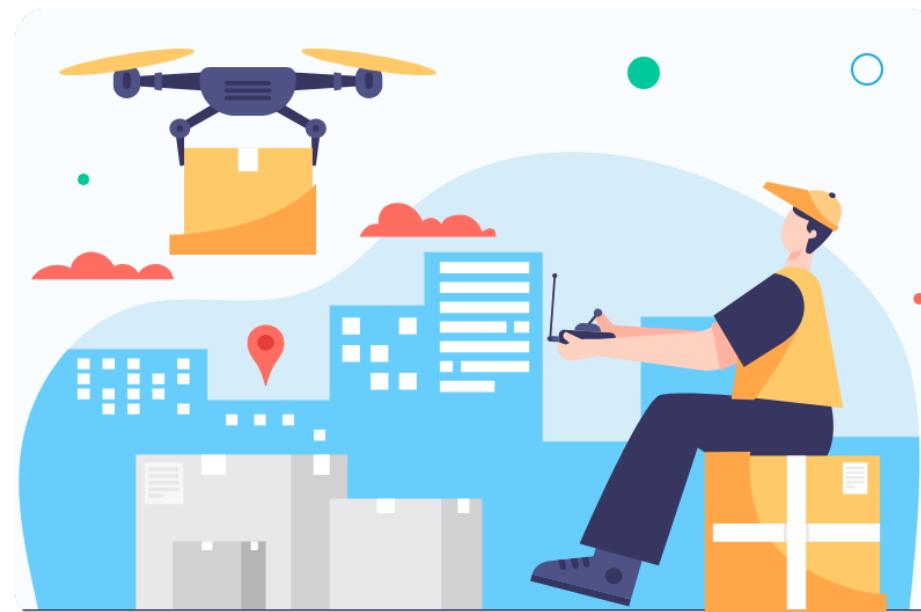
---



# CD's practices

## 1. Automated Deployments

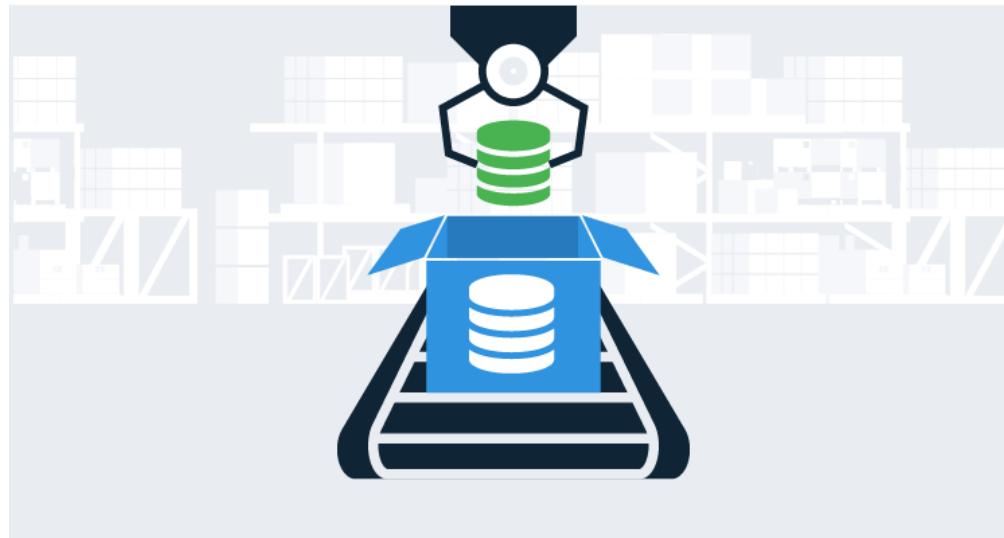
They emphasize that the deployment pipeline should be fully automated to ensure reliability and repeatability.



# CD's practices

## 2. Version Control for All Artifacts

The importance of versioning all components, including infrastructure and configuration, as part of continuous delivery.



# CD's practices

## 3. Deploy to Production-Like Environments Early

Need to test in production-like environments to reduce surprises in the actual production release.



# CD's practices

## 4. Zero-Downtime Deployments

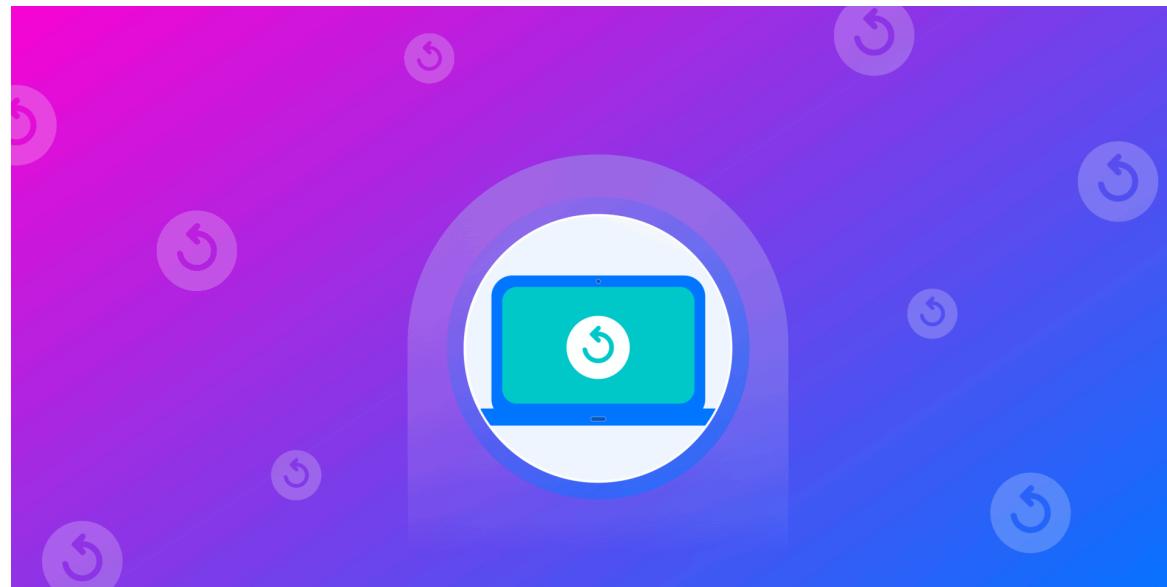
Techniques like blue-green deployment and canary releases as methods to avoid downtime and reduce risk during deployments



# CD's practices

## 5. Rollback Mechanism

The need for a robust rollback mechanism to handle any issues arising in production



# CD's benefits

---

- Automate the Software Release Process
- Reduced Manual Intervention, reduce human error
- Improve Developer Productivity
- Find and Address Bugs Quicker
- Deliver Updates Faster
- Easier Rollbacks



# Deployment strategies

# Deployment strategies

---

- Recreate deployment
- Ramped (Rolling deployment)
- Blue/green deployment
- Canary deployment
- A/B testing
- Shadow deployment



# Recreate deployment

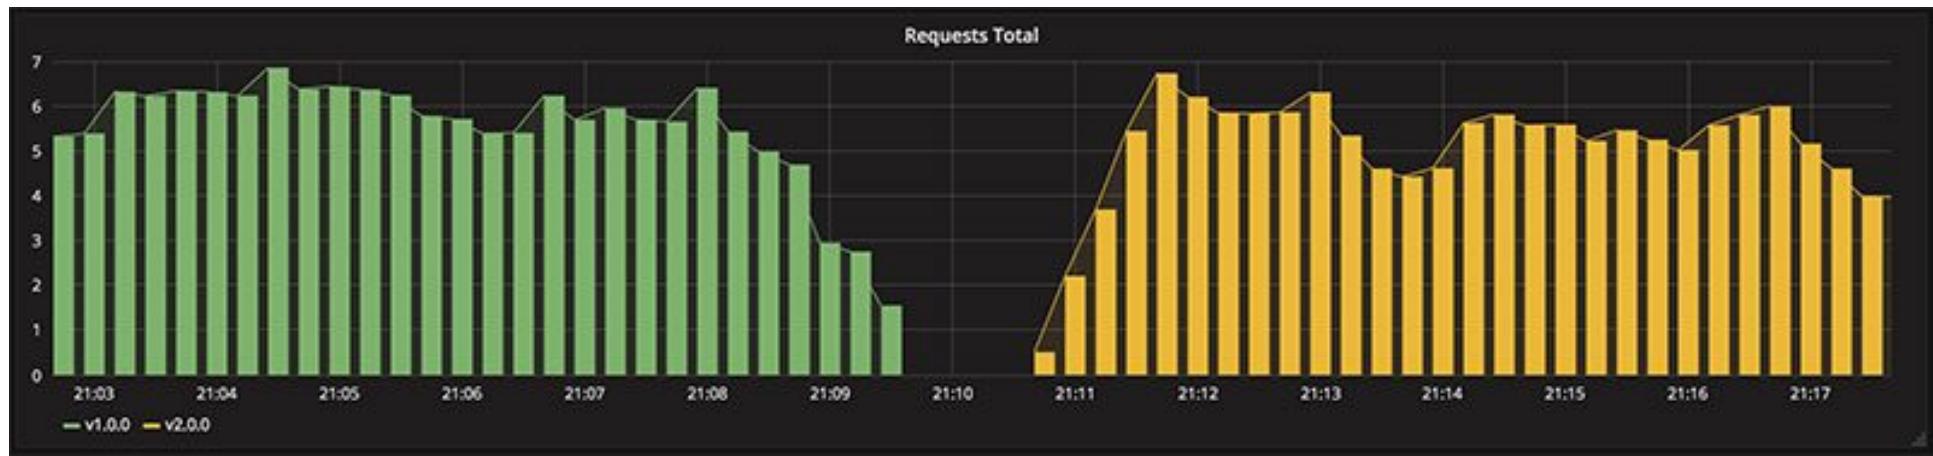
Terminating all the running instances then recreate them with the newer version

## Advantage

- Application state entirely renewed

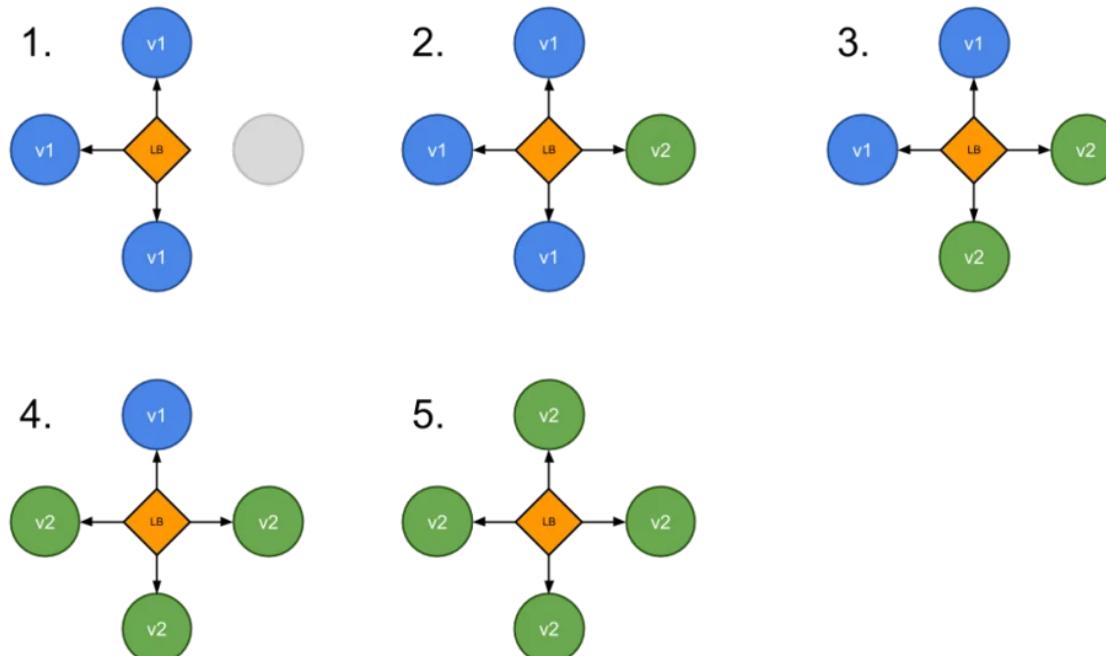
## Disadvantage

- Downtime that depends on both shutdown and boot duration of the application



# Ramped (Rolling deployment)

The rolling update is the **default** deployment strategy in Kubernetes that update the pods one by one without causing cluster downtime. the deployment objects create new pods using the new pod specifications and shut down the older ones when the application starts running successfully



# Ramped (Rolling deployment)

## Advantage

- Enables zero-downtime deployments
- Well-suited for large deployments and stateful applications

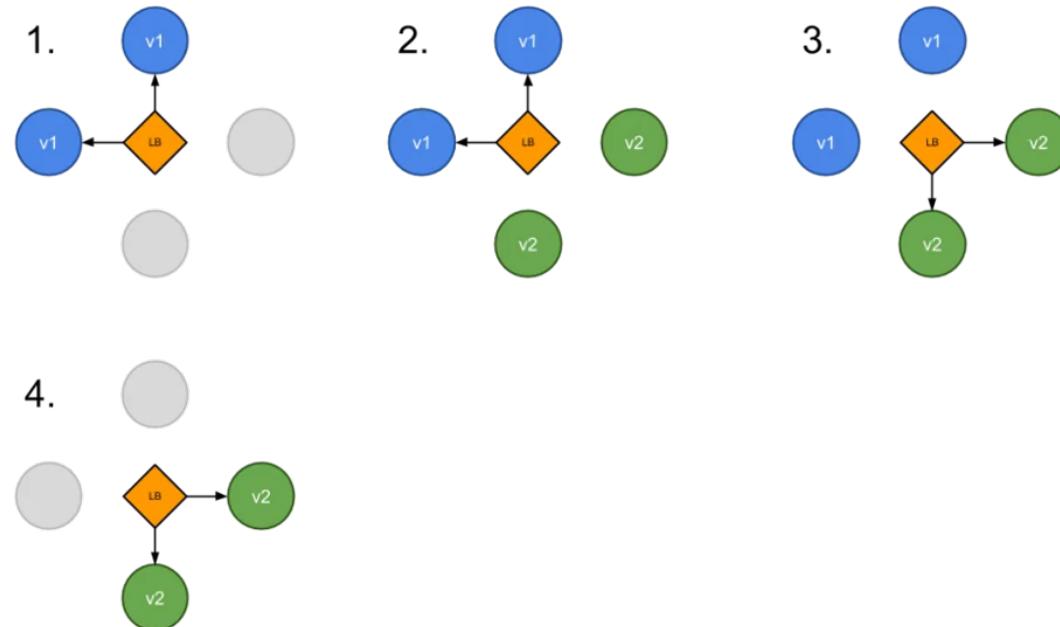
## Disadvantage

- Rollout/rollback can take time
- Supporting multiple APIs is hard
- No control over traffic



# Blue/green deployment

- While the “Blue” version is running, engineers deploy the “green” version in another deployment and run basic tests to ensure that the pods within the deployments are stable, then using services, you switch to the newer version of the application
- Best to avoid API versioning issues



# Blue/green deployment

## Advantage

- Instant rollout/rollback
- Clear separation between old and new versions

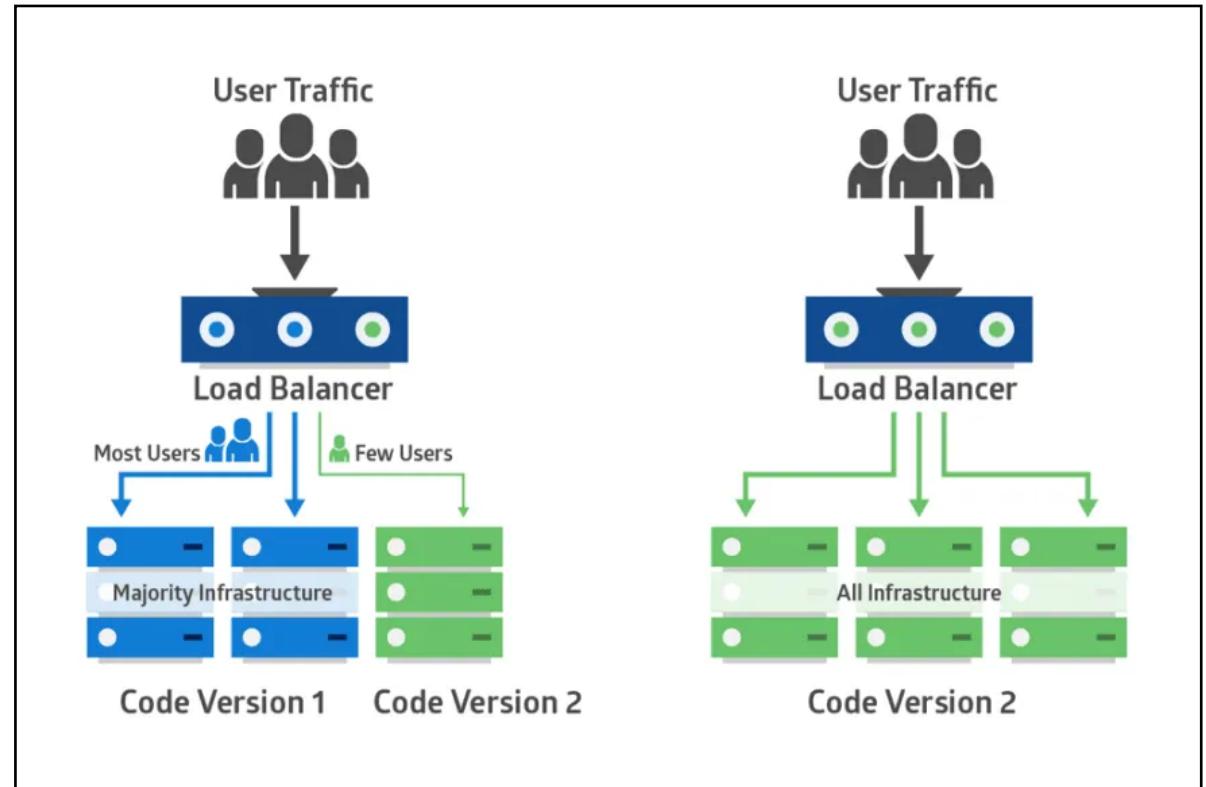
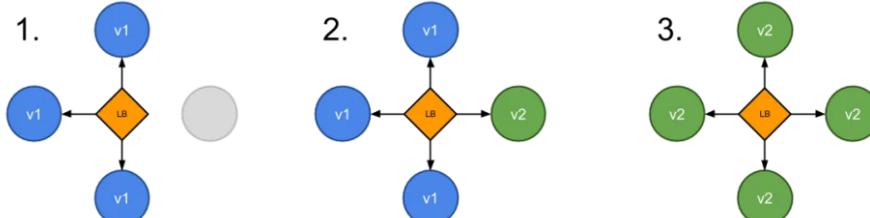
## Disadvantage

- Requires double the resources
- More complex configuration



# Canary deployment

- Similar to blue/green deployments, but only a small subset of traffic is directed to the new version (canary) initially. This allows for early detection of issues before a full rollout



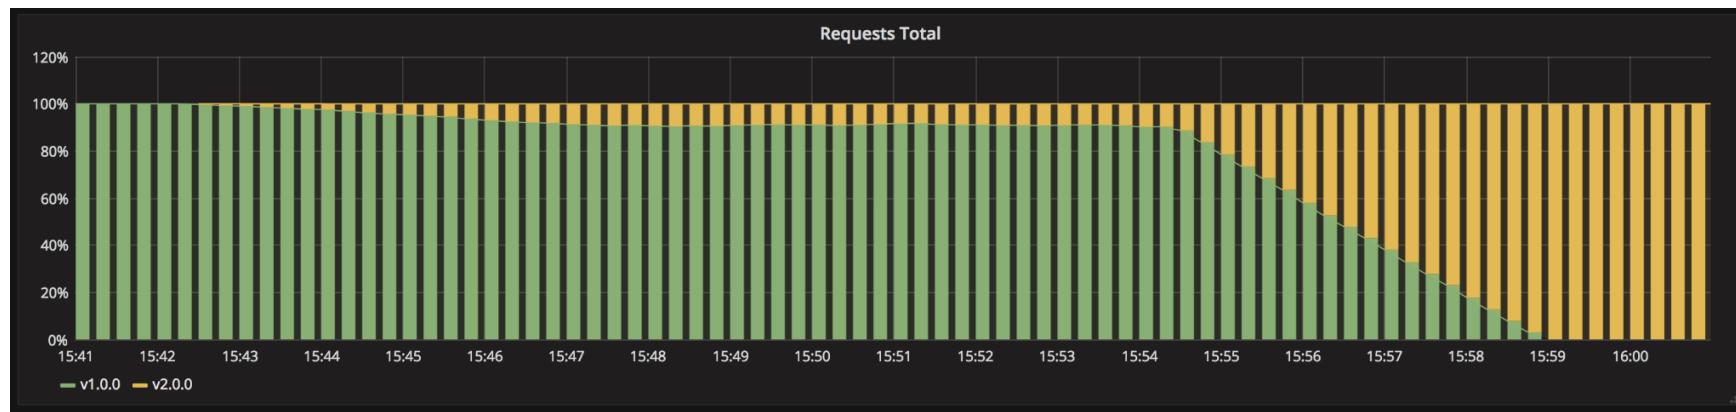
# Canary deployment

## Advantage

- Lowers risk by gradually rolling out new version to a small subset of users
- Enables early detection of issues

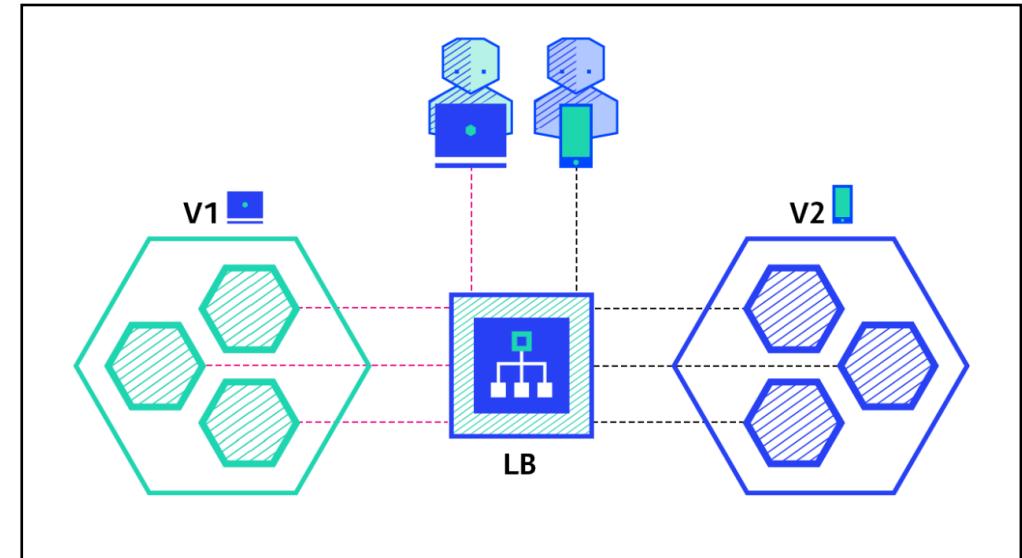
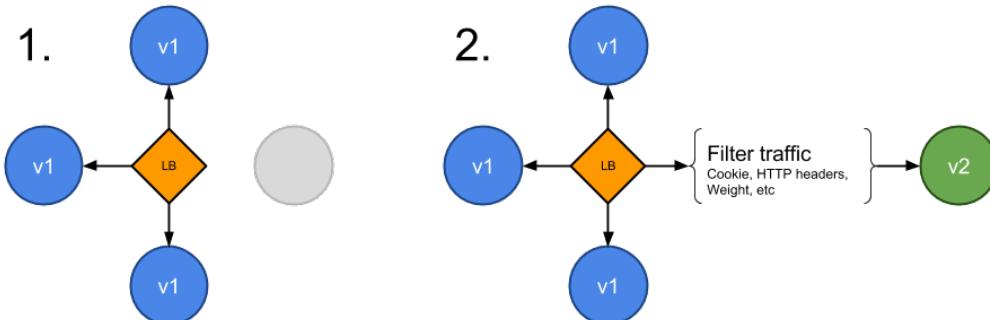
## Disadvantage

- Requires additional configuration for traffic routing
- Slow rollout, most of users cannot use new versions
- Application must support both of old & new versions



# A/B testing

- A/B testing deployments consists of routing a subset of users to a new functionality under specific conditions. It is usually a technique for making business decisions based on statistics, rather than a deployment strategy.
- Best for feature testing on a subset of users



# A/B testing

## Advantage

- Several versions run in parallel
- Full control over the traffic distribution

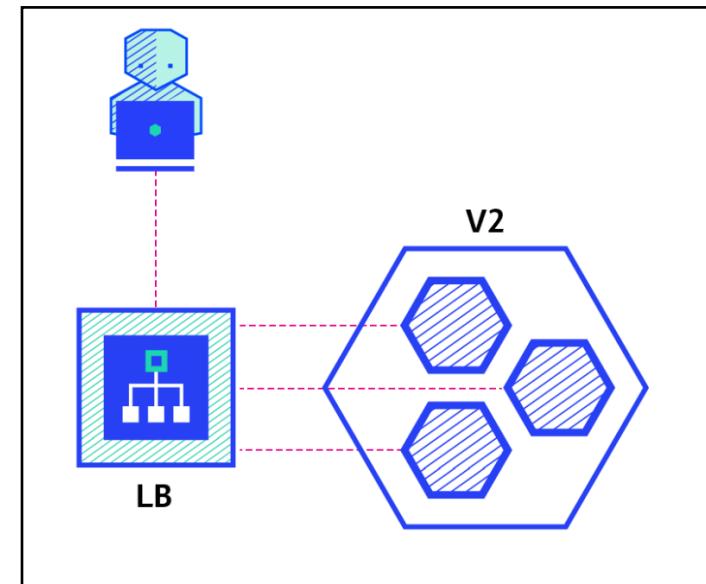
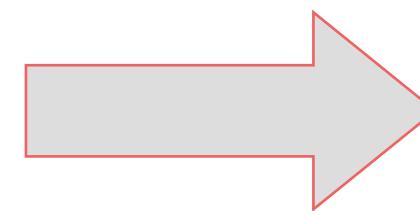
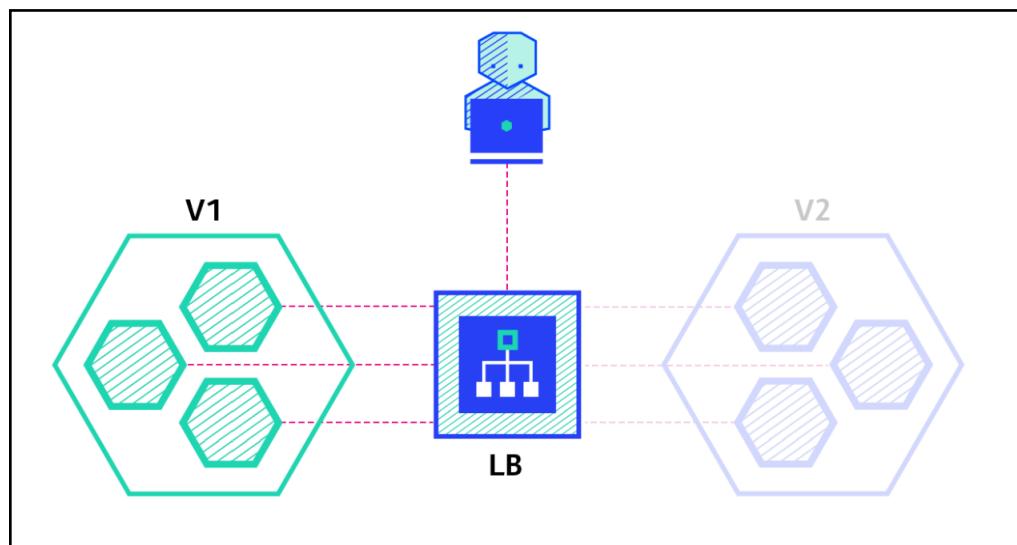
## Disadvantage

- Requires intelligent load balancer
- Not straightforward, you need to setup additional tools
- Hard to troubleshoot errors for a given session



# Shadow deployment

- Consists of releasing version B alongside version A, fork version A's incoming requests and send them to version B as well without impacting production traffic. This is particularly useful to test production load on a new feature



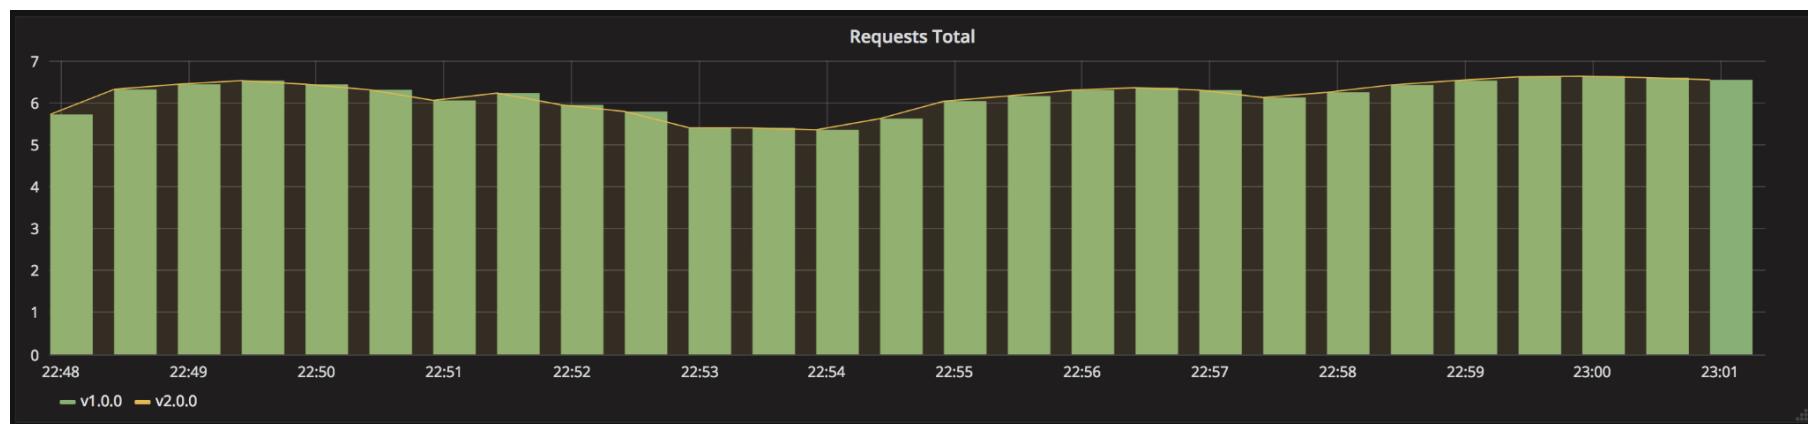
# Shadow deployment

## Advantage

- Performance testing, auto scale of the application with production traffic
- No impact on the user

## Disadvantage

- Expensive as it requires double the resources.
- Complex to setup



# DEPLOYMENT STRATEGIES

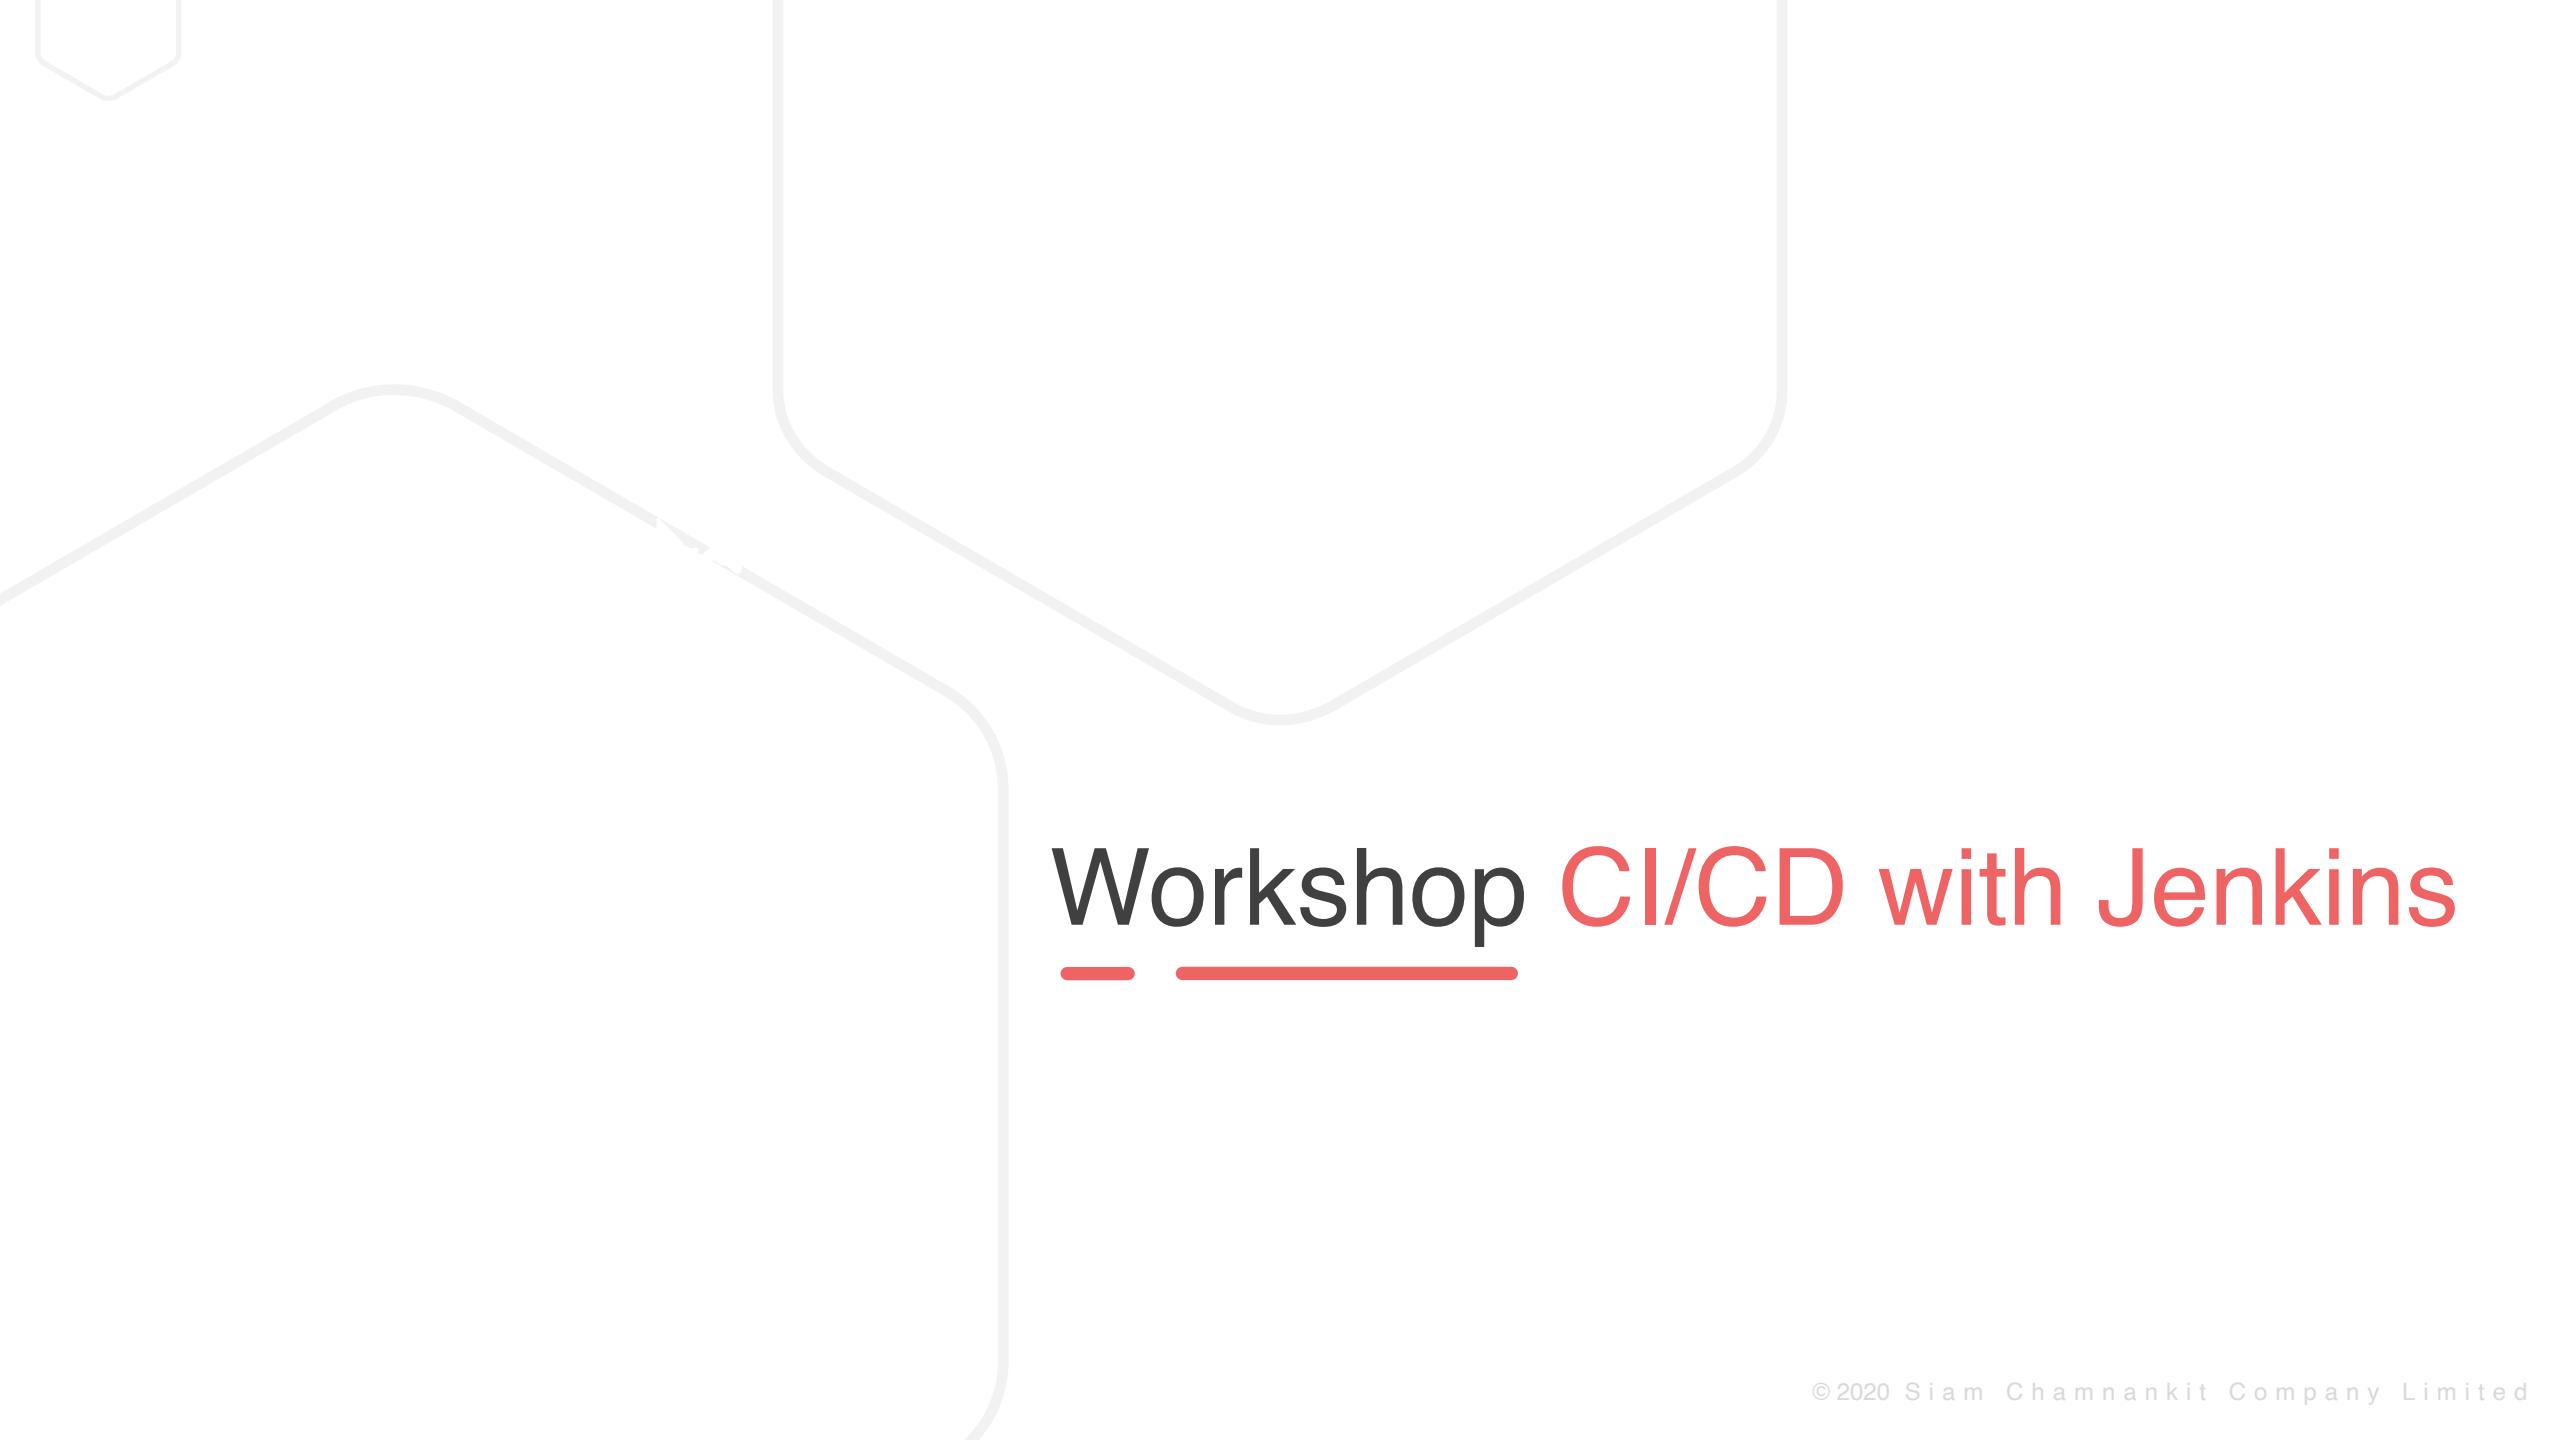
When it comes to production, a ramped or blue/green deployment is usually a good fit, but proper testing of the new platform is necessary.

Blue/green and shadow strategies have more impact on the budget as it requires double resource capacity. If the application lacks in tests or if there is little confidence about the impact/stability of the software, then a canary, a/b testing or shadow release can be used.

If your business requires testing of a new feature amongst a specific pool of users that can be filtered depending on some parameters like geolocation, language, operating system or browser features, then you may want to use the a/b testing technique.



Strategy	ZERO DOWNTIME	REAL TRAFFIC TESTING	TARGETED USERS	CLOUD COST	ROLLBACK DURATION	NEGATIVE IMPACT ON USER	COMPLEXITY OF SETUP
<b>RECREATE</b> version A is terminated then version B is rolled out	✗	✗	✗	■□□	■■■	■■■	□□□
<b>RAMPED</b> version B is slowly rolled out and replacing version A	✓	✗	✗	■□□	■■■	■□□	■□□
<b>BLUE/GREEN</b> version B is released alongside version A, then the traffic is switched to version B	✓	✗	✗	■■■	□□□	■■□	■■□
<b>CANARY</b> version B is released to a subset of users, then proceed to a full rollout	✓	✓	✗	■□□	■□□	■□□	■■□
<b>A/B TESTING</b> version B is released to a subset of users under specific condition	✓	✓	✓	■□□	■□□	■□□	■■■
<b>SHADOW</b> version B receives real world traffic alongside version A and doesn't impact the response	✓	✓	✗	■■■	□□□	□□□	■■■



# Workshop CI/CD with Jenkins

# Challenges of implementing CI/CD

- Require organization changes and mindset shifts
- Deep technical knowledge to automate the entire process
- Team is required to use various technologies depend on application stack
- Time Investment: setting up an effective CI/CD pipeline can require significant initial time investment
- Frequent Failures: When adopting CI/CD, teams may initially experience frequent build and test failures due to process changes

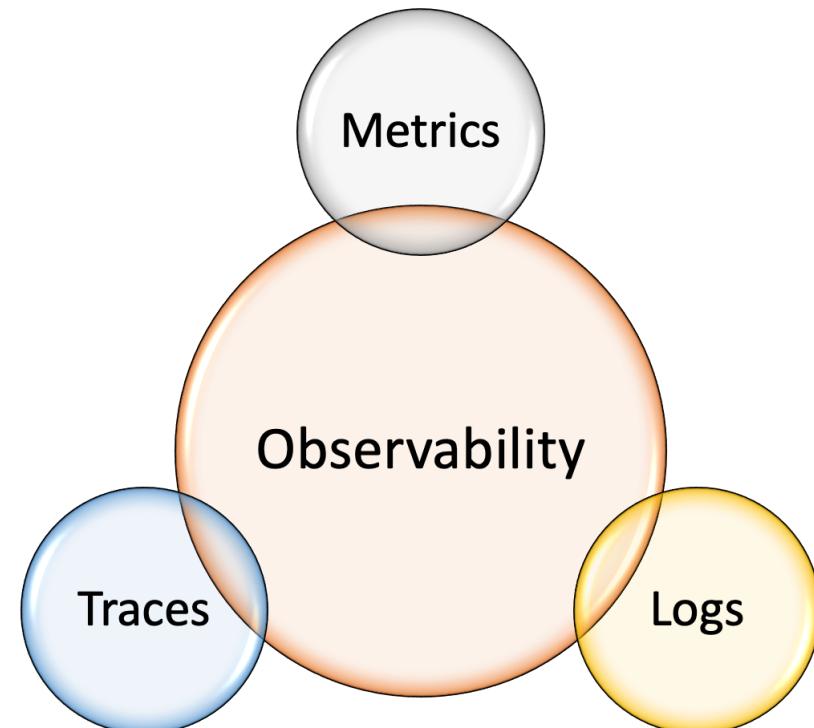


# Observability

# Three pillars of **Observability**

---

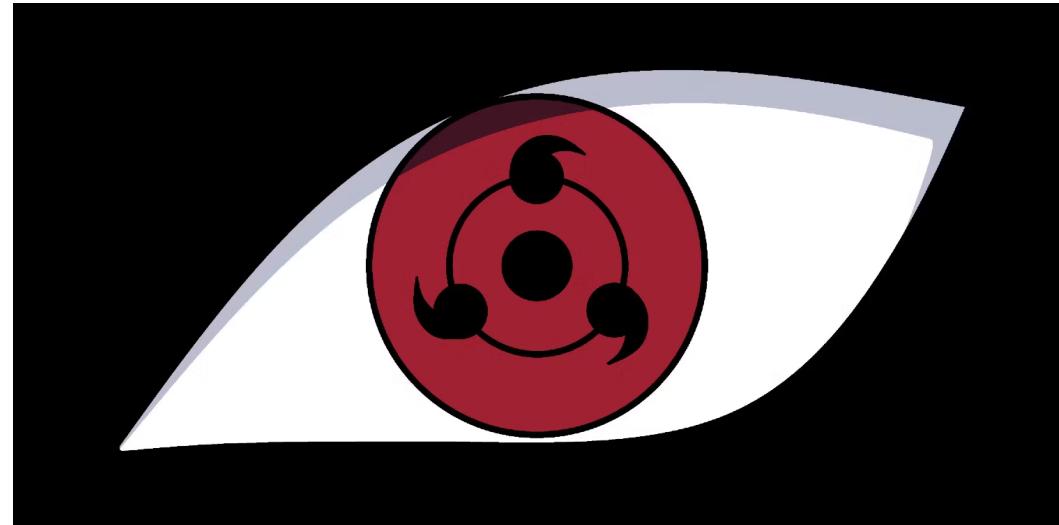
- Metrics
- Traces
- Logs



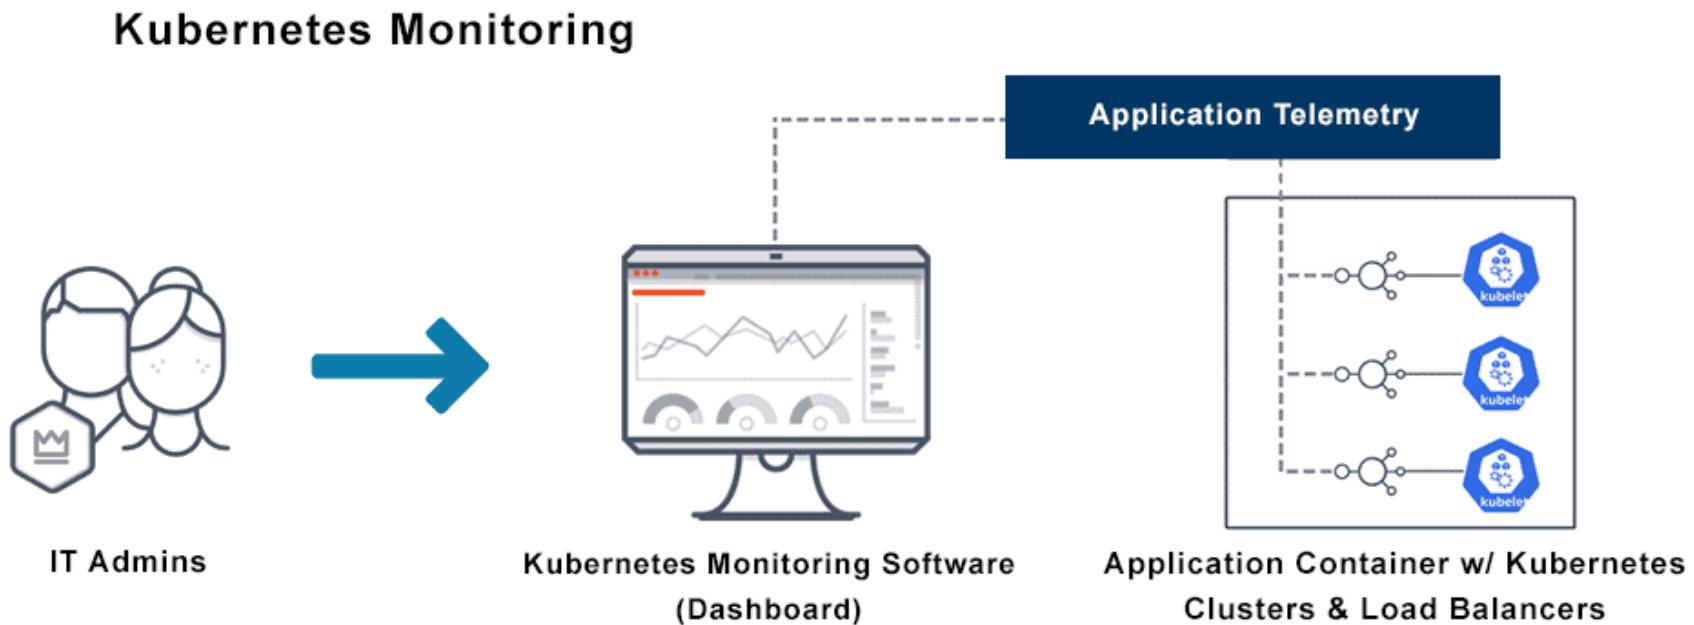
# Three pillars of Observability

---

- Metrics
- Traces
- Logs



# Metrics monitoring



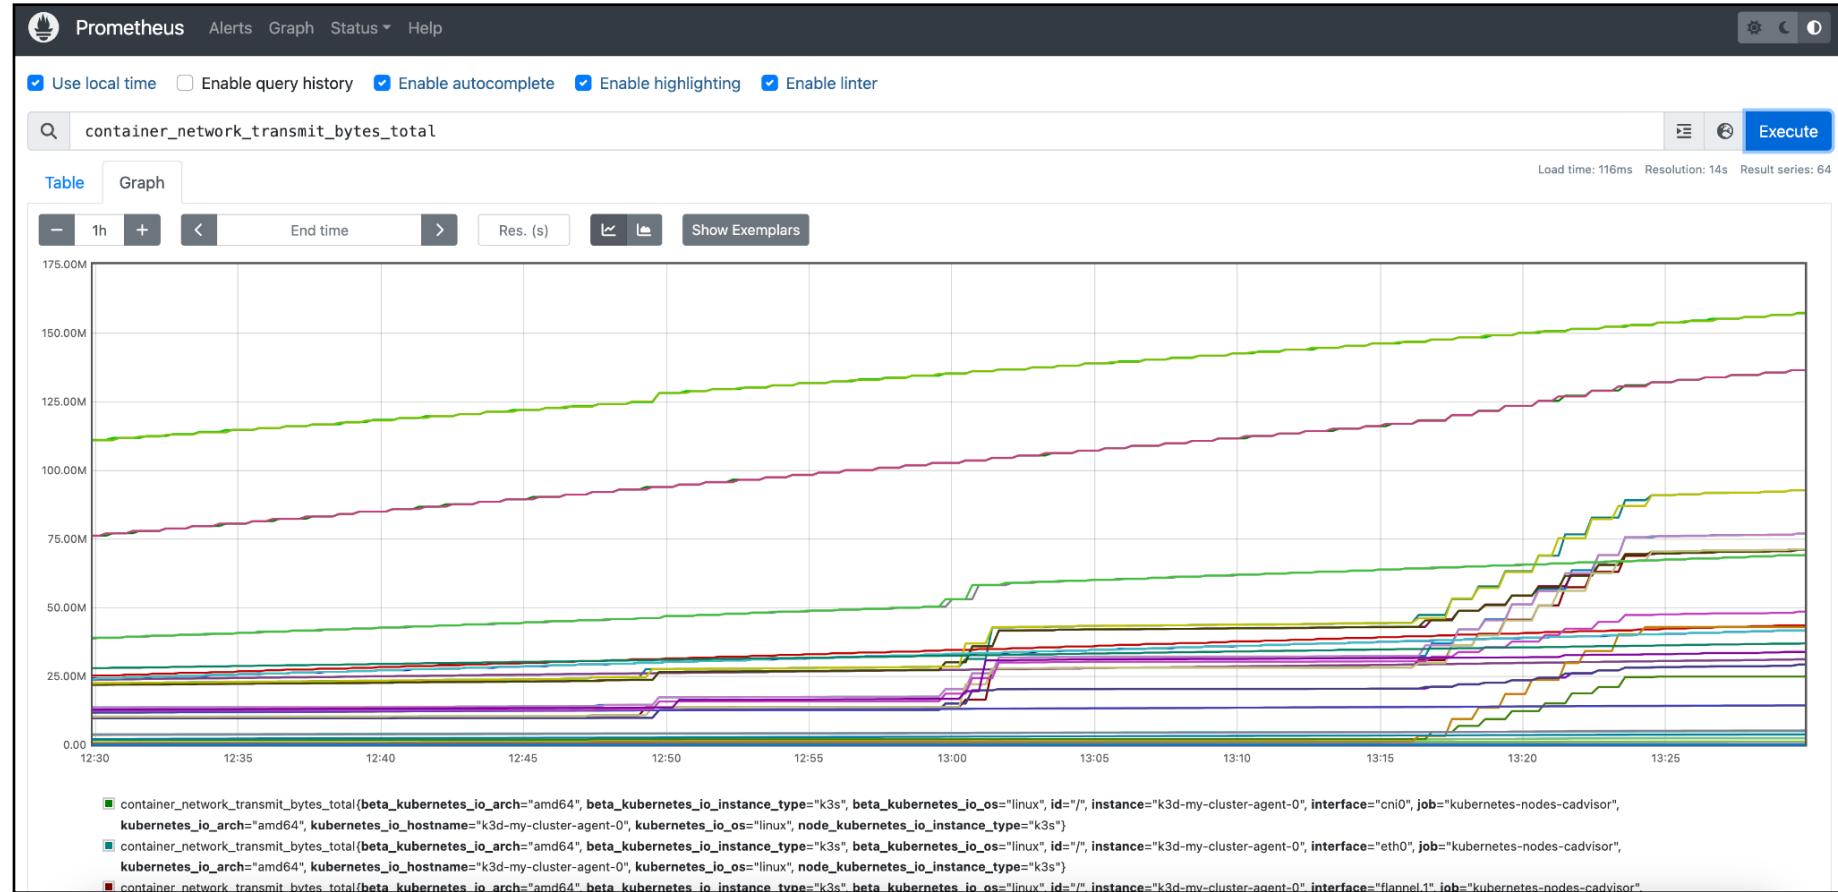
# Metrics monitoring

---

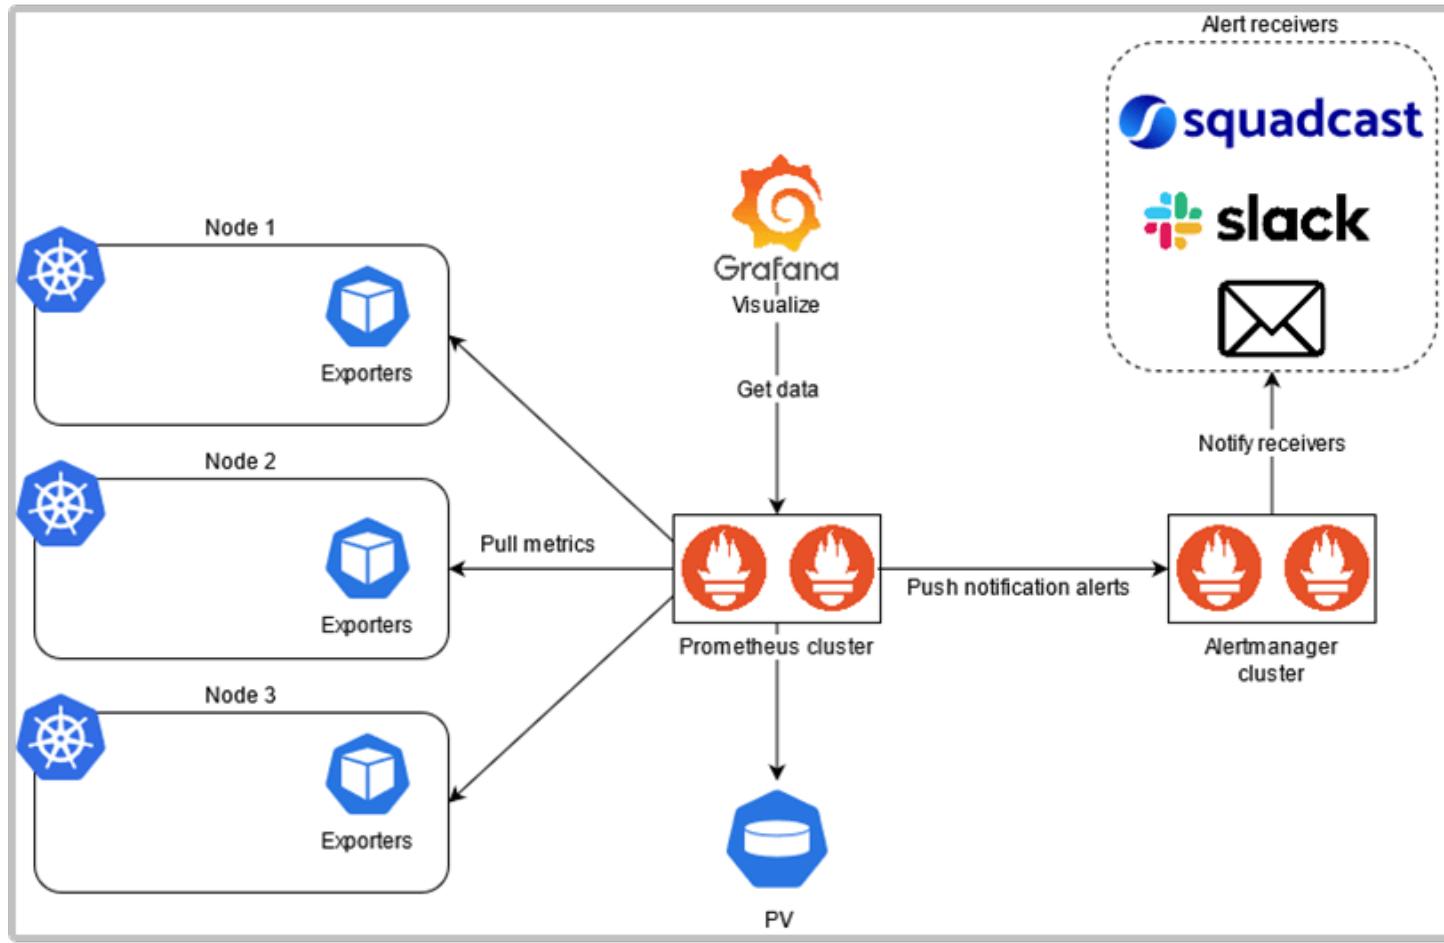
- Cluster state metrics, including the availability and health of pods
- CPU utilization
- Disk utilization
- Node status, including disk or processor overload, memory, network availability, and readiness
- Pod availability (unavailable pods may indicate poorly designed readiness probes or configuration issues)



# Prometheus



# Prometheus alert



# Prometheus alert example

---

- <https://samber.github.io/awesome-prometheus-alerts/rules>



# Grafana

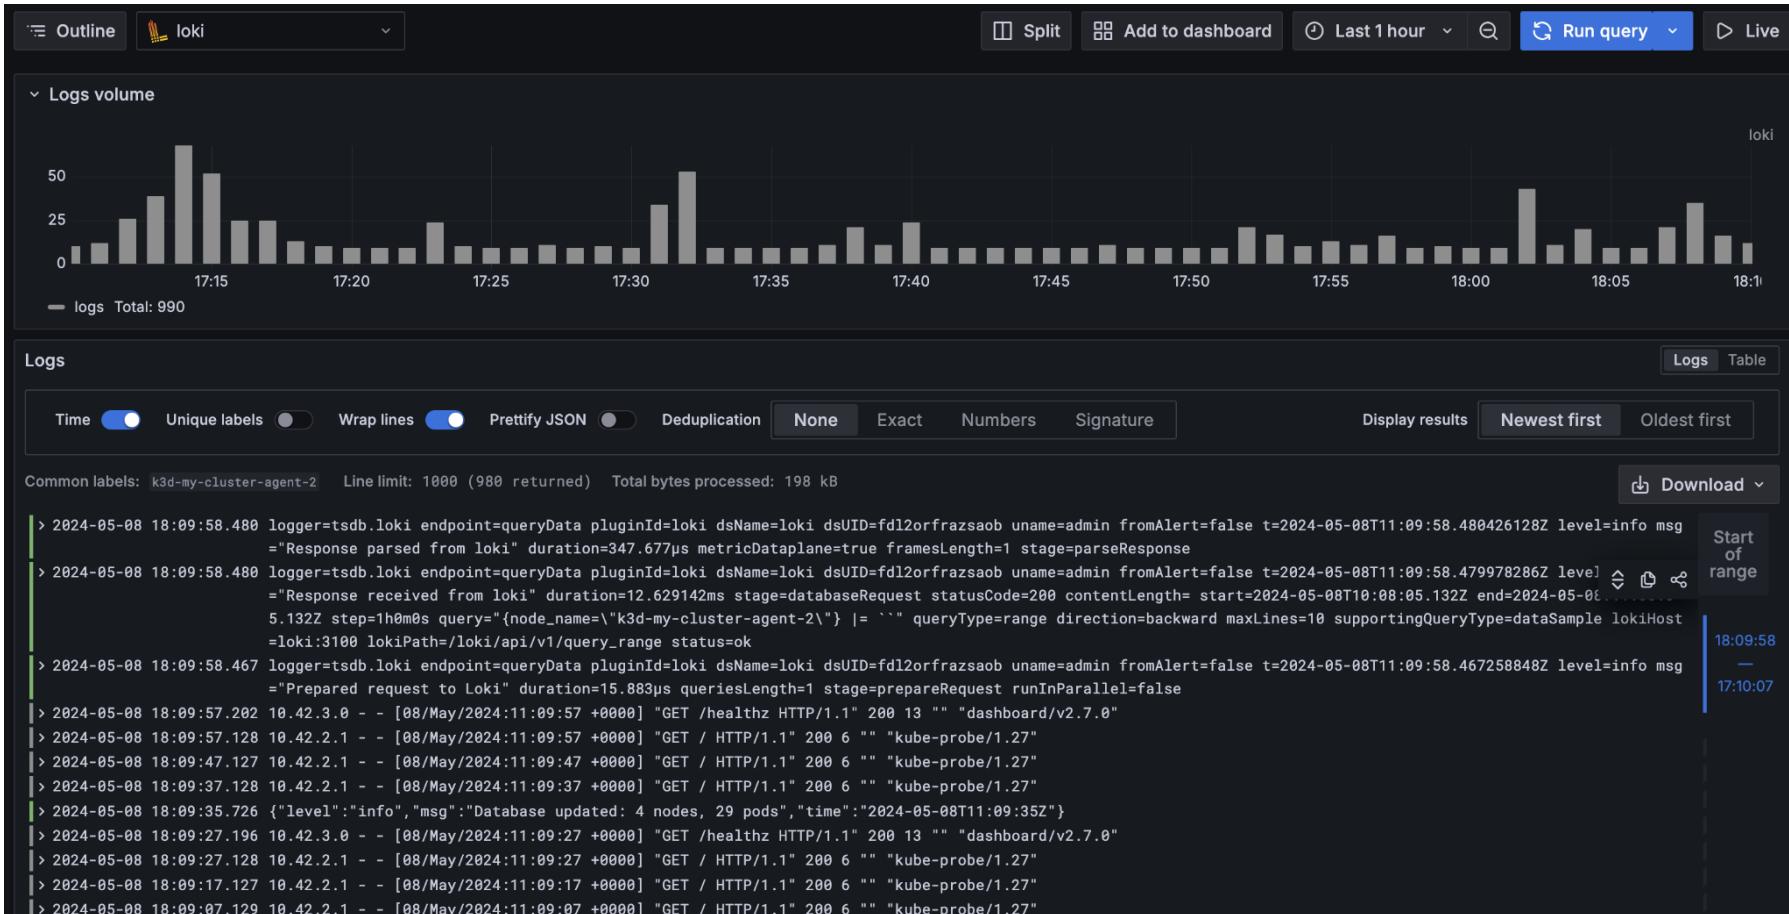


# Logs monitoring

- Cluster logs
- Node logs
- Container logs
- Kubernetes Events

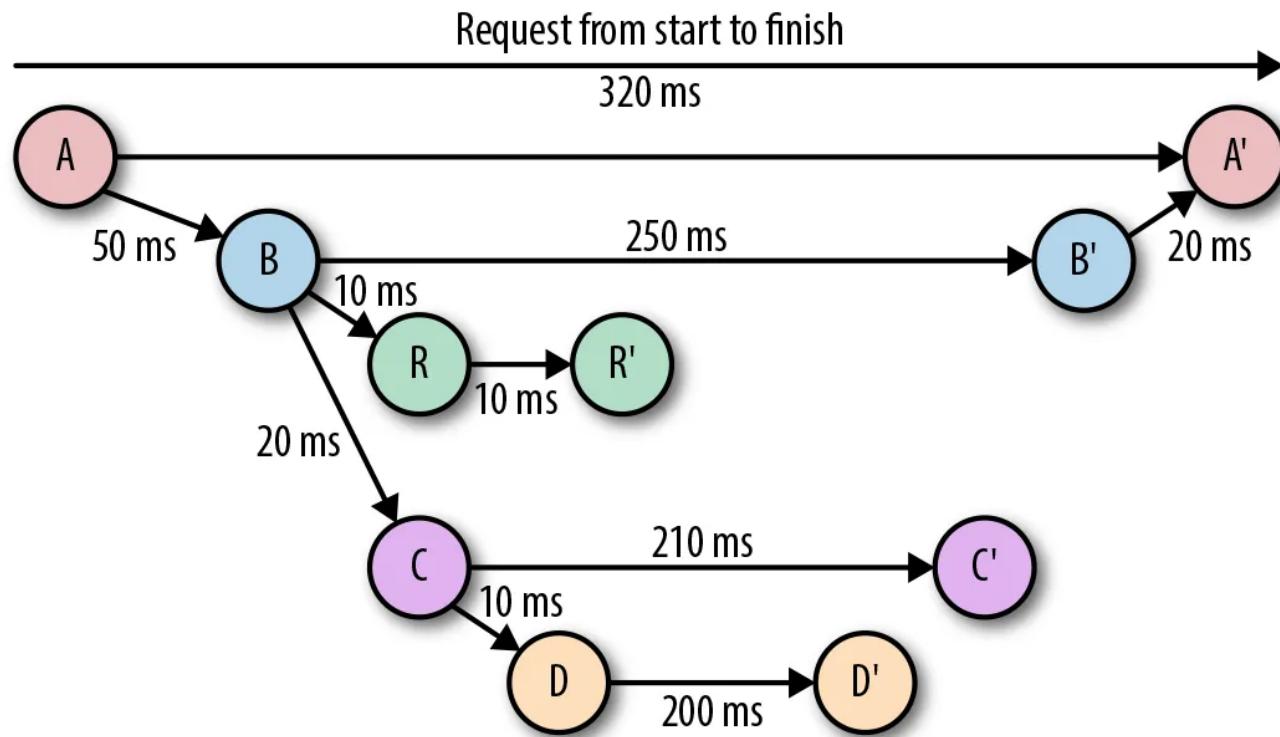


# Loki

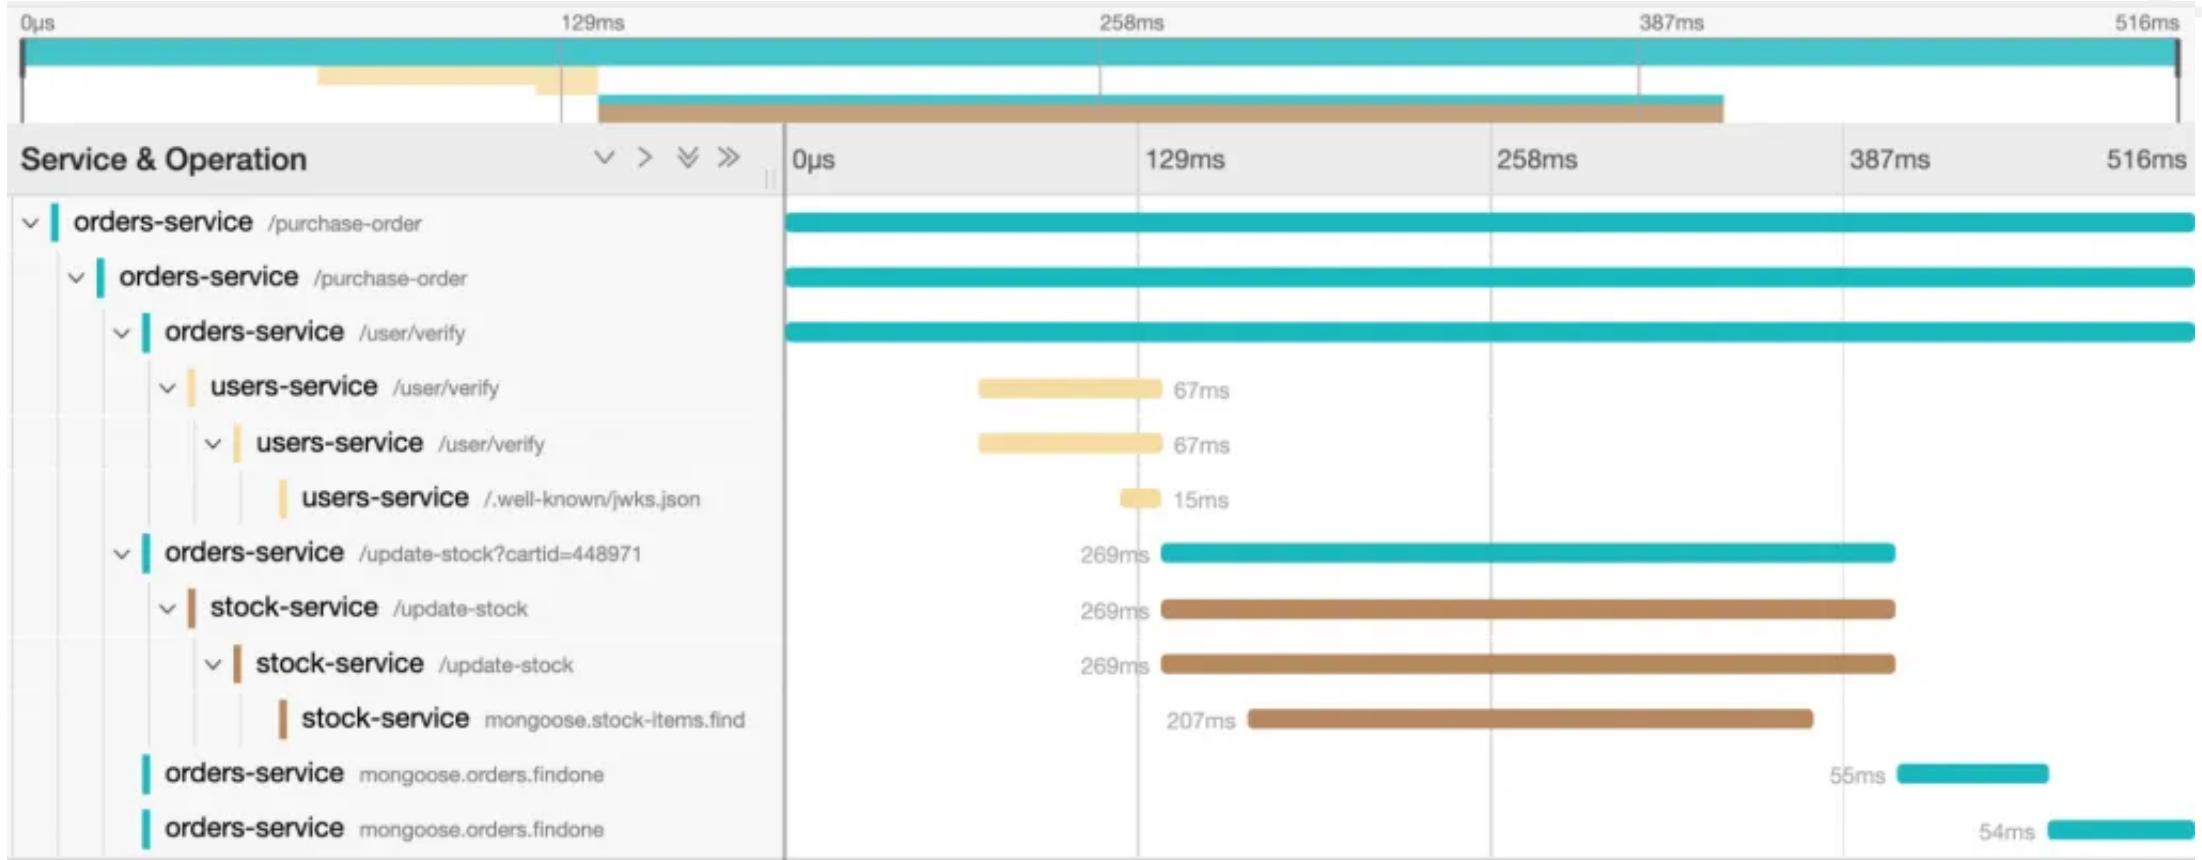


# Traces monitoring

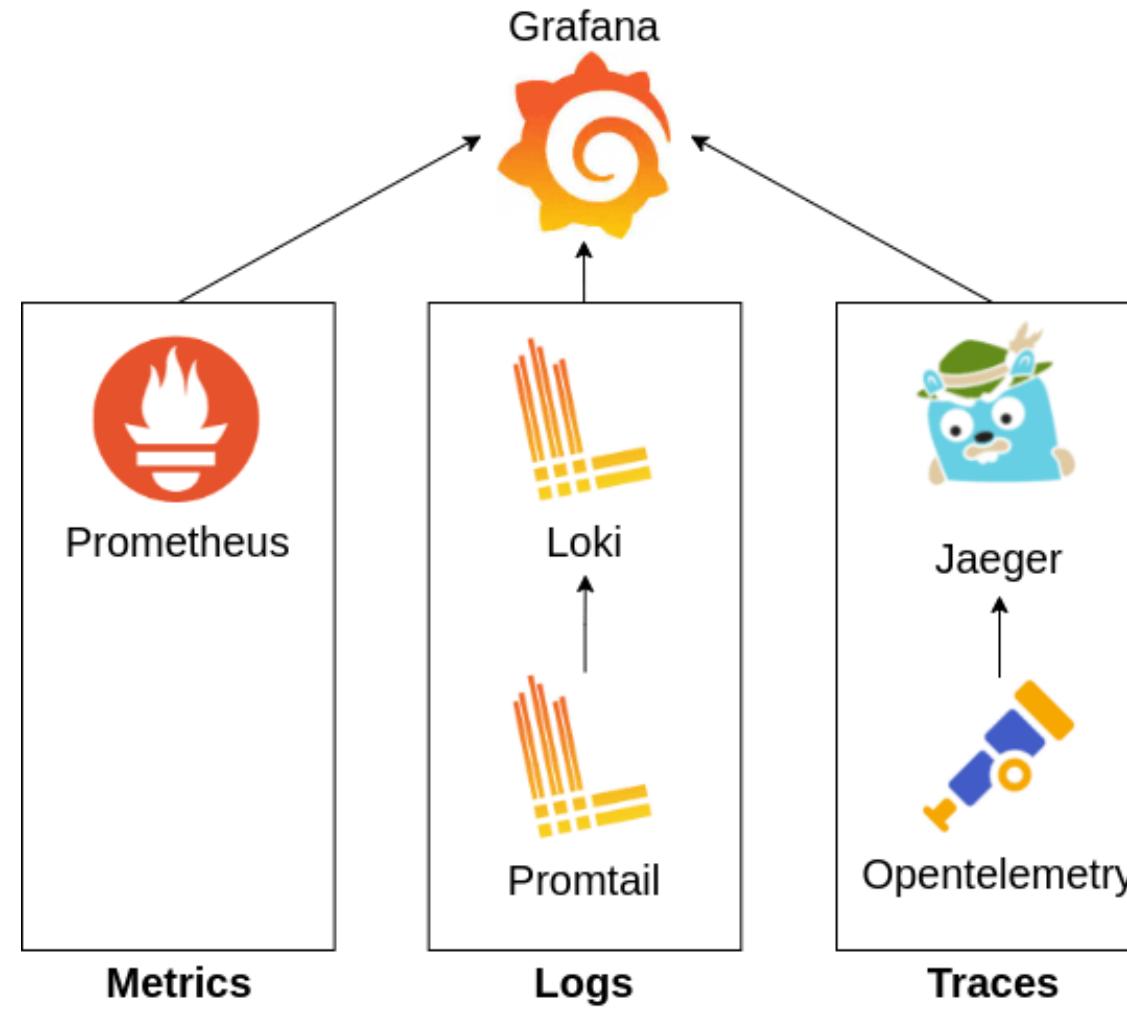
- A trace represents the entire journey of a request or action as it moves through all the nodes of a distributed system



# Jaeger



# Observability





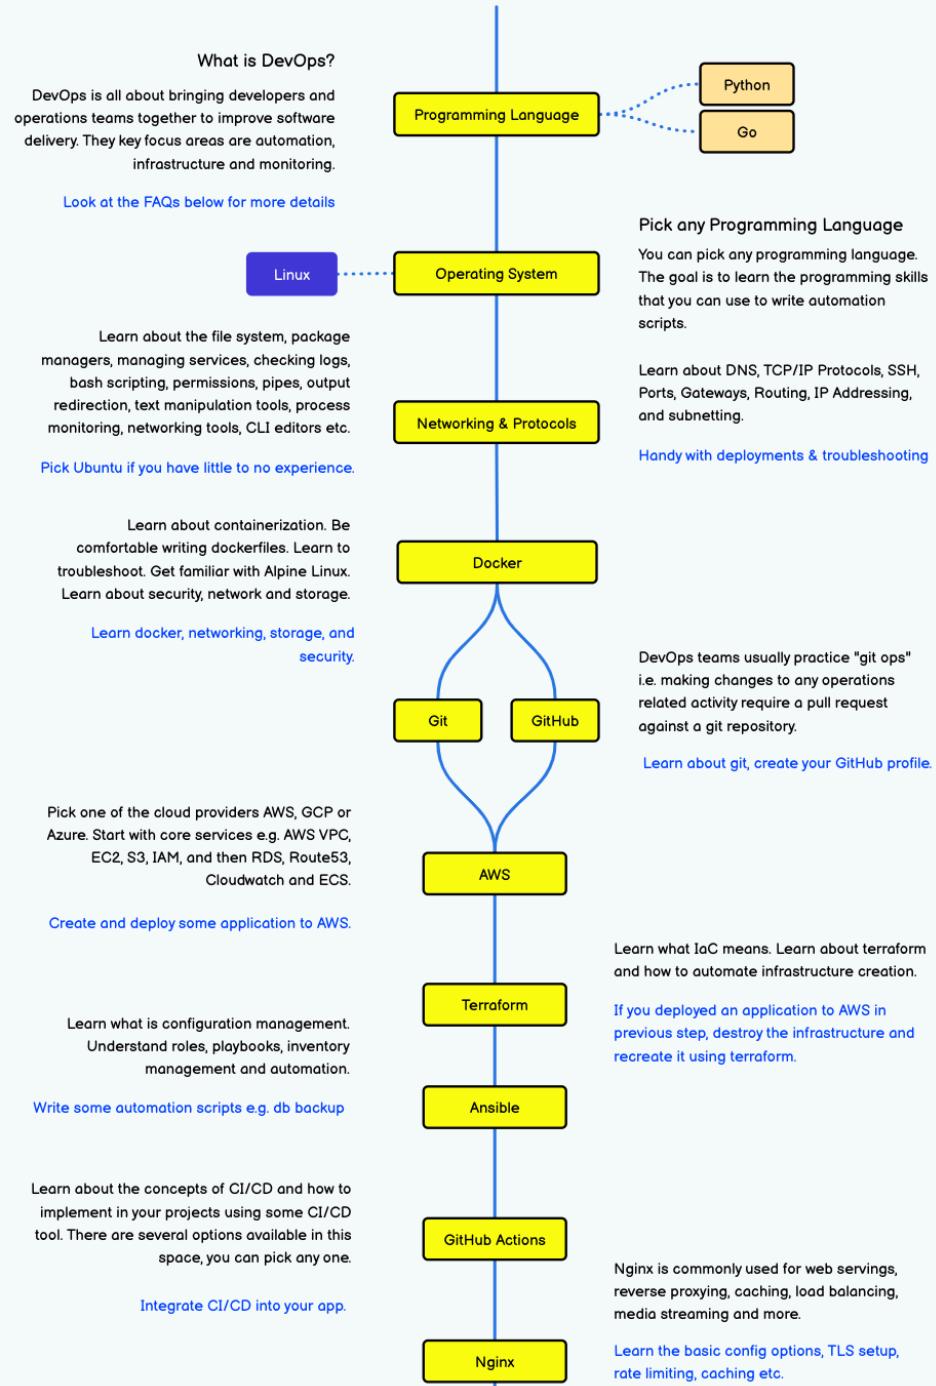
# Workshop Observability

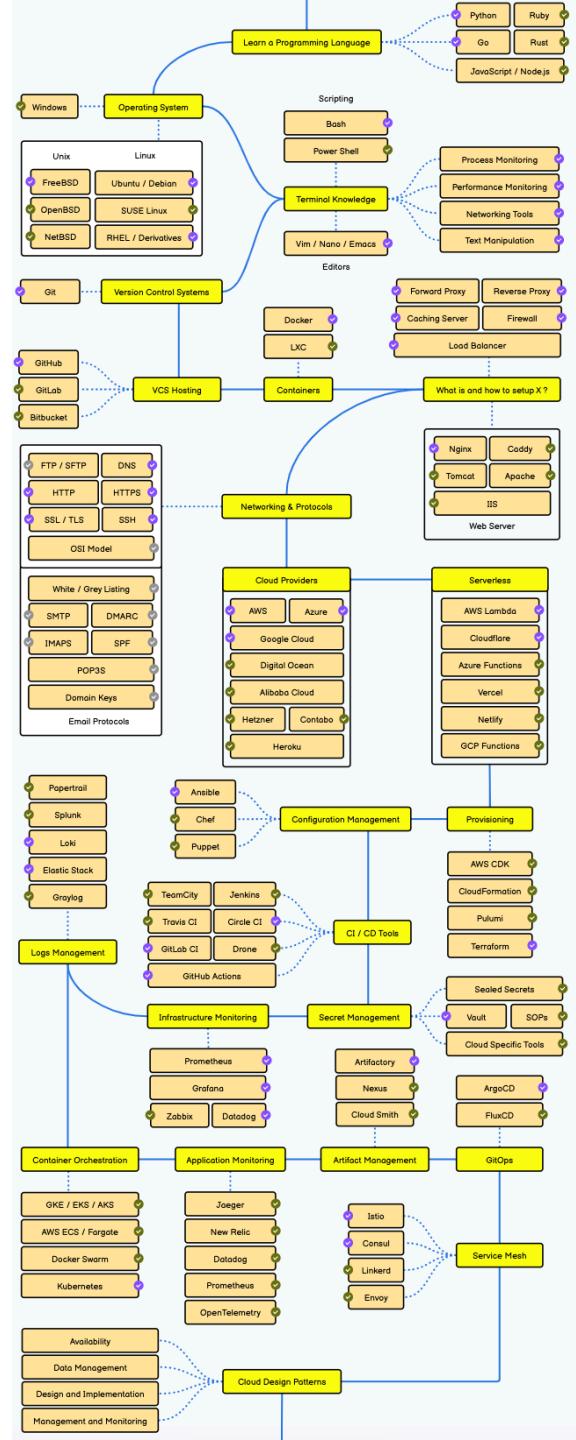
# What next in DevOps?

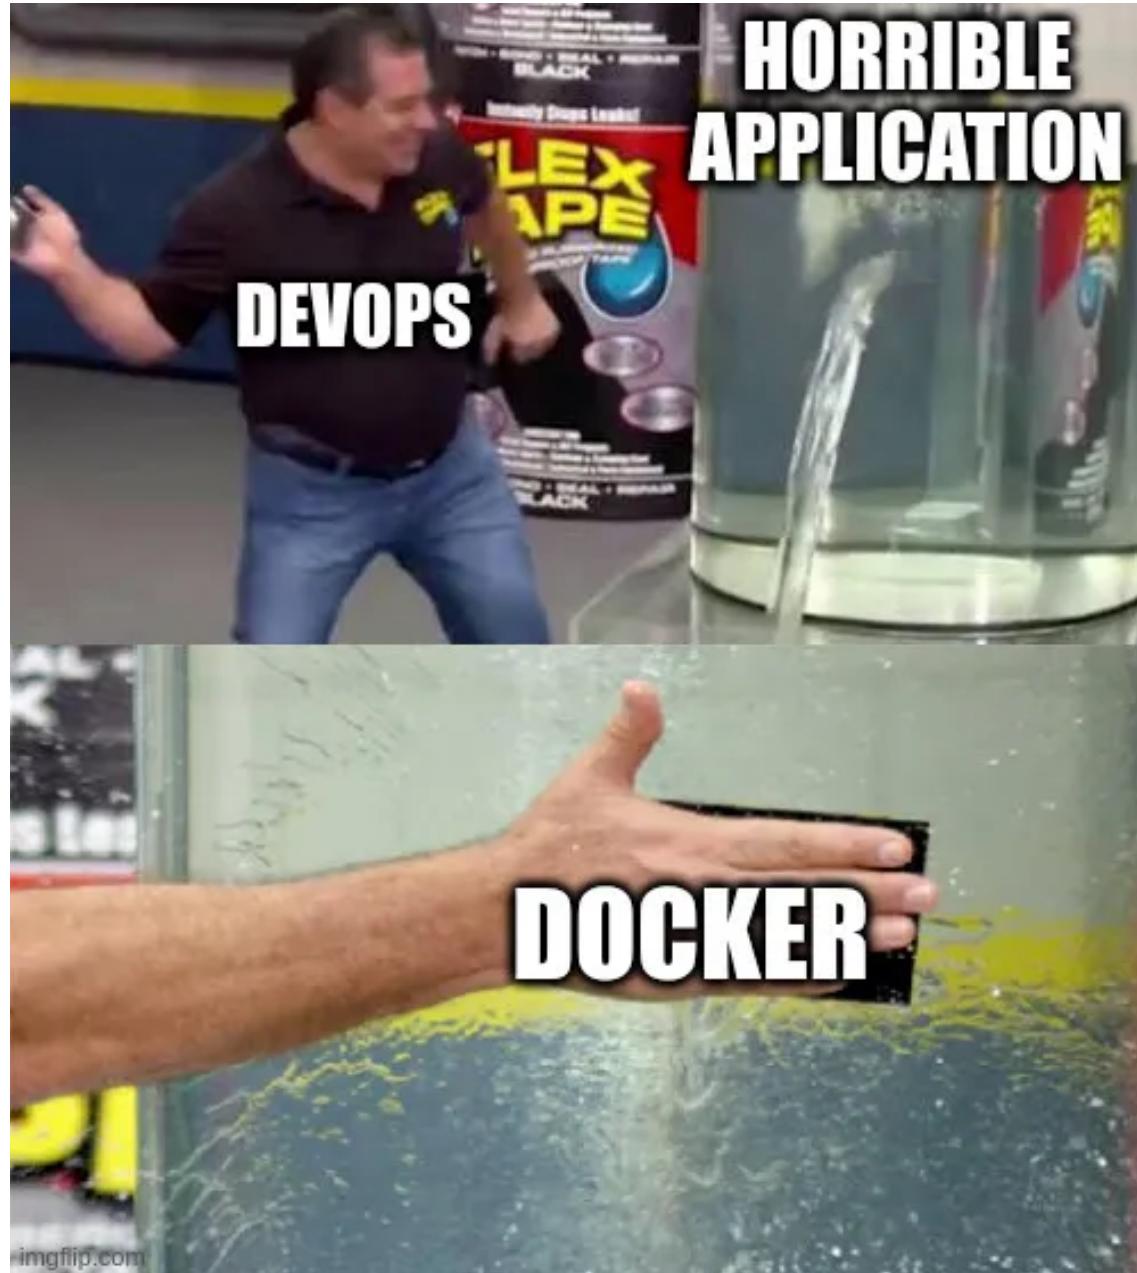
---

- DevSecOps
- GitOps
- Artificial Intelligence and Machine Learning
- Serverless Computing
- Site Reliability Engineering



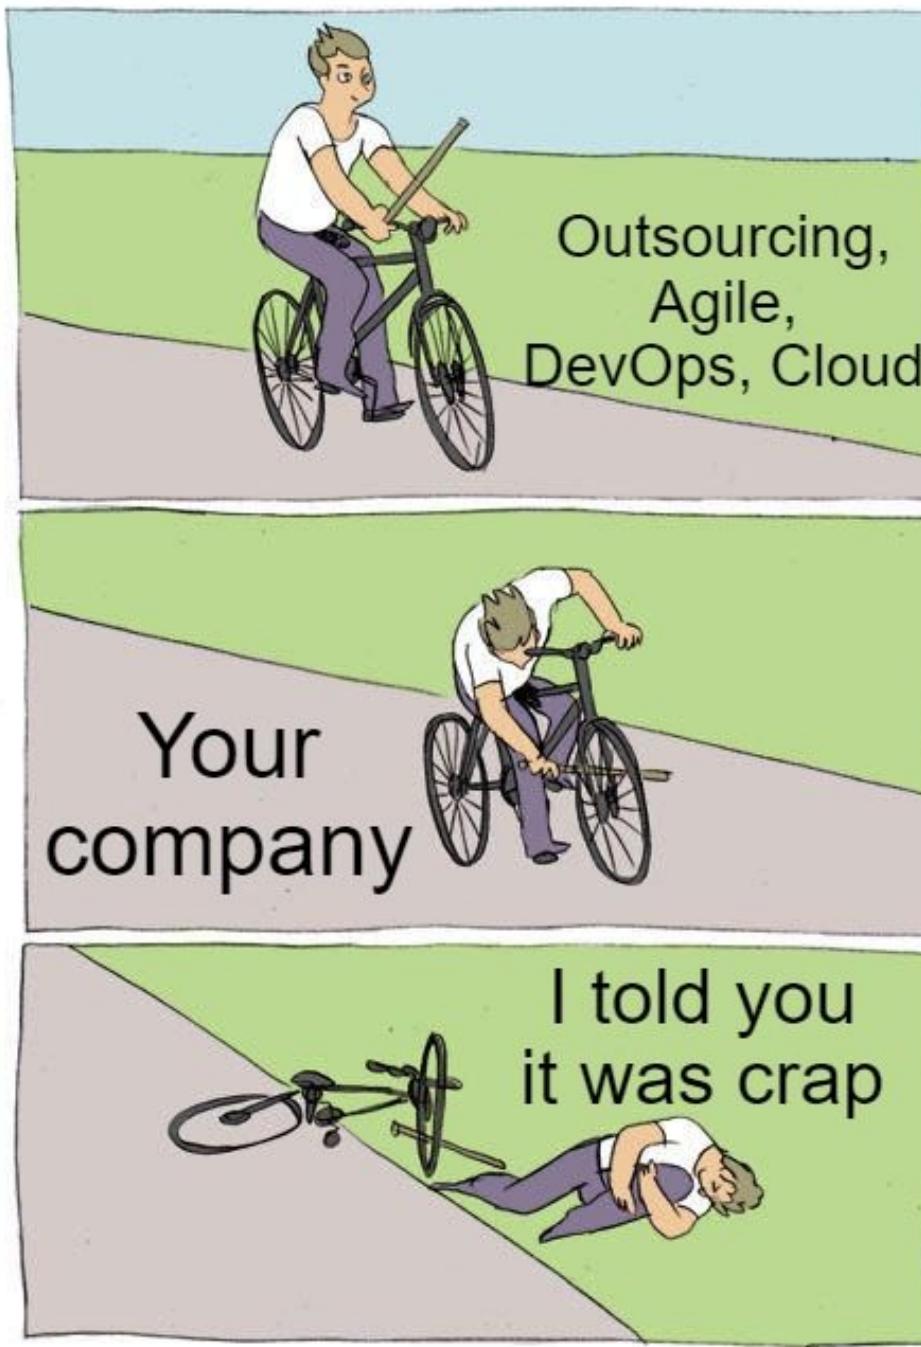






DevOps is not magic. It's about **continuous daily improvement**. If it's  
painful. You're doing it wrong. Do it again, and **bring the pain forward**





# 3 C's of Software Quality

- **Correctness:** perform its intended functions correctly and without errors or bugs
- **Completeness:** fulfills all its specified requirements. Ensures meets the needs of users
- **Consistency:** Ensures that the software is reliable and can be used in different situations.

